



Platform documentation

HOX project

Stuart J Mackintosh

Thursday 02 October 2025

At the bottom of the page, there are two large, overlapping curved shapes. The one on the left is teal and the one on the right is dark purple, matching the colors in the logo.

Contents

1	HPS system documentation	3
1.1	Overview	3
1.2	Quick start	3
1.3	System administration	3
1.4	Functions reference	3
1.5	Advanced configuration	4
1.6	Troubleshooting	4
1.7	Development	4
1.8	Appendices	4
1.9	Design decisions	4
2	Overview	5
2.1	Introduction	5
2.2	Terminology	5
2.3	Booting a host	5
2.4	First cluster setup	5
3	Quick start	6
3.1	Installation	6
3.2	Prerequisites	6
3.3	Verification	6
4	Installing HPS	7
4.1	Hardware	7
4.2	Obtaining HPS	7
4.3	Dependencies and prerequisites	7
4.4	Designing networking and numbering	7
4.5	Service verification	8
4.6	Upgrades and maintenance	8
4.7	Cluster configuration	8
4.8	Cluster creation	8
4.9	Distribution management	8
5	Deploying and configuring a cluster with nodes	9
5.1	Host profiles	9
5.2	Node provisioning	9
5.3	Service management	9

5.4	Verification	9
5.5	Deployment process	9
5.6	Purpose of the DR node	9
5.7	Failover and recovery testing	9
6	Deploying the disaster recovery node	10
6.1	Preparation	10
6.2	Synchronisation	10
7	System Administration	11
8	OpenSVC V3 Alpha Cheatsheet	12
8.1	Verified Commands & Patterns	12
8.2	Working Patterns	15
8.3	Known Limitations (V3 Alpha)	17
8.4	Help Commands	17
8.5	File Locations	17
9	HPS Storage Provisioning - System Administrator Manual	18
9.1	Overview	18
9.2	Prerequisites	18
9.3	Architecture	18
9.4	Manual Operations	18
9.5	Low-Level Operations	19
9.6	OpenSVC Integration	21
9.7	Troubleshooting	21
9.8	Manual Cleanup	22
9.9	Safety Notes	22
9.10	References	22
10	Developer documentation	24
10.1	Cluster management	24
10.2	Environment variables	25
10.3	Glossary	25
10.4	Glossary	25
11	Reference	28
11.1	Library functions	28
11.2	Paths	28
11.3	Reference configurations	28
11.4	External resources	28
11.5	Storage configuration	31
11.6	Troubleshooting	33

1 HPS system documentation

1.1 Overview

- **Project introduction** – Purpose of HPS, architecture, intended use cases.
- **Design decisions** – Records of key technical choices and rationale.
- **System components** – Description of major directories and modules.

1.2 Quick start

- **Prerequisites** – Required host operating system, packages, and network setup.
- **Installation** – Deploying `hps-container` with `hps-system` and `hps-config`.
- **First cluster setup** – Initialising a cluster with `cluster-configure.sh`.
- **Booting a host** – PXE boot process and selecting a host profile.
- **Verification** – Checking service status and logs.

1.3 System administration

- **Directory layout** – Locations for configuration, logs, distributions, and packages.
- **Cluster management** – Creating, switching, and editing clusters.
- **Host management** – Adding, removing, and updating host configurations.
- **Distribution and repository management** – Managing ISOs, PXE trees, and package repositories.
- **Service management** – Starting, stopping, and reloading `dnsmasq`, `nginx`, and `supervisord`.
- **Backup and restore** – Protecting and recovering configuration and repository data.

1.4 Functions reference

- Automatically generated from `lib/functions.d/` and other libraries.
- One function per page with purpose, arguments, usage examples, and related functions.

1.5 Advanced configuration

- **Kickstart and preseed templates** – Structure, variables, and customisation.
- **iPXE menus** – How menus are built and extended.
- **Storage provisioning** – SCH node disk detection, reporting, and configuration.
- **Integration points** – Hooks for OpenSVC, monitoring, and external systems.

1.6 Troubleshooting

- **PXE boot issues** – Common causes and fixes.
- **Service failures** – Diagnosing and restarting services.
- **Distribution and repository problems** – Checksums, GPG keys, and synchronisation errors.
- **Network problems** – DHCP conflicts, VLAN configuration, and firewall blocks.

1.7 Development

- **Code layout** – File structure and conventions for scripts and libraries.
- **Adding functions** – Naming, argument handling, logging, and documentation guidelines.
- **Testing** – Using `cli/test.sh` and other test harnesses.
- **Contributing** – Workflow, coding standards, and submission process.

1.8 Appendices

- **Glossary** – Definitions of terms and acronyms used in HPS.
- **Environment variables** – Description of exported variables and their use.
- **Decision records** – Full list of design decisions.
- **Reference configurations** – Example cluster, host, and service configuration files.

1.9 Design decisions

Stub: Summarise key design decisions. [Link to full decision records in the reference section.](#)

2 Overview

This chapter introduces the HPS system, its purpose, core architecture, and intended use cases.

It also records the major design decisions that shape the system and defines terminology used throughout.

2.1 Introduction

Stub: Provide a high-level explanation of HPS, its goals, and where it fits into the broader infrastructure platform.

2.2 Terminology

Stub: Define common terms, acronyms, and abbreviations used across the documentation.

2.3 Booting a host

Stub: PXE boot process overview and selecting a host profile.

2.4 First cluster setup

Stub: Use `cluster-configure.sh` to create the first cluster and set its parameters.

3 Quick start

A fast path to installing HPS, configuring a cluster, and booting a node.
This section covers the essentials and assumes no prior HPS experience.

3.1 Installation

Stub: Step-by-step guide to deploy `hps-container` with `hps-system` and `hps-config`.

3.2 Prerequisites

Stub: List hardware, OS, packages, network setup, and permissions needed before starting.

3.3 Verification

Stub: Confirming services are active and hosts are provisioned correctly.

4 Installing HPS

How to install the HPS system on the provisioning node, configure services, and verify readiness.

4.1 Hardware

Depends on the scenario, for example:

- evaluation
- home lab
- small organisation
- production-grade
- maximum performance

4.2 Obtaining HPS

How to acquire HPS source

- HPS Repo link
- Link to ISO downloads

4.3 Dependencies and prerequisites

- Operating system ISO's

4.4 Designing networking and numbering

HPS will manage hosts on a directly connected LAN as it uses lower-level protocols such as AMC address to manage key functions. If HPS is required on another indirectly connected network, then that should have it's own IPN.

When configuring the cluster, make sure to use a network address that doesn't conflict with anything else. HPS should offer ranges that do not conflict with anything that it can detect.

In almost every case, HPS will be implemented on a new network segment.

4.5 Service verification

Stub: Checking that dnsmasq, nginx, supervisord, and other components are running.

4.6 Upgrades and maintenance

Stub: Keeping hps-system updated without overwriting configuration files.

4.7 Cluster configuration

Stub: Setting DHCP interface, storage subnets, OS type, and other cluster settings.

4.8 Cluster creation

Stub: Running cluster-configure.sh and choosing cluster parameters.

4.9 Distribution management

Stub: Adding ISOs, extracting PXE trees, and maintaining package repositories.

5 Deploying and configuring a cluster with nodes

Creating a cluster, configuring its settings, provisioning nodes, and verifying the environment.

5.1 Host profiles

Stub: Assigning profiles such as SCH, TCH, DRH, and CCH to nodes.

5.2 Node provisioning

Stub: PXE/iPXE boot workflow and automated node configuration.

5.3 Service management

Stub: Controlling dnsmasq, nginx, supervisord, and other HPS services.

5.4 Verification

Stub: Checking that nodes are deployed correctly and services are operational.

5.5 Deployment process

Stub: Steps for installing and integrating the DR node into the cluster.

5.6 Purpose of the DR node

Stub: Explain the role of the DR node in ensuring service continuity.

5.7 Failover and recovery testing

Stub: Procedures for verifying DR readiness and simulating failover scenarios.

6 Deploying the disaster recovery node

Installing and configuring the DR node, synchronising data, and testing recovery procedures.

6.1 Preparation

Stub: Hardware, storage, and network prerequisites for the DR node.

6.2 Synchronisation

Stub: Methods for keeping DR node data in sync with the primary environment.

7 System Administration

Useful information for system administrators

8 OpenSVC V3 Alpha Cheatsheet

Version tested: v3.0.0-alpha87

8.1 Verified Commands & Patterns

8.1.1 Basic Object Management

Check version

```
om -v
```

Monitor cluster status

```
om mon
```

List nodes

```
om node ls
```

List services

```
om svc ls
```

List all objects

```
om all ls
```

8.1.2 Service Management

8.1.2.1 Create Service

Basic creation

```
om <service-name> create
```

With keywords

```
om <service-name> create --kw <section>.<key>=<value>
```

Example

```
om mysvc create --kw task#hello.type=host --kw  
↪ task#hello.command="echo hello"
```

8.1.2.2 Service Operations

View configuration

```
om <service-name> config show
```

```
# Delete service
om <service-name> delete

# View logs
om <service-name> logs

# View status
om <service-name> print status
```

8.1.3 Task Resources

8.1.3.1 Create Task

```
# Basic task
om mysvc create \
  --kw task#name.type=host \
  --kw task#name.command="<command>"

# Source bash functions and execute
om mysvc create \
  --kw task#name.type=host \
  --kw task#name.command=". /path/to/functions.sh && my_function"
```

8.1.3.2 Run Tasks

```
# Run specific task
om <service-name> run --rid task#name

# Run on all nodes
om <service-name> run --rid task#name --node=\*

# Get session ID for tracking
# Output shows: OBJECT NODE SID
```

8.1.3.3 View Task Output

```
# View logs with session filter
om <service-name> log --filter SID=<session-id>

# Or use journalctl (local only)
journalctl SID=<session-id>
```

8.1.4 Environment Variables Available in Tasks

Tasks automatically receive these environment variables:

```
OPENSVC_ACTION=run
OPENSVC_NAME=<service-name>
OPENSVC_SVCNAME=<service-name>
OPENSVC_KIND=svc
OPENSVC_SID=<session-id>
OPENSVC_ID=<object-id>
OPENSVC_LEADER=0|1
OPENSVC_SVCPATH=<service-path>
OPENSVC_RID=task#<name>
OPENSVC_NAMESPACE=root
```

8.1.5 Sync Resources

8.1.5.1 Create Sync Resource

```
# Rsync between nodes
om mysvc create \
  --kw sync#name.type=rsync \
  --kw sync#name.src="/source/path" \
  --kw sync#name.dst="/dest/path"
```

```
# Provision sync
om mysvc provision --rid sync#name
```

Note: Sync resources distribute files FROM the node running the service TO other nodes. Single-node setups will show “no target nodes”.

8.1.6 Configmaps

8.1.6.1 Create and Manage Configmaps

```
# Create configmap (note: cfg/ prefix required)
om cfg/name create
```

```
# List configmaps
om cfg ls
```

```
# Add key to configmap
om cfg/name key add --name=filename.ext --value='content'
```

```
# Add key from file
om cfg/name key add --name=filename.ext --from=/path/to/file
```

```
# List keys
om cfg/name key list
```

```
# View key content
om cfg/name key decode --name=filename.ext
```

```
# Delete configmap
om cfg/name delete
```

8.1.6.2 Configmap Limitations in V3 Alpha

- The configs parameter in tasks does NOT currently expose configmap data as environment variables or files
- Configmap data is stored base64-encoded in `/etc/opensvc/cfg/<name>.conf` under `[data]` section
- **Workaround:** Use direct file paths or inline functions instead

8.1.7 Naming Conventions

8.1.7.1 Valid Names

- Use hyphens (-), not underscores (_)
- Lowercase only
- Must comply with RFC952 (DNS naming rules)
- Examples: `my-service`, `test-functions`, `storage-mgmt`

8.1.7.2 Object Path Formats

- Services: `<name>` (no prefix)
- Configmaps: `cfg/<name>`
- Secrets: `sec/<name>` (assumed, not tested)
- Volumes: `vol/<name>` (assumed, not tested)

8.1.8 Resource Section Naming

```
# Format: <type>#<name>
task#hello
sync#files
app#webserver
disk#data
```

8.2 Working Patterns

8.2.1 Pattern 1: Simple Task Execution

```
om test create \
  --kw task#hello.type=host \
  --kw task#hello.command="echo 'Hello from task'"

om test run --rid task#hello
```


8.2.2 Pattern 2: Tasks with Bash Functions

```
# Create functions file on node(s)
cat > /opt/opensvc/functions.sh << 'EOF'
#!/bin/bash
my_function() {
    echo "Output from function"
}
EOF

# Create service
om mysvc create \
  --kw task#run.type=host \
  --kw task#run.command=". /opt/opensvc/functions.sh &&
    ↪ my_function"

om mysvc run --rid task#run
```

8.2.3 Pattern 3: Multi-Node Execution

```
# Create service with multiple nodes
om cluster-task create \
  --kw nodes="node1,node2,node3" \
  --kw task#check.type=host \
  --kw task#check.command="hostname; df -h"

# Run on all nodes
om cluster-task run --rid task#check --node=\*

# Returns SIDs for each node
# View specific node output
om cluster-task log --filter SID=<session-id>
```

8.2.4 Pattern 4: File Distribution (When Multiple Nodes Exist)

```
# Create sync resource
om distribute create \
  --kw nodes="node1,node2" \
  --kw sync#files.type=rsync \
  --kw sync#files.src="/local/file.sh" \
  --kw sync#files.dst="/remote/file.sh"

# Provision to distribute
om distribute provision --rid sync#files

# Then use in tasks
om distribute create --kw task#run.type=host \
  --kw task#run.command=". /remote/file.sh && function_name"
```

8.3 Known Limitations (V3 Alpha)

1. **Configmap Exposure:** `configs` parameter doesn't expose data to tasks
2. **Single Node Sync:** Sync resources need multiple nodes to function
3. **No Service Prefix:** Services use bare names, not `svc/<name>`
4. **Incomplete Documentation:** Many features undocumented in alpha

8.4 Help Commands

General help

```
om --help
```

Subsystem help

```
om svc --help
```

```
om cfg --help
```

```
om node --help
```

Command-specific help

```
om svc create --help
```

```
om cfg key add --help
```

8.5 File Locations

Service configs

```
/etc/opensvc/<service-name>.conf
```

Configmap configs

```
/etc/opensvc/cfg/<name>.conf
```

Service runtime data

```
/var/lib/opensvc/svc/<service-name>/
```

Cluster config

```
/etc/opensvc/cluster.conf
```

```
/etc/opensvc/node.conf
```

Note: This cheatsheet is based on V3 alpha87 testing. Features and syntax may change before GA release.

9 HPS Storage Provisioning - System Administrator Manual

9.1 Overview

This manual documents the storage provisioning system for debugging purposes. Under normal operation, OpenSVC handles all storage provisioning automatically. This guide is for system administrators who need to manually execute storage operations for troubleshooting or testing.

9.2 Prerequisites

All operations require the node functions to be sourced first:

```
# Source the functions library
. /srv/hps/lib/node_functions.sh
```

9.3 Architecture

9.3.1 Node Types

- **TCH** (Thin Compute Host) - Virtual machine hosts that request storage
- **SCH** (Storage Cluster Host) - Physical storage nodes that provide zvol/iSCSI resources

9.3.2 Components

- **ZFS zvols** - Block devices for storage volumes
- **LIO/iSCSI** - Network block device targets
- **OpenSVC** - Cluster orchestration layer

9.4 Manual Operations

9.4.1 1. Check Available Storage Capacity

Query available space on local storage pool:

```
# Get available bytes
storage_get_available_space
```

```
# Example output: 51504332800 (approximately 48GB)
```

Convert human-readable sizes:

```
# Parse capacity string to bytes
storage_parse_capacity "100G"
```

```
# Example output: 107374182400
```

9.4.2 2. Provision a Storage Volume

Create a complete storage volume with zvol and iSCSI target:

```
storage_provision_volume \
  --iqn iqn.2025-09.local.hps:vm-disk-001 \
  --capacity 100G \
  --zvol-name vm-disk-001
```

What happens: 1. Validates host type is SCH 2. Retrieves local zpool name 3. Checks available capacity 4. Creates ZFS zvol 5. Creates iSCSI target with backstore 6. Configures LUN mapping 7. Sets up demo mode (no authentication)

Requirements: - Must run on SCH host - Sufficient capacity in zpool - Valid IQN format - Unique zvol name

9.4.3 3. Remove a Storage Volume

Delete both iSCSI target and underlying zvol:

```
storage_deprovision_volume \
  --iqn iqn.2025-09.local.hps:vm-disk-001 \
  --zvol-name vm-disk-001
```

What happens: 1. Validates host type is SCH 2. Deletes iSCSI target configuration 3. Removes backstore 4. Destroys ZFS zvol

9.5 Low-Level Operations

9.5.1 ZFS Volume Management

Direct zvol operations via the storage manager:

```
# Create zvol
node_storage_manager zvol create \
  --pool ztest-pool \
  --name my-volume \
  --size 50G
```

```
# Delete zvol
node_storage_manager zvol delete \
  --pool ztest-pool \
  --name my-volume

# List zvols in pool
node_storage_manager zvol list --pool ztest-pool

# Check if zvol exists
node_storage_manager zvol check \
  --pool ztest-pool \
  --name my-volume

# Get zvol information
node_storage_manager zvol info \
  --pool ztest-pool \
  --name my-volume
```

9.5.2 iSCSI Target Management

Direct LIO/targetcli operations:

```
# Start target service
node_storage_manager lio start

# Stop target service
node_storage_manager lio stop

# Check service status
node_storage_manager lio status

# Create iSCSI target
node_storage_manager lio create \
  --iqn iqn.2025-09.local.hps:target-name \
  --device /dev/zvol/pool/volume

# Create with ACL (optional)
node_storage_manager lio create \
  --iqn iqn.2025-09.local.hps:target-name \
  --device /dev/zvol/pool/volume \
  --acl iqn.2025-09.initiator:client1

# Delete iSCSI target
node_storage_manager lio delete \
  --iqn iqn.2025-09.local.hps:target-name

# List all targets
node_storage_manager lio list
```

9.6 OpenSVC Integration

9.6.1 Service Structure

The storage-provision service provides two tasks:

Check capacity on all storage nodes

```
om storage-provision run --rid task#check-capacity --node=\*
```

Provision on specific node

```
om storage-provision instance run \  
  --rid task#provision \  
  --node=SCH-001 \  
  --env IQN=iqn.2025-09.local.hps:vm-disk-001 \  
  --env CAPACITY=100G \  
  --env VOLNAME=vm-disk-001
```

9.6.2 Viewing Logs

Get session ID from run output

```
om storage-provision run --rid task#check-capacity --node=\*
```

Output shows: OBJECT NODE SID

View logs for specific session

```
om storage-provision log --filter SID=<session-id>
```

9.7 Troubleshooting

9.7.1 Common Issues

“ERROR: This host type is ‘XXX’, not ‘SCH’” - Storage operations only allowed on Storage Cluster Hosts - Verify: `remote_host_variable TYPE`

“ERROR: Insufficient space” - Requested capacity exceeds available space - Check: `storage_get_available_space` - Reduce capacity or free up space

“Zvol already exists” - Volume name conflict - List existing: `node_storage_manager zvol list` - Choose different name or delete existing

“Failed to create backstore” - Backstore name already in use - List: `node_storage_manager lio list` - Delete conflicting target first

9.7.2 Debugging Steps

1. **Verify host configuration:**

```
remote_host_variable TYPE  
remote_host_variable ZPOOL_NAME
```

2. **Check available resources:**

```
storage_get_available_space
zpool list
node_storage_manager zvol list
node_storage_manager lio list
```

3. Test connectivity:

```
node_storage_manager lio status
systemctl status target
```

4. Review logs:

```
journalctl -u target -n 50
tail -f /var/log/messages
```

9.8 Manual Cleanup

If automated cleanup fails, manually remove resources:

```
# 1. Delete iSCSI target
```

```
targetcli /iscsi delete iqn.2025-09.local.hps:target-name
```

```
# 2. Delete backstore
```

```
targetcli /backstores/block delete backstore-name
```

```
# 3. Save configuration
```

```
targetcli saveconfig
```

```
# 4. Delete zvol
```

```
zfs destroy pool-name/volume-name
```

```
# 5. Verify cleanup
```

```
zfs list -t volume
```

```
targetcli ls
```

9.9 Safety Notes

- Always verify host type before provisioning
- Check capacity before creating large volumes
- Ensure unique IQN and volume names
- Use deprovision function for proper cleanup
- Monitor zpool space regularly

9.10 References

- Functions location: /srv/hps/lib/node_functions.sh
- OpenSVC service: storage-provision

- Target config: `/etc/target/saveconfig.json`
- Logs: `journalctl` and `om storage-provision log`

10 Developer documentation

This section provides support for someone who wants to customise their own HPS or contribute to the HPS development.

10.0.1 Load the HPS host libraries

to source the functions, run this on a node:

```
bash <(curl -fsSL "http://10.99.1.1/cgi-bin/boot_manager.sh?cmd=bootstrap_initialise_distro")
```

this calls a remote script which pulls down the functions, that are then sourcable

10.1 Cluster management

Clusters are managed by OpenSVC.

The central config file is generated on the IPS and downloaded on demand by cluster hosts. It is dynamically built based on the cluster config.

10.1.1 OpenSVC

we are using v3

Note:

Docs are incomplete.

- V3 docs: <https://book.opensvc.com/a>
- V2 docs: <https://docs.opensvc.com/latest/>

Thing you want to set	Where / How
Agent log path/level	opensvc.conf (log_file, log_level)
Agent TCP listener / Web UI ports	opensvc.conf (listener_port, web_ui*)

Thing you want to set	Where / How
Node local tags for default behavior	<code>opensvc.conf</code> (tags)
Cluster members / node IPs / names	<code>cluster.conf</code>
Service resources (zpool/zvol/f-s/ip/share)	<code>om ... create/set</code> → lives under <code>services/*</code>
Placement rules (tags=storage, nodename=...)	<code>om mysvc set --kw placement=...</code> (in service <code>cfg</code>)
Start/stop/provision services	<code>'om mysvc start stopprovision'</code>
Distribute service configs to other nodes	<code>om mysvc push / om mysvc sync</code>

10.1.1.1 Useful commands:

= the defined service name

om print config Prints the config for the service

om config validate Checks the config for the node and reports on errors

om purge `purge`, `unprovision` and `delete` are asynchronous and do things on all node with a object instance

10.1.1.2 References

10.2 Environment variables

Stub: Explanation of exported variables and their purpose.

10.3 Glossary

10.4 Glossary

Btrfs A modern Linux file system with support for snapshots, pooling, and checksumming. Considered for HPS storage but not selected due to lack of native block

device export.

CCH (Compute Cluster Host) A host profile type in HPS representing a node dedicated to compute workloads.

CIDR (Classless Inter-Domain Routing) A method for allocating IP addresses and routing, using a prefix length (e.g. /24) to indicate the network mask.

DHCP (Dynamic Host Configuration Protocol) Network protocol that automatically assigns IP addresses and other network settings to devices on a network.

DHCP interface The network interface on which HPS's DHCP service (dnsmasq) listens to respond to PXE boot and other client requests.

DHCP range The range of IP addresses available for assignment via DHCP in a given network segment.

DHCP reservation A mapping of a specific MAC address to a fixed IP address in the DHCP server configuration.

DHCP server A service that hands out IP addresses and network configuration to clients; in HPS, typically provided by dnsmasq.

DHCP subnet The network segment configuration for DHCP, usually defined by IP address and netmask/CIDR.

DHCP static lease A lease configuration mapping a specific client to a specific IP, without dynamic changes.

DHCP options Additional configuration data sent by DHCP server to clients (e.g., boot file name, domain name, DNS servers).

DHCP vendor class identifier A DHCP field that can identify the client's hardware or software type.

DHCP relay A service that forwards DHCP requests from clients in one network to a DHCP server in another network.

DHCP snooping A network switch feature that limits DHCP responses to trusted ports.

DHCP starvation An attack in which an attacker sends repeated DHCP requests to exhaust the available address pool.

DHCPDISCOVER DHCP message type sent by clients to find available DHCP servers.

DHCPOFFER DHCP message type sent by servers in response to a DHCPDISCOVER, offering an IP configuration.

DHCPREQUEST DHCP message type sent by clients to request offered configuration.

DHCPACK DHCP message type sent by the server to confirm an IP lease to the client.

dnsmasq Lightweight DNS forwarder and DHCP server used by HPS to provide PXE boot and local DNS services.

DR node (Disaster Recovery node) A dedicated node used to provide failover capability and data recovery in the event of a primary node failure.

DRH (Disaster Recovery Host) Host profile type in HPS representing a disaster recovery node.

HPS (Host Provisioning Service) The provisioning framework implemented in this project for automated PXE-based OS deployment and configuration.

HPS container The Docker container providing the HPS services.

HPS config The configuration directory structure for HPS, stored separately from the

core scripts to allow upgrades without overwriting site-specific settings.

HPS system The set of Bash scripts, functions, and service configurations that implement the HPS provisioning environment.

iPXE Open-source network boot firmware supporting protocols such as HTTP, iSCSI, and PXE. Used by HPS for dynamic boot menus and provisioning.

ISO (International Organization for Standardization image file) A disk image format commonly used to distribute operating system installation media.

Kickstart Automated installation method used by Red Hat-based distributions, configured via a `.ks` file.

MAC address Media Access Control address, a unique identifier assigned to a network interface.

NFS (Network File System) Protocol for sharing files over a network, not used for boot in HPS but sometimes relevant for Linux provisioning.

PXE (Preboot eXecution Environment) Network boot framework that allows a system to boot from a network interface before an OS is installed.

PXE boot menu The interactive menu shown to PXE/iPXE clients to select a boot option.

SCH (Storage Cluster Host) Host profile type in HPS representing a node dedicated to storage services.

syslog Standardised system logging protocol used to collect logs from services.

TCH (Thin Compute Host) Host profile type in HPS representing a lightweight compute node.

TFTP (Trivial File Transfer Protocol) Simple file transfer protocol used in PXE boot to transfer bootloaders and configuration.

ZFS Advanced file system with volume management, snapshots, and data integrity features. Selected in HPS for iSCSI exports due to native zvol support.

11 Reference

Static technical reference information for HPS, including function documentation, troubleshooting guides, and configuration details.

11.1 Library functions

There are two main libraries of functions, the hps functions, and host functions.

hps functions are used during the cluster build and configure process whereas the host functions are available on the running host.

11.2 Paths

11.2.1 hps provisioning node

11.2.2 hps operating node (TCN, CCN etc)

Storage under /srv/storage

scripts under /srv/scripts

/srv is iscsi mounted

everything else is transient.

11.3 Reference configurations

Stub: Example `cluster.conf`, `host.conf`, and service configuration files.

11.4 External resources

Below are external projects, code repositories, and documentation sources that are relevant to HPS.

These provide additional background, tooling, or dependencies that are either directly used or referenced in the design.

11.4.1 Btrfs

Link: <https://btrfs.readthedocs.io>

Summary: Linux file system with features such as snapshots, compression, and sub-volumes. Considered for HPS storage but not selected as the primary iSCSI export filesystem.

11.4.2 dnsmasq

Link: <http://www.thekelleys.org.uk/dnsmasq/doc.html>

Summary: Lightweight DNS forwarder and DHCP server used in HPS to provide PXE boot services, static leases, and local DNS.

11.4.3 Docker

Link: <https://docs.docker.com>

Summary: Containerisation platform used to run the HPS environment (hps-container) in a controlled, reproducible manner.

11.4.4 iPXE

Link: <https://ipxe.org>

Summary: Open-source network boot firmware supporting PXE, HTTP, iSCSI, and scripting. HPS uses iPXE binaries (`ipxe.efi`, `undionly.kpxe`, `snponly.efi`) for boot menus and network installs.

11.4.5 ISO images for Rocky Linux

Link: <https://download.rockylinux.org/pub/rocky/>

Summary: Official Rocky Linux distribution media. HPS uses these ISOs to populate PXE boot trees and perform installations.

11.4.6 Kickstart documentation

Link: <https://pykickstart.readthedocs.io>

Summary: Red Hat-based distributions' automated installation system. Kickstart files are used in HPS to perform unattended OS deployments.

11.4.7 OpenSVC

Link: <https://www.opensvc.com>

Summary: Cluster resource manager and service orchestrator considered for integration with HPS for managing ZFS-backed iSCSI storage and service failover.

11.4.8 Pandoc

Link: <https://pandoc.org>

Summary: Document converter used to compile HPS Markdown documentation into PDF, HTML, and other formats.

11.4.9 PXE specification

Link: https://en.wikipedia.org/wiki/Preboot_Execution_Environment

Summary: Standard network boot process for x86 systems. HPS extends PXE with iPXE for additional protocol support and boot menu scripting.

11.4.10 Rocky Linux

Link: <https://rockylinux.org>

Summary: Enterprise-grade Linux distribution used as a primary OS target in HPS deployments.

11.4.11 syslog protocol (RFC 5424)

Link: <https://datatracker.ietf.org/doc/html/rfc5424>

Summary: Standard for message logging in IP networks. HPS services can log via syslog for centralised collection.

11.4.12 TFTP

Link: <https://datatracker.ietf.org/doc/html/rfc1350>

Summary: Simple file transfer protocol used to serve PXE/iPXE bootloaders and configuration files to network boot clients.

11.4.13 ZFS on Linux (OpenZFS)

Link: <https://openzfs.org>

Summary: Advanced file system and volume manager selected for HPS storage exports due to native block device (zvol) support, snapshots, and data integrity features.

11.5 Storage configuration

This section describes how storage is provisioned and managed within HPS.

It covers the role of ZFS, iSCSI exports, and the functions used to configure storage cluster hosts (SCHs) and thin compute nodes (TCNs).

11.5.1 Overview

HPS provides storage to diskless compute nodes via **ZFS-backed iSCSI exports**.

Rather than replicating data at the storage back-end, redundancy is achieved through **client-side RAID** on the TCNs. This design keeps the storage back-end simple (“just a bunch of block devices”) and allows flexible RAID layouts on the client.

11.5.2 Functions and building blocks

Several library functions are provided to initialise and manage storage nodes:

- **remote_log**
Sends log output from a remote node back to the provisioning system.
- **remote_cluster_variable** and **remote_host_variable**
Used to read and update cluster-level and host-level configuration variables across the environment.
- **initialise_opensvc_cluster**
Prepares an OpenSVC cluster context for managing storage resources.
- **load_opensvc_conf**
Loads the OpenSVC configuration to ensure cluster services are consistent across nodes.
- **ZFS install and configure functions**
Automate installation of ZFS, creation of zpools, and export of zvols for use as iSCSI devices.

Additional functions can be added over time. These are distributed to remote nodes (e.g. TCNs and SCHs) and sourced locally so they can operate with the same provisioning logic as the HPS controller.

11.5.3 Host configuration variables

The iSCSI devices used by each thin compute node are defined in the host configuration. Two key variables are:

- **ISCSI_ROOT_0**

- **ISCSI_ROOT_1**

These identify the iSCSI targets that provide the root disks for the TCN.

When a TCN is first configured, these variables are created and stored in its host config.

Future expansions will allow additional variables such as **ISCSI_DATA_N** to define data disks for workloads running on the TCN.

11.5.4 Provisioning workflow

1. **SCH setup**

Storage cluster hosts are provisioned with ZFS and OpenSVC. ZFS zpools and zvols are created as needed for iSCSI targets.

2. **TCN request**

When a thin compute node is being installed, the provisioning system determines its storage requirements and allocates ISCSI_ROOT variables.

3. **OpenSVC orchestration**

OpenSVC is instructed to create the relevant iSCSI targets on one or more SCHs. These zvol-backed targets are exported over the storage network.

4. **Client RAID**

The TCN combines the iSCSI devices into a RAID set (e.g. RAID1 using mdadm) to provide redundancy. This allows failover if one SCH becomes unavailable.

5. **Installation**

The TCN OS is installed directly onto the iSCSI-backed RAID devices, making it fully diskless and resilient to back-end failures.

6. **Expansion**

Additional iSCSI targets can later be provisioned for data disks and attached to the TCN as needed.

11.5.5 Future extensions

- Automated provisioning of **ISCSI_DATA** targets for application and workload storage.
- Enhanced ZFS tuning (e.g. volblocksize, compression, logbias) to optimise for specific workloads.

- Integration of monitoring hooks so that OpenSVC and HPS can track zpool health and iSCSI target performance.
 - Support for flexible RAID scenarios where high-value services use multiple SCHs while low-cost services may only rely on a single disk.
-

11.5.6 Acronyms

- **SCH** – Storage cluster node (storage host).
- **TCN** – Thin compute node (diskless compute host).
- **ZFS** – Advanced file system and volume manager used for iSCSI backing.
- **iSCSI** – Protocol that exports block storage devices over a TCP/IP network.
- **OpenSVC** – Cluster manager and orchestrator used to control zvol exports and failover behaviour.

11.6 Troubleshooting

Stub: Common problems, causes, and fixes for HPS operation.

11.6.1 logging

See the log files / syslog

11.6.2 Function test

```
env QUERY_STRING="cmd=generate_opensvc_conf" bash -x /srv/hps-system/http/cgi-bin/boot_manager.sh
```

11.6.3 HPS system documentation

11.6.3.1 Overview

- **Project introduction** – Purpose of HPS, architecture, intended use cases.
- **Design decisions** – Records of key technical choices and rationale.
- **System components** – Description of major directories and modules.

11.6.3.2 Quick start

- **Prerequisites** – Required host operating system, packages, and network setup.
- **Installation** – Deploying `hps-container` with `hps-system` and `hps-config`.
- **First cluster setup** – Initialising a cluster with `cluster-configure.sh`.
- **Booting a host** – PXE boot process and selecting a host profile.
- **Verification** – Checking service status and logs.

11.6.3.3 System administration

- **Directory layout** – Locations for configuration, logs, distributions, and packages.
- **Cluster management** – Creating, switching, and editing clusters.
- **Host management** – Adding, removing, and updating host configurations.
- **Distribution and repository management** – Managing ISOs, PXE trees, and package repositories.
- **Service management** – Starting, stopping, and reloading `dnsmasq`, `nginx`, and `supervisord`.
- **Backup and restore** – Protecting and recovering configuration and repository data.

11.6.3.4 Functions reference

- Automatically generated from `lib/functions.d/` and other libraries.
- One function per page with purpose, arguments, usage examples, and related functions.

11.6.3.5 Advanced configuration

- **Kickstart and preseed templates** – Structure, variables, and customisation.
- **iPXE menus** – How menus are built and extended.
- **Storage provisioning** – SCH node disk detection, reporting, and configuration.
- **Integration points** – Hooks for OpenSVC, monitoring, and external systems.

11.6.3.6 Troubleshooting

- **PXE boot issues** – Common causes and fixes.
- **Service failures** – Diagnosing and restarting services.
- **Distribution and repository problems** – Checksums, GPG keys, and synchronisation errors.
- **Network problems** – DHCP conflicts, VLAN configuration, and firewall blocks.

11.6.3.7 Development

- **Code layout** – File structure and conventions for scripts and libraries.
- **Adding functions** – Naming, argument handling, logging, and documentation guidelines.
- **Testing** – Using `cli/test.sh` and other test harnesses.
- **Contributing** – Workflow, coding standards, and submission process.

11.6.3.8 Appendices

- **Glossary** – Definitions of terms and acronyms used in HPS.
- **Environment variables** – Description of exported variables and their use.
- **Decision records** – Full list of design decisions.
- **Reference configurations** – Example cluster, host, and service configuration files.

11.6.4 `_apkovl_create_resolv_conf`

Contained in `lib/functions.d/tch-build.sh`

Function signature: `7f68c373f65cdb9fd367f8e973f0778486a609eeeb348c3d3779336610f77261`

11.6.5 Function overview

The bash function `_apkovl_create_resolv_conf()` aims to create a new `resolv.conf` file in the specified directory with a passed nameserver. It takes two arguments - a temporary directory and a nameserver, creates a `resolv.conf` file in the given directory and logs the process. If the operation fails, it logs an error message and returns 1.

11.6.6 Technical description

- **Name:** `_apkovl_create_resolv_conf`
- **Description:** Creates a new `resolv.conf` file with a specified nameserver in a given directory.
- **Globals:** None
- **Arguments:**
 - `$1: tmp_dir`: The directory in which the `resolv.conf` file will be created.
 - `$2: nameserver`: The nameserver that will be written into the `resolv.conf` file.
- **Outputs:** A `resolv.conf` file in the specified directory; log messages of the process or an error message.
- **Returns:** 0 on success, 1 if it fails to create a `resolv.conf` file.
- **Example usage:**

```
_apkowl_create_resolv_conf "/temp/dir" "8.8.8.8"
```

11.6.7 Quality and security recommendations

1. Input validation: Even though not strictly necessary for this function to work, consider validating inputs, like verifying if a directory exists or if a nameserver is valid.
2. Error handling: Instead of just returning 1, consider throwing an exception or giving a more descriptive error output that includes more details of the issue.
3. Code comments: Including comments in the code would help other developers understand the function better, particularly explaining the purpose of the function and its arguments.
4. Logging: It would be beneficial to include more detailed logging, such as logging the successful creation of the `resolv.conf` file.
5. Security improvement: Be cautious with the use of `echo` with redirection `>` - if not used correctly, it might lead to security vulnerabilities. You should always make sure that user-provided inputs are properly escaped or cleaned up.

11.6.8 _apkowl_create_structure

Contained in `lib/functions.d/tch-build.sh`

Function signature: `de21a4a364feb0c001016f8d48f06a318c8c14aa1115120e71e72b590755ee99`

11.6.9 Function Overview

This function, `_apkowl_create_structure`, is responsible for creating a directory structure for Alpine Linux overlay (`apkowl`). It takes one argument, `tmp_dir`, which specifies the temporary directory where the structure is to be created. The function attempts to create two directories namely, `etc/local.d` and `etc/runlevels/default` within `tmp_dir`. It also sets up a symlink for the local service at boot. Any failure in the creation of directories or symlink results in an error message and halts the execution of the script.

11.6.10 Technical Description

- **Name:** `_apkowl_create_structure`
- **Description:** This function creates a directory structure for `apkowl` at a specified temporary location. It tries to set up the local service to be enabled at boot.
- **Globals:** None.
- **Arguments:**
 - **\$1 (tmp_dir):** The root directory where the `apkowl` structure is to be created.
- **Outputs:** Logs messages about operation status (debug and error).
- **Returns:**

- **1**: If failed to create directories or local service symlink.
- **0**: If execution completes without any error.

- **Example Usage:**

```
temp_directory="/tmp/my_directory"  
_apkovl_create_structure $temp_directory
```

11.6.11 Quality and Security Recommendations

1. Check if the argument input (tmp_dir) is not empty. This would prevent potential issues from attempting to create directories at the filesystem root.
2. Check if tmp_dir ends with a trailing slash and handle the scenario appropriately to prevent potential directory creation errors.
3. Consider using absolute paths for directory creation to prevent unexpected results due to relative paths.
4. Validate success of each operation, not just directory and symlink creations. This will improve error reporting and make troubleshooting easier in complex setups.
5. Make use of more secure and reliable logging mechanisms for output messages.
6. Employ appropriate file and folder permissions when creating the directories and setting up the symlink to avoid security risks.

11.6.12 _apkovl_create_bootstrap_script

Contained in lib/functions.d/tch-build.sh

Function signature: 6df90be762a78be7f74c156a3dabe8600615bcd697f73b84b3a65a3b3b4069de

11.6.13 Function overview

The _apkovl_create_bootstrap_script function is a bash function responsible for generating a bootstrap script for a temporary directory on an Alpine Linux distribution. It takes three arguments namely temporary directory, gateway IP and the Alpine version and creates a bootstrap script with these placeholders. The function also handles error scenarios robustly using logging infrastructure and returning relevant values to indicate success or failure.

11.6.14 Technical description

Name: _apkovl_create_bootstrap_script

Description: The function is defined to create a bootstrap script fulfilling a specific sequence of operations. It's a part of a larger set of operations, usually related to preparing an Alpine environment on a temporary directory including configuring Alpine repositories, updating the package index, installing essential packages, and sourcing other necessary functions.

Globals: None

Arguments: - \$1: Temporary Directory (tmp_dir) - \$2: Gateway IP (gateway_ip) - \$3: Alpine Distribution Version (alpine_version)

Outputs: This function does not directly output any value but it does create or manipulate several files in addition to printing logging statements.

Returns: This function returns 0 upon a successful execution and 1 on failure.

Example Usage:

```
_apkovol_create_bootstrap_script "/tmp/dir" "192.168.1.1" "v3.14"
```

11.6.15 Quality and security recommendations

1. Robust error handling: The function does well on this front, logging errors and returning failure immediately. It is also careful to validate its input and exit early in case of irregularities.
2. Secure placeholder substitution: In the step where placeholders are substituted, take into account special characters. Use `printf '%q'` to escape special characters to avoid accidental execution.
3. Validate input: Before accepting directory input or creating files inside unvalidated directories, ensure that path traversal attacks are not possible. Sensitizing path inputs helps avoid such scenarios.
4. Network-dependent operations should be segmented and adopt a fail-fast methodology. This is already in place for this function.
5. Variable localisation: Keep the use of local variables where they are known within, rather than using global variables.
6. Permission handling: Ensure that file permissions are properly managed throughout to prevent unprivileged access. The function does handle this, albeit at the last step, which could be too late in some scenarios.
7. Codify constants: For elements such as number of retries, instead of using magic numbers, their meaning could be codified in the form of well-named constants.

11.6.16 bootstrap_get_functions

Contained in `lib/functions.d/configure-remote-host.sh`

Function signature: `c437d9e5d7177b2e02fbafa34c1a8d3a19587a89ebceeda21ff24800c5beca8a`

11.6.17 Function overview

The `bootstrap_get_functions` function is designed to initialize local variables `gateway` and `distro` by invoking respective utility functions. The URL, created using these variables, is intended to reach a specific script on the server. The function uses `curl` to fetch the desired script and runs it within the script's current shell context. If script loading fails, an error message is printed, distinguishing between the error cases.

11.6.18 Technical description

- Name:
 - `bootstrap_get_functions`
- Description:
 - The function initializes `gateway` and `distro`, constructs a URL based on these variables along with a specific command, fetches and sources a script from the created URL via `curl`, validating successful sourcing. An error message is printed if unsourced.
- Globals:
 - No globals are used directly by this function.
- Arguments:
 - The function does not take any arguments.
- Outputs:
 - Success message indicating the sourced script URL, or
 - Error message indicating the failure to fetch or source functions from the URL.
- Returns:
 - 2 if the sourcing fails.
- Example usage:
`bootstrap_get_functions`

11.6.19 Quality and security recommendations

1. Consider outputting different messages or detailed exit codes to vividly distinguish between “Failed to fetch” and “Failed to source”.
2. Ensure that the `gateway` and `distro` returned by their respective functions are sanitized and do not contain any potentially harmful or unexpected characters.
3. Validate the existence and functional state of the specified `gateway` and the `boot_manager.sh` script before proceeding to fetch and source to avoid remote execution failures.
4. Use secure HTTP (HTTPS) for the URL to maintain secure communication while fetching the script.
5. Consider adding error handling and recovery logic for network interruptions.

11.6.20 `bootstrap_get_functions`

Contained in `lib/functions.d/configure-remote-host.sh`

Function signature: `77cd5f3eaa37754f85ce772c4240cd41a79e0a5ef9cc5efaaa33dffd9478e9e5`

11.6.21 Function Overview

The `bootstrap_get_functions` function's main objective is to fetch and source bootstrap functions from a URL generated using a given gateway and distro parameter. The URL is generated with a gateway caught from the `bootstrap_get_provisioning_node` function and a distro string from `bootstrap_initialise_distro_string` function. Curl is then used to retrieve and source the bootstrap functions. The function echoes a success or failure message, indicating whether the bootstrap functions were successfully loaded or not.

11.6.22 Technical Description

- **Name:** `bootstrap_get_functions`
- **Description:** This function fetches bootstrap functions from a specific server URL using curl, then sources them. The URL varies based upon the variables defined in the locally retrieved gateway and distro parameters.
- **Globals:** [None]
- **Arguments:** [None]
- **Outputs:** The function echoes whether the fetch and source operations were successful or failed.
- **Returns:** 2 if the fetch or source operations failed; nothing if successful.
- **Example Usage:**
`bootstrap_get_functions`

11.6.23 Quality and Security Recommendations

1. Ensure URL sanitation: For enhanced security and stability, it's crucial to guarantee that the URL used in the curl command is well-formed and devoid of harmful characters or injections.
2. Employ robust error handling: If the curl command fails to retrieve the data for any reason, the entire function will fail. To prevent this, incorporate more extensive error handling.
3. Add a timeout to the curl function: To prevent the script from hanging indefinitely if the server does not respond, it's advisable to implement a timeout feature.
4. Implement HTTPS: For the sake of security, it would be better to use a secure HTTPS connection instead of HTTP while making curl requests.
5. Verify SSL certificate: When using HTTPS, it would be beneficial to confirm the SSL certificate to prevent Man-in-The-Middle attacks.
6. Use explicit variable declarations: The script could cause unforeseen problems if global variables with similar names exist elsewhere in the script. Hence, it's important to make all variable declarations as local as possible.

11.6.24 bootstrap_get_provisioning_node

Contained in `lib/functions.d/configure-remote-host.sh`

Function signature: 510e95a7ea40a9fb7173fdeeb81b4d5e94033f079c1adf146dddc4040cf4b9c0

11.6.25 Function Overview

The function `bootstrap_get_provisioning_node` is designed to retrieve the default gateway IP which is also known as the provisioning node. It does so by using the `ip route` command piped into an `awk` command, which finds the line starting with 'default', and then prints the third field from that line before exiting. The resulting output is the IP address of the default gateway on the network.

11.6.26 Technical Description

Name: `bootstrap_get_provisioning_node`

Description: This function retrieves the IP address of the default gateway in a given network, otherwise referred to as the provisioning node.

Globals: None

Arguments: This function does not take any arguments

Outputs: The IP address of the default gateway (provisioning node)

Returns: Exit status of the command used within the function. If successful, the function will return the exit status 0, reflecting the successful execution of the `ip route` and `awk` commands. However, if either command fails, the function will return the non-zero exit status of the failed command.

Example usage:

```
provisioning_node=$(bootstrap_get_provisioning_node)
echo "The current provisioning node's IP is $provisioning_node"
```

11.6.27 Quality and Security Recommendations

To improve the quality and security of this function, consider the following:

1. Check the availability of the `ip` and `awk` commands before running them. This can prevent errors in environments where these commands might not be available.
2. Handle errors explicitly. Currently, the function assumes that the commands will execute without error. Handling possible errors can make the function more robust and easier to debug.
3. Validate the IP address obtained. Although unlikely, the `awk` command might return a non-IP-address string under unusual circumstances. Validating the output can prevent bugs later in the script.

4. Ensure correct routing table. The script assumes that the user intends to use the default IP routing table. If there are additional routing tables, this might not work as expected. This should be considered in more complex network setups.

11.6.28 bootstrap_initialise_distro

Contained in `lib/functions.d/configure-remote-host.sh`

Function signature: `1664d8a7eb2277600c48a3bc6974c34ea586df3fb5c9bfb5b7b8268285d91c63`

11.6.29 Function overview

This function, `bootstrap_initialise_distro()`, is utilised for initializing the process of bootstrapping on a specific distribution. It primarily utilizes a local variable `mac` which is passed as an argument. The function then runs a bash script, the output of which is printed out using the `cat` command.

11.6.30 Technical description

- **name:** `bootstrap_initialise_distro()`
- **description:** This function's main use is to initialize the bootstrapping process in a specific Linux distribution. The function takes MAC address as input which specifies the target machine for bootstrapping. The body of the function is a bash script that is implying an offline bootstrapping process from a provisioning server.
- **globals:** None
- **arguments:**
 - `$1`: `mac` - This argument represents the MAC address of the target machine.
- **outputs:** The output of the function is a printed bash script meant to conduct an offline bootstrap process from a provisioning server.
- **returns:** The function does not have an explicit return value since it's mainly a bash script output. The function's job is to print out a bash script.
- **example usage:**

```
bootstrap_initialise_distro "00:0a:95:9d:68:16"
```

11.6.31 Quality and security recommendations

1. Security can be improved by validating the MAC address input to ensure it's formatted correctly and exists in the network.
2. Validate the user's privilege level before running the script to safeguard against unauthorized access.
3. Better error handling could be implemented to account for any issues during the script's operation.
4. Avoid storing sensitive information such as passwords in the script to enhance security.

5. Include a logging functionality that keeps a record of changes made by the script for auditing and debugging purposes.

11.6.32 bootstrap_initialise_distro_string

Contained in `lib/functions.d/configure-remote-host.sh`

Function signature: `74d5bc0ea18318cdc4687791282138abdbb9349d49dd193ffda6f99d3f638705`

11.6.33 Function overview

The function `bootstrap_initialise_distro_string` is used to gather system information for a Linux machine and generate a string description of the system's architecture, manufacturer, operating system name and version number. This description string is outputted in a specific format: `cpu-mfr-osname-osver`.

11.6.34 Technical description

- **Name:** `bootstrap_initialise_distro_string`
- **Description:** This function collects system details including CPU architecture, manufacturer and operating system details (name and version). Uses this information to produce a string detailing these settings in the format `cpu-mfr-osname-osver`.
- **Globals:** [None]
- **Arguments:** [None]
- **Outputs:** The produced string of system settings (`cpu-mfr-osname-osver`) will get echo'ed out to stdout.
- **Returns:** No specific return value.
- **Example usage:**

```
$ bootstrap_initialise_distro_string
x86_64-linux-ubuntu-20.04
```

This example indicates a usage on an Ubuntu 20.04 system running on an x86_64 architecture.

11.6.35 Quality and security recommendations

1. Ensure the file `/etc/os-release` is reliably secure because the function reads from it. In case the file doesn't exist or is corrupted, the output string may not be accurate.
2. Handle edge cases where the CPU architecture, manufacturer or operating system details cannot be obtained. Currently, the function would default to "unknown", which may not be the desired output.
3. Consider checking and validating input before processing it to better handle any unexpected or malformed input.

4. Always keep the system and its software up-to-date to ensure that the `uname` command and `/etc/os-release` file return reliable and accurate information.
5. Run this program with the least privilege necessary to reduce potential damage in the event of a bug or breach.

11.6.36 `bootstrap_initialise_functions`

Contained in `lib/functions.d/configure-remote-host.sh`

Function signature: `f46a49b8c02ebe76f341c3e618292bab08779a5390717289e5fcb7b14072c5c5`

11.6.37 Function overview

The function `bootstrap_initialise_functions` reads a heredoc, EOF, which contains an offline bootstrap initializer from a provisioning server written in `sh`. It is designed to accept a wide variety of distributions (distro agnostic) and functions, even in environments where `bash` is not installed or supported. The initializer script defines a placeholder for functions.

11.6.38 Technical description

- **name:** `bootstrap_initialise_functions`
- **description:** The function initializes a series of offline functions from a provisioning server that are distro agnostic. It ensures the correct functions are used even if `bash` is not supported in the environment.
- **globals:** none
- **arguments:** none
- **outputs:** It outputs a shell script that includes an offline bootstrap initialization from a provisioning server.
- **returns:** The function does not have a specific return value. It returns the exit status of the last command executed which, in this case, is `cat`.
- **example usage:** To use this function, simply call it: `bootstrap_initialise_functions`

11.6.39 Quality and security recommendations

1. It is recommended to test this function with various Linux distributions to ascertain its distro-agnostic behavior.
2. Ensure the provisioning server can handle any sort of exception and downtime to avoid unexpected failures.
3. It is best to validate the contents of your heredoc.
4. Running scripts from a provisioning server can pose security risks. Verify the content and use checksums to ensure the integrity and authenticity of the scripts.
5. Ensure that the shell script handles errors appropriately, and reports problems in a way that is meaningful to the user.

11.6.40 build_dhcp_addresses_file

Contained in `lib/functions.d/dns-dhcp-functions.sh`

Function signature: `1e8517f7be98ca6ca266e569111938c53137ee61efbb1a3efc44cb18aa3f3d52`

11.6.41 Function overview

The Bash function `build_dhcp_addresses_file` is utilized to generate a DHCP addresses file which includes the Mac, IP, and hostname of each host in a cluster. If the `services` directory does not already exist, it is created. The function fetches a list of all hosts via `list_cluster_hosts`, and validates the IP address and hostname for each host before appending it to the DHCP addresses file.

11.6.42 Technical description

- **Name:** `build_dhcp_addresses_file`
- **Description:** This function produces a DHCP addresses file which lists the Mac, IP, and hostname for each of the hosts in the cluster. If the directory does not exist, the function will create one.
- **Globals:**
 - `HPS_CLUSTER_CONFIG_DIR`: The configuration directory for the cluster.
- **Arguments:** None
- **Outputs:** A file, `dhcp_addresses`, is created. If the directory does not exist, the function creates it. For every host in the cluster, the function adds entries into the file. The entry includes the formatted MAC address, IP, and hostname of the host.
- **Returns:** Returns 0 on executing successfully. Returns 1 in case of any error.
- **Example usage:**

To use this function, it should be called from the script in following manner: `bash build_dhcp_addresses_file ### Quality and security recommendations`

1. Ensure that the `HPS_CLUSTER_CONFIG_DIR` is set with the correct permissions, allowing only necessary privileges.
2. Ensure that the list of hosts being processed is reliable, avoiding the risk of processing malicious or compromised hosts.
3. Validate the `HOSTNAME` and IP that is fetched, to avoid possible injection attacks.
4. Avoid logging sensitive data in error messages or logs which can be read by users with lower privileges.
5. Consider using a temporary directory when creating temporary files to avoid potential security issues.
6. Cleaning up or securely erasing the temporary file after moving its content to destination file.

11.6.43 build_dns_hosts_file

Contained in `lib/functions.d/dns-dhcp-functions.sh`

Function signature: 202234c143dff920a5780db8420eb67121377864cab2f9cd898343ab9493ab

11.6.44 Function overview

The `build_dns_hosts_file` function in Bash is essentially a method that configures the details required for a DNS hosts file in a cluster. The function works by setting the specific directories and files, logging the process, analysing cluster configurations for domain name and IP address, validating the given IP, defining IPS service aliases, starting with an empty temporary file, and writing these entries into the file. It then moves this temporary file to the destined location. If any step fails, it logs an error and returns from the function.

11.6.45 Technical description

- **Name:** `build_dns_hosts_file`
- **Description:** The function `build_dns_hosts_file` sets up a DNS hosts file in a cluster. It first sets directories and goes on to fetch important details like DNS domain and IP address from the cluster configuration. It validates the received IP and then constructs the DNS hosts file. If any step fails, it logs the error message and returns from the function.
- **Globals:** [`HPS_CLUSTER_CONFIG_DIR`: An environmental variable defining the directory path of the cluster services]
- **Arguments:** [None: This function doesn't require any arguments]
- **Outputs:** This function writes entries into a temporary file and then moves this file to the final location. If any error occurs during the process, it logs the error message.
- **Returns:** The function will return '1' if the process fails at any step, and '0' if the function is executed successfully.
- **Example Usage:**

`build_dns_hosts_file`

11.6.46 Quality and security recommendations

1. Implement error message handling for each step, to avoid cascading failure effects.
2. Include input validation to the best extent possible.
3. Use security measures to protect IP and other sensitive details from being mishandled.
4. Workflow debugging might be needed to spot code imperfections to pre-empt errors.

5. Use file permissions correctly to prevent unauthorized access to sensitive files and directories. The DNS hosts file should have appropriate restrictive permissions.
6. Protect the function and its processes from interruptions and ensure graceful exits in case of errors.

11.6.47 build_yum_repo

Contained in `lib/functions.d/repo-functions.sh`

Function signature: 55fba2752d840b45670dcc10f8be084d3234f9c87a37a1959fcde6a9247be828

11.6.48 Function Overview

The `build_yum_repo` function checks for changes in RPM packages inside a specified repository. If changes are detected or if there's no previous state, it uses the `createrepo_c` command to create or update the metadata for the YUM repository. Regardless of the outcome, a checksum of current state is stored for future comparison. Outputs and potential errors are reported using `hps_log` helper function.

11.6.49 Technical Description

- **Name:** `build_yum_repo`
- **Description:** The function checks the RPM changes inside the provided repo path. If there are changes or no previous state, `createrepo_c` is used to create or update the repo metadata. Checksums of the current state are saved for future checks.
- **Globals:** None
- **Arguments:** `repo_path` (\$1): the path to the repo that needs to be checked and updated.
- **Outputs:** Associated logs with `hps_log` displaying info about events and potential errors.
- **Returns:** 0 if no changes were made or repo created successfully, 1 if the repo path is not provided or doesn't exist, 2 if `createrepo_c` command is not found.
- **Example usage:** `build_yum_repo "${HPS_PACKAGES_DIR}/${DIST_STRING}/Repo"`

11.6.50 Quality and Security Recommendations

1. Consider validating the checksum process and handling any exceptions it might throw. This can improve the function's quality and resilience.
2. Assert the existence of `'createrepo_c'` command at the start of the function to fail early if it's not installed, this can save execution time.
3. Look into the security of the checksum generation. Make sure it is robust against potential risks such as spoofing or collision attacks.

4. Validate the structure and content of the repo path argument to prevent potential bugs or security issues related to wrong or malicious inputs.

11.6.51 `build_zfs_source`

Contained in `lib/host-scripts.d/rocky.d/rocky.sh`

Function signature: `c64e26c5e886b1a0aded061414c66873630e97cae41cf2c7a37f800fd6866674`

11.6.52 Function Overview

`build_zfs_source` is a bash function which aims to download, build and install ZFS from the source files. It fetches source file index from a given source base URL and identifies matching source file for `build_zfs_source`. The function then attempts to download the source file and installs ZFS build dependencies solving automatic makefile generation, foreign function interface libraries, library for UUID generation etc. It extracts the source archive, moves to the build directory, and builds ZFS. The function ensures that ZFS module is found after installation and returns 0 or 1 based on the results of the implementation.

11.6.53 Technical Description

- **Name:** `build_zfs_source`
- **Description:** This function fetches, downloads, and builds ZFS from the source files, installing the necessary dependencies for the build, and ensuring the ZFS module is present post-installation.
- **Globals:** [gateway: gateway URL for the source]
- **Arguments:** [\$1: desc, \$2: desc]
- **Outputs:** Log messages indicating the process flow and possible errors during the execution.
- **Returns:** 0 if ZFS is successfully built and installed, 1 if an error occurs at any step of the process.
- **Example Usage:** `build_zfs_source`

11.6.54 Quality and Security Recommendations

1. It would recommend validating the URL before attempting to fetch or download files from it. This helps prevent potential security issues related to downloading malicious files.
2. Checking if the required dependencies are already installed before attempting to install them can improve the efficiency of the script.
3. Running `configure`, `make`, and `make install` operations could lead to potential security vulnerabilities. Therefore, it would be better to consider user privilege separation and avoid running the script as a superuser if not necessary.

4. Adding more logging could improve the traceability of errors and the overall debugging experience.
5. To make the function more robust, it is suggested to always quote variable expansions in strings to prevent word splitting and pathname expansion.

11.6.55 `cgi_auto_fail`

Contained in `lib/functions.d/cgi-functions.sh`

Function signature: `6b67dd57145386562143b5a27ad4c4f0a1e5170fc073f8d4e91738ade5779a4d`

11.6.56 Function overview

The `cgi_auto_fail` function is primarily used to detect the client type and based on the detection, it uses different failure methods. Messages are passed as arguments to the function, and based on the client type, these messages are processed differently. If the client type is `ipxe`, the function `ipxe CGI_fail` is called; for client types `cli`, `browser`, `script`, `unknown`, the function `CGI_fail` is called; and for all other scenarios, the function simply logs the error and echoes the message.

11.6.57 Technical description

- **Name:** `cgi_auto_fail`
- **Description:** This function is made to detect the client type and handle failure messages differently based on the client type. The client type is first detected by the function `detect_client_type`, and then care branches deal with the different cases accordingly.
- **Globals:** None
- **Arguments:**
 - `$1`: msg The failure message to be processed.
- **Outputs:** Depending on the client type and the corresponding branch the execution enters, the output could be the execution of the `ipxe CGI_fail` function, the `CGI_fail` function, or simply echoing the error message and logging it.
- **Returns:** No specific return value.
- **Example usage:** `CGI_auto_fail "Failure message"`

11.6.58 Quality and security recommendations

1. Always sanitize user inputs, especially if these inputs are incorporated into messages or logs.
2. Add comprehensive error handling and logging to help with problem detection and troubleshooting.
3. Regularly update your logging methods to take advantage of whatever logging resources are currently available on the system.

4. Separate the logic of error detection from the error messaging to provide a cleaner and more maintainable code.
5. Consider employing a standard and internationalized method for handling error messages.

11.6.59 `cgi_fail`

Contained in `lib/functions.d/cgi-functions.sh`

Function signature: `5bc8c57b97b640ef4a194c0065da11fec44b1f1bb0525bf46e8163d2a50cabd5`

11.6.60 Function overview

This function `cgi_fail()` is used within a CGI script to handle errors. It takes an error message as an argument, `cmsg`, logs this message to some error tracking system via the `hps_log()` function, and then responds to the HTTP request with a plain header and the same error message echoed back. It's an in-built function for managing errors in CGI scripts.

11.6.61 Technical description

- **Name:** `cgi_fail`
- **Description:** This function handles errors in CGI scripts by logging errors and sending responses with plain headers and echoed error messages.
- **Globals:** None.
- **Arguments:**
 - `$1`: `cmsg` – Error message that will be logged and sent back as response.
- **Outputs:** The function outputs an error message.
- **Returns:** This function does not return anything.
- **Example Usage:** New error message can be passed directly to function as follows

```
bash      cgi_fail "An error has occurred."
```

11.6.62 Quality and security recommendations

1. Validate the input: You should consider validating your function's input. This way it ensures that the error message passed is a string so that it can be correctly processed.
2. Error handling and logging: You should ensure proper error handling and logging is in place. This is vital for auditing and rectifying the faults in the system. Implement a standardized error logging method.
3. Escape special characters: In the echoed error message, make sure to escape any HTML special characters for security purpose.
4. Usage of local variables: Use local variables where possible. Usage of global variables may lead to the risk of variable pollution.

5. Use appropriate HTTP status codes: Always use appropriate HTTP status codes in conjunction with error messages to help client understand the problem better.

11.6.63 `cgi_fail`

Contained in `lib/functions.d/cgi-functions.sh`

Function signature: `5bc8c57b97b640ef4a194c0065da11fec44b1f1bb0525bf46e8163d2a50cabd5`

11.6.64 Function Overview

The function `cgi_fail` depicted above performs the essential role of logging an error message and outputting it through the server's CGI header. It takes the error message as an argument and performs two main actions: logging the error message using `hps_log` function and outputting the error message. This function can come handy in cases where you want to keep track of all the errors that occur on your server, and also display them on the server for debugging purposes.

11.6.65 Technical Description

In a more detailed technical perspective, the function's definition is structured as follows:

- **Name:** `cgi_fail`
- **Description:** This function logs an error message to the system and outputs it through the server's CGI header. It operates by invoking the `cgi_header_plain` function and the `hps_log` function.
- **Globals:** None
- **Arguments:**
 - `$1`: An error message to log and output through the CGI header.
- **Outputs:** This function debugs the provided message and prints it to the standard output.
- **Returns:** Does not return any explicit value.
- **Example Usage:** To use `cgi_fail`, you would generally pass it an error message as such: `cgi_fail "Unknown error encountered"`.

11.6.66 Quality and Security Recommendations

1. **Input Validation:** Given that the function is accepting user-provided strings, the function should perform proper input validation on the error message to prevent possible injection attacks.
2. **Error Handling:** The function does not contain any explicit error handling. Implementing checkpoints to verify successful suppression of errors might be useful.
3. **Logging:** Consider adding timestamps and additional context to the logged errors for ease of troubleshooting.

4. **Fallbacks:** It may be beneficial to add a fallback or default error message if a blank error message is provided.
5. **Code Documentation:** Make sure to document what the function does and how to use it. Include details of what each argument should be and what the function returns. This can help others to use the function correctly and avoid mistakes.

11.6.67 `cgi_header_plain`

Contained in `lib/functions.d/cgi-functions.sh`

Function signature: `cce7eb1544681966ec11d5f298135158497fc2ac56c89b00c63c84b8e1bc733a`

11.6.68 Function Overview

The `cgi_header_plain` function is a Bash shell function that is designed to produce the HTTP header for a plain text response. In a CGI (Common Gateway Interface) context, this function can be used to clearly define the content type of the output as plain text.

11.6.69 Technical Description

```
def cgi_header_plain:
  - name : cgi_header_plain
  - description : A Bash function that prints the HTTP header for
  ↪ a plain text response, typically used in a CGI context.
  - globals : None
  - arguments : None
  - outputs :
    - "Content-Type: text/plain"
    - A line break
  - returns : None
  - example usage :
    ```
 #!/bin/bash
 cgi_header_plain
    ```
```

11.6.70 Quality and Security Recommendations

1. Consider checking the status of the `echo` commands to ensure that they successfully sent the outputs.
2. If this function is used in a larger script, make sure that it's being called appropriately and that the output is being properly used.
3. Ensure secure execution with proper user privileges. Running the script with unnecessary admin rights can be a security risk.
4. Audit and keep a record of function usage in order to trace and troubleshoot any issues when they occur.

5. Include error handling to respond to failed operations and cleanly exit the script when necessary. This will prevent any misuse or security vulnerabilities.

11.6.71 `cgi_log`

Contained in `lib/functions.d/cgi-functions.sh`

Function signature: `9f2c2cf7c0d57e85a08611717b5d691eddf235f096bbc311bf9d58541f0c77b3`

11.6.72 Function overview

The function `cgi_log()` is designed to output log messages with a timestamp into a file. It takes a string as an argument and appends it to the log file with a timestamp for traceability purposes.

11.6.73 Technical description

- **name:** `cgi_log`
- **description:** The function accepts a string as an input and writes it to the log file (`/var/log/ipxe/cgi.log`) with a timestamp. This provides a chronological record of all the logging information.
- **globals:** None
- **arguments:**
 - `[$1: msg]` Log message as string
- **outputs:** Appends the log message with a timestamp to the `/var/log/ipxe/cgi.log`.
- **returns:** Not applicable.
- **example usage:** `cgi_log "This is a test message"`

11.6.74 Quality and security recommendations

1. **Input Validation:** Always validate the input (`msg`) before using it. This will prevent log injection attacks.
2. **Log Rotation:** To manage the size of the logs properly, implement some type of log rotation either via a Bash script or a system utility.
3. **Permissions:** Ensure appropriate permissions are set on the log file to prevent unauthorised access or alteration of the logs.
4. **Sensitive Information:** Be cautious while logging messages as it should not include any sensitive data like passwords which can be exposed through logs.
5. **Error Handling:** Consider adding error handling to log any potential errors when trying to write to the log file.

11.6.75 `cgi_param`

Contained in `lib/functions.d/cgi-functions.sh`

Function signature: 5007f95c313c04a01df7ba39bff0241f44511cd570d795bc11a890edf032f323

11.6.76 Function overview

The bash function `cgi_param` is designed to decode and retrieve parameters from the `QUERY_STRING` in a CGI context. The function uses a query-string from the CGI context and processes it to detect commands and key-value pairs. The function is triggered to process the query string only once, with subsequent calls to commands (`get`, `exists`, `equals`) retrieving or comparing the saved values. If an unknown command is passed, an error message is returned.

11.6.77 Technical description

- **Name:** `cgi_param`
- **Description:** This function parses a query string into parameters, then provides a way to retrieve a parameter value, check if a parameter exists, or see if a parameter's value matches a provided value by using the `'get'`, `'exists'` or `'equals'` commands respectively.
- **Globals:**
 - `QUERY_STRING`: The string to extract parameters and their values from.
 - `__CGI_PARAMS_PARSED`: Global flag to detect if the query string has been parsed.
 - `CGI_PARAMS`: An array to save decoded keys and their values.
- **Arguments:**
 - `$1`: The command to run: `'get'`, `'exists'`, `'equals'`.
 - `$2`: The name of the parameter.
 - `$3`: Optional. The value to compare against when the command is `'equals'`.
- **Outputs:** If the `'get'` command is called, the value of the named parameter. If not, a success status or an error message in case of an invalid command.
- **Returns:** The function can return different values depending on the command given. If `'get'` command, will print value to stdout. If `'exists'`, 0 is returned if parameter exists, 1 if not. If `'equals'`, returns 0 if parameter equals given value, 1 if not. Returns 2 if command is invalid.
- **Example usage:**

```
cgi_param get username
cgi_param exists page_count
cgi_param equals user_role admin
```

11.6.78 Quality and security recommendations

1. Use stricter validation on parameter keys while parsing. Right now, only alphanumeric characters plus underscores are allowed. However, consider more restrictive set.

2. Be sure to escape all variable expansions to prevent code injection.
3. Consider ways of handling or communicating parse failures more explicitly - it may prove difficult to debug if the QUERY_STRING is not formatted as expected.
4. Implement a more robust error handling mechanism. For example, when the user provides an invalid command, an exception should be handled that doesn't allow further execution of the script.
5. Add more guard clauses for blank, null, or unexpected inputs to avoid unexpected behaviour.
6. Always try to keep your function's behavior predictable and documented.

11.6.79 `cgi_success`

Contained in `lib/functions.d/cgi-functions.sh`

Function signature: `3c468a0d5bf1a1d432f4efcb2a108889f0d0e762f542f50c29b92350f02de3bc`

11.6.80 Function overview

The `cgi_success` function is a Bash function used within CGI (Common Gateway Interface) scripting to return a plain HTTP header and output a message. This function is common in web development contexts where Bash scripting is used to interact with websites.

11.6.81 Technical Description

Name:

`cgi_success`

Description:

The function ``cgi_success`` calls the ``cgi_header_plain`` function
↪ and then uses ``echo`` to print the first argument passed to it.

Globals:

None

Arguments:

- `$1`: This is a message that is outputted after the HTTP header. It
↪ can be any string message that the user want to display.

Outputs:

- The function outputs an HTTP header followed by the content of
↪ the string stored in `$1`.

Returns:

- No explicit return value.

Example usage:

```
```bash
message="Your process was successful."
cgi_success "$message"
```

### ### Quality and Security Recommendations

1. Validate input: Ensure the input string (`\$1`) is sanitized and validated before being used.
2. Error handling: Consider adding error handling for the `cgi\_header\_plain` function.
3. Documentation: Document the purpose and usage of the function in the codebase for ease of use.
4. Return codes: This function currently does not use explicit return codes. Though it's not a requirement, it would be helpful to have a return code to indicate success or failure.

```
`cgi_success`
```

Contained in `lib/functions.d/cgi-functions.sh`

Function signature: 3c468a0d5bf1a1d432f4efcb2a108889f0d0e762f542f50c29b92350f02de3b

### ### Function Overview

The `cgi\_success` function is a utility function in Bash designed for use in CGI scripts.

### ### Technical Description

The details of the `cgi\_success` function are as follows:

- **Name**: `cgi\_success`
- **Description**: This function outputs a plain header and then echoes the first input argument.
- **Globals**: None.
- **Arguments**:
  - `\$1`: The text to be displayed in the body of the CGI response.
- **Outputs**: Prints out the contents of `\$1` following a plain header in CGI response.
- **Returns**: Nothing since `echo` does not have a return value.
- **Example usage**:

```
```bash
cgi_success "The CGI Script executed successfully."
```

11.6.82 Quality and Security Recommendations

1. Always validate and sanitize the input passed to the function to avoid a potential injection attack.

2. Implement error checking for the function calls inside `cgi_success` and handle them appropriately.
3. Document the expected values for `$1` clearly for users.
4. To protect against unintended side effects, make sure to quote variables that are referenced to prevent word splitting or pathname expansion.
5. Always use the function in conjunction with proper Header declaration in CGI scripting.

11.6.83 `check_and_download_latest_rocky`

Contained in `lib/functions.d/iso-functions.sh`

Function signature: `ac86f8c78c7149b4fbc3125a2f52dd65b160ad01f437b5eeb5b504bd1a90c6b1`

11.6.84 Function overview

This shell function `check_and_download_latest_rocky`, written in bash script, is designed to check for the latest version of Rocky Linux for a specific architecture and download the ISO file if it is missing. It sets up the target base directory and checks whether the latest version ISO file is present. If it is not present, it downloads the ISO from the base url which is initially defined.

11.6.85 Technical description

- **Name:** `check_and_download_latest_rocky`
- **Description:** This bash function checks and downloads the latest version of Rocky Linux for a specific architecture if it is missing.
- **Globals:** `HPS_DISTROS_DIR`: Directory for storing ISOs.
- **Arguments:** None.
- **Outputs:** Logs and messages about the process and final results.
- **Returns:** None in this function but could return 1 if erred in practical usage.
- **Example Usage:** `check_and_download_latest_rocky`. As it doesn't require any arguments, the function can be called without providing any.

11.6.86 Quality and security recommendations

1. **Validating inputs:** Input variables, especially those sourced from outside the function such as `HPS_DISTROS_DIR`, should be checked and validated to ensure they are correct and safe.
2. **Error trapping:** Consider adding more error trapping to handle situations where commands within the function fail.
3. **Use More Robust Code:** String operations on paths can be replaced with more robust code using `dirname` and `basename`.

4. **Use `https://`:** Always use secure protocol (`https`) while downloading files for ensured security.
5. **Checksum Verification:** For further security, the downloaded ISO could be verified against a known checksum to ensure its integrity. Any mismatches should trigger a warning or error.
6. **Permission and Access:** Ensure proper permission schemes for the directories and files mentioned in the function.

11.6.87 `check_available_space`

Contained in `lib/functions.d/repo-functions.sh`

Function signature: `106f18dba6188914b70a8af6d1ddf900c6e5702a9f782fb29e86c3906bc878f6`

11.6.88 Function Overview

The `check_available_space` function is used to determine whether there is enough available disk space at a given path. The path to be checked is passed to the function as an argument, along with a specified amount of required space in megabytes. If no amount is specified, the function defaults to checking for at least 500MB of available space.

11.6.89 Technical Description

- **Name:** `check_available_space`
- **Description:** This function verifies if enough disk space is available at a given path. The amount of required disk space can optionally be passed as an argument. If not provided, it defaults to checking for 500MB.
- **Globals:** None
- **Arguments:**
 - `$1`: Path to the directory for disk space check. This is a mandatory argument.
 - `$2`: The required disk space in MB. This is an optional argument, if not provided defaults to 500MB.
- **Outputs:**
 - Outputs available disk space in MB at the given path to `stdout`.
 - Logs error messages to `stderr` using `hps_log` if there's an error.
- **Returns:**
 - Returns 0 if the available space is greater than or equal to the required space OR if the path exists and available space can be determined.
 - Returns 1 if the available space is less than the required space OR if the path does not exist or available space cannot be determined.
- **Example usage:**
 - `check_available_space /path/to/directory 1000`: Checks if the directory at `/path/to/directory` has at least 1000MB of available space.

11.6.90 Quality and Security Recommendations

1. The function currently relies on the availability and behavior of external commands such as `df`, `awk`, and `sed`. The reliance on these commands might introduce potential vulnerabilities and room for behavior inconsistencies across different systems. It is recommended to reduce this dependence by using built-in Bash features wherever possible.
2. Consider adding more comprehensive error handling to deal with situations like the absence of utilities the function depends on.
3. To enhance readability and maintainability, consider refactoring lengthy operations into separate, smaller functions.
4. The function could benefit from input validation. Currently, there is no validation that the required parameters (path and required space) are of the correct format or within plausible ranges.
5. Ensure the function is thoroughly tested with various inputs, including edge and failure cases.
6. Document any assumptions made in the code and any system dependencies.
7. Consider using a static analysis tool to detect potential security vulnerabilities and to enforce a coding standard.

11.6.91 `check_iscsi_export_available`

Contained in `lib/host-scripts.d/common.d/iscsi-management.sh`

Function signature: `747cf8b13f4b0b09cc62797344d4a6efa48f7dd4e42c1431fbc4771d9a4058f5`

11.6.92 Function overview

The `check_iscsi_export_available` function verifies if all necessary software and hardware components are present and correctly set up in order to export iSCSI targets on a Unix-based system. It does this by validating the presence of the `targetcli` package, checking for a properly mounted `configfs` filesystem, and ensuring the availability of Light-weight Input/Output (LIO) kernel target modules.

11.6.93 Technical description

- **Name:** `check_iscsi_export_available`
- **Description:** This function checks for the presence and configuration of various prerequisites necessary for iSCSI target exporting.
- **Globals:** None
- **Arguments:** No arguments are expected by this function.
- **Outputs:** This function outputs various error or success messages outlining the state of the export environment.

- **Returns:** Returns 0 if all checks pass and the iSCSI export environment is ready (both ‘targetcli’ and LIO kernel modules are available). Returns 1 and echoes an error message if any of the checks fail.
- **Example usage:**

```
check_iscsi_export_available
```

11.6.94 Quality and security recommendations

1. To improve soundness and maintainability of the function, it is recommended to implement a more robust error handling mechanism. This might include handling for potential issues during the execution of the embedded shell commands.
2. For security purposes, consider controlling permissions (via ‘chown’, ‘chmod’, etc.) on /sys/kernel/config to ensure only the required processes and users can access and modify it.
3. To further enhance function reliability, consider adding checks for specific versions of the ‘targetcli’ package and LIO kernel modules, as different versions may not behave identically.

11.6.95 check_latest_version

Contained in lib/functions.d/iso-functions.sh

Function signature: 3b874dd0168548a5363b9c63f357dab1016d452e1155213b1190930b50bb44c5

11.6.96 Function overview

The bash function `check_latest_version()` is designed to check the latest version of an operating system provided by a manufacturer for a specified CPU architecture. The function currently supports operating systems hosted on `rockylinux.org`. Inputs include the CPU architecture, the manufacturer, and the operating system name. Outputs will be the latest version of the specified operating system or appropriate error messages.

11.6.97 Technical description

- **Name:** `check_latest_version()`
- **Description:** This function checks the provided URL for the latest version of an operating system. Specifically tailored for `rockylinux.org`, the function parses the fetched webpage to find OS version numbers and returns the latest version found.
- **Globals:** No global variables are used in this function.
- **Arguments:** \$1: CPU architecture (Not currently used in function), \$2: Manufacturer (Not currently used in

function), \$3: Operating System name (Used to form URL and output appropriate version or error messages)

- **Outputs:** Latest version of the operating system or appropriate error messages.
- **Returns:** 0 if the latest version of the operating system is successfully found, 1 otherwise
- **Example Usage:** `check_latest_version x86_64 Intel rockylinux`

11.6.98 Quality and security recommendations

1. Input Validation: Implement input validation for CPU architecture and manufacturer parameters, which are currently unused.
2. Error Handling: Additional error handling could be implemented if the curl command fails due to network issues.
3. Expand OS Support: The function's functionality could be expanded to support other OS providers beyond just `rockylinux.org`.
4. Secure HTTP Transfers: Consider enforcing HTTPS when fetching the HTML to enhance the security of the function.
5. Scrutinize Regular Expressions: Regular expressions used for matching versions could be reviewed and made more robust to prevent potential errors in output.

11.6.99 `check_zfs_loaded`

Contained in `lib/host-scripts.d/common.d/zpool-management.sh`

Function signature: `e83726e842477cf79c67cdcf1de046b56eacfa7d7f97ba195d70b1f2bccfab2`

11.6.100 Function Overview

The `check_zfs_loaded` function checks whether the ZFS (Zettabyte File System) is installed and operational on the machine where the script is running. If the ZFS command isn't found or the ZFS kernel module isn't loaded, the function attempts to perform the necessary actions and gives feedback on the actions performed. The function ultimately returns 0 if the system passes all checks; otherwise it returns 1.

11.6.101 Technical Description

11.6.101.1 Name

`check_zfs_loaded`

11.6.101.2 Description

This bash function checks two conditions: if the ZFS command is found on the system, and if the ZFS kernel module is loaded. If any of these conditions is not met, it prompts the

user with an error message. Moreover, if the ZFS module is not loaded, it attempts to load it using the `modprobe` function. This operation might require superuser privileges.

11.6.101.3 Globals

None.

11.6.101.4 Arguments

No arguments taken.

11.6.101.5 Output

The function outputs to `stdout`:

- A successful or unsuccessful message indicating ZFS command installation status.
- A status update on whether the ZFS module is loaded, whether a load attempt was made, and whether that was successful.

11.6.101.6 Returns

- 1: if either the ZFS command is not found OR the ZFS module failed to load.
- 0: if the ZFS command is found AND the ZFS module is successfully loaded.

11.6.101.7 Example Usage

```
if check_zfs_loaded; then
    echo "ZFS is ready to use."
else
    echo "ZFS is not properly set up."
fi
```

11.6.102 Quality and Security Recommendations

1. Authentication Check: The function should check if the user has required permissions to perform actions such as loading a kernel module using `modprobe`. Otherwise, it could misleadingly signal a failure when the underlying issue is a lack of permission.
2. Error Catching: It might be beneficial to add additional error handling or checking. For example, catching and further diagnosing errors that result from the command or `modprobe` functions could provide more clarity as to what caused a failure.
3. Reliable Checking: Checking presence of ZFS only through presence of `zfs` command and loading status of module might not fully confirm its operational status. Integrate a more reliable way of verifying whether ZFS is properly working, such as creating a test file on a ZFS filesystem.

4. Logging: Consider adding logging for better troubleshooting capabilities. Rather than just echoing the status, also write it into a log file.
5. Use Safe Bash Flags: Consider setting safe bash flags (`set -euo pipefail`) at the beginning of your scripts to avoid certain common bash pitfalls.

11.6.103 check_zfs_loaded

Contained in `lib/host-scripts.d/common.d/zpool-management.sh`

Function signature: `e83726e842477cf79c67cdcf1de046b56eacfa7f97ba195d70b1f2bccfab2`

11.6.104 Function Overview

The `check_zfs_loaded` function is used to ensure that the ZFS filesystem is installed and enabled in the kernel. It first checks if the `zfs` command is available. If the command is not found it will notify the user and return an error code. If the command is found, the function then checks if the ZFS kernel module is loaded using the `lsmod` and `grep` commands. If the module is not loaded, the function will attempt to load it using `modprobe`. If successful, it notifies the user and returns a success code. Otherwise, it notifies the user of the failure and returns an error code.

11.6.105 Technical Description

Definition Block:

- **Name:** `check_zfs_loaded`
- **Description:** A bash function to check whether the ZFS filesystem is installed and enabled in the kernel on the current system.
- **Globals:** None.
- **Arguments:** None.
- **Outputs:** Echoes information to the console regarding the state of the ZFS filesystem and its kernel module.
- **Returns:** 1 if the ZFS command is not found or the kernel module fails to load; 0 otherwise.
- **Example Usage:**

```
bash      check_zfs_loaded      if [ $? -eq 0 ];  
then      echo "ZFS is installed and loaded"      else  
echo "ZFS is either not installed or not loaded"      fi
```

11.6.106 Quality and Security Recommendations

1. The function should validate the return status of the command `-v zfs` and `modprobe zfs` calls directly to avoid false positives in case of error situations.
2. The function could handle more error scenarios, such as failing to run `modprobe`.
3. Other tuning parameters could be checked to ensure optimal ZFS performance, such as the `zfs` module option values.

4. To enhance security, consider ensuring that the function is run with the right permissions, as loading/unloading kernel modules typically requires root access.
5. This function could be supplemented by implementing additional checks, such as whether or not a given ZFS pool or filesystem exists.

11.6.107 `cidr_to_netmask`

Contained in `lib/functions.d/network-functions.sh`

Function signature: 9da9925a8a1bbe4ed1131bde6d5e428aa8c11ff5f8a692b1049fc1dd9078c9d4

11.6.108 Function Overview

The function `cidr_to_netmask()` is a network utility function in Bash which converts a CIDR prefix into a netmask address. This is useful because the CIDR format is commonly used for network configuration and the function provides an easy conversion to a more traditional netmask configuration.

11.6.109 Technical Description

- **Name:** `cidr_to_netmask`
- **Description:** Converts CIDR prefix to netmask address.
- **Globals:**
 - `mask`: A string that builds and holds the final netmask value
- **Arguments:**
 - `$1`: This argument represents the CIDR prefix, which should be an integer between 0 and 32.
- **Outputs:** Prints the resulting netmask value.
- **Returns:** Does not return a value since the result is printed directly.
- **Example Usage:**

```
{  
  echo "CIDR to netmask conversion:"  
  
  echo $(cidr_to_netmask 24)  
}
```

The example above will output: 255.255.255.0.

11.6.110 Quality and Security Recommendations

1. Ensure correct input values: The function currently does not perform any checks on the input to ensure it is a valid CIDR prefix. Implementing this would increase its robustness.

2. Use more specific variable names: While `mask`, `prefix`, `full_octets`, and `partial_octet` are informative, including their intent in names could make the function easier to understand.
3. The function could be modified to output to a variable instead of directly printing the result. This would allow its output to be captured and used elsewhere in a script rather than simply displaying it.
4. Avoid command substitution and parsing echo output for determining the octet count (`$(echo "$mask" | tr -cd '.' | wc -c) < 3`). It's better to count internally using an integer to prevent any erroneous output or potential code injection.

11.6.111 `cluster_config`

Contained in `lib/functions.d/cluster-functions.sh`

Function signature: `ea561bb0141a57e976a8d270b800ed25635238facc99a0dcb1b5ee58af4ad106`

11.6.112 Function overview

The Bash function `cluster_config()` provides an interface to manage the configuration of a cluster. The function takes three arguments: an operation, a key, and optionally a value. The operation can be 'get', 'set', 'exists', or any other string (which will be treated as an error). The function attempts to find the active cluster configuration file. If it doesn't find one, it returns an error. If it does find one, it performs the desired operation on the key-value pair.

11.6.113 Technical description

- **Name:** `cluster_config`
- **Description:** A Bash function to manage the key-value pairs in the configuration of an active cluster.
- **Globals:** [`cluster_file`: a file containing the active cluster configuration]
- **Arguments:** [`$1`: operation to perform (get, set, exists), `$2`: a key to operate on, `$3(Optional)`: a value corresponding to the key]
- **Outputs:** Outputs depend on operation. If 'get', it prints the value corresponding to the key. If 'set', it modifies or creates a key-value pair. If 'exists', it searches for the key. Any other string will result in an error message.
- **Returns:** 1 if no active cluster config is found, 2 if an unknown operation is passed to the function.
- **Example usage:** To set the value of a key 'key1' in the active cluster's config to 'value1', command will look like `cluster_config set key1 value1`.

11.6.114 Quality and security recommendations

1. The function currently writes error messages to `stderr`, which is a good practice. However, no positive messages are provided for successful operations. Providing feedback for success can improve usability and debuggability.
2. The ‘set’ operation makes use of `echo` to append to a file, which can be unsafe if the script does not control the content being echoed. A safer alternative would be to use `printf`.
3. The ‘get’ operation uses `cut` and `grep` to parse the key-value pairs. If the values contain special characters, or if another ‘=’ appears in the line, this could have unintended consequences. A more robust parsing mechanism could improve this.
4. The function may consider input validation to ensure that provided keys and values are valid before attempting operations.
5. It could be beneficial to check if the operation succeeded (i.e., the value was actually set or got), and return appropriate errors if not.

11.6.115 `cluster_has_installed_sch`

Contained in `lib/functions.d/cluster-functions.sh`

Function signature: `befb7e577a31bf8e64c1179ffa1c6cd4d2ec9a30913e1c6b26986c37fd0762cc`

11.6.116 Function overview

This function, `cluster_has_installed_sch()`, checks a directory containing configuration files by reading and processing each file one at a time. Specifically, it checks for the presence of files ending with `.conf` extension. Within each file, it reads key-value lines and processes the lines containing “TYPE” and “STATE”. It then checks if `type` is “SCH” and `state` is “INSTALLED”. If it finds a match, it returns 0, indicating success. If no match is found after reading all the files, it returns 1, indicating failure.

11.6.117 Technical description

- **name:** `cluster_has_installed_sch`
- **description:** The function checks a directory of configuration files to find if a file contains `TYPE=SCH` and `STATE=INSTALLED`.
- **globals:** [`HPS_HOST_CONFIG_DIR`: The directory containing the configuration files (`.conf`) to be checked.]
- **arguments:** None.
- **outputs:** Outputs nothing.
- **returns:** Returns 0 if a match is found; otherwise it returns 1.
- **example usage:**

```
$ cluster_has_installed_sch
$ echo $?
0
```

11.6.118 Quality and security recommendations

1. Add a check at the beginning of the function to ensure `HPS_HOST_CONFIG_DIR` is set and is a valid directory.
2. Provide a graceful exit or report an error message if the configuration directory does not exist.
3. Consider using stricter checks when processing each file's content. This can help prevent potential issues related to unexpected data.
4. Add more comments throughout the script to explain the function of each part.
5. In terms of security, be aware of potential "Path Traversal" vulnerabilities if the directory contains symbolic links pointing outside of the intended directory tree. You might want to add a check to ensure symbolic links are not processed.

11.6.119 `_collect_cluster_dirs`

Contained in `lib/functions.d/cluster-functions.sh`

Function signature: `21a54dd1fed816cbbecd01967b98b68192e6453200d8c224a543b3e6b074b633`

11.6.120 Function Overview

The function `_collect_cluster_dirs()` is used to collect all directories from a base directory (excluding symbolic links). The base directory is specified by the global variable `HPS_CLUSTER_CONFIG_BASE_DIR`. The names of directories are stored in an array, which is passed to the function via a reference variable. If the base directory doesn't exist, the function prints an error message and exits with status zero.

11.6.121 Technical Description

- **Name:** `_collect_cluster_dirs()`
- **Description:** This bash function collects all the directory entries from the base directory specified by the `HPS_CLUSTER_CONFIG_BASE_DIR` global variable and stores them in the referenced array. It skips symbolic links and any non-directory entries.
- **Globals:**
 - `HPS_CLUSTER_CONFIG_BASE_DIR`: Description not provided in the sample function. Presumably, this global variable specifies the base directory in which the function searches for directories.
- **Arguments:**
 1. `$1`: This is a reference to an array. The names of directories found will be added to this array.
- **Outputs:**
 - If the base directory doesn't exist, an error message is written to the standard error output.

- The function modifies the array passed to it via the reference variable, adding names of directories found.

- **Returns:**

- The function always returns zero.

- **Example Usage:**

```
# Assume that the global variable HPS_CLUSTER_CONFIG_BASE_DIR is
# already set to some valid directory path
declare -A my_array
_collect_cluster_dirs my_array
```

11.6.122 Quality and Security Recommendations

1. Document the purpose and usage of global variables used by the function.
2. Consider having the function return non-zero status when an error is encountered, such as when the base directory doesn't exist.
3. Avoid polluting the global scope by unsetting local variables at the end of the function.
4. It would be a good security practice to sanitize inputs to the function or ensure they are of the expected type.
5. Include error handling for cases where the argument passed is not a valid reference to an array.

11.6.123 `configure_kickstart`

Contained in `lib/functions.d/configure_kickstart.sh`

Function signature: `b9c974579a4cf0918b1f523ab5546c0fabf4f4fe0d6a0a4ab77a23765ef1cbca`

11.6.124 Function Overview

The Bash function, `configure_kickstart()`, is used to automate the generation of a Kickstart configuration file for a new Linux cluster. This function requires a cluster name as the argument, which is then used to name and generate a Kickstart file in the designated path. This function includes a configuration set up that ensures system requirements, including network configuration, root password, partitioning, and package selection, among others, are met. An error message is displayed when a cluster name is not provided as the argument.

11.6.125 Technical Description

- **Name:** `configure_kickstart()`
- **Description:** This function automates the process of generating a Kickstart configuration file for a Linux cluster using a provided cluster name.

- **Globals:** [`CLUSTER_NAME` : Stores the name of the cluster, `KICKSTART_PATH` : Path where the kickstart file will be generated]
- **Arguments:** [`$1`: Cluster name]
- **Outputs:** Prints status messages during the function execution.
- **Returns:** Returns an error message if the required argument (cluster name) is not provided.
- **Example Usage:** `configure_kickstart test_cluster`

11.6.126 Quality and Security Recommendations

1. Make sure the Kickstart file path is a secure location and has the right permission settings to avoid unauthorized access or alterations.
2. Ensure proper validation is performed on the input cluster name to avoid possible command injection vulnerabilities.
3. The root and user passwords are currently set with placeholder values. Change these values to secure ones before using on production systems.
4. Ensure you handle the potential issue where the cluster name given might already exist, overwriting an existing kickstart file.
5. Consider including a verification step to confirm the successful generation of the Kickstart file or to catch any possible errors during the file creation process.
6. It's recommended to ensure the installation log (`/root/ks-post.log`) is properly secured or even disabled in a production environment to protect system information.

11.6.127 `configure_kickstart`

Contained in `lib/functions.d/configure_kickstart.sh`

Function signature: `b9c974579a4cf0918b1f523ab5546c0fabf4f4fe0d6a0a4ab77a23765ef1cbca`

11.6.128 Function Overview

The function `configure_kickstart()` is used to generate a Kickstart configuration file for a specified cluster. Kickstart files are used by the CentOS/Fedora/RHEL boot process to configure new installations. It's a shell script that automates the post-installation process of setting up a customized system.

11.6.129 Technical Description

- **Name:** `configure_kickstart`
- **Description:** This function accepts a cluster name as an argument. It takes this name and creates a Kickstart file customized for that specific cluster. If no cluster name is provided, it will return an error message and exit. After successfully creating the Kickstart file, it prints the location of the newly generated file.

- **Globals:** [`CLUSTER_NAME` : The name of the cluster, `KICKSTART_PATH` : The location where the Kickstart file will be created]
- **Arguments:** [`$1` : The name of the cluster]
- **Outputs:** Outputs logs descriptive of the processing, including an error message if the cluster name is not provided, a status log regarding kickstart file generation, and the location of the generated Kickstart file.
- **Returns:** Does not return a value and will stop execution if cluster name is not provided.
- **Example Usage:** `configure_kickstart cluster1`

11.6.130 Quality and Security Recommendations

1. Make sure to validate the input parameter to avoid any unwanted results.
2. Always check for existing files before generating a new file to prevent data loss.
3. Ensure the generated files have the correct permissions and owner to prevent potential security breaches.
4. Consider encrypting sensitive data, such as passwords, to enhance security.
5. Develop comprehensive error handling routines to catch and handle any issues during execution.
6. Avoid using hardcoded values (such as `"/srv/hps-config/kickstarts"`) as much as possible.

11.6.131 `configure_supervisor_core`

Contained in `lib/functions.d/configure-supervisor.sh`

Function signature: `85bc4baf9a5f406b060f9e65c003dcd4719029b002e68c90befff7936eb18994`

11.6.132 Function Overview

The `configure_supervisor_core` function generates a Supervisor core configuration file at a location defined by the `CLUSTER_SERVICES_DIR` environment variable and logs its location for further usage. It validates required environment variables, creates required directories, writes the configuration file, checks for errors, and ensures the config file is readable and contains the required sections.

11.6.133 Technical Description

- **Name:** `configure_supervisor_core`
- **Description:** Generates a Supervisor core configuration file.
- **Globals:**
 - `CLUSTER_SERVICES_DIR`: Directory for cluster services.
 - `HPS_LOG_DIR`: Directory for logging.
- **Arguments:** None

- **Outputs:** Path to the generated configuration file through stdout.
- **Returns:**
 - 1 if CLUSTER_SERVICES_DIR or HPS_LOG_DIR is not set.
 - 2 if creation of the configuration or log directories fails.
 - 3 if writing supervisord configuration file fails.
 - 4 if configuration file does not exist, is not readable, appears to be empty, or is missing a required section after writing.
 - 0 if successful.
- **Example Usage:**

`configure_supervisor_core`

11.6.134 Quality and Security Recommendations

1. Ensure to validate and sanitize all environment variables in the beginning of the script execution to avoid potential security vulnerabilities.
2. Write proper error messages in the case of application failure. This will help in better debugging and understanding of what is happening inside your code.
3. Handle all possible exceptional cases, such as checking if the necessary directories exist before proceeding.
4. Make sure to secure the configuration files and restrict the permissions to only necessary users to prevent unauthorized modifications.
5. Regularly update and maintain the application to protect from potential security threats.

11.6.135 `configure_supervisor_services`

Contained in `lib/functions.d/configure-supervisor.sh`

Function signature: `deb84b582dac93ac1bafb35b941997e641067f625a53e65feb8a6cdd938cc887`

11.6.136 Function overview

The `configure_supervisor_services` function in Bash creates a configuration file for handling Supervisor services. This function ensures the core header and defaults exist. If the core configuration file does not exist or can't be read, logger function `hps_log` logs an error message. The function forms several service configurations (`dnsmasq`, `nginx`, `fcgiwrap`, `opensvc`) by invoking helper function `*supervisor*append_once` which checks if a service (stanza) exists already, if not, it appends a new service configuration block.

11.6.137 Technical description

- **Name:** `configure_supervisor_services`

- **Description:** This function is primarily used for configuring Supervisor service settings and managing associated directories and files. It helps in setting up different services and their environments and handling errors along the way by logging them.
- **Globals:** HPS_LOG_DIR, CLUSTER_SERVICES_DIR, HPS_SERVICE_CONFIG_DIR.
- **Arguments:** This function does not accept any runtime arguments.
- **Outputs:** Logs informative, debug and error messages related to the creation, existence and validation of supervisor services configuration.
- **Returns:** 0 if success, 1 if failed to create supervisor core configuration, 2 if failed to create directory, and 3 if failed to write service block or Supervisor configuration file validation failed.
- **Example usage:** `configure_supervisor_services`

11.6.138 Quality and security recommendations

1. The function is missing strict mode (`set -euo pipefail`). Using strict mode makes scripts halt execution on the first error, making debugging easier and limiting data corruption.
2. More extensive error handling could be implemented or use of `set -e` to stop script on error.
3. Consider handling signals and traps for more graceful exits or cleanup routines.
4. Comments and sections could be better formatted for easier understanding.
5. Avoid hardcoding paths and user names. Consider parameterizing more aspects of the function to increase reusability.
6. Some parts of the function are redundant and could likely be simplified.
7. Conceptually, this function may do too many things. It may be beneficial to split it into multiple subfunctions for a cleaner code structure.

11.6.139 `count_clusters`

Contained in `lib/functions.d/cluster-functions.sh`

Function signature: `35e810345753544ce73618109f16ad97d33b5185c8ff2d374a8837aa569e5960`

11.6.140 Function overview

The `count_clusters` function is a shell command used primarily to gather a count of directories that are considered “clusters”. More specifically, it stores these cluster directories into an array and returns the count. If the array is empty, it outputs an error message, and return an exit code of 0 along with a count of 0 to illustrate that no clusters were found.

11.6.141 Technical description

Name: count_clusters

Description: This function is used to obtain a count of “cluster” directories.

Globals: [HPS_CLUSTER_CONFIG_BASE_DIR: The base directory where clusters are located]

Arguments: [None]

Outputs: If it cannot find any clusters, it outputs an error message stating “No clusters found in ‘HPS_CLUSTER_CONFIG_BASE_DIR’”. If it finds clusters, it outputs the count of clusters.

Returns: It returns 0 regardless of whether it finds any clusters. If it finds clusters, it still returns the number of clusters.

Example Usage:

```
# Count the directories in the base directory.  
count_clusters
```

11.6.142 Quality and security recommendations

1. Based on the provided function, there’s an implication that some global variables are used across multiple functions. This could lead to obscure bugs if one function modifies the global variable in a way that another function doesn’t expect. In a safer, clearer, and easier-to-maintain option, those global variables should be turned into arguments to isolate side effects.
2. There is a lack of argument validation in the function. For robust and error-free functioning, it is advisable to check if the calculated directories exist and are indeed directories before counting them.
3. The use of echo to communicate errors is not following best practices. Instead, we should consider switching to an logger or exposing error messages through the use of special return codes. This would give users more context on how to resolve the error than just presenting it as a string.

11.6.143 create_config_dnsmasq

Contained in lib/functions.d/create_config_dnsmasq.sh

Function signature: 22bc6c19aa391d4d58e332e373400aad4f02035bedd99d8d94127662dcf30360

11.6.144 Function Overview

The create_config_dnsmasq function is a helper script designed to create the dnsmasq.conf and related configuration files required for dnsmasq functions, like binding to an IP, DHCP reservations, DNS settings, TFTP, and PXE. This function creates and writes the configuration files, and ensures they exist in the specified paths. The function will generate an error message if there is no DHCP_IP given.

11.6.145 Technical Description

- **Name:** `create_config_dnsmasq`
- **Description:** This function creates the `dnsmasq.conf`, `dns_hosts`, and `dhcp_addresses` configuration files for `dnsmasq`. It also logs the configuration process and checks for the presence of `DHCP_IP`. If `DHCP_IP` is absent, it returns an error message and exits.
- **Globals:**
 - `DHCP_IP`: description TBD
 - `DNSMASQ_CONF`, `DNS_HOSTS`, `DHCP_ADDRESSES`: Paths for the configuration files
- **Arguments:**
 - Not applicable for this function.
- **Outputs:** Creates the configuration files and logs the process. If `DHCP_IP` is not set, it outputs an error message.
- **Returns:** Technically, as a function, it doesn't return anything.
- **Example Usage:** `create_config_dnsmasq`

11.6.146 Quality and Security Recommendations

1. It is always a good practice to validate input parameters before proceeding with the rest of the function. For this function, more checks could be included to ensure the validity of the supplied IP addresses.
2. Error handling should be improved - currently, if `DHCP_IP` is not set it echos an error line and then exits. It would be better to implement more robust error handling.
3. Avoid using relative paths for configuration files, as this increases the risk of path injection vulnerabilities.
4. Add specific file permissions to the configuration files created using this function. This would restrict unauthorized access to these configuration files.
5. The function currently creates files and logs the process - logging should be considered for all significant steps of the function, not restricted to just the configuration section.
6. The script could benefit from some comments explaining what each section of the created configuration file does.

11.6.147 `create_config_dnsmasq`

Contained in `lib/functions.d/create_config_dnsmasq.sh`

Function signature: `5c19d2840751f093e5d235e25b027f015ff8c3397bf8e3295c8f9678f9127911`

11.6.148 Function Overview

This Bash function, named `create_config_dnsmasq`, is used to generate and set up the dnsmasq configuration file for a given active cluster. The generated configuration file supports DHCP and TFTP functionalities, and can be used in a PXE boot environment. The configuration is generated based on active cluster details and is used to set up the working environment, especially when it involves processes such as network booting and assigning IP addresses within the network.

11.6.149 Technical Description

In deeper details, the function can be broken down as below:

- **Name:** `create_config_dnsmasq`
- **Description:** This function generates and sets up a dnsmasq configuration file for running DHCP and TFTP services.
- **Globals:** These include:
 - `HPS_SERVICE_CONFIG_DIR`: Description of this variable can be as: The directory path where service config files are stored.
 - `DHCP_IP`: This variable contains the DHCP IP address to bind the dnsmasq service.
 - `NETWORK_CIDR` & `DNS_DOMAIN`: They provide the required network and DNS details for the configuration.
- **Arguments:** The function doesn't take any explicit arguments (\$1, \$2, etc.)
- **Outputs:** This function writes out configurations to the `dnsmasq.conf` file in the `HPS_SERVICE_CONFIG_DIR` directory. Any errors or issues encountered are logged and signaled.
- **Returns:** No explicit return value, but will exit with 0 in case of an error i.e., when the `DHCP_IP` is not set.
- **Example usage:** It can be called without any arguments like `create_config_dnsmasq`.

11.6.150 Quality and Security Recommendations

1. The function should include error checking for every step that might fail, not just checking if `DHCP_IP` is set.
2. Explicit security measures like directory and file permission setting and verification should exist to ensure that only authorized users can modify or manage these configurations.
3. It will be better if there's a version history or change tracking for the configuration changes generated by this function.
4. Instead of exiting immediately upon encountering an error, it might be better if the function handled errors more gracefully, allowing for cleanup tasks to complete despite errors.

5. There might be a need for a parameter validation mechanism to ensure the accuracy and security of supplied values.
6. Any sensitive information, like IP addresses and domain names, should be handled securely and be secured inside the function to prevent leaks.

11.6.151 create_config_nginx

Contained in `lib/functions.d/create_config_nginx.sh`

Function signature: `d3d22d399af4c7c80fb8449bf421c864014acf615622669621bd6f5a2999ef5e`

11.6.152 Function overview

The `create_config_nginx` shell function is designed to create and set up a configuration file for the nginx server. It begins by sourcing the active cluster file (if it exists) and defining the path for nginx's configuration file. The function then uses a heredoc (`<<EOF`) to write the server configurations, including worker processes, user and events into the nginx configuration file.

11.6.153 Technical description

Name: `create_config_nginx`

Description: This function creates an nginx configuration file with appropriate server settings.

Globals: There is one global variable that this function interacts with: `HPS_SERVICE_CONFIG_DIR`: The directory in which the nginx configuration file is located.

Arguments: This function accepts no arguments.

Outputs: This function outputs an `info` level log message about the configuration of nginx server.

Returns: The function doesn't return any explicit value, since its primary task is writing to a file. If this operation is successful, it will implicitly return `0`. Else, it will return whatever error code is thrown by the failing command inside the function.

Example usage:

```
create_config_nginx
```

11.6.154 Quality and security recommendations

1. Implement a feature to validate the nginx configuration file after it's been written. The `nginx -t` command could be used for this.
2. Consider using strict mode (`set -euo pipefail`) to handle potential errors.
3. Add descriptive comments to the script for easier understanding and debugging.

4. Configure the function to accept inputs such as `worker_processes`, `user`, `worker_connections` as arguments so it can be used more flexibly.
5. Use ShellCheck (or a similar linter tool) to analyze the script for potential issues related to robustness, portability, and maintainability.
6. Implement error handling. For example, check whether the `HPS_SERVICE_CONFIG_DIR` directory exists before trying to write to it.
7. Consider encrypting sensitive information that might appear in logs or configuration files.

11.6.155 `create_config_nginx`

Contained in `lib/functions.d/create_config_nginx.sh`

Function signature: 581497653f8d51b9a4cbf2d4bf39d79e8c580671467d2b61a5c8dfd8654b0888

11.6.156 Function Overview

The Bash `create_config_nginx()` function is clearly involved in the creation of an Nginx configuration file. It obtains a filename from the `get_active_cluster_filename` function (silencing error output), stores it within a local variable `NGINX_CONF`, and then writes a configuration template to it.

11.6.157 Technical Description

- **Name:** `create_config_nginx`
- **Description:** This function creates an Nginx configuration file.
- **Globals:** `HPS_SERVICE_CONFIG_DIR`: This global variable holds the path to the service configuration directory.
- **Arguments:** None
- **Outputs:** Produces an Nginx configuration file in a specific location.
- **Returns:** Not explicitly defined in the code.
- **Example Usage:** Call the function without any arguments like so - `create_config_nginx`

11.6.158 Quality and Security Recommendations

1. Handle errors from the `source` command: The function sources a file determined by the `get_active_cluster_filename` function but doesn't handle possible errors. To improve, check the file exists before sourcing it.
2. Manage permissions: Ensure that the script runs with enough, but not excess, privileges to write to `HPS_SERVICE_CONFIG_DIR`.
3. Log or handle errors from `cat`: The function writes to a file, but any errors (e.g., from insufficient permissions, out-of-disk-space errors) are ignored. These should be logged or handled.

4. Validate variables: Validate the `HPS_SERVICE_CONFIG_DIR` variable before using it. This could include checks to ensure it's a directory and has appropriate permissions.
5. Hardcoded configuration: The Nginx configuration is hardcoded into this function. It would offer more flexibility to load configuration from an external source, allowing for easy modification without requiring script changes.

11.6.159 create_config_opensvc

Contained in `lib/functions.d/create_config_opensvc.sh`

Function signature: `1e0006940e609bdc531da8856a0e4150fe29d39bfd65037dae44fc9ba0ec892f`

11.6.160 Function Overview

The function `create_config_opensvc()` generates configuration for Open Service Controller (OpenSVC) and ensures it has single cluster agent key policy. This involves creating specific directory paths for configuration, logs, and variable data. Temporary files are created and written to before being atomically moved to their permanent locations to ensure file creation does not fail halfway through. The agent key is checked against any existing keys and overwritten under specific circumstances, providing error handling if disk key and cluster key do not match.

11.6.161 Technical Description

11.6.161.1 Name

`create_config_opensvc()`

11.6.161.2 Description

A function to generate configuration files for OpenSVC and apply a single cluster agent key policy.

11.6.161.3 Globals

- `conf_dir`: The path to the configuration directory.
- `log_dir`: The path to the log directory.
- `var_dir`: The path to the variable data directory.
- `conf_file`: The location of the configuration file.
- `key_file`: The location of the agent key file.

11.6.161.4 Arguments

- `$1 (ips_role)`: Role of the IP Service for OpenSVC configuration.

11.6.161.5 Outputs

- Write configuration, create directories and files, log error/success messages.

11.6.161.6 Returns

- 1: If mktemp fails or generate_opensvc_conf() fails.
- 2: If disk_key and cluster_key do not match.

11.6.161.7 Example Usage

```
create_config_opensvc "database"
```

11.6.162 Quality and Security Recommendations

1. Add more comments and documentation within the function for better maintainability.
2. Implement more robust error checking mechanisms throughout the function to handle any potential failures. For example, check the success of mkdir, mv, chmod, chown commands.
3. Consider variable sanitization for the safety of the script. Always ensure paths and file names are secure before using them.
4. Avoid suppressing errors (2>/dev/null) to quickly identify and fix issues.
5. Enforce the use of secure file permissions and ownership, especially for key files that store sensitive data. Use least privilege principle.
6. Use cryptographic functions from trusted libraries instead of implementing them in house. This function uses openssl rand and in its absence, a fallback method for generating a key which would be less secure. Make openssl a requirement instead.

11.6.163 create_iscsi_target

Contained in lib/host-scripts.d/common.d/iscsi-management.sh

Function signature: 31b1d82c8fe3715f33280abbeccda64305c1dccc58d1fc70c3df63585ebba7d1e

11.6.164 Function Overview

The create_iscsi_target is a bash function primarily used for creating an iSCSI target on a remote host. iSCSI (Internet Small Computer System Interface) is an Internet Protocol-based storage networking standard for linking data storage facilities. The function receives three parameters: remote hosts, IP (default is 0.0.0.0), and port (default is 3260). It uses these to create an iSCSI target which exposes block-level storage for use by other networked servers, leveraging the targetcli command to perform the actual operations. It also checks for required iSCSI prerequisites before proceeding.

11.6.165 Technical Description

- **Name:** `create_iscsi_target`
- **Description:** Creates an iSCSI target at the specified remote host using the provided IP and port.
- **Globals:** None
- **Arguments:**
 - `$1`: `remote_host` - The hostname of the remote machines
 - `$2`: `ip` - The IP to create the iSCSI target on. Default is 0.0.0.0.
 - `$3`: `port` - The port to use for the iSCSI target. Default is 3260.
- **Outputs:** Prints out messages about the status of the iSCSI target creation.
- **Returns:** 1 if the remote host is not specified or iSCSI target prerequisites are not met; otherwise, the exit status of the last command executed.
- **Example usage:** `create_iscsi_target my-host 192.168.1.100 3260`

11.6.166 Quality and Security Recommendations

1. Since this function does privileged operations, ensure that it's only run by authorized users or services. Implement proper authorization checks to prevent unwanted usage.
2. Validate input parameters to ensure they are of correct format and within expected ranges. This helps avoid possible command injection vulnerabilities.
3. Handle errors and exceptions appropriately. Right now, if anything fails, the function would just continue with the remaining commands which might lead to inconsistent results.
4. Consider the implications of not using any encryption in your iSCSI traffic. Use IPSec or similarly secure tunneling protocols if this traffic traverses untrusted networks.
5. Consider limiting access to the iSCSI targets from trusted initiator IPs only, to prevent unauthorized accesses.

11.6.167 `create_supervisor_services_config`

Contained in `lib/functions.d/supervisor-functions.sh`

Function signature: `b804c13eebed28fe3f1dd8879efd00699a1641a6e35eada4d27d16679d7a2abd`

11.6.168 Function overview

The `create_supervisor_services_config` function in Bash scripting is responsible for creating the necessary configuration for three crucial services: `nginx`, `dnsmasq`, and `opensvc`. It automatically triggers the configuration routines for `nginx` and `dnsmasq`. In case of `opensvc`, it invokes the configuration routine by flagging the node as an IPS node specifically.

11.6.169 Technical description

The following provides a more technical detailing of the `create_supervisor_services_config` function:

- **Name:** `create_supervisor_services_config`
- **Description:** This is a Bash function that calls three other functions – `create_config_nginx`, `create_config_dnsmasq`, and `create_config_opensvc`. The main objective of all these functions is preparing the running environment for these services via respective configuration.
- **Globals:** None
- **Arguments:** No arguments for `create_supervisor_services_config` function itself. But, it internally passes `IPS` as a parameter to the `create_config_opensvc` function.
- **Outputs:** No explicit output. However, the environment gets prepared and the respective configurations for the services are made as a result of the function execution.
- **Returns:** Nothing specific is returned from the function execution.
- **Example usage:**
`create_supervisor_services_config`

11.6.170 Quality and security recommendations

1. Consider making the services for which configurations are to be made customizable through arguments. This will make the function more flexible and adaptable for different use-cases.
2. Check for necessary permissions before attempts to create configurations or making service alterations. This can prevent unexpected errors and enhances security by confirming proper access levels.
3. Always test the new configurations in a safe environment before sliding them into production. This precaution could prevent possible service downtime.
4. Include error-catching mechanisms and create logs where appropriate. This would help in debugging when something goes wrong.
5. Check backward compatibility if you intend to make alterations to this function. This ensures the function remains useful in diverse environments.

11.6.171 `create_supervisor_services_config`

Contained in `lib/functions.d/configure-supervisor.sh`

Function signature: `000144298c724dfbd327b7354be11dd901d1726a3f593104f6a6bbdb96329588`

11.6.172 Function overview

The function, `create_supervisor_services_config`, is responsible for invoking three other functions: `create_config_nginx`, `create_config_dnsmasq` and `create_config_opensvc`. These functions presumably generate configuration files for `nginx`, `dnsmasq` and `opensvc` respectively.

11.6.173 Technical description

The following is a block definition for `create_supervisor_services_config` function:

- **Name:** `create_supervisor_services_config`
- **Description:** This function is designed to trigger three other functions, namely `create_config_nginx`, `create_config_dnsmasq` and `create_config_opensvc`. Presumably, each of these functions will create a configuration file for their respective service.
- **Globals:** None
- **Arguments:** This function does not take any arguments.
- **Outputs:** This function does not have a return output, but it is implied by the function names it calls that configuration files for `nginx`, `dnsmasq`, and `opensvc` will be created as a result of this function.
- **Returns:** Nothing
- **Example usage:** `create_supervisor_services_config`

11.6.174 Quality and security recommendations

1. Each function called (`create_config_nginx`, `create_config_dnsmasq`, `create_config_opensvc`) should incorporate error handling to ensure the process of creation indeed occurs without fail.
2. Ensure that the configuration files created by these functions have tight file permissions to prevent unauthorized access or modifications.
3. Implement logging inside each function to have an audit trail of the operations performed and to allow for efficient debugging in case of any issues.
4. Validate the configuration files for `nginx`, `dnsmasq` and `opensvc` after they are created to ensure they do not contain any misconfigurations or security vulnerabilities.
5. Always handle sensitive data (passwords, secret keys, etc.) securely if such data is being written into any of these configuration files. Avoid hardcoding secrets in the scripts or configs and use secure methods to fetch such information.

11.6.175 `detect_call_context`

Contained in `lib/functions.d/system-functions.sh`

Function signature: f167df149452fd670d9f40bd87d52442f2f5d5153026bdfcab5f9c60997d7f96

11.6.176 Function Overview

The bash function `detect_call_context` is designed to identify and echo the current context in which the script is being executed. It handles three main contexts: when the script is sourced instead of directly executed, when it gets invoked as a common gateway interface (CGI), and finally, when it gets directly executed in a shell or reading from stdin without CGI environment variables. If none of these contexts apply, the function defaults to “SCRIPT”.

11.6.177 Technical Description

- **Name:** `detect_call_context`
- **Description:** This function identifies the context in which the script is running, which could be either “SOURCED”, “CGI”, or “SCRIPT”. It prints out the current context and then returns.
- **Globals:** [`BASH_SOURCE[0]`: Describes the source of the bash script, `GATEWAY_INTERFACE`: A required variable to detect CGI, `REQUEST_METHOD`: Another required variable to detect CGI, `PS1`: Helps in explicitly detecting the script]
- **Arguments:** None
- **Outputs:** Prints one of the four possible states - “SOURCED”, “CGI”, “SCRIPT”, or a fallback “SCRIPT”.
- **Returns:** `null`
- **Example Usage:**

```
source your_script.sh
detect_call_context
```

This script would output “SOURCED” if `your_script.sh` contains a call to this function.

11.6.178 Quality and Security Recommendations

1. It is essential that global variables are well-defined and adequately protected in a function. Use local variables or provide default values to prevent empty or undefined global variables issues.
2. Bash lacks some advanced features like well-defined namespaces, classes, or functions. For better security and efficiency, consider using a more powerful scripting language like Python or Perl for complex scripts.

3. Make sure that the script running the function has appropriate permissions. Bash scripts can be a significant security risk if they are writable by any user. Therefore, secure your script by limiting access.
4. Implement error handling and fallbacks for unexpected behaviour or exceptions.
5. The function implicitly trusts that certain global environment variables are not maliciously set. Ensure that these environment variables are validated before use.
6. Regularly update your system and software to prevent security vulnerabilities.

11.6.179 detect_client_type

Contained in `lib/functions.d/network-functions.sh`

Function signature: `a5996dd77379049ae4564d5c7eaab09a5189d239c610eea12c0d452cd96097e7`

11.6.180 Function Overview

The function `detect_client_type()` is designed to determine the type of client making a request and echo it back. The function looks at both the query string and the user agent to make this determination. If the client type cannot be determined, it echoes back “unknown”.

11.6.181 Technical Description

- **Name:** `detect_client_type()`
- **Description:** Determines the type of the client from `$QUERY_STRING` or `$HTTP_USER_AGENT`. This function echoes the client type: ‘ipxe’, ‘cli’, ‘browser’, ‘script’ or ‘unknown’ if it can’t determine the client type.
- **Globals:**
 - `QUERY_STRING`: Utilized to determine the client type based on the presence of certain keywords. If not set, default value is an empty string.
 - `HTTP_USER_AGENT`: Utilized to determine the client type based on the presence of certain keywords. If not set, default value is an empty string.
- **Arguments:** The function does not accept any arguments.
- **Outputs:** Echoes the type of client making the request.
- **Returns:** Always returns 0 as the function is designed to not fail, falling back to a default output of “unknown” when necessary.
- **Example Usage:**

```
$ export QUERY_STRING="via=cli"
$ detect_client_type
cli
```

11.6.182 Quality and Security Recommendations

1. To reduce potential errors or misuse, explicitly document the environments and contexts in which this function should be used or not used.
2. Consider adding error handling for undesired or unexpected input to improve stability.
3. Make sure to sanitize user-generated inputs such as query strings to prevent injection attacks.
4. If feasible, add type checks for input values in cases where the function starts accepting arguments.
5. To prevent leakage of potentially sensitive information, use discretion with echoing data, especially if it includes data from HTTP headers in a web server environment.

11.6.183 detect_storage_devices

Contained in `lib/functions.d/storage_functions.sh`

Function signature: 26407abf67962b945819ec70c2d91c7a26b60c144eeb209e22298b013ba307ed

11.6.184 Function Overview

The function `detect_storage_devices` is used to identify all available block devices on a system and gather important details such as device model, vendor, serial number, type, bus, size, usage, and speed. This information is collected in a structured format for easy analysis and troubleshooting.

11.6.185 Technical Description

- **Name:** `detect_storage_devices`
- **Description:** This function detects all block devices in a Linux system and retrieves information of each device including the device path, model, vendor, serial number, bus type, device type, size, usage, and speed.
- **Globals:** None
- **Arguments:** None
- **Outputs:** The function generates a formatted string containing details about all detected storage devices.
- **Returns:** The function doesn't return a specific result – it echoes the output directly, making the output available in the standard output stream.
- **Example Usage** `detect_storage_devices` The function will deliver an output listing all the storage devices and their relevant details.

11.6.186 Quality and Security Recommendations

1. Always sanitize input, if any, to prevent any potential code injection attacks.

2. Incorporate error handling to make the function more robust. This can include scenarios wherein the queried device information is not available.
3. Develop a unit test case for the function to ensure its accuracy and validity.
4. Avoid potential command injection by checking names of the block devices before executing commands.
5. Assign meaningful names to the variable to make the code more readable.
6. Document the function usage and its arguments properly for clarity and future reference. If the function's behavior changes, update the documentation timely.
7. Instead of directly accessing hardware related information, consider using system APIs or other safer methods, if available.
8. The function currently prints information to standard output (via echo). Instead, consider returning status codes indicating success or failure for greater control over function reactions to specific scenarios.

11.6.187 disks_free_list_simple

Contained in `lib/host-scripts.d/common.d/zpool-management.sh`

Function signature: `23bd3ac7877a7c90a934c9f11fb7a056bf39f675410585baaaa84e107157944c`

11.6.188 Function overview

The function `disks_free_list_simple()` is a Bash script that lists all available block storage devices in a system. It singles out regular, non-removable hard drives, discounting specific types of block devices such as loop, ram, md, dm devices, or those with partitions. Moreover, the function ignores devices that are mounted, belong to LVM2, Linux RAID, ZFS storage pool, or exist in a zpool. Whenever possible, it prefers World Wide Name (WWN) identifier for device representation.

11.6.189 Technical description

Below is a technical block describing the function:

- Name: `disks_free_list_simple()`
- Description: This function scans and lists the available block storage devices, taking into consideration specific exclusions.
- Globals: [None]
- Arguments: [No arguments are used with this function]
- Outputs: The function will output a list of free storage devices based on the criteria and filters it applies.
- Returns: No value returned.
- Example usage: After defining the function in bash, simply call `disks_free_list_simple` to use. Note this should be used in a Bash environment such as a script or the command line.

11.6.190 Quality and security recommendations

1. Ensure that you have proper permissions before running this function. It requires access to low-level disk information which might be restricted.
2. Extend error-handling mechanisms to include a fallback or feedback in case the function is not able to execute the commands correctly.
3. Validate that the output of each command is as expected, dealing particularly with edge cases where there might be unique devices or mounting schemes at play.
4. Carefully consider the security implications of exposing low-level hardware information. Apply necessary restrictions when and where needed, such as on a multi-user system.
5. Implement testing methodologies to maintain the function's integrity and efficiency.

11.6.191 `download_alpine_release`

Contained in `lib/functions.d/tch-build.sh`

Function signature: `f1438a9265c0b60b2947c704f30ee5980245150741b8578a6cb28e185b7cd407`

11.6.192 Function Overview

The bash function `download_alpine_release()` aims to download a specific version of the Alpine Linux distribution. If no version is specified, it will try to determine and download the latest available version. The specific requirements of this function include a specified Alpine version input and a path location under the `HPS_RESOURCES` environment variable where the downloaded file will be stored.

11.6.193 Technical Description

- **Name:** `download_alpine_release()`
- **Description:** This shell function is designed to download the specified version of the Alpine iso and save it under a defined path. If no version is specified, it downloads the latest version. If the file already exists in destination path, it does not download but returns the file path.
- **Globals:**
 - `HPS_RESOURCES`: The destination directory for the downloaded Alpine ISO. If not set, function will return error.
- **Arguments:**
 - `$1` (optional): Version of the Alpine Linux distribution to download.
- **Outputs:**
 - Outputs log messages via `hps_log()` function.

- Prints the file path of the downloaded ISO.

- **Returns:**

- 0 for successful downloads or if the file is already available.
- 1 for missing HPS_RESOURCES or for failure in detecting the latest Alpine version.
- 2 for failure in downloading the file.

- **Example Usage:**

```
download_alpine_release 3.20.2
```

11.6.194 Quality and Security Recommendations

1. Add input validation to ensure correct format of the `alpine_version` if provided.
2. Introduce a verbose mode to provide additional information during the download process.
3. Implement checksum validation after download to ensure the integrity of the downloaded iso.
4. Make the function more general by allowing the user to define which architecture (x86_64, armv7, etc.) they would like to download instead of always downloading the x86_64 version.
5. The HPS_RESOURCES variable should be verified not only as a non-empty string but also as a valid writable directory.

11.6.195 `download_file`

Contained in `lib/functions.d/system-functions.sh`

Function signature: 77275d6db2a730d2f09e1f7111242d9607b215be97269a81302557f4e047a820

11.6.196 Function Overview

`download_file` is a Bash function designed to download a file from the provided URL to the specified destination path. The function uses either `curl` or `wget` based on availability and supports resuming interrupted downloads. If provided, the function verifies the checksum of the downloaded file using `sha256sum`. If an error is encountered (missing arguments, failure to create the destination directory, download failure, or checksum mismatch), the function will log the error and return a non-zero exit code.

11.6.197 Technical Description

- **Name:** `download_file`
- **Description:** Bash function to download a file from a URL to a specified destination path, with optional SHA256 checksum verification.

- **Globals:** None.
- **Arguments:**
 - \$1: `url` - The URL of the file to be downloaded.
 - \$2: `dest_path` - The destination path where the downloaded file will be placed.
 - \$3: `expected_sha256` - (Optional) The expected SHA256 checksum of the downloaded file.
- **Outputs:** Information, warning and error logs.
- **Returns:** 0 if the file is successfully downloaded and the checksum (if provided) matches. 1 if necessary tools are missing or required arguments are not provided, 2 if any operation (e.g., directory creation or file download) fails, and 3 if the checksum does not match.
- **Example Usage:**

```
download_file "https://example.com/test.txt" "/path/to/test.txt"  
↪ "e3b0c44298fc1c149afbf4c8996fb92427ae41e4649b934ca495991b7852b855"
```

11.6.198 Quality and Security Recommendations

1. The `download_file` function currently lacks input validation. Malformed or malicious values for the URL or destination path could lead to unexpected behavior or security risks.
2. The function silently falls back to `wget` if `curl` is not available. Explicitly specifying the desired tool or alerting the user to the fallback could improve transparency and control.
3. If `sha256sum` is not available, the function logs a warning but continues the download. It might be preferable to fail outright to support secure environments where checksum verification is required.
4. The function only supports SHA256 for checksums. Supporting additional or newer checksum methods could improve versatility and security.
5. It may be worth considering a timeout for the download operation, to prevent hanging in the case of a slow or unresponsive server.

11.6.199 `download_iso`

Contained in `lib/functions.d/iso-functions.sh`

Function signature: `a769efc53df917e64e3dbdfb8acb70dff4b4cb4a89efd3b55a689c47cde86e91`

11.6.200 Function overview

The function `download_iso` retrieves a specific version of an Operating System (OS) installation ISO file based on certain parameters such as the CPU architecture, the manufacturer, the name and version of the OS.

11.6.201 Technical description

The function is defined as follows:

- **Name:** `download_iso`
- **Description:** Downloads an ISO file for a certain OS from a base URL, builds the proper download URL and the filename of the downloaded ISO file, checks if the ISO file already exists, and if not, downloads the ISO file and saves it to a designated directory. In case of download failure, it cleans up the failed download file. Currently, only “rockylinux” is supported as the OS name.
- **Globals:** No global variables used.
- **Arguments:**
 - \$1: CPU architecture
 - \$2: Manufacturer
 - \$3: OS name
 - \$4: OS version
- **Outputs:** Prints status messages about the download process, such as whether the ISO file already exists or if the OS variant is unsupported, the status of ISO download process, etc.
- **Returns:** 0 if the ISO file already exists or if it was successfully downloaded; 1 in case of any failures, such as unsupported OS variant or a failed download.
- **Example usage:** `download_iso x86_64 intel rockylinux 10`

11.6.202 Quality and security recommendations

1. Consider adding more OS support, not just Rocky Linux.
2. Catch and properly handle potential problems, such as issues with the directory creation (`mkdir -p "$iso_dir"`), to prevent unexpected results or vulnerabilities.
3. Implement checksum validation after download to ensure the downloaded ISO file is correct and wasn't tampered with during transit.
4. Refactor the case block for OS names into separate functions for better readability and maintainability.
5. Use secure coding practices and constant code reviews to avoid potential security issues.
6. Test the function thoroughly under different conditions and scenarios to ensure it behaves properly.

11.6.203 `enable_iscsi_target_reload_on_boot`

Contained in `lib/host-scripts.d/common.d/iscsi-management.sh`

Function signature: `c18133b271784d139d5b0a951926f1dac4277b5134a96a54ef074be5dfe03d5d`

11.6.204 Function overview

This bash script function, `enable_iscsi_target_reload_on_boot`, checks for the existence of the required iSCSI configuration files, verifies if the `target.service` is installed, and if everything is in place, it enables the iSCSI target to reload its configuration at the system boot. If any of these conditions are not met, it returns an error message and exits with a non-zero status code.

11.6.205 Technical description

- Name: `enable_iscsi_target_reload_on_boot`
- Description: This Bash function checks for the required iSCSI configuration files and services, and if found, enables the iSCSI target configuration to reload at system boot.
- Globals: None.
- Arguments: None.
- Outputs: Success or failure message depending on whether the operations were successful or not.
- Returns: The function will return 1 in case of failure (when required files or services are not found) else it will successfully run without explicit return value.
- Example usage: Simply calling the function without arguments is how it is to be used:

```
enable_iscsi_target_reload_on_boot
```

11.6.206 Quality and Security recommendations

1. Input sanitization: Since the function is not currently handling any inputs, this issue doesn't exist in the present case. However, if any arguments are introduced in future versions, be sure to sanitize them.
2. Error handling: It is recommended to handle other potential errors and exit conditions as well. The function currently checks for two potential issues, but other problems may occur.
3. Logging: To see the complete progress of function, you should add more echo statements at different stages. This will offer better debugging ability if something goes wrong in future.
4. Security enhancements: Validate the required privileges for running the function. In this case, running the script with necessary privileges is crucial for successful operation.
5. You can use more precise tools instead of `grep` to prevent false positives when searching for the 'target.service'. Do make sure that such tools are installed though.

11.6.207 `export_dynamic_paths`

Contained in `lib/functions.d/system-functions.sh`

Function signature: a6b2a47e08ed460524c491151fd944d515572f0fe9d26a1a2d5d310ad72b99b5

11.6.208 Function Overview

This Bash function, `export_dynamic_paths`, is designed to set and export paths dynamically within a cluster server environment. It makes use of local cluster names to base its operation while providing an alternative default directory path. This function is significant for managing multiple active clusters and ensuring that the active cluster is properly recognized. The function also sets and exports environment variables representing various paths in the cluster's configuration.

11.6.209 Technical Description

- **Name:** `export_dynamic_paths`
- **Description:** A Bash function designed to set and export cluster configuration paths dynamically using the provided cluster name as a reference. This includes the active cluster, the cluster's configuration directory, and the hosts configuration directory. The function also considers the case where an active cluster has not been specified.
- **Globals:** [`HPS_CLUSTER_CONFIG_BASE_DIR`: This global variable provides the root directory for storing cluster configs. By default, its value is set to `/srv/hps-config/clusters`]
- **Arguments:**
 - `$1`: This argument is the string value that represents the cluster name. If it is not provided, the function will use the currently active cluster (default to empty string).
- **Outputs:** Outputs a warning message “[x] No active cluster and none specified.” to `stderr` if no active cluster exists and none is specified by user.
- **Returns:** Returns 1 if there is no active cluster and none has been specified by the user, or 0 if execution was successful.
- **Example usage:** `export_dynamic_paths 'cluster_name'`

11.6.210 Quality and Security Recommendations

1. Validate inputs at the start of the function. Be sure that the supplied cluster name does not contain unsafe characters (e.g., slashes, backticks, etc.) that could potentially lead to command or path injection attacks.
2. Handle all error or exceptional scenarios. Improve error handling so that more specific messages are returned based on the failure's nature.
3. Make sure that proper permissions are set for the directories and files involved, especially when the function is handling paths and using these to access potentially sensitive data or system configurations.

4. Incorporate a logging mechanism to trace the function's behavior when debugging is required to aid in future troubleshooting.
5. Use more descriptive variable names for readability and maintenance. Ensure the variable and function names accurately describe their purposes or behavior.

11.6.211 `extract_iso_for_pxe`

Contained in `lib/functions.d/iso-functions.sh`

Function signature: 291539ccd925935805dd8c7d2a52eeb587e5f13f95dab12edd40bdbe8cf2a210

11.6.212 Function overview

This function `extract_iso_for_pxe` is primarily used to extract ISO files to a specific directory for PXE (Preboot eXecution Environment). The function is provided with the architecture, manufacturer, operating system details and version, and it tries to extract the corresponding ISO file if it hasn't been extracted before. If anything goes wrong during the extraction process, an error message is printed and the function returns with a failure code.

11.6.213 Technical description

- **name:** `extract_iso_for_pxe`
- **description:** Extracts an ISO file under a specified directory for PXE booting.
- **globals:** [`HPS_DISTROS_DIR`: Directory where the ISO file is to be extracted]
- **arguments:** [\$1: CPU architecture, \$2: Manufacturer, \$3: Operating system name, \$4: Operating system version]
- **outputs:** Prints status messages on console, indicating the progress of the operation.
- **returns:** Returns status code 0 if the operation is successful, or 1 if otherwise.
- **example usage:**

```
extract_iso_for_pxe "x86_64" "dell" "ubuntu" "20.04"
```

11.6.214 Quality and security recommendations

1. Ensure that the function parameters, especially `iso_dir` and `extract_dir` are properly sanitized to avoid directory traversal attacks.
2. Implement checks for the existence of the `bsdtar` utility before proceeding with the extraction. If `bsdtar` is not found, exit the function early with a suitable return code and message.
3. Perform more robust error handling. For example, it would be helpful to check whether the directory creation was successful before proceeding with the extraction.

4. Consider adding additional checks or accept only whitelisted inputs to reduce possibilities of misuse or error.
5. Input validation can go beyond the file system. For instance, validate the architecture and OS version parameters to confirm they're valid before proceeding.

11.6.215 `extract_rocky_iso_for_pxe`

Contained in `lib/functions.d/iso-functions.sh`

Function signature: `52e8a2e1e3a65096a2e0e1ac3696f4c21521a7cc3c3b006a7056ede3446b2ada`

11.6.216 Function overview

The function `extract_rocky_iso_for_pxe()` is created to extract a Rocky Linux ISO for PXE (Preboot Execution Environment). This extraction is required to prepare the PXE boot environment. The function takes in three arguments: the path to the ISO, the version of the Linux OS, and the CPU architecture information. It then prepares a location to extract the ISO contents and performs the operation using `bsdtar` or `fuseiso`. If neither utility is available, the function logs an error message and exits with a non-zero status code.

11.6.217 Technical description

- **name:** `extract_rocky_iso_for_pxe`
- **description:** The function aims to extract a Rocky Linux ISO for PXE (Preboot Execution Environment).
- **globals:**
 - `VAR: HPS_DISTROS_DIR`: a variable holding the base directory for Linux OS distributions.
- **arguments:**
 - `$1: iso_path`: Path to the Rocky Linux ISO file.
 - `$2: version`: The version number of the Linux OS.
 - `$3: CPU`: The required CPU architecture.
- **outputs:**
 - Logs about the extraction process.
 - An error message if neither `bsdtar` or `fuseiso` commands are installed.
- **returns:**
 - 0 if the Rocky Linux ISO extraction is successful.
 - 1 if neither `bsdtar` or `fuseiso` commands are installed.
- **example usage:**

```
extract_rocky_iso_for_pxe "/path/to/iso" "8.4" "x86_64"
```

11.6.218 Quality and security recommendations

1. Always sanitize inputs to the function to prevent any possible malicious or unintended actions.
2. Make sure to check the existence of the ISO file passed as an input, and fail fast, if it doesn't exist.
3. Use full paths to commands like bsdtar, fuseiso, etc. to avoid dependency on PATH and prevent potential "command not found" errors or command injection vulnerabilities.
4. Be mindful of returning appropriate status codes on failure scenarios, it helps to debug faster.
5. Consider providing more informative and meaningful log messages for clarity about each operation.
6. Always handle potential errors in filesystem operations like mkdir, cp, etc.
7. In the "globals" section, important potential side effects of using global variables in a function like potential conflicts with other scripts/environment variables should be kept in mind.

11.6.219 fetch_and_register_source_file

Contained in `lib/functions.d/prepare-external-deps.sh`

Function signature: 97650ca173393457a3e6151b2fbb0a997e425c3010fc3c13018364e050618d70

11.6.220 Function Overview

The function `fetch_and_register_source_file()` is designed to first download a file from a specified url and then register this file using a given handler. The function accepts three arguments; the url from which the file will be downloaded, the handler which will be used to process the file after download and the filename of the downloaded file. If the filename is not provided, the function will take the base name from the url.

11.6.221 Technical Description

- Name: `fetch_and_register_source_file`
- Description: This function is used to download a file from a given url and then register it using a provided handler.
- Globals: None.
- Arguments:
 - \$1: url from which the file will be downloaded.
 - \$2: handler used to process the file after the download.
 - \$3: filename of the downloaded file. If not provided, the function will derive it from the given url.
- Outputs: The function performs the task of downloading and registering a file. There are no specific output returns on the console.

- **Returns:** The function will return the status of the last command executed within it. Thus, if both the fetch and register operations are successful, the function will return 0. If either operation fails, the function will return the corresponding error status.
- **Example usage:**

```
fetch_and_register_source_file "http://example.com/file.txt"  
↪ "myHandler"
```

11.6.222 Quality and Security Recommendations

1. **Validate inputs:** For robustness and security, validate the input parameters to ensure they are in the correct format and have valid values.
2. **Handle download issues:** Consider adding more granular error handling for the `fetch_source_file` function to allow the function to easily troubleshoot issues related to file download.
3. **Handle registration issues:** Similarly, address possible failure points in `register_source_file` with adequate error handling and messaging.
4. **Check handler validity:** Before invoking the handler on the downloaded file, ascertain if the handler exists and can be executed safely.

11.6.223 `fetch_source_file`

Contained in `lib/functions.d/prepare-external-deps.sh`

Function signature: `9851dec43ce9f3e94856800874741f8ff28deda9346709daaa88282763e49999`

11.6.224 Function overview

The `fetch_source_file` function is a utility to download a file from a given URL and save it on a specified destination directory. It first checks if the file already exists in the destination directory. If the file does not exist, it downloads the file using `curl` command. If the `filename` argument is not provided, it infers the filename from the URL. The default directory to save the file is `/srv/hps-resources/packages/src`, nevertheless, it can be overwritten through `HPS_PACKAGES_DIR` environment variable.

11.6.225 Technical description

- **Name:** `fetch_source_file`
- **Description:** Fetches a file from a provided URL and saves it in a specific location.
- **Globals:** [`HPS_PACKAGES_DIR`: Specifies the root directory for saving the downloaded files]

- **Arguments:** [\$1: URL of the file to be downloaded, \$2: Name of the downloaded file]
- **Outputs:** Logs to stdout showing the progress and result of the download.
- **Returns:** 0 if the file is successfully downloaded or already exists on the server, 1 if the download fails.
- **Example usage:**

```
fetch_source_file "https://example.com/file.zip" "myFile.zip"
```

11.6.226 Quality and security recommendations

1. Input validation: As the function downloads content from a URL, it is advisable to add checks for ensuring that the URL is properly formatted and secure (uses `https://`, belongs to trusted domain etc).
2. Error handling: While the function checks if the file is successfully downloaded, it would be better to add error handling for other operations as well like creating directory.
3. Sanity checks: It would be safer to add some sanity checks on the downloaded file, like checking its size, verifying its checksum and more.
4. Avoid using global variables: The use of global variables makes code hard to predict and debug, it is better to pass them as parameters to the functions.
5. Logging: consider redirecting error logs to stderr consistently. In the current case, some logs are written to stdout, some - to stderr, which might be cumbersome to debug if the function is used in a script.

11.6.227 `find_hps_config`

Contained in `lib/functions.sh`

Function signature: `2b1359b9639780889fba67b113809b966b3326fe6600551c16100c71c64f76ef`

11.6.228 Function Overview

This function, `find_hps_config()`, is a simple Bash script function that iterates over an array `HPS_CONFIG_LOCATIONS`. It checks for existing files in the specified locations. Once it comes across the first existing file, it assigns the location to a local variable `found` and breaks the loop. The function then outputs the value of `found` and returns a success status code 0. If no files are found in any of the locations in the array, the function returns a failure status code 1.

11.6.229 Technical Description

- **Name:** `{find_hps_config()}`

- **Description:** This function iterates over an array `HPS_CONFIG_LOCATIONS`, checking each location for existing files. The first existing file found is assigned to the local variable `found`, which is then echoed out. The function returns 0 when a file is found and 1 when not.
- **Globals:** [`HPS_CONFIG_LOCATIONS`: An array of file paths where the function will look for existing files]
- **Arguments:** [None]
- **Outputs:** The filepath of the first existing file found in the array `HPS_CONFIG_LOCATIONS`. If no file is found, it does not produce any output.
- **Returns:** 0 if an existing file location is found, 1 if not.
- **Example Usage:**

```
bash    HPS_CONFIG_LOCATIONS=("/path/to/file1"
"/path/to/file2" "/path/to/file3")    find_hps_config
```

11.6.230 Quality and Security Recommendations

1. Ensure the array `HPS_CONFIG_LOCATIONS` only contains paths to trusted, secure locations to prevent any unintended access to malicious files.
2. Check whether files found at any of the locations are not just existing but also readable before assigning to the variable `found`.
3. Document this function on usage and expected inputs and outputs, especially since it does not have any named input arguments. This will help avoid any potential misuse of the function.
4. Implement error handling to capture scenarios where the `HPS_CONFIG_LOCATIONS` is either not an array or is an empty array. This will improve the function reliability.
5. Validate and sanitize the file paths to ensure they are not misused to access unintended parts of the file system, adding an extra layer of security.

11.6.231 `format_mac_colons`

Contained in `lib/functions.d/network-functions.sh`

Function signature: `95ebd5c4198ab3943c5263e6ec4e563812216afc2059562fef0d3c1357dc53ca`

11.6.232 Function Overview

The `format_mac_colons` function is designed to validate a MAC address input, convert it to lowercase, and insert colons in the MAC address format. This simplifies the process of formatting MAC addresses to adhere to standard MAC address representation.

11.6.233 Technical Description

Function name: `format_mac_colons`

Description: This function validates a supplied MAC address to ensure it consists of exactly 12 hexadecimal characters. After validation, it converts the MAC address to

lowercase and inserts colons in the appropriate position to adhere with the standard representation of a MAC address.

Globals: None

Arguments: - \$1: This argument should be a MAC address of 12 hexadecimal number.

Outputs: If valid, a MAC address with colons inserted would be produced. If invalid, the function will output an error message directed to the standard error output (stderr).

Returns: - 1: if supplied MAC address is invalid. - Formatted MAC address: if the supplied MAC address is valid.

Example usage:

```
format_mac_colons "123456abcdef"
```

11.6.234 Quality and Security Recommendations

1. Input Sanitization: Ensure that the input to the function is properly sanitized before it is processed by the function.
2. Error Handling: The function should have plans to handle unexpected inputs (like a null input or an input with non-hexadecimal characters) and not just invalid length.
3. Include a clear and descriptive comment on what the function does at the beginning of the function. This includes listing out all assumptions, preconditions and postconditions.
4. Add more detailed output messages that can guide the user on the type and format of input the function requires whenever invalid inputs are provided.
5. It would be beneficial to provide more return codes to cover other potential error situations, such as an empty input string. Making return codes more descriptive can make debugging simpler.
6. Always ensure to follow the least privilege principle - only give the minimum amount of permissions necessary for the function to execute.

11.6.235 generate_dhcp_range_simple

Contained in `lib/functions.d/network-functions.sh`

Function signature: `0b96ef1d9a801dba7584c3c03ea9f1327da1cd2685d343519ee3ae72aa66ecd8`

11.6.236 Function Overview

This function, `generate_dhcp_range_simple()`, is used to generate a DHCP range for a given local network provided in CIDR notation. The range generated is based on the network's gateway IP and its broadcast IP, as well as an optional count parameter that specifies the number of IP addresses in the range.

11.6.237 Technical Description

Function Name: `generate_dhcp_range_simple()`

Description: This function generates a range of IP addresses suitable for use as a DHCP dynamic address pool. The IP range is calculated using the provided local network (in CIDR notation), a gateway IP, and an optional count parameter, which defaults to 20 when not specified.

Globals: None

Arguments: - \$1: `network_cidr` (e.g. `192.168.50.0/24`) - \$2: `gateway_ip` (e.g. `192.168.50.1`) - \$3: `count` (default: 20)

Outputs: A range of IP addresses for DHCP purposes is generated.

Returns: The function doesn't return anything, it just executes side effects.

Example Usage:

```
generate_dhcp_range_simple "192.168.50.0/24" "192.168.50.1"
```

11.6.238 Quality and Security Recommendations

1. For improved data validation, consider adding checks to ensure the CIDR and gateway IP are valid before the function uses them.
2. To avoid potential information leak, consider muting `ipcalc` command internal verbose logs if not necessary for active debugging.
3. Test the function thoroughly using different network setups and gateway IPs to ensure it functions correctly under different attribute values.
4. Consider integrating a logging system for easier debugging.
5. Make sure the script has appropriate permissions, limiting the exposure of the function to avoid unauthorized use. Remember, script that alters network configuration could be a potential security risk if misused.

11.6.239 `generate_initramfs_script`

Contained in `lib/functions.d/tch-build.sh`

Function signature: `d63236f298e94334bdf196a44afa8f541d4e405dc6d1a0e18890a7520df8f974`

11.6.240 Function overview

The `generate_initramfs_script` function creates a bootstrap script that is executed post-boot. It first ensures that the `/sysroot` directory exists, then writes the RC script content dynamically, using the `generate_rc_script` function, into the file `/sysroot/etc/local.d/hps-bootstrap.start`. It finally gives execution permissions to `hps-bootstrap.start` script and informs the user that the bootstrap script has been installed.

11.6.241 Technical description

Function Name: generate_initramfs_script

Description: This function generates a script that executes after boot. It checks for the presence of /sysroot directory, then, it dynamically writes the RC script content into a file /sysroot/etc/local.d/hps-bootstrap.start, gives it execute permissions and outputs a success message.

Globals Variables: None.

Arguments: - \$1: Represents the gateway IP address.

Outputs: 1. Success message informing the user of the installed bootstrap script. 2. Error Message if /sysroot directory is not found.

Returns: Does not return a value. If /sysroot directory is not found, the script execution stops with an exit status of 1.

Example Usage:

```
generate_initramfs_script "192.168.1.1"
```

11.6.242 Quality and security recommendations

1. Check validity of the input argument: the function should not just accept any argument as IP address. The function should incorporate data sanity checks to ensure that the passed argument is a valid IP address.
2. Use more descriptive error message: when the /sysroot directory is not available, consider providing suggestions or possible steps that a user could take to resolve the error.
3. Ensure Secure Coding Practices: practice secure coding to protect against injection and other attacks. For instance, don't blindly trust the input passed into generate_rc_script function.
4. Consider better error handling: rather than just printing output to stdout, consider logging errors to a dedicated error log for better diagnostic and debugging capability.
5. Validate successful execution of commands: rather than just displaying a completion message, ensure that commands were successfully executed before displaying the message or proceed to next command.
6. Keep system and programming practices up-to-date. Regularly check for and apply updates and patches to your systems and coding tools, to maintain pace with evolving threats.

11.6.243 generate_ks

Contained in lib/functions.d/kickstart-functions.sh

Function signature: b973477b6d7653d80151f03150fadcf0d8f1f85a77f353025a267d5257c8529

11.6.244 Function overview

The function `generate_ks()` handles the generation of a kickstart file, used to automate the installation process of an operating system. This function exports different variables containing configuration parameters needed for the OS installation and finally prepares the installation script for rendering.

11.6.245 Technical description

- **Name:** `generate_ks()`
- **Description:** This function accepts two parameters. It exports different variables like `macid`, `HOST_TYPE`, `HOST_NAME`, etc., used for host configuration. It uses helper functions like `hps_log()`, `cgi_header_plain()`, `host_config()`, `cluster_config()`, and `script_render_template()` to log information, get and set configurations, and render the installation script respectively.
- **Globals:** [`macid`: first argument, represents MAC ID / Unique Host ID of a host machine, `HOST_TYPE`: second argument, represents the type of the host machine]
- **Arguments:** [`$1`: MAC ID / Unique Host ID of a host machine, `$2`: Type of the host machine]
- **Outputs:** Logs about the function calls and configurations, Along with the final render of the installation script.
- **Returns:** Does not return any value. However, the function could be interrupted and denote failure with a return of 1.
- **Example Usage:**

```
generate_ks "00:0c:29:c0:94:bf" "CentOS"
```

11.6.246 Quality and security recommendations

1. Avoid usage of global variables as much as possible to reduce side effects and improve function modularity.
2. Input validation and error handling should be more rigorous. Check for the validity of MAC addresses before using them.
3. Preferably restrict access to root or privileged users to reduce potential security risks.
4. Avoid inlining large scripts. Use a separate file to manage large scripts.
5. Implement a logging mechanism to record errors and important events during installation.
6. If possible, encrypt sensitive data like IP addresses, hostnames, and other network details.

11.6.247 `generate_opensvc_conf`

Contained in `lib/functions.d/opensvc-functions.sh`

Function signature: `5c8ee3ce32d8ed63bb644f40505a6e6f8d4d0abf12ff9bc2de331baa6e0db1f7`

11.6.248 Function Overview

The Bash function `generate_opensvc_conf` is designed to generate a configuration file for the OpenSVC v3 agent node. It starts by setting some local variables and fetching data from the host configuration using the `host_config` function and from the cluster configuration using the `cluster_config` function. It checks the type of service and assigns corresponding tags, then establishes some cluster-scoped variables. Finally, it assigns static paths to certain directories and files before generating and printing a configuration section by section.

11.6.249 Technical Description

- Name: `generate_opensvc_conf`
- Description: This function is used to generate the conf file for an OpenSVC v3 node agent by fetching and setting various necessary settings and variables originating from both host and cluster configurations.
- Globals: None
- Arguments: `$1 (ips_role)`: The role of the IPS, if different it will be set to 'provisioning'
- Outputs: A generated OpenSVC v3 Node Configuration.
- Returns: None
- Example Usage: This function can be used without arguments, as follows:
`generate_opensvc_conf`

11.6.250 Quality and Security Recommendations

1. It would be beneficial to include validation checks for each variable being set or fetched from the host & cluster configurations.
2. Error handling could be improved by including more explicit failure messaging and by using exit status values to define specific error types.
3. In terms of security, be careful when echoing variables directly into output, especially if they contain untrusted input, as this could lead to command injection. To avoid this, consider using quoted arguments wherever possible.
4. It would be useful to add comments explaining the purpose and workings of more complex parts of the script. This enhances maintainability, making it easier for future developers to understand the function.
5. Finally, comprehensive unit testing will ensure that changes to the function don't inadvertently break the intended behavior or introduce new bugs.

11.6.251 `generate_rc_script`

Contained in `lib/functions.d/tch-build.sh`

Function signature: 9113ea3e96cf28a1e93526bcfc049491e6848d350162947286fddb23bf7dd405

11.6.252 Function Overview

The `generate_rc_script()` is a shell function that dynamically generates a shell script which will update the package repository list of an Alpine Linux system, update its package index, installs `bash` and `curl`, and execute a `boot_manager.sh` shell script fetched from a specified gateway. This function is typically utilized in a network boot environment.

11.6.253 Technical Description

- Name: `generate_rc_script()`
- Description: This function generates a shell script that modifies the package repository list to point to a particular gateway IP address, updates the system's package index, installs `bash` and `curl`, and executes a shell script located on the gateway.
- Globals: None.
- Arguments:
 - `$1`: The IP address of the gateway. This IP address is used to point the package repositories and retrieve the `boot_manager.sh` script.
- Outputs: Outputs a shell script that performs the aforementioned actions.
- Returns: None.
- Example usage:

```
generate_rc_script 192.168.1.1
```

11.6.254 Quality and Security Recommendations

1. Consider using a secure protocol such as `HTTPS` for the repository and the `boot_manager.sh` script to prevent man-in-the-middle attacks.
2. The function could be improved by making it more flexible and allowing the Alpine version and the package list to be defined as variables.
3. Sanity checks for the input argument (IP Address) could be added to ensure it is in the correct format and the targeted server is reachable.
4. Error handling can improve the resilience of the script - for instance, actions could be taken if `apk` update or the script download fails.
5. Sensitive operations such as changing repositories and executing scripts fetched from the network should be run with minimal privileges to limit potential damage.

11.6.255 `get_active_cluster_dir`

Contained in `lib/functions.d/cluster-functions.sh`

Function signature: `bea86c9b08f47ce60a2602b8a8391990588b8ca0b96c0252ba0070cc3623cfce`

11.6.256 Function overview

The `get_active_cluster_dir` function is fundamental to resolving symlinks and retrieving the path of the currently active cluster directory in a system. It defines multiple local variables to find and check the authenticity of this directory, subsequently echoing its path if satisfactory conditions are fulfilled.

11.6.257 Technical description

- **Name:** `get_active_cluster_dir`
- **Description:** This function works to resolve the symlink of the active cluster and retrieve its directory. It does this by storing the symlink and its full path in local variables, verifying their validity, checking if the stored full path is a directory, and if so, outputs the directory. It returns an error if the symlink fails to resolve or the target isn't a directory.
- **Globals:** None.
- **Arguments:** None.
- **Outputs:** If successful, the function will print the directory of the active cluster. Error messages will be printed to `stderr` in any of the following cases:
 - The symlink can't be resolved.
 - The active cluster target is not a directory.
- **Returns:** 0 if the function is successful. Returns 1 if the symlink can't be resolved or the target is not a directory.
- **Example Usage:**
`get_active_cluster_dir`

11.6.258 Quality and security recommendations

1. Verify if the `get_active_cluster_link` function and the `get_cluster_dir` function used within `get_active_cluster_dir` are secure and optimized. The efficiency of `get_active_cluster_dir` is highly dependent on the performance of these two.
2. Error messages are redirected to `stderr`, which is a best practice and should be continued.

3. Ensure that this function is running with the right privileges - it should not have more permissions than what is required to read the link and possibly traverse directories.
4. Sanitize the output of the `readlink` and `basename` command to prevent potential path manipulation or symbolic link attacks.
5. Implement unit tests for this function to safeguard it from potential edge cases and bugs.

11.6.259 `get_active_cluster_file`

Contained in `lib/functions.d/cluster-functions.sh`

Function signature: `6461e5f5814cbc510b1bad68822a5b0d8eb3c58618d50b5418200c62ea6dff01`

11.6.260 Function overview

The bash function `get_active_cluster_file()` is designed to retrieve the name of the active cluster saved in a file and output its content.

11.6.261 Technical description

Definition:

- **name:** `get_active_cluster_file`
- **description:** This function retrieves the name of the “active” cluster from the method `get_active_cluster_filename`, assigns it to a local variable `file` and outputs its content, i.e., the active cluster’s data. If the `get_active_cluster_filename` method fails, the function will exit and return 1.
- **globals:** None
- **arguments:** None
- **outputs:** The contents of the file retrieved from `get_active_cluster_filename`.
- **returns:** Content of the file if successful, 1 if the `get_active_cluster_filename` fails.
- **example usage:**

```
get_active_cluster_file
```

11.6.262 Quality and security recommendations

1. Include more error handling for situations where file does not exist or it fails to be read by `cat` command.
2. Ensure that the file reading process is secure and its contents are not accessible to unauthorized users. This can be achieved by setting appropriate permissions on the file.

3. Sanitize output to minimize the potential impact of malicious data.
4. The function should not trust the file's content blindly, it should validate the input before processing it. This is important to prevent possible code injection flaws.
5. Include a more comprehensive documentation, specifying what the function expects as an input and output. This will be beneficial for users of the function. Keep improving the unit tests aiming at improved code coverage.

11.6.263 `get_active_cluster_filename`

Contained in `lib/functions.d/cluster-functions.sh`

Function signature: 20bf828d972eed107690831358c6e9148058e88842050f23b8944d669bfab00a

11.6.264 Function Overview

The `get_active_cluster_filename` function is designed to find the configuration file for the currently active cluster. It performs the following steps: 1. It obtains the path to the active cluster's directory. 2. It extracts the name of the active cluster from the directory path. 3. It retrieves the cluster configuration file based on the cluster's name. 4. If the file is not found, it prints an error message and returns 1. 5. If the file is found, it prints the file.

11.6.265 Technical Description

- **Function Name:** `get_active_cluster_filename`
- **Description:** This function gets the configuration file name for the currently active cluster.
- **Globals:** None
- **Arguments:** None
- **Outputs:**
 - The path to the configuration file for the active cluster.
 - An error message if the configuration file does not exist.
- **Returns:**
 - Returns 1 when the active directory or the configuration file of the cluster does not exist.
 - Returns 0 otherwise.
- **Example Usage:**

```
filename=$(get_active_cluster_filename)
```

11.6.266 Quality and Security Recommendations

1. To improve readability and maintainability, consider adding comments explaining what each command does.

2. It's good practice to test return values directly in `if` statements rather than relying on `|| return 1` expressions.
3. We should validate user inputs where possible, and handle errors explicitly.
4. Consider using a variable to capture the error message string, which would allow you to change the message in just one place if needed.
5. Make sure to use secure coding practices, such as not making assumptions about input data, checking for null or unexpected values, and handling errors and exceptions as much as possible.

11.6.267 `get_active_cluster_info`

Contained in `lib/functions.d/cluster-functions.sh`

Function signature: `51d7c68169c79cb4271c76d4914a13afce322dd3dab4f93ccfeba936798070ea`

11.6.268 Function Overview

The `get_active_cluster_info` function is a bash function designed to gather and display information about currently active clusters. This function first calls on another function, `_collect_cluster_dirs`, in order to get a list of cluster directories. Then the function checks if there are any active clusters. If there are none, it will print an error message and end the program. If there are active clusters, the function will proceed to print specific information about each cluster.

11.6.269 Technical Description

- **Name:** `get_active_cluster_info`
- **Description:** This function gathers and displays information about currently active clusters.
- **Globals:** `HPS_CLUSTER_CONFIG_BASE_DIR`: It holds the base directory path that the function will check for active clusters.
- **Arguments:** None
- **Outputs:** Error message if no clusters are found, otherwise list of directories stored in `dirs`.
- **Returns:** Returns 1 if no clusters are found, otherwise returns 0.
- **Example Usage:** `get_active_cluster_info`

11.6.270 Quality and Security Recommendations

1. The function should include validation for the global variable `HPS_CLUSTER_CONFIG_BASE_DIR` to ensure it is correctly defined and it points to a valid directory.
2. Use descriptive error messages to allow for easier debugging, and in those error messages include potential solutions for the issues.

3. Ensure that the `_collect_cluster_dirs` function is properly securing and validating the data that it is returning.
4. Check directories for appropriate read permissions before attempt to collect directories.
5. It would be recommended to also provide some form of error handling, for scenarios when the function `_collect_cluster_dirs` isn't defined or fails to execute.
6. Consider adding logging functionality for tracking warnings or errors. This will help in maintaining the system and diagnosing problems.

11.6.271 `get_active_cluster_link`

Contained in `lib/functions.d/cluster-functions.sh`

Function signature: 8840658f2d249224284adc89b84feddc787683b794880b06d7f3d566dc412130

11.6.272 Function overview

The `get_active_cluster_link` function is a Bash function used to get the active cluster link in a system. It first assigns the output of `get_active_cluster_link_path` function to the `link` variable and then checks if this link is a symbolic link. If not, it prints an error message and returns 1. If the link is a symbolic link, it simply echoes the `link` - printing the link.

11.6.273 Technical description

- **Name:** `get_active_cluster_link`
- **Description:** This function is used to retrieve the link to the active cluster in a system. It returns an error if the link does not exist or isn't a symbolic link.
- **Globals:** None
- **Arguments:** None
- **Outputs:** If successful, prints the active cluster link to stdout. If not, an error message is printed to stderr.
- **Returns:** The function returns 0 if the active cluster link is retrieved successfully, and 1 if either no link exists or the link isn't a symbolic link.

- **Example usage:**

```
get_active_cluster_link
```

11.6.274 Quality and security recommendations

1. The function does not have any arguments. As such, it would run with any number of arguments provided. To improve quality, error checking can be added to enforce

that no arguments are expected.

2. It assumes that `get_active_cluster_link_path` function has already been defined and works correctly. To improve maintainability, there should always be a check if a function exists before it's called.
3. Logging level of error could be standardized. Instead of directly printing to `stderr`, a logging function can be used to control the logging levels.
4. Error messages printed are currently hardcoded strings. To increase maintainability and readability, these error messages can be converted to constants at the top of the script or inside a configuration file.
5. For security improvements, input validation is a must. For this specific function, checking that the path points to an expected location can prevent symbolic link attacks. For instance, the bash function “`realpath`” could be used.
6. Lastly, the function needs to handle permissions errors gracefully. It can use defensive programming practices to ensure that the user running the script has the authority to access the link.

11.6.275 `get_active_cluster_link_path`

Contained in `lib/functions.d/cluster-functions.sh`

Function signature: 00422606e36c805412e226764ea91ac8d9620442c34f9c3bf83c8de949497756

11.6.276 1. Function Overview

The function `get_active_cluster_link_path` is used to echo an active cluster link path existing in the `${HPS_CLUSTER_CONFIG_BASE_DIR}` directory. It is specifically designed to get the path for the current active cluster configuration file within a High-Performance Computing (HPC) cluster environment. This function does not accept any arguments or modify any global variables.

11.6.277 2. Technical Description

11.6.277.1 Name

`get_active_cluster_link_path`

11.6.277.2 Description

The `get_active_cluster_link_path` function is used to output the path of the active cluster config file. It does this by appending the string “active-cluster” to the `HPS_CLUSTER_CONFIG_BASE_DIR` environment variable using forward slash (/) as the separator to build the file path for the active cluster config file.

11.6.277.3 Globals

[HPS_CLUSTER_CONFIG_BASE_DIR: The directory containing the cluster config files]

11.6.277.4 Arguments

This function does not require any arguments.

11.6.277.5 Outputs

The output is a string indicating the path of the active cluster config file. Example: /path/to/HPS_CLUSTER_CONFIG_BASE_DIR/active-cluster

11.6.277.6 Returns

This function will return 0 on successful execution.

11.6.277.7 Example Usage

```
# Get the path of the active cluster config file
active_cluster_path=$(get_active_cluster_link_path)
echo ${active_cluster_path}
```

11.6.278 3. Quality and Security Recommendations

1. This function does not perform any error checking. Therefore, it's recommended to add error handling to deal with the situation where HPS_CLUSTER_CONFIG_BASE_DIR might not be set or could be set to an invalid directory.
2. Confirm that the active-cluster file actually exists before attempting to return its path.
3. Protect against Command Injection attacks by sanitizing HPS_CLUSTER_CONFIG_BASE_DIR if it's externally controlled data.
4. Always quote your variables - in this case \${HPS_CLUSTER_CONFIG_BASE_DIR} - to avoid word splitting and globbing. This is important if there are spaces or special characters in the names.
5. Consider adding comments to explain what the function is doing a bit more clearly, especially if other people who are not familiar with the code will be reading or maintaining it.

11.6.279 get_active_cluster_name

Contained in lib/functions.d/cluster-functions.sh

Function signature: 56eb8b1d52908fb62f61b716d0ffccedd95bf6c229314f23d9a10385163e0cfd

11.6.280 Function overview

The function `get_active_cluster_name()` is a shell function that retrieves the name of the currently active cluster by utilizing other helper functions. It sets the `dir` variable as the active cluster directory obtained from `get_active_cluster_dir` function. Once the directory is obtained, it retrieves only the last segment from the pathname that denotes the active cluster's name.

11.6.281 Technical description

Name: `get_active_cluster_name()`

Description: This function retrieves the name of the active cluster from the cluster directory's path. It uses the `get_active_cluster_dir` function to get the active directory first, then leverages the `basename` utility to retrieve the active cluster's name.

Globals: None

Arguments: None

Outputs: The active cluster's name.

Returns: It returns 0 on successful execution and 1 if the `get_active_cluster_dir` function fails to execute.

Example Usage:

```
active_cluster=$(get_active_cluster_name)
echo "Active cluster is: $active_cluster"
```

11.6.282 Quality and security recommendations

1. Always quote your variable expansions. For instance, `basename -- "$dir"`.
2. In this function, the `get_active_cluster_dir` function is used, but it cannot handle the case if the function isn't defined. Error handling for this use case should be taken into account.
3. Avoid globals as much as possible as they can produce unpredictable side effects, which can be difficult to debug and maintain.
4. Validate user-defined inputs for security concerns to prevent injection attacks.
5. Write comments for your functions and complex code sections for better readability and maintainability.
6. Always consider edge cases while writing the function. For example, if there is no active cluster, the function should handle this scenario gracefully.
7. Writing unit tests for the functions is a good practice to ensure that they work as expected. It helps to detect function errors, gaps, or missing requirements.

11.6.283 `get_all_block_devices`

Contained in `lib/functions.d/storage_functions.sh`

Function signature: add7deb6a087d72238984be839fb5488e13aff3ae7251f93f2c1814d79162625

11.6.284 Function overview

The `get_all_block_devices` function is used to retrieve all the block devices in a system where their type is 'disk'. It iterates over all block devices in `/sys/block` directory and which uses a helper function named `get_device_type` to get the type of the device. We get the basename, or the least significant part of the device's path in this instance, to identify the device. If it is a disk, then the device name is printed to the standard output.

11.6.285 Technical description

- **Function Name:** `get_all_block_devices`
- **Description:** This function retrieves all the block devices whose type is 'disk' from a system.
- **Globals:** `devname`: contains the name of the device.
- **Arguments:** None.
- **Outputs:** The function outputs to stdout the names of all block devices which are of type 'disk'.
- **Returns:** No explicit return value. Success or failure can be inferred from the lack or presence of output.
- **Example Usage:** `get_all_block_devices`. It needs no arguments.

11.6.286 Quality and security recommendations

1. Validation of the device path: The function directly accesses the `/sys/block/` path. The command `basename` can potentially fail if the path does not exist. Hence, validation of the actual device path before processing can improve robustness.
2. Error handling: There is no explicit error handling in case the `get_device_type` returned value is not 'disk' or some other error occurs. It would be beneficial to add error handling mechanisms to improve the robustness of the code.
3. Defensive programming: There are no checks to ensure that the function operates as expected in an abnormal or unanticipated scenario. Therefore, enhancing the function with more checks would increase its resilience.
4. Documentation: There is no comment in the function explaining what it does and how it works. Good documentation makes it easier to maintain and debug the code in the future.

11.6.287 `get_alpine_bootstrap`

Contained in `lib/functions.d/tch-build.sh`

Function signature: 6d5c7768966553e101da81968a5d1d06f7a887f781e76db0807a69dcaef396ac

11.6.288 Function overview

The `get_alpine_bootstrap` function provides a mechanism for generating a bootstrap script for a given stage within an Alpine Linux environment. The function does this by retrieving the gateway IP from the cluster configuration and using it as part of creating either an `initramfs` or `rc` bootstrap script, depending on the stage specified.

11.6.289 Technical description

- **name:** `get_alpine_bootstrap`
- **description:** This function is intended to generate a bootstrap script for a specified stage within an Alpine Linux environment. The specific stage (either 'initramfs' or 'rc') is indicated by the passed argument. If no stage is passed, the default is 'initramfs'. The gateway IP gathered from the cluster configuration is utilized in generating the appropriate bootstrap script.
- **globals:** [`gateway_ip`: This variable holds the gateway IP retrieved from the cluster configuration.]
- **arguments:** [`$1`: Stage indicator, it determines if an 'initramfs' or 'rc' bootstrap script is to be generated. The default is 'initramfs' if no argument is passed.]
- **outputs:** It outputs an error message with `hps_log` function in case of failure in getting the gateway IP from cluster configuration or if an invalid stage parameter is passed.
- **returns:** If an invalid stage parameter is passed, it returns with an exit code of 1.
- **example usage:**

```
get_alpine_bootstrap initramfs
```

11.6.290 Quality and security recommendations

1. Secure the IP retrieval by validating the response from the `cluster_config get DHCP_IP` command. For instance, check if the returned IP is valid to prevent potential security vulnerabilities.
2. Incorporate more error checks to ensure the reliability of the script. For instance, check if the `generate_initramfs_script` and `generate_rc_script` functions get executed successfully.
3. Validate the stage argument before using it to ensure it is properly formatted and as expected, rather than after. This helps prevent potential security risks from improperly formatted or unexpected arguments.
4. Implement logging for audit and troubleshooting purposes. Logs can facilitate problem identification and can be valuable when troubleshooting an issue.
5. Avoid using globals where possible, by passing variables as arguments to functions, as globals may lead to name clashes and hard to debug issues.

11.6.291 `get_client_mac`

Contained in `lib/functions.d/network-functions.sh`

Function signature: `8cf7b608ec665ff47f16996d542d5d3ed0c9392116e453e8abfcc00eab616973`

11.6.292 Function Overview

The function `get_client_mac` is used to extract the MAC address corresponding to a given IP address in a local network. It sends a ping to trigger an ARP update, which is then parsed for the required MAC address. If the initial method does not succeed, it uses the `arp` command as a fallback and returns a normalized MAC address.

11.6.293 Technical Description

- **Name:** `get_client_mac`
- **Description:** This function uses IP address to get its corresponding MAC address from the Address Resolution Protocol (ARP) cache.
- **Globals:** None
- **Arguments:**
 - `$1`: The IP address of the client.
- **Outputs:** Normalized MAC address related to the IP address if found.
- **Returns:**
 - 1 if the IP address is not valid or MAC address is not found.
 - MAC address corresponding to given IP address in normalized form.
- **Example Usage:**
 - `get_client_mac 192.168.1.1`
 - `MAC_ADDRESS=$(get_client_mac 192.168.1.1)`

11.6.294 Quality and Security Recommendations

1. Ensure the user has required permissions to run the `ping`, `ip neigh` and `arp` commands to avoid permission denied errors.
2. Include error handling for cases where ARP does not contain an entry for the given IP address, emitting appropriate error messages.
3. Consider including a validation check for the normalized MAC address before returning it.
4. Use secure command execution to prevent injection attacks.
5. Consider returning a standardized error value, such as a null address or a specific error string, if the MAC cannot be found.

11.6.295 `get_cluster_conf_file`

Contained in `lib/functions.d/cluster-functions.sh`

Function signature: `fc7feb39024383e4cb59d1bec6341e1f7248d7dd978aa33bac047e372cc4613e`

11.6.296 Function overview

The `get_cluster_conf_file` function takes a cluster name as an argument and returns the path to its configuration file. If no cluster name is provided, it prints out an error message and exits the function.

11.6.297 Technical description

Name:

`get_cluster_conf_file`

Description:

This Bash function takes a cluster name as an argument and returns the path of the configuration file for that cluster. It first checks if a cluster name has been provided, if not an error message is printed and the function exits. It then gets the path to the cluster using the `get_cluster_dir` function.

Globals:

No global variables are used in this function.

Arguments:

- \$1: The name of the cluster

Outputs:

- If no cluster name is provided, an error message is printed to stderr: `[ERROR] Usage : get_cluster_conf_file <cluster-name>`
- If successful, it prints the path to the cluster's configuration file

Returns:

- 1 if no cluster name is provided or if `get_cluster_dir` function fails - 0 if the function executes successfully

Example usage:

```
get_cluster_conf_file my-cluster
```

11.6.298 Quality and security recommendations

1. Consider using more descriptive error messages. Instead of `[ERROR] Usage : get_cluster_conf_file <cluster-name>`, something like `[ERROR] Missing required argument: cluster_name` might be more helpful to users.
2. This function trusts that the `get_cluster_dir` function is well-implemented and doesn't validate the return value other than checking for errors. If `get_cluster_dir` could potentially return a harmful or malicious path, then this function will pass it along to the caller.
3. This function assumes that there is always a `cluster.conf` file inside the directory returned by `get_cluster_dir`. Ensure proper error handling if the file does not exist.

4. Avoiding using hardcoded strings (e.g. “/cluster.conf”) which could potentially cause problems in the future if there is a change in the configuration filenames or directory structure. Consider using a configuration or properties file to manage these.

11.6.299 `get_cluster_dir`

Contained in `lib/functions.d/cluster-functions.sh`

Function signature: `f5e9f656e463f433ab0e3bfed0da73d9166cc4e5802421827e55bef022685a0a`

11.6.300 Function Overview

The `get_cluster_dir()` function is a simple bash function designed to fetch and echo the directory of a specified cluster. The function takes the name of a cluster as an input, constructs the path to the directory, and outputs the path. If no cluster name is given or if the cluster name is empty, the function outputs an error message and returns with an error status.

11.6.301 Technical Description

- **Name:** `get_cluster_dir`
- **Description:** This function generates the path to a specified cluster directory by concatenating a base directory string with the specified cluster name.
- **Globals:** `HPS_CLUSTER_CONFIG_BASE_DIR`: This global variable is used as the base path to create the full cluster directory path.
- **Arguments:** `$1`: This argument is expected to be the name of a cluster. It is used to construct the full cluster directory path.
- **Outputs:** The function will output the constructed path string to stdout. If no cluster name or an empty string is provided, it will output an error message to stderr.
- **Returns:** Returns 0 if it successfully outputs the full cluster path; returns 1 if no cluster name or an empty string is provided.
- **Example usage:**

```
echo $(get_cluster_dir example_cluster)
```

```
# Output: path/to/existing/HPS_CLUSTER_CONFIG_BASE_DIR/example_cluster
```

11.6.302 Quality and Security Recommendations

1. A proper validation of the cluster name should be added before further processing to make sure the input is not malicious and adheres to appropriate naming conventions.

2. To avoid confusion or flawed operations, add checks to ensure that the “HPS_CLUSTER_CONFIG_BASE_DIR” and the constructed directory path actually exist in the file system.
3. To enhance clarity of the function behavior, include a clear and verbose error message to indicate when the function encounters any problem.
4. Be sure to make use of proper permission settings for the directories and files to prevent unauthorized access. This is particularly crucial if this function is part of a larger system that has security implications.
5. Consider defining strings like the error message and base directory as constants at the top of the program or in a configuration file. This enhances maintainability especially when changes need to be made in the future.

11.6.303 `get_device_bus_type`

Contained in `lib/functions.d/storage_functions.sh`

Function signature: `71e77c52eb2ba1d481c9ef51c928878b08e2c1088eeb0bf2fbe433c62633a476`

11.6.304 Function overview

The function `get_device_bus_type()` takes one argument, which represents a device, and returns the device bus type based on the device name or device property defined in the local environment. If the device name starts with `/dev/nvme`, it will echo “NVMe”. If not, it then checks if the device is rotational, if it’s not rotational, it echoes “SSD”, otherwise, it echoes “HDD”.

11.6.305 Technical description

- **Name:** `get_device_bus_type`
- **Description:** This function identifies the type of bus a provided device is using. It checks whether the input is a Non-Volatile Memory Express (NVMe) device, a Solid State Drive (SSD), or Hard Disk Drive (HDD) by analyzing its name or its rotational property.
- **Globals:** None
- **Arguments:**
 - `$1`: `dev` A string that represents the device.
- **Outputs:**
 - If the device name starts with `/dev/nvme`, it outputs “NVMe”.
 - Otherwise, if the device is not rotational, it outputs “SSD”.
 - Otherwise, it outputs “HDD”.
- **Returns:** None
- **Example usage:** `get_device_bus_type /dev/nvme0n1`

11.6.306 Quality and security recommendations

1. Always validate the input to ensure that the device provided exists in the system.
2. Consult the device properties from a trusted source, or directly from the system if possible, instead of purely relying on the device name pattern.
3. Maintain the single-responsibility principle. The function may benefit from being split into multiple smaller functions, each with its own responsibilities: one for checking if the device is NVMe, another for checking if the device is an SSD, and another for checking if the device is an HDD.
4. Always handle potential errors or exceptional cases to avoid unexpected behaviors. In this function, an else-case would be beneficial to handle situations where the device is neither an NVMe device, SSD, nor HDD.

11.6.307 `get_device_model`

Contained in `lib/functions.d/storage_functions.sh`

Function signature: `ec59456e4d8546a811f0f4533daf51da08474fb4ada4f1efb264e30d2fea7091`

11.6.308 Function overview

The `get_device_model()` function in Bash is used to get the model of a specific device. It takes a device identifier as an argument and then accesses the corresponding system information to return the model of the device. If the function cannot find the specified device, or if any other error occurs, it will return a string saying “unknown”.

11.6.309 Technical description

- **Name:** `get_device_model`
- **Description:** This function retrieves the model of a device given its identifier. It looks in the “`/sys/block`” directory, removes any whitespace, and then returns the model of the device. If the system cannot find the model for any reason, it reports “unknown”.
- **Globals:** None.
- **Arguments:**
 - `$1`: `dev` - This is the identifier of the device whose model is being retrieved
- **Outputs:** Model of the device or “unknown” if the device model can’t be found.
- **Returns:** The function returns the model of the device or “unknown” if there is an error or if the device can’t be found.
- **Example usage:**

```
model=$(get_device_model sda)
echo $model
```


11.6.310 Quality and security recommendations

1. Implement error checking for the `cat` command.
2. Check the validity of the input device identifier before processing.
3. Provide a more descriptive error message.
4. Sanitize the input to avoid command injection vulnerabilities.
5. Handle device names that contain spaces correctly.
6. Implement proper logging to understand the behavior of the function in case of failures.

11.6.311 `get_device_rotational`

Contained in `lib/functions.d/storage_functions.sh`

Function signature: 1002aec4163a82a48171eea735ad8700633333695a4b75dcb9ecc3317d14222f

11.6.312 Function overview

The `get_device_rotational()` function in Bash is designed to fetch the rotational status of the specified device. This status is derived from the `sysfs` filesystem and signals whether the device uses rotational storage media (like a traditional hard drive) or not (like an SSD). If the function is unable to retrieve this information, it defaults to returning `"1"`.

11.6.313 Technical description

- **Name:** `get_device_rotational()`
- **Description:** This function fetches and prints the rotational status of a specific block device. '1' implies the device uses rotational media, while '0' indicates use of non-rotational media.
- **Globals:** None.
- **Arguments:** `$1: dev` — The device whose rotational status is to be fetched.
- **Outputs:** The rotational status ('1' or '0') of the specified device.
- **Returns:** Nothing.
- **Example Usage:** `get_device_rotational sda`

In the example above, 'sda' is the block device for which the rotational status is desired.

11.6.314 Quality and security recommendations

1. Consider using more descriptive variable names to improve code readability.

2. Remember to properly quote your variables to avoid word-splitting or pathname expansion.
3. Always use `#!/bin/bash` or `#!/usr/bin/env bash` for writing bash scripts, not `#!/bin/sh`, because it may not actually link to bash on many systems, leading to unexpected behavior.
4. You should check if the device path exists in the `sys` filesystem as a preliminary step before attempting to fetch the rotational status.
5. Consider error handling for cases where the provided device name does not exist.

11.6.315 `get_device_serial`

Contained in `lib/functions.d/storage_functions.sh`

Function signature: `8e1b210627973a0cc629e646a16460819b61ae08954c5dba162358f2b6cb4532`

11.6.316 Function overview

The function `get_device_serial` is used for extracting the serial number from a specific device in a Linux system. This is accomplished by querying the Udev device manager using the `udevadm info` command with the particular device and parsing the output to get the serial number.

11.6.317 Technical description

- **name:** `get_device_serial`
- **description:** This function takes in a device (a local variable `dev`) as an argument, runs a query for its properties using `udevadm`, and extracts its `ID_SERIAL` value (the device's serial number). If no serial is found, it returns "unknown".
- **globals:** No global variables are used in this function.
- **arguments:**
 - `$1`: `dev`, the name of the device to get the serial number from.
- **outputs:** Either the device's serial number or the string "unknown" if no serial number is found.
- **returns:** The function doesn't have a `return` statement, its output is a side-effect of the function.
- **example usage:** `get_device_serial /dev/sda`

11.6.318 Quality and security recommendations

1. Add error checking and handling for non-existent devices or permission issues when running the `udevadm` command.
2. Currently, the function silently defaults to "unknown" when the device serial cannot be retrieved. This could be enhanced by incorporating a verbose mode or a warning message to inform the user about possible issues.

3. To avoid potential command injection, ensure that the provided input is properly validated and sanitized before it is used.
4. Prefer using the `printf` function over `echo` for compatibility and predictability reasons.
5. Always use double quotes around variable substitutions to avoid word splitting and pathname expansion. For instance, `--name="$dev"` instead of `--name=$dev`. This is already correctly done in the provided function.

11.6.319 `get_device_size`

Contained in `lib/functions.d/storage_functions.sh`

Function signature: `1d4e207d13b80b5424975cd10a3ed7197f78e6832fbbf6662e417abdd6def6d2`

11.6.320 Function Overview

The `get_device_size` function is a bash function that accepts a device name as an argument and then obtains the size of the device using the `lsblk` command. If the device size cannot be determined, it will output “unknown”.

11.6.321 Technical Description

- **Name:** `get_device_size`
- **Description:** This function retrieves the size of a specified device. It utilizes the `lsblk` command and, in the event the device size cannot be determined, it outputs the string “unknown”.
- **Globals:** None
- **Arguments:**
 - `$1`: The device whose size is to be determined. It is usually a block device-like “`/dev/sda`”.
- **Outputs:** The size of the device. If the size can’t be determined, it outputs the string “unknown”.
- **Returns:** The status of execution of the last command executed, typically the `lsblk` command.
- **Example Usage:** `get_device_size "/dev/sda"`

11.6.322 Quality and Security Recommendations

1. Since the function performs operations on block devices, it should have proper error checking and handling. This would help in maintaining the integrity of the device and prevent any data loss.
2. Validate the input to the function to ensure that it is a valid device name.
3. The function echoes “unknown” when it can’t determine the device size. It might be better to return a specific error code for this situation.

4. Consider adding checks to ensure that the required utilities (`lsblk` in this case) are available on the system. This could prevent errors from occurring due to missing utilities.
5. It's always a good idea to run scripts as a non-root user when possible. If this function needs root privileges, make sure you handle that securely.

11.6.323 `get_device_speed`

Contained in `lib/functions.d/storage_functions.sh`

Function signature: `a3925bc33db6a5252197ab449a78edcdfb4d9117f6e067cedb4b22c2eaea3a3e`

11.6.324 Function overview

The function `get_device_speed` measures the read speed of a particular device.

11.6.325 Technical description

- **name:** `get_device_speed`
- **description:** This bash function measures and outputs the read speed of a given device. It reads data from the specified device using the `dd` command and pipes the output to the `grep` command to filter the speed data. In case the speed data can't be retrieved, it prints "N/A".
- **globals:** None
- **arguments:**
 - `$1`: This refers to the device for which the read speed is being measured.
- **outputs:** It outputs the read speed of the provided device in the format '[0-9.]+ MB/s', else if it can't fetch the data, it prints "N/A".
- **returns:** It won't return any standard exit codes, but the output of the speed of the device.
- **example usage:**

```
get_device_speed "/dev/sda1"
```

11.6.326 Quality and security recommendations

1. Add proper error handling: We should have a way to handle situations where the device does not exist, is not readable or the `dd` command is not available.
2. Return error codes: This function doesn't have a return value, it just emits the output. To improve its usability in scripts, it would be better to return distinct codes for different error conditions.
3. Check for needed utilities: It would be great if function checks for presence of `dd` and `grep` commands at the start of execution.
4. Input sanitation: Inputs should be sanitized or validated to prevent potential security risks. Validate the device name accordingly.

5. Improve the output: To increase the usability of the function, it could return both the raw speed data and a human-readable string.

11.6.327 `get_device_type`

Contained in `lib/functions.d/storage_functions.sh`

Function signature: `c04dccee20f8398f1cc00c1f339464f78c0ea09416025a9627dcbd45477ea68d`

11.6.328 Function Overview

This function, `get_device_type()`, utilizes the `udevadm` command to fetch specific device properties. Given a device ID as a parameter, it queries for properties and filters out the device type. By design, if no correct device ID is given or if the `udevadm` fails, the function defaults to returning “disk”.

11.6.329 Technical Description

- Name: `get_device_type`
- Description: Fetches and returns the type of a device by utilizing the device’s unique ID. If a type is not found or the function fails, it defaults to returning “disk”.
- Globals: None
- Arguments: [`$1`: The unique ID of the device for which the type is to be found]
- Outputs: The type of the given device. Prints “disk” if no correct ID is given or if the function fails.
- Returns: 0 on successful execution, non-zero on error.
- Example Usage:

```
device_type=$(get_device_type /dev/sda1)
echo $device_type
```

11.6.330 Quality and Security Recommendations

1. It is recommended to add formal error handling to help diagnose potential issues or incorrect inputs more easily. Relying solely on a default return value can mask actual errors.
2. The usage of `grep` and `cut` to parse the output of `udevadm` assumes a specific output format and might break if the output or format changes in the future. A more robust parser could be considered.
3. All inputs, even if they are supposed to be device IDs, should be sanitized to prevent potential Bash command injection issues.
4. A detailed comment describing the function, its parameters, and its return values can improve maintainability and readability of the code.

11.6.331 `get_device_usage`

Contained in `lib/functions.d/storage_functions.sh`

Function signature: 859cf1265955125053e13b61c6c4bfa14f3715f652ae3ee914eb7e3927870233

11.6.332 Function overview

The function `get_device_usage()` is used to fetch the usage details about a device. The function accepts a device identifier as an argument and returns a comma-separated string of mount points where the device is in use. If the device is not used anywhere, it returns “unused”.

11.6.333 Technical description

- **Name:** `get_device_usage`
- **Description:** This function utilizes Linux command line utilities to decipher the usage of a given device. The device identifier is passed as an argument and usage details are then obtained with the `lsblk` command. The list of places where the device is in use is compiled into a comma-separated string. If the device is not currently in use anywhere, “unused” is returned.
- **Globals:** None
- **Arguments:**
 - `$1`: Device identifier (e.g., `/dev/sda1`)
- **Outputs:** Comma-separated string of device usage locations or “unused” if the device has no usage records.
- **Returns:** 0 on success.
- **Example Usage:**

```
usage=$(get_device_usage "/dev/sda1")
echo $usage
```

11.6.334 Quality and security recommendations

1. Input validation: This function currently performs no validation on input. It would be beneficial to add checks to ensure that the argument passed is actually a valid device identifier.
2. Error handling: Updates could be done to the function to make it handle, recover from, or report any errors that may occur during the execution of command line utilities used.
3. Use of unassigned variables: In the current format, if `lsblk` fails to execute, `usage` will be unset causing an ‘unbound variable’ error to be thrown. Consider setting a default value for `usage` to prevent this.
4. Secure handling of command substitution: Use the `“$()”` construct for command substitution to avoid potential security issues with backticks. The current implementation already follows this recommendation.

11.6.335 `get_device_vendor`

Contained in `lib/functions.d/storage_functions.sh`

Function signature: `e635f4842a790d321a2d2bf3371ad26f62f2881ced6effdff7eaea8887f48cf1`

11.6.336 Function overview

The `get_device_vendor` function aims to find out and return the vendor information of a given device specified by the argument passed to the function. It is a bash function implemented to dive deep into system files to extract this information. If the information is not available or if the system encounters any error while trying to find out the information, it simply returns an “unknown” string.

11.6.337 Technical description

- **Name:** `get_device_vendor`
- **Description:** This function retrieves the vendor information of a given device in a Linux-based operating system.
- **Globals:** None.
- **Arguments:** [`$1`: The device (e.g., `/dev/sda`) for which to retrieve the vendor information.]
- **Outputs:** Prints vendor information to the standard output. If the vendor information could not be obtained, prints “unknown”.
- **Returns:** None.
- **Example usage:** `get_device_vendor /dev/sda` *This will return the vendor information of the sda disk.*

11.6.338 Quality and security recommendations

1. Validate the input argument as a valid device before proceeding with the command to prevent unexpected behavior or potential security breaches.
2. Adding error checks and handling could improve the function further. Right now, if anything goes wrong, the function will simply print “unknown” which might not be the most informative way to handle errors in some cases.
3. Check for the appropriate permissions before trying to access system files. Your script might not work without the right permissions or in a restricted environment.
4. Include more detailed comments in the code to explain decision reasons and to clarify hard-to-understand parts. Even though the code looks simple, good commenting is a strong marker of software quality.
5. You should consider setting a stricter error handling policy, like `set -euo pipefail`, to make the script exit on the first error it encounters.

11.6.339 `_get_existing_zpool_name`

Contained in `lib/host-scripts.d/common.d/zpool-management.sh`

Function signature: `cc41e1bc5b39a7a3e09eef4edc9166c5f5783d182d1cd860936367173015ae84`

11.6.340 Function Overview

This function, `_get_existing_zpool_name`, is used for obtaining the existing ZFS pool name from a remote host. If the ZFS pool name is present, the script prints the name to stdout and returns a zero exit status. If the ZFS pool name is not found, the function returns a non-zero exit status.

The code processes options (e.g. `strategy`, `mountpoint`, `force`, `dry-run`, `no-defaults`) and calls helper functions like `zpool_name_generate` and `zfs_get_defaults` for specific functionalities. It also checks the `ZPOOL_NAME` before proceeding, validates the policy rules, generates the pool name and selects a disk based on the strategy.

Subsequently, the script performs the creation of the pool, persists the host variable `ZPOOL_NAME` and provides appropriate logging and error handling throughout the execution.

11.6.341 Technical Description

- Function name: `_get_existing_zpool_name`
- Description: This function is responsible for obtaining the name of an existing ZFS pool from a remote host.
- Globals:
 - `ZPOOL_NAME`: The name of the ZFS storage pool.
- Arguments:
 - `$1`: Command line arguments and options.
 - `$2`: Value of argument where it applies.
- Output: Prints the ZFS storage pool name to stdout if exists.
- Returns: Exit status 0 when pool name is found, 1 or 2 when there is an error.
- Example Usage: `_get_existing_zpool_name --strategy first --mountpoint /tmp --dry-run`

11.6.342 Quality and Security Recommendations

1. The function could benefit from additional input validation. Checking whether the provided arguments are recognizable before their first usage could prevent unexpected behaviors.
2. The error messages can be improved to be more descriptive. For example, stating precisely why a specific helper function is missing could allow an easier troubleshooting experience.

3. The function relies heavily on other functions (helpers), their availability and their output. As such, this function could be prone to break if any of the helper functions are modified. Robust integration tests could help catch these issues.
4. Security-wise, the script does not seem to sanitize inputs when making system calls. This could potentially lead to command injection vulnerabilities. Performing rigorous input validation and sanitization throughout the script would help mitigate this risk.
5. Consider implementing a more comprehensive logging functionality, such as logging every action and result, which can help in debugging and can be crucial in live environments.
6. If possible, make the function idempotent. The function should give the same output and have the same side effects, regardless of how many times it's run with the same inputs. This would ensure predictable behavior.

11.6.343 `get_external_package_to_repo`

Contained in `lib/functions.d/repo-functions.sh`

Function signature: `a00cf7324f6795dbbcefc882d9a0346242d31cdeaef0ed00e85c79720244ebc9`

11.6.344 Function overview

The `get_external_package_to_repo` function is designed to download a package from a specified URL and save it to a specified repository path. It performs various checks to ensure that the package URL and repository path are valid, that the repository directory exists, and that the package file has a `.rpm` extension. It logs various information and error messages during its operation, and it uses the `curl` command to download the package.

11.6.345 Technical description

- **Name:** `get_external_package_to_repo`
- **Description:** This function downloads an RPM package from a specified URL to a specified repository path. It checks for correct usage and the existence of the target directory, and it verifies that the URL points to an RPM file. It logs information about its operation and any errors it encounters.
- **Globals:** None.
- **Arguments:**
 - \$1: URL from which to download the RPM package.
 - \$2: Path to a directory that will hold downloaded package.
- **Outputs:** Logs information and error messages to standard output.
- **Returns:**
 - 0 if the package is successfully downloaded
 - 1 if incorrect usage

- 2 if target directory does not exist
- 3 if URL does not point to an RPM file
- 4 if it fails to download package URL.
- **Example usage:** `get_external_package_to_repo "http://example.com/package.rpm"`
`"/path/to/repo"`

11.6.346 Quality and security recommendations

1. Use absolute paths for the repository path to avoid potential confusions or errors with relative paths.
2. Validate the URL more thoroughly. Currently, it only checks if the URL ends with `.rpm`. More comprehensive validation would be beneficial.
3. Use the `-s` (silent) option with `curl` to suppress unnecessary output and only display important information.
4. Consider using a more secure protocol, such as HTTPS, for downloading the package to ensure the integrity and security of the download.
5. Add more detailed logging, including timestamps and the full file path of the downloaded files, to allow easier troubleshooting.
6. Rather than relying on return codes, consider propagating more detailed error information to give invokers more context of failures.

11.6.347 `get_host_type_param`

Contained in `lib/functions.d/cluster-functions.sh`

Function signature: `2aa086f62cbb99876d8c7312d176dfce1f88af4d006d3f9a64537fa96e3f9008`

11.6.348 Function overview

The `get_host_type_param()` function in Bash is designed to retrieve a specific value from an associative array, based on the provided key. This function takes in two parameters: the name of the associative array, and the key for which the value should be retrieved. The value corresponding to the given key is then echoed out, allowing the function's output to be captured and used elsewhere in the script.

11.6.349 Technical description

- **Name:** `get_host_type_param()`
- **Description:** This Bash function retrieves a specific value from an associative array. The name of the associative array and the key are provided as input arguments. The function uses Bash's variable reference (`declare -n`) to reference the associative array, and then uses array indexing (`${ref[$key]}`) to fetch the value corresponding to the provided key.
- **Globals:** None

- **Arguments:**
 - \$1 (type) : The name of the associative array from which to retrieve the value.
 - \$2 (key) : The key for which to retrieve the value from the associative array.
- **Outputs:** Echoes out the value from the associative array that corresponds to the provided key.
- **Returns:** Does not return value.
- **Example usage:** `get_host_type_param "server_param"`
 `"ip_address"`

11.6.350 Quality and security recommendations

1. It's important to ensure the arguments passed into this function (the name of the associative array and the key) are not controlled by untrusted user input, as this could potentially lead to unintended behavior.
2. Consider adding error checks in the function to handle cases where the provided associative array or key might not exist. Currently, if either the array or the key doesn't exist, the function won't return an error or warning, which might lead to bug diagnosis challenges.
3. It might be beneficial to enclose all variable references, including array indices, within double quotes. This will prevent word splitting and pathname expansion, which could lead to unexpected results in some cases.

11.6.351 `get_iso_path`

Contained in `lib/functions.d/iso-functions.sh`

Function signature: `c9a43a62bf1aba8fb474972e5a14edb51e00a6b3ad7b99ced97edcfedfb00929`

11.6.352 Function Overview

The `get_iso_path` function is used in Bash to construct the path to the ISO file in a distribution directory. It verifies if the `HPS_DISTROS_DIR` global variable is set and is a directory, and if so, it concatenates the string `"/iso"` to it. If `HPS_DISTROS_DIR` is either not set or not a directory, it outputs an error message and returns an exit status of 1.

11.6.353 Technical Description

- **Name:** `get_iso_path`
- **Description:** This function constructs the path to the ISO file within a specified directory by appending `"/iso"` to the `HPS_DISTROS_DIR` global variable. If `HPS_DISTROS_DIR` is not a directory or isn't set, it reports an error and exits with a status of 1.

- **Globals:**
 - HPS_DISTROS_DIR: The directory where distributions are stored.
- **Arguments:**
 - None.
- **Outputs:**
 - The full constructed ISO path if HPS_DISTROS_DIR is set and is a directory.
 - An error message directed to stderr if HPS_DISTROS_DIR is not set or is not a directory.
- **Returns:**
 - 0 if the function successfully prints the path.
 - 1 if the HPS_DISTROS_DIR is not set or not a directory.
- **Example usage:**

```
path=$(get_iso_path)
if [[ $? -eq 0 ]]; then
    echo "The ISO path is: $path"
else
    echo "Unable to get ISO path."
fi
```

11.6.354 Quality and Security Recommendations

1. Check if the HPS_DISTROS_DIR is an absolute path. Relative paths can be manipulated in unexpected ways.
2. Check if the directory is readable and not empty.
3. Apply proper error-handling to ensure meaningful error messages are returned to help in debug efforts.
4. Document the expected behavior of this function thoroughly to avoid confusion or misuse.
5. Use `set -u` to catch unset variables in the script, which may help catch bugs or errors.
6. Use `set -e` to stop script execution upon encountering an error, which can be beneficial to prevent script from continuing in an unpredictable state.

11.6.355 get_latest_alpine_version

Contained in `lib/functions.d/tch-build.sh`

Function signature: `f1404bc114d5e831d9237d9abb99e8b1e61a8b81eef840b822e47765fdb7591`

11.6.356 Function Overview

This function `get_latest_alpine_version()` is designed to fetch the latest version number of Alpine Linux from the official Alpine website. It uses either `curl` or `wget` to download the page, extracts the version number using a regular expression and

sorting, and falls back to a predefined version number if it fails. If the extraction fails, the function provides a warning log message with the fallback version. Finally, the function returns this version number.

11.6.357 Technical Description

For the function `get_latest_alpine_version()`:

Name: `get_latest_alpine_version`

Description: This function fetches the most recent version number of Alpine Linux from its official website using either `curl` or `wget` for the process. If the methods fail in fetching, the function resorts to a fallback version defined within it. The function provides a warning log message with the fallback version in case the extraction fails and finally returns the version number.

Globals: None

Arguments: None

Outputs: - The version number of the latest Alpine Linux. - Warning message in case of failure fetching with the version falling back to 3.20.2

Returns: - 0 if the function is successful - Warning message with fallback version if the function fails

Example Usage: Run the function like so: `get_latest_alpine_version`

11.6.358 Quality and Security Recommendations

1. Error handling can be better: Currently, the function falls back to a hard-coded version when it fails to fetch the latest version via `curl` or `wget`. However, the reason might be transient network issues, and a simple retry might work. Implementing a retry mechanism with exponential backoff would improve the quality of this script.
2. Robustness: Consider checking the returned status codes of the `curl` or `wget` commands, to know if the fetch was successful or not. The result does not always indicate the command's success.
3. For security reasons, consider using a more secure protocol like secure https to access the website compared to http which is currently in use. It can protect the data being exchanged from lurking security threats.
4. The function logs a warning message if it fails to get the versions, it can also provide some debugging information like the status code, error messages, etc., which could be helpful while troubleshooting.
5. The function could have a mechanism to periodically check for an update after some interval automatically so that the fetched version is always current without the need to invoke the function manually.

11.6.359 `get_provisioning_node`

Contained in `lib/host-scripts.d/common.d/common.sh`

Function signature: `1bbd682413f1c35f8f147b598c84f60028eea1205d56405e50b38d7d551b5b01`

11.6.360 Function overview

The `get_provisioning_node` function uses the built-in commands of `ip` and `awk` to retrieve the IP address of the default gateway in a Unix system. If successful, it prints the IP address of the default gateway to the standard output. This function is generally used for obtaining the IP needed for system provisioning.

11.6.361 Technical Description

11.6.361.1 Name

- `get_provisioning_node`

11.6.361.2 Description

- The function `get_provisioning_node` retrieves the IP address of the default gateway in a Unix system.

11.6.361.3 Globals

- None

11.6.361.4 Arguments

- None

11.6.361.5 Outputs

- The IP address of the default gateway.

11.6.361.6 Returns

- The default gateway IP address.

11.6.361.7 Example usage

```
gateway_ip=$(get_provisioning_node)
echo $gateway_ip
```

11.6.362 Quality and Security Recommendations

1. Validation: At its current state, the function does not validate the IP address it retrieves. This could pose a problem if the command fails for some reason. Implementing basic validation checks would improve the quality of the function.
2. Error handling: In case the command fails to obtain the IP address, the function currently offers no method of detecting and managing such failure. It is recommended to implement error handling to improve robustness.
3. Comments: More explanatory comments should be added in the function to improve maintainability and readability of the function.
4. Security: It's a good practice to limit the privileges of the script that uses this function to avoid potential security risks.

11.6.363 `__guard_source`

Contained in `lib/functions.sh`

Function signature: `e8beb9b32cabb9e73dba64f7b102ad5f1112590b455a914144bd65e5d3001168`

11.6.364 Function overview

The function `__guard_source()` is a bash script guard, designed to prevent sourcing a script more than once. This function takes the name of a previously-sourced script and stores it in a guard variable. The function returns '1' if the script has already been sourced (indicating an error and preventing further sourcing) and '0' if not.

11.6.365 Technical description

- **name:** `__guard_source()`
- **description:** This function is used to avoid multiple sourcing of the same bash script for the prevention of redundant operations, potential recursion and unexpected behavior. It generates a guard variable for the sourced script and checks if the script has been previously executed or sourced, and if so, the function will return 1 preventing the script from re-executing.
- **globals:** `[_guard_var]`: Guard variable created for each sourced script to mark its execution. It uses the script name with nonalphanumeric replaced with '_'
- **arguments:** `[$1]`: Unused in this function., `[$2]`: Also unused in this function.]
- **outputs:** No output is returned to stdout/stderr.
- **returns:** 0 if the script or source has not been run before, and 1 if it has and encountered an error
- **example usage:**

```
if ! __guard_source; then
    echo "Script already sourced once"
fi
source script.sh
```

11.6.366 Quality and security recommendations

1. Clear documentation: Each function, global variable and return value should be clearly documented for future developers.
2. Error handling: The function should handle possible errors gracefully and clear, meaningful messages should be returned.
3. Input Validation: Currently, the function does not take any arguments, but for future enhancement if it does, it should validate the input before processing it.
4. Use Strict Mode: To catch errors and undefined variables sooner, consider enabling a 'strict mode' in your scripts with `set -euo pipefail`.

11.6.367 handle_menu_item

Contained in `lib/functions.d/ipxe-menu-functions.sh`

Function signature: `dfa36a6a8b5f9bae0068c79fb19d4ebec3ef157d60edde5cb1b1475ef092face`

11.6.368 Function Overview

The function `handle_menu_item()` is designed to handle a wide variety of menu items within the iPXE interface. The function takes two parameters - an `item`, which is the menu item to be handled, and a `mac`, which is the MAC address of the host for whom the menu item is being processed. The function uses a case statement to handle a variety of possible menu items, and will log an error and fail if an unknown menu item is entered.

11.6.369 Technical Description

11.6.369.1 Name

`handle_menu_item()`

11.6.369.2 Description

This function handles the various menu options in the iPXE menu system.

11.6.369.3 Globals

VAR: `item` - The menu item to be handled.

11.6.369.4 Arguments

\$1: `item` - The item to be processed by the function. \$2: `mac` - The MAC address for the host for which the menu item is to be processed.

11.6.369.5 Outputs

Various outputs depending on the `item` input.

11.6.369.6 Returns

Varies depending on the `item` input.

11.6.369.7 Example usage

```
handle_menu_item host_install_menu 00:11:22:33:44:55
```

11.6.370 Quality and Security Recommendations

1. The function should have a more standardized output mechanism, rather than simply echoing to the console in some cases.
2. It would be beneficial to add further error checking to ensure that the `item` and `mac` parameters are properly formed and valid before the function attempts to process them.
3. Certain case options may produce unintended side effects - for example, `unconfigure` and `reboot` both trigger an `ipxe_reboot` which could potentially interrupt ongoing processes. Further consideration should be given to these potential issues.
4. Security could be improved by ensuring that only permitted users are allowed to call the `handle_menu_item()` function, or by adding additional security checks within the function to prevent unauthorized users from executing certain menu options.

11.6.371 `has_sch_host`

Contained in `lib/functions.d/host-functions.sh`

Function signature: `379b6704d9af414555293ef691b70449d39306ef567e44aabbf9ef70f91fa692d`

11.6.372 Function Overview

The `has_sch_host` function is designed to check if there is any configuration file in the specified directory (`$HPS_HOST_CONFIG_DIR`) that contains a certain type 'SCH'. It first checks if the specified directory exists. If not, it outputs an error message and returns 1 to indicate an error. If the directory exists, it checks every `.conf` file in it for the

presence of 'TYPE=SCH'. If at least one file with 'TYPE=SCH' is found, it returns 0 (True). If no such file is found, it returns 1 (False).

11.6.373 Technical Description

- **name:** `has_sch_host`
- **description:** This function checks if there is any configuration file in the specific directory that contains "TYPE=SCH".
- **globals:**
 - `HPS_HOST_CONFIG_DIR`: Directory to search for configuration files.
- **arguments:** None
- **outputs:** An error message if the specified directory does not exist.
- **returns:** 1 if the specified directory doesn't exist or no file containing "TYPE=SCH" is found. 0 if at least one file containing "TYPE=SCH" is found.
- **example usage:** `if has_sch_host; then echo "SCH host config found"; else echo "SCH host config not found"; fi`

11.6.374 Quality and Security recommendations

1. It would be safer to use an absolute path for `$HPS_HOST_CONFIG_DIR` to avoid ambiguity and confusion.
2. It is recommended to standardize the configuration files' extensions, such as `.conf`, for better searchability.
3. Consider improving the error-handling mechanism to make it more robust, for example by detailing which scenarios result in which errors.
4. Further sanitize the search pattern `('^TYPE=SCH')` to prevent potential command injections.
5. Utilize built-in shell checks to validate that `"${HPS_HOST_CONFIG_DIR}"` is, in fact, a directory, not a file.

11.6.375 `host_config_delete`

Contained in `lib/functions.d/host-functions.sh`

Function signature: `e05b6467ee0958047fe16e2eb5a0519a4ceab375f01a2f8a27e8ee6f35cf7400`

11.6.376 1. Function Overview

The function `host_config_delete()` is designed to delete a specified host configuration file. The function accepts the MAC address of a host as an argument. It determines the configuration file corresponding to that MAC address in a predefined directory and deletes it. It logs an informational message if the file was deleted successfully, or a warning message if the file was not found.

11.6.377 2. Technical Description

name: `host_config_delete()`

description: This function deletes the configuration file of a host identified by a specific MAC address.

globals: - `HPS_HOST_CONFIG_DIR`: This is the directory where the host configuration files are stored.

arguments: - `$1`: This is the MAC address of the host whose configuration file is to be deleted.

outputs: Informational or warning logs are output depending on whether the operation is successful or not.

returns: - `0`: if the host configuration file is deleted successfully. - `1`: if the host configuration file was not found.

example usage:

```
host_config_delete "00:11:22:33:44:55"
```

It will delete the configuration file mapped to MAC address "00:11:22:33:44:55".

11.6.378 3. Quality and Security Recommendations

1. Make sure the function correctly handles special characters in the MAC address to avoid unexpected behavior or security issues.
2. The function should check if the 'rm' command succeeded before logging a success message. This would improve the reliability of the logging message.
3. The `HPS_HOST_CONFIG_DIR` directory should have appropriate permissions set to prevent unauthorized access or modifications.
4. It would be beneficial to sanitize inputs to reduce the risk of code injection or Directory Traversal vulnerabilities.
5. To reduce the possibility of any potential logs forgery, it would be wise to incorporate a secure logging mechanism.
6. Lastly, check the existence of the `$HPS_HOST_CONFIG_DIR` directory before performing any operation to prevent any unnecessary errors.

11.6.379 `host_config_exists`

Contained in `lib/functions.d/host-functions.sh`

Function signature: `93698d3265775cdeb0581fb007522cdb74832e2c705812027e67b512cd33964b`

11.6.380 Function Overview

The `host_config_exists` function is a Bash function in that checks for the existence of a specific configuration file based on a provided *mac* address. It takes a *mac* address as

an argument, and constructs a file path to the configuration file using the *mac* address as the file name and a pre-defined directory as its path. It then uses a conditional statement to check if the file exists, returning a boolean indicating the result.

11.6.381 Technical Description

- **Name:** `host_config_exists`
- **Description:** The function checks for the existence of a specific configuration file that correlates to the provided *mac* address.
- **Globals:**
 - `HPS_HOST_CONFIG_DIR`: This global variable defines the directory where the configuration files are stored.
- **Arguments:**
 - `$1: mac`: This argument is the *mac* address which will be used as the base for configuration file name.
- **Outputs:** This function does not directly output any value.
- **Returns:** The function returns a boolean value which indicates whether the sought configuration file exists.
- **Example Usage:**

```
if host_config_exists "00:11:22:33:44:55"
then
    echo "Configuration file exists."
else
    echo "Configuration file does not exist."
fi
```

11.6.382 Quality and Security Recommendations

1. Make sure the `HPS_HOST_CONFIG_DIR` is always defined and set to a valid directory to prevent the function from searching in the incorrect location, leading to incorrect results.
2. Consider validating the *mac* argument to ensure that it is a valid *mac* address. This could prevent potential errors down the line or potential security vulnerabilities if malicious or incorrectly formatted *mac* addresses are provided.
3. The function should handle or communicate any errors it encounters during execution in a secure manner. Currently, it does not communicate any issues outwardly.
4. To better adhere to Unix philosophy, consider making the function output informative messages or codes to standard error when it encounters problems (e.g. when the directory does not exist).
5. Always remember to mark your variables as local where possible to prevent them from leaking into the global script environment, which can be a security risk and a source of bugs. This function does a good job of this by declaring `mac` and `config_file` as local variables.

11.6.383 host_config

Contained in `lib/functions.d/host-functions.sh`

Function signature: 389b01155f4ed6b074bab989a189c3e33a5dfb44e6ade66966a2b1966e95960d

11.6.384 Function Overview

The function `host_config()` is an imperative program in bash scripting language. This function accepts four arguments, namely `mac`, `cmd`, `key`, and `value`. Using these parameters, a configuration file is processed for the host machine identified by the MAC address. The function can execute commands (`get`, `exists`, `equals`, `set`) on the HOST configuration map, utilizing the `key` and `value` parameters. If the configuration file doesn't exist or if configuration has not been parsed for the current MAC address, the configuration file is first parsed and the `HOST_CONFIG` associative array is updated with key-value pairs from this file.

11.6.385 Technical Description

Name: `host_config()`

Description: Mainly, this function is for configuring host configuration files which store settings for the host machine identified by the MAC address. Commands to process settings include `get`, `exists`, `equals`, and `set`. An associative array `HOST_CONFIG` is used for this purpose.

Globals: [`HOST_CONFIG`: An associative array storing key-value pairs from the host configuration file, `HPS_HOST_CONFIG_DIR`: Directory where host configuration files are stored]

Arguments: - `$1`: MAC address identifier for host machine - `$2`: Command to process host settings. Valid commands are `get`, `exists`, `equals`, `set`. - `$3`: Key of setting to process. - `$4`: (Optional) Value to update the setting with if the 'set' command is used.

Outputs: Based on the command: - `get`: Prints the value of the setting identified by the key - `exists`: Checks if the key exists. No visual output. - `equals`: Checks if the setting identified by the key matches the value. - `set`: Logs info about the new setting and update. - A warning if an invalid command is used.

Returns: - No specific integer return values. For 'get', 'exists', 'equals' commands the function effectively returns 0 (true) if successful and 1 (false) if unsuccessful. - For 'set' command no explicit return values are declared. - Returns 2 if invalid command is supplied.

Example Usage:

```
mac_addr="00:11:22:33:44:55"
setting_key="Test_key"
```

```
# Returns the value of Test_key  
host_config $mac_addr get $setting_key
```

11.6.386 Quality and Security Recommendations

1. Strong input validation: Although the function performs some input validation, it could be improved. For example, MAC address format could be validated.
2. Error handling: Currently, if an invalid command is given, the function simply outputs a message and returns 2. Comprehensive error handling here would improve the robustness of the function.
3. Unknown globals: Globals like HPS_HOST_CONFIG_DIR are used in the function without prior validation. It's good practice to check if these are set before usage.
4. Logging: Use a proper logging utility instead of using command-line echo for important operations. Monitoring could also be implemented for security-sensitive operations.

11.6.387 host_config_show

Contained in `lib/functions.d/host-functions.sh`

Function signature: 6105ece190686264b3985f4c2de384ff48edc977456deeb1f70f8a863bb065a8

11.6.388 Function Overview

The `host_config_show` function in Bash is used to read a host configuration file that corresponds to a certain MAC address. If a configuration file doesn't exist for a given MAC address, an informational log message is displayed. If found, the function reads through the file, trims any leading or trailing quotes from each value, escapes any embedded quotes and backslashes, and then echoes each key-value pair.

11.6.389 Technical Description

- **Name:** `host_config_show`
- **Description:** Reads a host configuration file that matches a given MAC address. The function will output key-value pairs from the file, with necessary characters escaped for safety. If no such file exists, it logs an informative message.
- **Globals:**
 - `HPS_HOST_CONFIG_DIR`: This global points to the directory where host configuration files are stored.
 - `hps_log`: This global logs messages based on the application's events.
- **Arguments:**
 - `$1`: This is the MAC address of the device. It's supposed to match with a configuration file within the directory specified by `HPS_HOST_CONFIG_DIR`.
 - `$2`: This argument is not used by the function.

- **Outputs:** Key-value pairs from the configuration file, if it exists. Otherwise, an information log message is output.
- **Returns:** Returns 0- indicating successful execution of the function.
- **Example Usage:** `host_config_show "04:0E:3F:A1:B2:C3"`

11.6.390 Quality and Security Recommendations

1. Implement argument validation: `host_config_show` could fail unexpectedly (or silently) if it receives unexpected data. Implement checks for the MAC address format and whether `HPS_HOST_CONFIG_DIR` is set.
2. Use clearer env var names: `HPS_HOST_CONFIG_DIR` and `hps_log` could be renamed to more descriptive names for better readability of the code.
3. Handle failure condition: Does the user need to know when `host_config_show` can't find a particular MAC address? If so, consider changing the return code from 0 in case a corresponding configuration file doesn't exist for the mac address passed.
4. Sanitize file reading: The `read` command used in the while loop might encounter issues dealing with special character sequences. Use `-r` option to prevent this.
5. Protect against variable shadowing: By using `local` for `mac` and `config_file`, this function avoids shadowing variables in the parent scope. Make sure to follow this good practice in the rest of your code as well.

11.6.391 `host_initialise_config`

Contained in `lib/functions.d/host-functions.sh`

Function signature: `9cbdb13ea8edf79033b1a96edf52ebe54e082910b6a8a403b6b2fd86a4d5b486`

11.6.392 Function Overview

The function `host_initialise_config()` is an initial set up tool designed to generate and assign a configuration file for a host, identified by its thus passed MAC address. It ensures that the host configuration directory exists. Then, it sets the state of the host config file to "UNCONFIGURED". Additionally, it logs the initialization of the host configuration.

11.6.393 Technical Description

- **Name:** `host_initialise_config`
- **Description:** This function initialises host configuration file within `HPS_HOST_CONFIG_DIR`. It takes in the MAC address of the host, creates a configuration file uniquely associated with the host, sets the initial state to "UNCONFUGURED", and logs the initialization.
- **Globals:** `[{ HPS_HOST_CONFIG_DIR: The directory to store host configuration files }]`

- **Arguments:** [{ \$1: MAC address of the host }]
- **Outputs:** Logs the initialization process with the file name
- **Returns:** N/A
- **Example Usage:**

```
host_initialise_config "00:0a:95:9d:68:10"
```

11.6.394 Quality and Security Recommendations

1. The function should include error handling for potential failures while creating directories or setting the state.
2. Consider integrating data validation procedures to ensure that the MAC address provided as an input matches the anticipated format.
3. Remnants of deprecated operations should be removed from the function to prevent future confusion or unintentional uncommenting.
4. To prevent unexpected errors or unauthorized manipulation, it is suggested to place appropriate file and directory permissions on the created configuration file and directory.
5. It would be more secure to generate unique names for configuration files rather than using the MAC address, which can be predictable or easily available. This would make them less susceptible to targeted attacks.

11.6.395 host_network_configure

Contained in `lib/functions.d/host-functions.sh`

Function signature: `2d391e7910d04552e8ca1284295323649c38c44df9d4c9f8e335c080e4d38dac`

11.6.396 Function overview

The `host_network_configure` is a Bash function that configures network settings for a given host. It accepts the MAC address and host type as arguments, and uses these variables to customize the network setup.

11.6.397 Technical description

Name: `host_network_configure`

Description: The function takes in the identifier and type of the host, retrieves the Dynamic Host Configuration Protocol (DHCP) IP and CIDR block from the cluster configuration, and configures the network settings accordingly. If the DHCP IP or CIDR block is not present, it logs a debug message and returns 1. It validates the presence of the `ipcalc` tool required for the configuration process, and if not present, logs another debug

message and returns 1. The function calculates the netmask and network base using the `ipcalc` tool and stores them in local variables.

Globals: - `VAR: dhcp_ip, dhcp_cidr` The DHCP IP address and CIDR block respectively, retrieved from the cluster configuration.

Arguments: - `$1: mac_id` The MAC address of the host. - `$2: host_type` The type of the host.

Outputs: Debug messages logged when DHCP IP or CIDR is missing or `ipcalc` is not installed.

Returns: 1 if either DHCP IP or CIDR block is missing from the cluster configuration, or if the `ipcalc` tool is not installed.

Example Usage: `host_network_configure MAC_ADDRESS HOST_TYPE`

11.6.398 Quality and security recommendations

1. Make sure to keep sensitive network details in secured and protected configurations.
2. Regularly check the status and availability of `ipcalc` tool.
3. Handle the failure case with a meaningful error message instead of simply logging a debug message.
4. Quaternion further error handling for unexpected return values from the `ipcalc`.
5. Check the validity of the received arguments (MAC address and host type).
6. The function should consider validating the retrieved DHCP IP and CIDR block before proceeding with the rest of the function execution.

11.6.399 hps_configure_opensvc_cluster

Contained in `lib/functions.d/opensvc-functions.sh`

Function signature: `844f546ce501e16cbbb4c4680bc897ae12e2512aa91c458bd9674d5a949bd6dc`

11.6.400 Function overview

The function `hps_configure_opensvc_cluster` is designed to configure an OpenSVC cluster. It waits for the daemon's socket to be available and then attempts to configure the cluster's identity. In case of failure configuring the cluster, it logs a warning message. If the daemon socket is not available after repeated checks, cluster configuration is skipped.

11.6.401 Technical description

- **Name:** `hps_configure_opensvc_cluster`

- **Description:** This function attempts to configure the identity of an OpenSVC cluster by periodically checking for the availability of the daemon's socket.
- **Globals:** None.
- **Arguments:** No arguments required for this function.
- **Outputs:** Logs information about the successful configuration of the OpenSVC cluster or a failure to do so.
- **Returns:** Always returns 0.
- **Example usage:** `hps_configure_opensvc_cluster`

11.6.402 Quality and security recommendations

1. Set a maximum number of attempts or a timeout for attempting to configure the cluster to prevent the script from being stuck in an infinite loop if the daemon socket is unavailable.
2. Include additional error handling for specific errors that could be returned when the identity configuration fails.
3. Ensure all log messages include the date and time to assist with potential troubleshooting.
4. Validate that the socket file or path does not contain any unsafe characters or sequences to avoid potential security vulnerabilities.

11.6.403 `hps_log`

Contained in `lib/functions.d/hps_log.sh`

Function signature: `7097e5f6dbf7eb35af6d5c1fe86b7ea40ce7ed8d95259fe409cfc4e8eb117559`

11.6.404 Function overview

The `hps_log` function is a custom log function that logs system events in a file named `hps-system.log` located in the directory specified by `HPS_LOG_DIR`. The function creates logs with timestamps and a log level which can be set from the command input. It also helps organize logs by supporting an identifier input. Lastly, if this function is called without the `HPS_LOG_IDENT` variable having been set, it defaults the identifier as 'hps'.

11.6.405 Technical description

- **name:** `hps_log`
- **description:** This function logs data with the specified level, message, and identity into a file defined by environment variable `HPS_LOG_DIR`. Default identity is 'hps' when `HPS_LOG_IDENT` is not set.
- **globals:** [`HPS_LOG_DIR`: Directory to store the log files, `HPS_LOG_IDENT`: Identity for the logs]

- **arguments:** [\$1: Log level, \$... : Log message]
- **outputs:** Writes logs to `hps-system.log` file in the location specified by `HPS_LOG_DIR`.
- **returns:** Not applicable since logging functions typically do not return a value.
- **example usage:** `hps_log "error" "This is a sample error message"`

11.6.406 Quality and security recommendations

1. Protect the log file by setting appropriate permissions to prevent unauthorized access or tampering.
2. Regularly rotate and archive log files to avoid them becoming too large.
3. Implement checks to ensure that `HPS_LOG_DIR` is set to a valid directory.
4. Provide error handling mechanisms if the log file cannot be written to.
5. Assign a default value to `HPS_LOG_DIR` that points to a secure and writeable directory if it has not been set.
6. Avoid logging sensitive information to maintain user data privacy.

11.6.407 `hps_origin_tag`

Contained in `lib/functions.d/system-functions.sh`

Function signature: `34e636eb4b7c0c9bf98b49ce8416227a48485bd990954c5b7f74feba9f1c472a`

11.6.408 Function Overview

The `hps_origin_tag` function attempts to generate a unique tag based on the origin of a given process. The function considers several aspects such as an override option, user, host and process ID in the context of an interactive terminal, and also client IP/MAC in the context of a non-interactive terminal.

11.6.409 Technical Description

- **Name:** `hps_origin_tag`
- **Description:** This function generates a unique tag indicating the origin of a process. It first checks if an explicit override is provided. If the script is running from an interactive terminal, the function captures the user, host, and process ID. If the script is running from a non-tty environment (e.g., a web server), it attempts to use client IP/MAC information to generate the tag.
- **Globals:** `REMOTE_ADDR`: Internet protocol address of remote computer
- **Arguments:** `$1`: Overrides the need for automatic origin determination

- **Outputs:** Prints a string that stands as the unique origin tag. The format could be process ID, user-host data, IP or MAC address.
- **Returns:** 0 on successful execution
- **Example Usage:** tag=\$(hps_origin_tag)

11.6.410 Quality and Security Recommendations

1. Ensure proper validation and sanitation of command outputs like `id -un` and `hostname -s` to prevent any potential command injection attacks.
2. Use stringent error handling and check the return codes of executed commands as much as possible.
3. Avoid putting sensitive data like MAC addresses within origin tags as they can leak data by exposing it in logs or other output. If it is necessary, make sure logs/output storing these tags are adequately secured.
4. When printing out the tag, consider using an appropriate log level.
5. If the function fails to create a tag, it would be advisable to include fallback methods or return a standard error code.

11.6.411 hps_services_post_start

Contained in `lib/functions.d/system-functions.sh`

Function signature: `afd33703649ac97f669eda6f6bc20af42f3aaae2584d16fd61e145be4387d5bc`

11.6.412 Function overview

The function, `hps_services_post_start()`, is a Bash function coded to configure the OpenSVC cluster if applicable. The configuration is completed by running the `hps_configure_opensvc_cluster` function.

11.6.413 Technical description

Here is a breakdown of this function:

- **Name:** `hps_services_post_start`
- **Description:** This function configures the OpenSVC cluster if needed by executing a sub function named `hps_configure_opensvc_cluster`.
- **Globals:** None
- **Arguments:** None
- **Outputs:** The outcome of the `hps_configure_opensvc_cluster` function. Outputs depend entirely on this function and could range from simple print statements to full configuration results.

- **Returns:** The function does not return any specific value. Returns will depend entirely on the sub function that is executed.
- **Example usage:**
`hps_services_post_start`

11.6.414 Quality and security recommendations

1. Add logging for monitoring purposes: It's better to add a logging utility for debugging and maintenance. Specific events to log could include successful configuration, error during configuration, status pre-configuration and post-configuration.
2. Parameterize the function: If the configurations vary by OpenSVC cluster, it may make more sense to parameterize the function.
3. Error handling: Any failure in the `hps_configure_opensvc_cluster` function could lead to the OpenSVC cluster being unconfigured, potentially leading to errors downstream. Therefore, error handling is crucial in this function.
4. Function documentation: Be sure to document what exactly `hps_configure_opensvc_cluster` is doing, what configurations/settings are being changed, and any potential side effects.
5. Security: If configuring OpenSVC requires credentials, ensure those are not being hardcoded in the script and are stored safely. In addition, ensure that any necessary access control is implemented.

11.6.415 `hps_services_restart`

Contained in `lib/functions.d/system-functions.sh`

Function signature: `7b93e7411974a37f633379e3621dbdd656b539da01549e95ae3697ceb23cb5bc`

11.6.416 Function Overview

The function `hps_services_restart` is used for restarting the HPS services by configuring, creating and reloading supervisor services. After these operations it logs the restart status and executes post start steps.

11.6.417 Technical Description

- **Name:** `hps_services_restart`

- **Description:** This function restarts the HPS services. It starts by configuring and creating supervisor services. Then, it reloads the supervisor configuration. It logs an information message regarding the restart status of all supervisor services, using the configuration file located at “`${CLUSTER_SERVICES_DIR}/supervisord.conf`”. Finally, it performs post start tasks for the HPS services.

- **Globals:** [`CLUSTER_SERVICES_DIR`: This global variable holds the directory where the supervisor services' configurations are stored]
- **Arguments:** This function does not take any arguments.
- **Outputs:** Logs the outcome of the restart command to the standard output.
- **Returns:** Does not return a value.
- **Example usage:**

```
hps_services_restart
```

11.6.418 Quality and Security Recommendations

1. Validation checks should be implemented at the start of the function to ensure that the `CLUSTER_SERVICES_DIR` global variable is set and refers to a valid directory.
2. Error handling should be implemented to catch unsuccessful operations, such as in case of failed reload of supervisor configuration or if logging returns an error.
3. Consider using more specific logging levels (e.g., debug, warning, error) instead of using the 'info' level for all types of logs. This makes the system easier to debug and monitor.
4. Carefully manage file and directory permissions for `CLUSTER_SERVICES_DIR` and `supervisord.conf`, ensuring that only authorized users/services can modify them. This can help to prevent unauthorized modification of services, which could lead to security breaches.
5. Evaluate the necessity of executing `hps_services_post_start` at the end of the function in terms of security. If this function is not necessary, or could potentially be exploited, consider removing it.

11.6.419 `hps_services_start`

Contained in `lib/functions.d/system-functions.sh`

Function signature: `706bef3ec7c7c1616f4ef6bc0f1386604ca17cbbc784f1302e9fc230eda5eee3`

11.6.420 Function Overview

The Bash function `hps_services_start()` is responsible for command-line execution of a series of operations. These operations involve configuring, reloading, starting, and post-start processes of supervisor services. In essence, `hps_services_start()` is a higher-level function which provides an interface to manage various supervisor services within a cluster environment.

11.6.421 Technical Description

- **Name:** `hps_services_start`

- **Description:** This function encapsulates a sequence of operations which work together to manage supervisor services. Initially, the supervisor services are configured by executing the `configure_supervisor_services` function. Post that, supervisor config is reloaded through `reload_supervisor_config` function. Next, all the supervisor services are started with `supervisorctl` using a specific configuration file. Finally, post-start functions of the services are executed by calling `hps_services_post_start`.
- **Globals:** [`CLUSTER_SERVICES_DIR`: This global variable defines the directory where the supervisor configuration files are located]
- **Arguments:** [No direct arguments are passed to this function]
- **Outputs:** Depending on the functions called within `hps_services_start`, it outputs success messages of the functions. This might include the success of configuring, reloading, and starting supervisor services, as well as the output of any post-start operations.
- **Returns:** It does not return any explicit value.
- **Example Usage:** `hps_services_start`

11.6.422 Quality and Security Recommendations

1. Error handling should be implemented to catch any potential issues during the execution of the associated functions. This will result in a more robust and fail-safe function.
2. Security check should be implemented before executing any operations, specifically before accessing or modifying any directory or file.
3. Variables should be properly sanitized and validated to protect from code injection.
4. Any potential race conditions should be addressed to prevent unexpected behavior.

11.6.423 `hps_services_stop`

Contained in `lib/functions.d/system-functions.sh`

Function signature: `b43f68b5092581de60e6451048da5e3181daccfff13eee5943191ac9d98eab43`

11.6.424 1. Function overview

The function `hps_services_stop()` is used to stop all services managed by `supervisord`, a process control system. This is done by referring to the configuration file named `supervisord.conf` within the `CLUSTER_SERVICES_DIR` directory.

11.6.425 2. Technical description

- **Name:** `hps_services_stop()`
- **Description:** This function stops the running services governed by the `supervisord` located in the `CLUSTER_SERVICES_DIR` directory by using the `supervisorctl` command with `-c` flag and `stop all` argument.
- **Globals:** [`CLUSTER_SERVICES_DIR`: This is the directory path where the `supervisord.conf` configuration file lives.]
- **Arguments:** [None]
- **Outputs:** This function does not explicitly outputs anything. However, command line output from `supervisorctl` command will be visible which likely includes status information about stopping the services.
- **Returns:** It returns nothing. However, the exit status of the function will be the same as the `supervisorctl` command's exit status.
- **Example usage:**

```
hps_services_stop
```

11.6.426 3. Quality and security recommendations

1. The function relies on the global variable `CLUSTER_SERVICES_DIR`. It is advisable to validate that this variable is set and points to the right directory for robustness.
2. If necessary, logging should be implemented within the function to capture the status and any potential errors during the execution for debugging and traceability purposes.
3. Ensure proper permissions are set for the scripts containing this function to prevent unauthorized execution or modification.
4. Use secure coding practices like escaping any variables used in the function call to avoid command injection vulnerabilities. In this case, using quotes around `${CLUSTER_SERVICES_DIR}` ensures this variable is safely used even if it contains spaces or other special characters.
5. Make sure `supervisord` and `supervisorctl` are installed, configured correctly, and updated regularly for reliable and secure function execution.

11.6.427 `_ini_get_agent_nodename`

Contained in `lib/functions.d/opensvc-functions.sh`

Function signature: `3e6393764b80ab51638ba55763f7c3bac53a45bc4f47d68347a05c657c90996b`

11.6.428 1. Function Overview

The function `_ini_get_agent_nodename` is used to extract the nodename from a given INI configuration file under the `[agent]` section. This function uses the tool `awk`

for parsing the file. The nodename declaration should follow the format `nodename = value`, where `value` is the name of the node. This value is then trimmed of leading and trailing spaces before being output.

11.6.429 2. Technical Description

- **Name:** `_ini_get_agent_nodename`
- **Description:** This function is used to extract the value of `nodename` from the `[agent]` section of a given INI file. Uses `awk` for parsing.
- **Globals:** None.
- **Arguments:** [`$1`: The file to parse and extract the `nodename` from]
- **Outputs:** The `nodename` value from the parsed file.
- **Returns:** Nothing. The function exits after printing the `nodename`.
- **Example usage:** Follows the format `_ini_get_agent_nodename config.ini`

```
agent_nodename=$( _ini_get_agent_nodename config.ini )
echo $agent_nodename
```

11.6.430 3. Quality and Security Recommendations

1. Handle the scenario when the input file is not an INI file or does not exist. Right now, the function will not provide useful feedback if this occurs, which can make it difficult to debug.
2. Clearly document the format of the INI file expected by this function to set the right expectation for the users and maintainers.
3. Consider sanitizing the value read from the INI file. If malicious values are present in the field `nodename`, it could lead to security vulnerabilities.
4. Consider failure scenarios such as what should happen when the `nodename` field is missing, or when the file can't be read, and provide appropriate error messages for these cases.
5. Multiple agent sections in one INI file may lead to problems as the current implementation prints out the `nodename` from the first agent section it encounters. This edge case should be communicated or addressed.

11.6.431 `initialise_cluster`

Contained in `lib/functions.d/cluster-functions.sh`

Function signature: `011e924e7a556253209f90291e5c11cc6d353538a17821169829f72b6d16fe0d`

11.6.432 Function overview

The function `initialise_cluster()` is a Bash function intended to initialize a new server cluster configuration within a given directory. If provided a cluster name, it

will create a new directory with that name and set up necessary subdirectories and configuration files. If the cluster name is not provided, or if a cluster with the provided name already exists, the function will return an error message. If the initialization ends successfully, it will export dynamic paths for this cluster.

11.6.433 Technical description

- Name: `initialise_cluster`
- Description: This function creates a new cluster configuration based on the given cluster name. It creates relevant directories, initial configuration file and exports dynamic paths for further use.
- Globals: `HPS_CLUSTER_CONFIG_BASE_DIR`: The base directory in which the cluster configuration will be created
- Arguments:
 - `$1`: The name of the cluster to initialize
- Outputs:
 - Various status messages, indicating whether configuration has been successfully completed or not.
 - Error messages, when an error occurred (e.g., missing argument, directory already exists).
- Returns:
 - 0: on successful initialization and cluster paths exported
 - 1: if the cluster name was not provided
 - 2: if the cluster directory already exists
 - 3: if exporting cluster paths fails
- Usage Example: `bash initialise_cluster "my-cluster"`

11.6.434 Quality and Security Recommendations

1. Validate user input: Currently the function does not validate that the input is safe, or whether the name is reasonable for a directory name.
2. Handle or report errors from `mkdir`: If the the directories cannot be created, the function will still try to create the files, which will always fail.
3. Reduce scope of variables: Currently the function uses many local variables, which increase complexity and can potentially clash with identically named variables outside the function.
4. Execute least privilege: Ensure that the script is run with as few privileges as possible to avoid potential security risks.
5. Securely handle potential failures in file writing operations: Always check the status code after writing to a file. Failure to write to a file should be treated as an error and should be handled properly.

11.6.435 initialise_distro_string

Contained in `lib/functions.d/configure-remote-host.sh`

Function signature: `c440af9a86acdde30570b73ae6d52c7bddf38d765f513ea26ebf15ee4805200`

11.6.436 Function overview

The `initialise_distro_string` function is a Bash function that creates a string giving a basic description of the system's distribution. It works by checking for the system's architecture, manufacturer, operating system, and version. If it can't find the OS name and OS version, it will list them as "unknown."

11.6.437 Technical description

- **Name:** `initialise_distro_string`
- **Description:** This function generates a string containing a description of the system's distribution. It takes no arguments and will list "unknown" if it cannot find the OS name and version.
- **Globals:** [None]
- **Arguments:** [None]
- **Outputs:** The function outputs a string in the format "[cpu]-[mfr]-[osname]-[osver]".
- **Returns:** Doesn't return.
- **Example usage:**

```
distro_string = $(initialise_distro_string)
echo ${distro_string}
```

11.6.438 Quality and security recommendations

1. Include error handling: If the `/etc/os-release` file doesn't exist or can't be accessed, it would be prudent to add error handling logic to the function and notify the user.
2. Validate variables before usage: To ensure only expected values are used, add sanity checks for the variables.
3. Use more strict condition checks: The function currently considers any existing `/etc/os-release` file as valid. Not just its presence, but also the correctness of its format should be checked.
4. Document the function: The function would benefit from inline comments explaining the logic, and a block comment giving a brief overview of the function and its usage.
5. Use underscore in function name: Using underscore (`_`) between words instead of camel case (`camelCase`) make function names more readable in bash script.

11.6.439 initialise_host_scripts

Contained in `lib/functions.d/configure-remote-host.sh`

Function signature: `f3a19023175709bb6341f62f9bd565ec7430df7cae722dcd18cf2a20800ee478`

11.6.440 Function Overview

The function `initialise_host_scripts` is a Linux bash script function that primarily fetches and sources a host functions script from a URL. It asks an external system (referred to as the provisioning node) for the proper distro string and fetches the appropriate host function script bundle, which it then sources for use in the current shell process. If this fetch and source operation fails, it raises an error and returns a non-zero exit status.

11.6.441 Technical Description

- **name:** `initialise_host_scripts`
- **description:** A bash script function which fetches and sources a script for host functions from a given URL.
- **globals:** None
- **arguments:** None
- **outputs:** This function outputs various status messages to indicate progress of the operation. It also modifies the ongoing shell environment by sourcing the fetched host functions scripts.
- **returns:** When successful, the function completes silently with a zero exit status. If the fetch and source operation failure occurs, it gives an error message and returns a non-zero exit status.
- **example usage:**

`initialise_host_scripts`

11.6.442 Quality and Security Recommendations

1. The function should validate the fetched script before sourcing it. The validity check could include size and content checks or even a signature check if the script is provided with a reliable signature.
2. The curl statement doesn't set any timeouts or retry strategies, leaving this function susceptible to hanging indefinitely. Appropriate timeouts and retry intervals need to be set.
3. It's a good practice to enclose the entire function logic within a try/catch block to ensure that the function fails gracefully in case of any unexpected issues.
4. For enhanced security, we recommend using HTTPS instead of HTTP while fetching any scripts or code.

5. It's advisable to standardize log messages and provide verbose and quiet command-line options to control stdout outputs. Generating logs can help in better tracking and debugging of code.

11.6.443 initialise_opensvc_cluster

Contained in `lib/host-scripts.d/rocky.d/opensvc-management.sh`

Function signature: `5ae00b0ed52acd2fc297bb42596c4d4828b6218e0a73be9e8e50fce1f927287d`

11.6.444 Function Overview

The Bash function `initialise_opensvc_cluster()` is responsible for setting up an OpenSVC cluster. It retrieves necessary values such as the cluster name, cluster secret and tags from the host configuration. It handles the setup logic pertaining to different conditions, whether the parameters are not found or the heartbeat type is not set. It applies the cluster configuration using the OpenSVC commands and sets the node tags (if any are found). The function concludes with logging the completion of cluster initialization.

11.6.445 Technical Description

- name: ``initialise_opensvc_cluster()``
- description: A function to initialize an OpenSVC cluster by
 - ↪ reading config values, configuring the cluster, and setting
 - ↪ node tags.
- globals:
 - ``VAR: cluster_name``: The name of the cluster.
 - ``VAR: cluster_secret``: The secret key associated with the
 - ↪ cluster.
 - ``VAR: ips_addr``: The IP addresses associated with the nodes of
 - ↪ the cluster.
 - ``VAR: node_tags``: Tags associated with the node types in the
 - ↪ cluster.
- arguments:
 - ``$1``: None: This function doesn't take any arguments.
- outputs: Logs detailing steps of the OpenSVC cluster
 - ↪ initialization and potential errors.
- returns: ``1`` if any error occurs during cluster initialization,
 - ↪ otherwise no explicit return.
- example usage:
 - `initialise_opensvc_cluster()`: This function doesn't take any
 - ↪ arguments. You can simply call it in your bash script to
 - ↪ initialize a cluster.

11.6.446 Quality and Security Recommendations

1. For security purposes, avoid logging the `cluster_secret` variable as it might expose sensitive information.
2. Add error handling to the remote cluster variable calls to handle scenarios where these calls fail.
3. Return explicit values for success scenarios, currently, the function does not return anything upon success.
4. Consider making the function more reusable by providing arguments instead of using global variables.
5. Make sure the global variable `node_tags` is not manipulated outside of the function as this function directly depends on its value.
6. Document the expected format and values for the `node_tags` variable for easier debugging and usage.

11.6.447 `install_opensvc_foreground_wrapper`

Contained in `lib/functions.d/opensvc-functions.sh`

Function signature: `044d09f5864508cf501804626eabf1e2282f7b5c3a9c290f490b470dcca0b251`

11.6.448 Function Overview

The function `install_opensvc_foreground_wrapper` is used to install a bash wrapper for the OpenSVC agent to run in the foreground. Its output is directed to a log directory, either specified by the user or defaulted to `/srv/hps-system/log`. If the required directories don't exist, the function will create them. This function performs a preflight check to ensure an agent key exists and is not empty. If the agent key is missing or empty, the function echoes an error message and exits with a status of 2. When the content of the bash wrapper is prepared, it then checks if the target file exists and only replaces the target if the contents therein differ from the intended set of contents.

11.6.449 Technical Description

- **Name:** `install_opensvc_foreground_wrapper()`
- **Description:** A Bash function that safely installs an OpenSVC agent bash wrapper script for running the agent in the foreground. Handles pre-flight checks and facilitates logging output of the agent.
- **Globals:** `[HPS_LOG_DIR: User defined log directory or defaulted to /srv/hps-system/log]`
- **Arguments:** None.
- **Outputs:** Writes a bash wrapper script with specific contents at a target location. It sends output to either user defined `HPS_LOG_DIR` or default path `/srv/hps-system/log`.

- **Returns:** Returns 1 if temporary file creation fails during execution. Returns 0 if the function completes successfully.
- **Example Usage:** Install OpenSVC wrapper in the foreground: `install_opensvc_foreground_wrapper`

11.6.450 Quality and Security Recommendations

1. To prevent file path injection issues, it is recommended to further validate the `HPS_LOG_DIR` global variable value before use.
2. The preflight check on the presence and emptiness of `/etc/opensvc/agent.key` is a good practice. You can enhance it by additionally checking the validity and correct format of this key file.
3. To prevent a full disk from causing a complete system halt, implement disk usage checks and automatic clean-ups of old log files in the log generation process. Avoid writing logs directly to critical locations like `/var/log`.
4. Consider adding shell option `set -u` at the start of the functions to prevent the script from running with unset variables, as this can cause unexpected behavior.
5. Utilize more detailed logging methods to capture more comprehensive logs for complex error scenarios. Log outputs of critical operations can help with debugging later on.

11.6.451 `install_virtualization`

Contained in `lib/host-scripts.d/alpine.d/install-virtualization.sh`

Function signature: `b9a1188f0d12d3fa0ad84887b0e02d3674d7c38cb2ce5084b540ca488b735579`

11.6.452 Function overview

The function `install_virtualization` is used to install various packages related to virtualization on a system. This function primarily uses the `apk` command to install necessary packages. It also enables and starts the `libvirtd` service, usually required for virtualization.

11.6.453 Technical description

- **Name:** `install_virtualization`
- **Description:** This function installs virtualization related packages and starts the required services.
- **Globals:** None
- **Arguments:** None
- **Outputs:** Displays a message about the installation of virtualization packages.
- **Returns:** None

- **Example usage:**

```
install_virtualization
```

11.6.454 Quality and security recommendations

1. Always run this function with appropriate permissions. Using it with root privileges might create security risks.
2. Add error handling for package download and installation steps to ensure that the function does not fail silently.
3. Include logging to track function execution and failures.
4. Make sure the system does not already have these services running to avoid redundancy and potential conflicts.
5. Check the availability of the packages before attempting to install them.
6. Encapsulate all commands that may fail in a try-catch block or equivalent to provide meaningful errors and fail gracefully.

11.6.455 int_to_ip

Contained in `lib/functions.d/host-functions.sh`

Function signature: `1acb712577aa9213ca920a051e80825c73d7775714d43d10ccad322458fc5946`

11.6.456 Function Overview

The function `int_to_ip()` is a bash function that converts an integer to an IP address. It's primarily used in a larger code base where IP addresses are calculated and used dynamically. This function is part of a script that attempts to assign a unique IP address and host name, from a defined range, to a host type within a network base setting.

11.6.457 Technical Description

- **Name:** `int_to_ip()`
- **Description:** This function receives an integer as an argument, does bitwise shifting on it and creates an IP address string by placing dots in between the numbers.
- **Globals:** [`start_ip`: A local variable to keep track of the start IP, `end_ip`: A local variable to keep track of the end IP, `try_ip`: A local variable to keep track of the tried IP, `max`: A local variable representing the max value to use in the loop]
- **Arguments:** [`$1` (local ip): This integer argument is used in computations to generate an output IP address]
- **Outputs:** An IPv4 address built from the input integer.
- **Returns:** None.

- **Example usage:**

```
int_to_ip 16843009  
#=> output: 1.1.1.1
```

11.6.458 Quality and Security Recommendations

1. Consider validating the input data, if it is an integer and in the acceptable range.
2. Implement error handling for improper input or unexpected output.
3. Be cautious about information leakages or unwanted side effects due to globally scoped variables.
4. Consider implementing some form of logging for debugging and traceability of function.
5. Adopt defensive coding practices across the board to ensure that potential misuse of code does not result in compromising the system.

11.6.459 ip_to_int

Contained in `lib/functions.d/host-functions.sh`

Function signature: `7a7e2ac879b38f493155a8d7ebe4e0938b6b76af6dd4451359927f5afb697e52`

11.6.460 Function overview

The Bash function `ip_to_int()` provided here is designed to convert an IPv4 address into an integer. This can be used for a variety of purposes, such as network calculations or storing IP addresses in an efficient manner. The function expects a string representing a standard IPv4 address as input and outputs the corresponding integer representation.

11.6.461 Technical description

- **Name:** `ip_to_int`
- **Description:** Converts an IPv4 address from the typical dot-decimal format (e.g., `192.0.2.146`) to an integer value.
- **Globals:** No global variables used or modified.
- **Arguments:**
 - `$1`: A string representing the IPv4 address to be converted into an integer.
- **Outputs:** This function simply prints the integer representation of the provided IP address to standard output.
- **Returns:** Not applicable as the function does not return a value, only outputs it.
- **Example usage:**

```
ip_to_int "192.0.2.146"
```

11.6.462 Quality and security recommendations

1. **Input validation:** It's strongly recommended to validate that the input is a properly formatted IPv4 address. If improperly formatted input is provided, behavior is undefined.
2. **Error handling:** The function currently doesn't handle errors or unexpected conditions. Consider adding error checking and a proper error message.
3. **Documentation:** Include clear documentation/comments within the function that describe what it does and how it works.
4. **Security considerations:** Ensure that the Bash script doesn't run with elevated privileges unless it's necessary. Remember that inputs should not be trusted and should be sanitized appropriately.
5. **Testing:** Various forms of testing should be done to ensure the function behaves as expected under different conditions and input values.

11.6.463 ipxe_boot_alpine_tch

Contained in `lib/functions.d/ipxe-menu-functions.sh`

Function signature: `6f0d1b191c44f3fce6bc1279370eb316b3e6133af20e28bc5501dc43f55ba4cf`

11.6.464 Function Overview

The function `ipxe_boot_alpine_tch()` initiates a booting sequence for Alpine TCH over Internet Protocol Extensible Firmware Interface (IPXE). It does this using a target's MAC address, and configuration parameters like IP, gateway, CIDR, and hostname. If the necessary Alpine `apkowl` file is not found, the function generates it. The boot arguments are constructed and passed along with kernel and `initramfs` to IPXE to launch the system.

11.6.465 Technical Description

- **Name:** `ipxe_boot_alpine_tch`
- **Description:** This function initiates a boot sequence of the Alpine TCH operating system instance via IPXE.
- **Globals:** [`$HPS_DISTROS_DIR` is directory path where distros are stored, `$mac` is the MAC identification of the target]
- **Arguments:** [`$1` is not directly used in this function, but it is implied that it would normally be MAC address for necessary MAC operations]
- **Outputs:** Kernel logs with details about the booting process. Script will output steps in its boot sequence to `stdout`. It will also output logs in case of errors.
- **Returns:** The function does not return any specific value but it exits with `0` status code upon successful booting.
- **Example usage:** It's typically used in system bootstrapping and not invoked manually, but for testing purposes could be invoked like: `ipxe_boot_alpine_tch`

11.6.466 Quality and Security Recommendations

1. Input Validation: Validate the inputs including the MAC address and CIDR block for correctness and presence.
2. Error Handling: Increase error handling and logging especially around critical operations like apkovl file creation.
3. Secure File Operations: File operations around the network should be secured to prevent snooping and manipulation.
4. Permissions: Script should be run with only necessary permissions. Excessive permissions could lead to abuse if a vulnerability is present.
5. Dependencies: There seems to be a reliance on other functions. Ensuring the robustness and security of those functions is key to the secure performance of this function.
6. Anonymity: Consider masking or replacing key identifiers in logs to ensure privacy.
7. Code Review: Frequent code reviews should be conducted to ensure that current security and quality standards are maintained.

11.6.467 ipxe_boot_from_disk

Contained in `lib/functions.d/ipxe-menu-functions.sh`

Function signature: `f26dc79cd8c9ea270e2758702b63d4a607fd3fd40737c5d947de898af37ff1fe`

11.6.468 Function overview

The `ipxe_boot_from_disk` function is designed to facilitate the process of booting from a local disk via BIOS by exiting the iPXE stack. The purpose of this function is to invoke the iPXE header and issue necessary commands to send control back to the BIOS, thus triggering it to boot from the local disk. A brief pause is also incorporated before exiting.

11.6.469 Technical description

- **name:** `ipxe_boot_from_disk`
- **description:** This function is used to boot from local disk via BIOS by exiting iPXE stack. It involves invoking iPXE header and issuing commands to hand control back to BIOS, along with a brief pause before completing its operation.
- **globals:** No global variables.
- **arguments:** No arguments.
- **outputs:** Outputs are the commands that are echoed- “Echo Handing back to BIOS to boot”, “sleep 5” and the exit command.
- **returns:** This function does not return any value.
- **example usage:** After defining the function, it can be invoked just by calling its name as follows:

ipxe_boot_from_disk

11.6.470 Quality and security recommendations

1. **Error Handling:** To improve upon this function, it is advisable to include error handling to account for any issues that might arise during the execution of the commands.
2. **Logging:** Adding logging statements could be beneficial for troubleshooting purposes.
3. **Security:** The function should have checks in place to ensure only authorized personnel can invoke a boot from the disk, reducing potential security risks.
4. **Function Validation:** Validate the outcome of each echo command to confirm it executed successfully.
5. **Commenting:** More comments can be included to clarify the role of each function step. This will make the code easier to evaluate and maintain.

11.6.471 ipxe_boot_from_disk

Contained in `lib/functions.d/ipxe-menu-functions.sh`

Function signature: `f26dc79cd8c9ea270e2758702b63d4a607fd3fd40737c5d947de898af37ff1fe`

11.6.472 Function overview

The `ipxe_boot_from_disk` function is used to boot from a local disk via BIOS by exiting the iPXE stack. It first calls the `ipxe_header` function, prints a message to the console indicating that control is being handed back to BIOS for booting, sleeps for 5 seconds, and then exits.

11.6.473 Technical description

- **Name:** `ipxe_boot_from_disk`
- **Description:** This function facilitates booting from a local disk via BIOS by exiting the iPXE stack.
- **Globals:** None.
- **Arguments:** None.
- **Outputs:** Prints a message to the console indicating that control is being handed back to BIOS for booting.
- **Returns:** Nothing.
- **Example Usage:**

`ipxe_boot_from_disk`

11.6.474 Quality and Security Recommendations

1. As a good practice, it would be beneficial to add error handling or exceptions for certain operations such as checking if the BIOS processes are executing correctly when handling back control from iPXE.
2. It would also be useful to log the operations for debugging and auditing purposes.
3. From a security perspective, ensure that the appropriate permissions are set for the function to prevent unauthorized access or changes.
4. Avoid hardcoding values like sleep time. Instead, allow these to be configured through variables or configuration files.

11.6.475 ipxe_boot_installer

Contained in `lib/functions.d/ipxe-menu-functions.sh`

Function signature: `1437c52de005d7eb1a9211f7f5c6da2aec0f19a08013b218e09e6cac1561ba14`

11.6.476 Function overview

The `ipxe_boot_installer` function is used to install a new host through ipxe boot. The host type and profile are defined as arguments, and these provide the essential configuration for the installation process. Different kernel and initrd files are utilized based on the host type. If the host type is defined as “TCH”, then the host is configured for network boot. Else, the host gets configured based on the loaded cluster host type profiles.

The function checks if the host is already installed and aborts the installation if that’s the case. After setting the configurations, the function specifies the distribution path and URL, mounts the distro ISO and prepares for PXE Boot for non-interactive installation. The function handles different OSNAME cases and tosses an error if the configuration for the particular OSNAME does not exist. In the end, the state of the host is set as “INSTALLING”.

11.6.477 Technical description

- **Name:** `ipxe_boot_installer`
- **Description:** This function checks the host type and based upon that, it configures and prepares the host for the installation through PXE Boot. It performs various tasks such as configuring network boot, loading cluster host profiles, setting host configuration, etc.
- **Globals:** None
- **Arguments:**
 - `$1`: `host_type`: The type of host that is going to be installed.
 - `$2`: `profile`: The profile that contains the configuration settings for the host.
- **Outputs:** Log messages about the installation process and status.

- **Returns:** N/A
- **Example usage:** `ipxe_boot_installer "TCH" "profile1"`

11.6.478 Quality and security recommendations

1. Input validation - Ensure that the supplied `host_type` and `profile` are valid and not malicious. This is especially important if the inputs are mentioned by an untrusted user or from an untrusted source.
2. Error handling - Add more robust error handling. Currently, if the `$state` fetch from `host_config` fails for any reason, the script continues to execution. This may lead to undesired outcomes or misconfiguration.
3. Log improvements - Including more detailed logging could assist in diagnosing problems or errors during installation.
4. Return codes - Although the function does not currently return a specific code, adding return codes to indicate successful execution or error conditions can be helpful in larger scripts or systems.
5. Code redundancy - There is some repeated code in the `ipxe_boot_installer` function, such as the configuration setting for `host_config`. This could be consolidated to reduce redundancy and improve code maintainability.
6. Configuration check - Additional checks can be configured for verifying the validity of `DIST_PATH` and `DIST_URL`.

11.6.479 ipxe_boot_installer

Contained in `lib/functions.d/ipxe-menu-functions.sh`

Function signature: `ad6c7317bbb7a71cfcd8f86037b9440faaacb0b1a309a4e03c807b567e589b21`

11.6.480 Function overview

The `ipxe_boot_installer` function is part of a Bash script designed to perform automated installations of operating systems on target hosts over networks using the iPXE boot firmware. The function takes in two positional parameters: `host_type`, and `profile`. The function interacts with global configurations to identify the host's hardware and corresponding OS requirements. Then, it begins the process of mounting the necessary boot images and preparing the host for installation. If the function detects a previously installed OS on the host, it will abort the installation to prevent damage. The function currently supports installations for `rockylinux` and has placeholder cases for `debian` and other operating system types, although these are not yet implemented.

11.6.481 Technical description

- **Name:** `ipxe_boot_installer`

- **Description:** This function prepares a target host for a network-based OS installation using the iPXE boot firmware.
- **Globals:**
 - **HPS_DISTROS_DIR:** Directory where the necessary boot files for each OS-/hardware setup are stored.
 - **mac:** Global unique mac address of the host.
 - **CGI_URL:** URL for CGI services.
- **Arguments:**
 - **\$1: host_type:** Specifies the type of the target host.
 - **\$2: profile:** Describes a specific configuration profile for the operating system to be installed.
- **Outputs:** Installation script steps to standard output, as well as debug or error logs.
- **Returns:** Varies based on multiple conditional branching. Often, failure cases will terminate script execution.
- **Example usage:**

```
ipxe_boot_installer rockylinux default
```

11.6.482 Quality and security recommendations:

1. Implement error handling for critical steps to prevent script failure.
2. Integrate exception handling to ensure clean termination during a fail state, including unmounting of ISO or other resources.
3. Clean up or secure temporary files used to store iPXE boot install scripts. These may contain sensitive data or configuration details.
4. Implement timeouts or checks to confirm a boot image loads successfully.
5. Hash-check downloaded distros to verify their integrity.
6. As functionality grows, consider creating multiple functions for different OS installations rather than adding more conditional branches to this function.

11.6.483 ipxe CGI fail

Contained in `lib/functions.d/ipxe-menu-functions.sh`

Function signature: 98eb961ae3ed7dffc7f4ae68cd1c88aa5f47672ee53b71b29f0428922e8efaa

11.6.484 Function overview

The function `ipxe CGI fail()` is utilized to handle failure during the Implicit PXE (iPXE) process. This function generates an iPXE header, sends an error message to the hp's log, reports the error in the iPXE code to be interpreted by iPXE supporting software, waits for 10 seconds, then reboots the system.

11.6.485 Technical description

- **Name:** `ipxe_cgi_fail`
- **Description:** This function is intended to handle failures during the iPXE process. Upon invocation, it creates an iPXE header, logs an error message, alerts the user about the failure, waits a bit, then reboots the system.
- **Globals:** None.
- **Arguments:**
 - `$1`: `cfmsg`: This is the error message to be logged and displayed.
- **Outputs:** An iPXE formatted error message, including the input error message.
- **Returns:** Nothing. The function does not have a return statement, but it does call `exit`, terminating the script it's within.
- **Example usage:** `ipxe_cgi_fail "Network boot failed"`

11.6.486 Quality and security recommendations

1. **Input Validation:** Implement input validation on `$1` to check if it is set and isn't an empty string before proceeding with the rest of the function. This would make the function more robust against erroneous invocations.
2. **Error handling on `exit`:** The script terminates abruptly with `exit`. It's recommended instead to return an error code, and let the main section of your script decide how to handle the error.
3. **Message Standardization:** It's suggested to use standardized error codes and messages to make troubleshooting more efficient.
4. **User Instructions:** Since the error causes the system to reboot, provide more information to the user regarding what they should do after the reboot.
5. **Securing Logging:** Ensure only authorized and authenticated applications or services can write to the log file. This prevents unauthorized modifications which could lead to misinterpretation of system states.

11.6.487 `ipxe_cgi_fail`

Contained in `lib/functions.d/ipxe-menu-functions.sh`

Function signature: `98eb961ae3ed7dffc7f4ae68cd1c88aa5f47672ee53b71b29f0428922e8efaa`

11.6.488 Function Overview

The function `ipxe_cgi_fail` is designed to convey an error message if something goes wrong during the operation of the iPXE. It first calls the `ipxe_header` function, logs the error message, and then displays an error in the iPXE shell. This includes the error message passed to the function, a sleep command to pause execution for 10 seconds, and then a reboot. The function then ends using the `exit` command.

11.6.489 Technical Description

- **Name:** `ipxe_cgi_fail`
- **Description:** Generates and displays an error message when IPXE encounters an issue.
- **Globals:** None.
- **Arguments:** [\$1: Error message to be displayed in IPXE and logged]
- **Outputs:** An error message in IPXE shell, and a log entry.
- **Returns:** None. The function ends with an `exit` command after displaying the error message, pausing for 10 seconds, and rebooting.
- **Example usage:**

```
ipxe_cgi_fail "Unable to connect to the server"
```

11.6.490 Quality and Security Recommendations

1. Always ensure that the function is supplied with a meaningful and useful error message to help with debugging.
2. It might be beneficial to include some environment information in the log message or the error output to provide additional debugging context.
3. Exercise caution when executing from an untrusted context or with unsanitized inputs, as the log message could potentially expose sensitive data.
4. Any usage of this function results in a reboot. Ensure that this is an acceptable behavior and that it does not unexpectedly disrupt other processes or tasks. If not, consider modifying the function to provide alternative execution paths.
5. Regularly review the logs for errors to inform improvements to the system as a whole and increase overall robustness.

11.6.491 `ipxe_configure_main_menu`

Contained in `lib/functions.d/ipxe-menu-functions.sh`

Function signature: `1f7c5adf70ab3b0d8ce1dc04122eff4ae6c8141a95956ef26ee00221876e111c`

11.6.492 Function Overview

The function `ipxe_configure_main_menu` is used for displaying the main menu options when a host is not configured in a cluster. The function logs and delivers a configuration menu. The menu contains options for a host to install, view configuration, recover from Disaster Recovery Host (DRH), enter rescue shell, boot from local disk, reboot, and set advanced options. Another significant feature is Forced installation which is set by checking global `'FORCE_INSTALL'` and can be enabled or disabled from the menu.

11.6.493 Technical Description

- **Name:** `ipxe_configure_main_menu`

- **Description:** This bash function delivers a configuration menu with several options to manage a host that is not configured within a cluster. The options include various actions such as installing, viewing configurations, enabling forced installations, etc.
- **Globals:** `FORCE_INSTALL`: An option that forces the installation process on next boot if set to “YES”.
- **Arguments:** None
- **Outputs:** Outputs the configuration menu to the terminal
- **Returns:** Doesn't return a value, executes commands based on the user's choice in the menu
- **Example usage:** This function does not require any arguments. An example of usage would be simply calling the function as `ipxe_configure_main_menu`.

11.6.494 Quality and Security Recommendations

1. Ensure that the host config and all command line arguments are properly sanitized to prevent command injection vulnerabilities.
2. It would be useful to check the statuses of operations like 'host_config' and short circuit the execution of the function, in the event of an error.
3. Ensure that the logging mechanism in 'hps_log' properly sanitizes and escapes any string input to prevent log injection attacks.
4. As a quality improvement, among the menu items, 'Recover from Disaster Recovery Host (DRH)' is mentioned as 'NOT YET IMPLEMENTED'. Consider implementing this feature or removing it from the menu to avoid user confusion.
5. For enhancing security while fetching the log, handle the failure case with more than just an echo statement. Prompt the user, or retry fetch.

11.6.495 `ipxe_configure_main_menu`

Contained in `lib/functions.d/ipxe-menu-functions.sh`

Function signature: `1f7c5adf70ab3b0d8ce1dc04122eff4ae6c8141a95956ef26ee00221876e111c`

11.6.496 Function overview

The bash function `ipxe_configure_main_menu` is used to generate a main menu structure for a firmware-level preboot eXecution Environment (iPXE), which gives computers the capability to load other software over network. This menu is shown when the cluster is configured but the host is not yet configured.

11.6.497 Technical description

11.6.497.1 Function Detail:

- **Name:** `ipxe_configure_main_menu`

- **Description:** This function is used to generate a menu structure for a preboot eXecution Environment. If the cluster is configured but the host is not, this menu is delivered. Based on various conditions like whether forced installation is enabled or not, it provides different options in the menu.
- **Globals:** None
- **Arguments:** None
- **Outputs:** Prints a menu structure for configuring a host. This menu contains various options such as enabling forced installation, entering rescue shell, booting from local disk, or rebooting the host, among several others.
- **Returns:** It doesn't return an explicit value, but generates output to stdout.
- **Example usage:** `ipxe_configure_main_menu`

11.6.498 Quality and Security Recommendations

1. To avoid shell injection, never pass untrusted input to `eval`, the system, or others unless you've properly sanitized it.
2. Use "\$@" instead of "\$*" to correctly handle parameters with spaces or special characters.
3. You should ensure that any local variables that should not be exported to global scope are not inadvertently exported.
4. Check for the existence and the type of commands before invoking them.
5. Include error checking for every command that can feasibly fail.
6. Use a consistent style in your code to ensure better readability and maintainability.

11.6.499 ipxe_header

Contained in `lib/functions.d/ipxe-menu-functions.sh`

Function signature: `f724cbc120f1a1eae5cf985238d6bc44a932f4521832fa8959c47118c5068ed5`

11.6.500 Function Overview

The `ipxe_header()` function transmits a pxe header to prevent boot failures. It establishes a series of variables for use in IPXE scripts, then outputs these to the console.

11.6.501 Technical Description

- **Name:** `ipxe_header`
- **Description:** The function initially transmits a pxe header to prevent any boot failures. It sets a couple of global variables (`CGI_URL`, `TITLE_PREFIX`) for use in IPXE scripts. The function constructs and prints a specific message structure.
- **Globals:** [`CGI_URL`: The URL to the boot manager script running on the selected server, `TITLE_PREFIX`: The title prefix combining the cluster name, mac address, and network IP of the server]

- **Arguments:** [None]
- **Outputs:** Console output which includes a log message, specifics about the server, and information about the client.
- **Returns:** No explicit return value.
- **Example usage:** It's frequently used in the context of IPXE scripting and will typically be invoked without arguments, like so:

```
ipxe_header
```

11.6.502 Quality and Security Recommendations

1. It is crucial to double-check all inputs, specifically while fetching images over the network from a specified URL. This safeguard prevents possible vulnerabilities associated with potential malicious content.
2. All user data or potentially sensitive information output to the console should ideally be sanitized or selectively displayed, as it may expose critical information to malicious actors or risk leaking sensitive data.
3. It may be useful to add error handling to address potential failures that may occur throughout the execution of the function. For instance, if `cgi_header_plain` or `imgfetch` fail, there should be appropriate error messages and failure handling.
4. Increase the readability of the code by adding further comments regarding the functionality and working of different code blocks within the function.

11.6.503 ipxe_header

Contained in `lib/functions.d/ipxe-menu-functions.sh`

Function signature: `f724cbc120f1a1eae5cf985238d6bc44a932f4521832fa8959c47118c5068ed5`

11.6.504 Function Overview

The function `ipxe_header()` is a Bash script function that is primarily used to generate a dynamic IPXE script. This function initially sends a PXE (Preboot eXecution Environment) header via a function `'cgi_header_plain'`, followed by setting some variables such as `'CGI_URL'` and `'TITLE_PREFIX'` for later use within the IPXE scripts. The main task of this function involves the construction and execution of an embedded `'heredoc'` (EOF) shell script. The script is defined to set a log message, fetch an image, and echo relevant cluster and client information.

11.6.505 Technical Description

- **name:** `ipxe_header()`

- **description:** This function sends a PXE header, sets the vital variables for later use in iPXE scripts, and concatenates an embedded here document which primarily aims to set a log message, fetch an image and echo important information involving the cluster and client.
- **globals:** [CGI_URL: URL sequence to be used in iPXE scripts, TITLE_PREFIX: Prefix string appended before {mac:hexraw} {net0/ip}]
- **arguments:** No arguments required
- **outputs:** The function will output a dynamically generated iPXE script
- **returns:** Function does not return any particular value
- **example usage:**
`ipxe_header`

11.6.506 Quality and Security Recommendations

1. Sanitize Input and Output: Ensure that all the variables used in the function are properly sanitized to prevent injection attacks.
2. Error Handlers: Add error handlers to catch failures and ensure that the function behaves as expected in all cases, particularly when fetching the image via 'imgfetch'.
3. Keep URLs and Ports Secure: The function's 'CGI_URL' that is obtained from 'next-server' must be kept secure, as any lapses could potentially open doors to malicious attacks.
4. Be Cautious of Command Injection: Since shell scripts are vulnerable to command injection attacks, it is advisable to use built-in language features that ensure variable expansion does not permit arbitrary command execution.
5. Code Maintainability: Comment the code in a structured manner for increased readability and ease of future maintenance.

11.6.507 ipxe_host_install_menu

Contained in `lib/functions.d/ipxe-menu-functions.sh`

Function signature: `0627b23e1b7a33e58451ad40a8e31ebc0e9911be4bc534e1afb8dc54dea050c4`

11.6.508 Function overview

The function `ipxe_host_install_menu` is primarily responsible for the creation and display of an interactive menu within an iPXE environment. The generated menu provides a list of options to the user for configuring and installing different host solutions such as Thin Compute Host, Storage Cluster Host, Disaster Recovery Host, and Container Cluster Host.

11.6.509 Technical description

- **name:** `ipxe_host_install_menu`
- **description:** This function creates an interactive menu for host installation. It uses heredoc (`cat <<EOF ... EOF`) to print an instalment menu list. Depending on user selection, a specific command chain is fetched and replaced in the iPXE environment.
- **globals:** [`TITLE_PREFIX`: Title prefix for the menu, `CGI_URL`: Base URL to use for interaction]
- **arguments:** [None]
- **outputs:** It displays an interactive menu to stdout, containing various options for host installation in the iPXE environment.
- **returns:** None. But calls various other commands depending on user selection.
- **example usage:** `ipxe_host_install_menu`

11.6.510 Quality and security recommendations

1. Set your Bash scripts to fail on unhandled errors (`set -e`) for better error handling and stability.
2. Always use double quotes around variables to prevent word splitting.
3. Prefer declaring function-specific variables with `local` to avoid scope issues.
4. Make sure to sanitize and validate user inputs to prevent potential security risks.
5. As the function interacts with external URLs, it is advisable to implement adequate measures to ensure a secure HTTPS connection.

11.6.511 `ipxe_host_install_menu`

Contained in `lib/functions.d/ipxe-menu-functions.sh`

Function signature: `0627b23e1b7a33e58451ad40a8e31ebc0e9911be4bc534e1afb8dc54dea050c4`

11.6.512 Function overview

The function `ipxe_host_install_menu` is used to create an iPXE installation menu for different types of hosts. Each host type corresponds to a different installation option, and users are given a choice between these options. The function sends a log message containing the user's selection and triggers the necessary installation process.

11.6.513 Technical Description

- **name:** ``ipxe_host_install_menu``
- **description:** Generates an iPXE installation menu and handles user selection for o
- **globals:**
 - ``TITLE_PREFIX``: Used for creating the menu title.

- ``CGI_URL``: The base URL for the cgi scripts.
 - ****arguments****: None
 - ****outputs****
 - The function outputs an installation menu.
 - Upon making a selection, a log message is printed, and an installation process is initiated.
 - ****returns****: The function doesn't explicitly return a value.
 - ****example usage****
- ```
```bash
ipxe_host_install_menu
```
- This will activate the function and display the installation menu to the user.

Quality and Security Recommendations

1. Always sanitize user inputs to prevent any potential security threats.
2. Consider implementing error checking for network calls (imgfetch and chain).
3. Add meaningful comments to improve readability and maintainability.
4. Use meaningful names for global variables, maintain consistent naming conventions.
5. Implement checks to ensure that the necessary global variables are set before the function is called.
6. Handle edge cases and expect the unexpected in the user's input. Validate the `selected` variable.
7. Encapsulate all the static string values used in the function as constants at the top of the file.
8. Make sure that adequate logging is in place. This not only helps in debugging but also provides a record of user actions.

``ipxe_host_install_sch``

Contained in ``lib/functions.d/ipxe-menu-functions.sh``

Function signature: `76d8d63df923bfeaa3d9b18625e12e7909180d679baee564a7d5760de6c84ba`

Function Overview

The ``ipxe_host_install_sch`` function is used to provide a menu for installing a Storage

Technical Description

- ****name****: ``ipxe_host_install_sch``
 - ****description****: This bash function creates an interactive iPXE menu for configuring
 - ****globals****: ``TITLE_PREFIX``: represents the prefix of the menu title, ``CGI_URL``: represents the base URL for the CGI scripts.
 - ****arguments****: This function does not require any arguments.
 - ****outputs****: The function prints an iPXE menu for user interaction.
 - ****returns****: The function does not have a return value but alters the control flow of the script.
 - ****example usage****
- ```
```bash
```

ipxe_host_install_sch

11.6.514 Quality and Security Recommendations

1. Check for the availability of global variables before using them. This will prevent unexpected outputs.
2. Verify your endpoint `${CGI_URL}` is secure and trusted to prevent possible injection attacks.
3. Avoid using plaintext for critical log messages or consider adding encryption for log messages to increase security.
4. It's recommended to handle possible failed fetch requests when `imgfetch` is unable to fetch data.
5. Consider adding validation for user input to avoid potential security risks.

11.6.515 ipxe_host_install_sch

Contained in `lib/functions.d/ipxe-menu-functions.sh`

Function signature: 76d8d63df923bfeaa3d9b18625e12e7909180d679baee564a7d5760de6c84baa

11.6.516 Function Overview

The `ipxe_host_install_sch` function is part of an iPXE script used for deploying storage cluster hosts. When called, it displays a menu to the user that allows them to initiate the installation of a storage cluster host. Different options are available depending on the user's storage infrastructure, such as single-disk or RAID configurations. The function also logs selections made from the menu and chainloads other scripts as per the user's choice.

11.6.517 Technical Description

Name: `ipxe_host_install_sch`

Description: This function displays an iPXE-based menu for installing a Storage Cluster Host. The user can initiate installation of a Single-Disk or RAID storage cluster.

Globals: [`TITLE_PREFIX`: prefix for the installation menu title, `CGI_URL`: URL of CGI scripts for logging and processing menu items]

Arguments: [None]

Outputs: The function outputs an interactive menu for the user. Logging and processing of user actions are facilitated by fetching and chaining CGI scripts.

Returns: The function does not return a value.

Example Usage:

`ipxe_host_install_sch`

11.6.518 Quality and Security Recommendations

1. To enhance security, ensure that input from the user is properly sanitized before it is fed into the `imgfetch` and `chain` commands to avoid potential command injection attacks.
2. Since this function relies on external scripts hosted at `CGI_URL`, make sure that this URL is secure and the scripts are reviewed for potential vulnerabilities.
3. The function currently does not handle errors gracefully. Adding some error handling logic would likely improve the function, for example by informing the user if fetching or chaining the CGI scripts fail.
4. Avoid storing any sensitive information such as keys or passwords in global variables. Use secure methods for authentication.

11.6.519 `ipxe_init`

Contained in `lib/functions.d/ipxe-menu-functions.sh`

Function signature: `b99f8e1adfe13d429e91c39da6aeab71b263f99c8d73aa5f35dd7351988760e2`

11.6.520 Function Overview

The function `ipxe_init` is designed to configure a network boot environment when the host isn't known yet because its MAC address hasn't been retrieved. It uses iPXE, an open-source boot firmware, to request the boot configuration from a server specified by the `CGI_URL` global variable. Once the configuration has been loaded, the iPXE shell executes it. If no configuration exists for the host, the function includes commented out code that handles this scenario by sending a message to the console and rebooting though this should never be the case as the boot manager is expected to create it.

11.6.521 Technical Description

- **Name:** `ipxe_init`
- **Description:** This function initializes iPXE, a network boot firmware, and fetches and executes the boot configuration from a remote server. This function is utilized in situations where the cluster is configured but the host is not yet known due to a missing MAC address.
- **Globals:**
 - `CGI_URL`: This global variable holds the URL of the remote server from which the boot configuration will be requested.
- **Arguments:** This function doesn't require any arguments.
- **Outputs:** The function produces a console output of the actions taking place including the fetching and loading of the boot configuration and its execution status.
- **Returns:** This function doesn't return any value because it's operations are all about booting a system with iPXE.

- **Example Usage:** `ipxe_init`

11.6.522 Quality and Security Recommendations

1. The function contains commented out code which handles a situation where the configuration for the given MAC address is not found on the remote server. It would be better to uncomment these lines to account for this exception.
2. The global variable `CGI_URL` should be sanitized before using it in the function to prevent potential command injection.
3. For recovery and fault tolerance, consider adding a retry mechanism for fetching the configuration from the remote server if it fails at the first attempt.
4. Make sure to use secure transport (HTTPS not just HTTP) for fetching the configuration to prevent potential man-in-the-middle attacks.
5. Enhance logging with error levels, and make use of a logging system that supports searching and filtering, which would help in troubleshooting and incident response.

11.6.523 `ipxe_init`

Contained in `lib/functions.d/ipxe-menu-functions.sh`

Function signature: `b99f8e1adfe13d429e91c39da6aeab71b263f99c8d73aa5f35dd7351988760e2`

11.6.524 Function Overview

The `ipxe_init` function is primarily used during the network booting process. Its function is to load the iPXE configuration for a client machine or host within a network. If the cluster is configured and it cannot identify the host yet (since it does not have the host's MAC address), this function gets utilized. The function requests the boot configuration from a specific URL, fetches, loads, and executes the configuration. The function also takes care of scenarios where the host configuration could not be found.

11.6.525 Technical Description

- **Name:** `ipxe_init`
- **Description:** The function is used to initialize ipxe, which includes fetching, loading, and executing the iPXE configuration for a host within a network.
- **Globals:** [`CGI_URL` : URL from where the iPXE configuration for hosts is fetched]
- **Arguments:** None
- **Outputs:** Initiate a request for fetching the configuration, fetch, load, and execute the config or an error message if no host config found.
- **Returns:** None, as the function does not return a value but executes certain operations.

- **Example usage:** `ipxe_init`

11.6.526 Quality and Security Recommendations

1. Proper error handling should be implemented. In the current structure, there are lines that have been commented out which are supposed to handle cases where no host configuration is found. These lines should be uncommented and ensure they are working correctly.
2. There should be validation and sanitization of the `CGI_URL` global variable given that it introduces potential security risks.
3. Consider cases where the `imgfetch`, `imgload`, or `imgexec` commands might fail. Ensuring these commands are successful before proceeding will increase the robustness of the script.
4. Better management and usage of global variables. Use of them can result in side effects, if they are modified in other parts of the scripts unknowingly. Consider passing `CGI_URL` as a parameter to the function.
5. Input validation should be included for safer and more reliable script execution. This becomes crucial especially if this script is expected to run in different environments with different inputs. This will prevent potential code injection and other related security risks.
6. Ensure the use of the script in a secure and encrypted communications environment. Since the fetching of the iPXE configuration happens over the web, using a secure HTTP protocol (HTTPS) is recommended. This helps to safeguard against potential man-in-the-middle attacks.

11.6.527 `ipxe_network_boot`

Contained in `lib/functions.d/ipxe-menu-functions.sh`

Function signature: `12978609d623d739eeb7768560e3b7ba0c3cd5d711ae46c9949271f4abc20919`

11.6.528 Function Overview

The function `ipxe_network_boot()` is primarily designed to determine the type of host system and perform a network boot accordingly. The function fetches the host type and logs it for debugging purposes, before proceeding to inspect the host type and perform operations correspondingly. If the host type is 'TCH', it validates the Alpine repository before attempting to boot. If the validation fails, it logs an error message, sends a failure message, and returns 1, demonstrating unsuccessful execution. If validation is successful, a boot operation specific to 'TCH' systems is performed. For any other type of host, the method logs a message notifying that network boot is not supported for that host type.

11.6.529 Technical Description

- **Function Name:** `ipxe_network_boot`
- **Description:** Responsible for determining the host type and carrying out a network boot operation accordingly.
- **Globals:** `mac`: the mac address of the host system.
- **Arguments:** None
- **Outputs:** Logs messages for debugging, errors, and host type support. May execute a failure response if the Alpine repository is not prepared.
- **Returns:** 1 if the Alpine repository validation fails, otherwise no explicit return.
- **Example Usage:**

```
ipxe_network_boot
```

Note: This function could only be invoked without any parameters.

11.6.530 Quality and Security Recommendations

1. Consider including an input validation to check if a `mac` global is defined before performing any other operations.
2. Additional error handling could be implemented to account for any potential pitfalls or unexpected issues during the booting process.
3. Providing additional support for other host types beyond 'TCH' could enhance flexibility and universality of the function.
4. Regularly review and update the function to ensure compatibility with updated versions of the Alpine repository and host configurations.
5. Uphold good practice of ensuring the sensitive information, such as MAC addresses, used within the logs and outputs is secured appropriately.

11.6.531 `ipxe_reboot`

Contained in `lib/functions.d/ipxe-menu-functions.sh`

Function signature: `c2e2fb28eaa8f9898d8b5cb181b2b278c884005d1a98813c38797c3225f12b52`

11.6.532 Function overview

The function, `ipxe_reboot`, accepts a single argument, `MSG`, logging that a reboot was requested with `MSG` as a parameter, outputs specific headers using the `ipxe_header` function, and then echoes `MSG`. If `MSG` is not an empty string, it echoes that the system is rebooting, puts the system to sleep for 5 seconds, and then reboots the system.

11.6.533 Technical description

- **Name:** `ipxe_reboot`
- **Description:** This function logs an info message showing that a reboot is requested, outputs headers using another function `ipxe_header`, checks if `MSG` (first input argument) is not a null string and echoes it if true. It then echoes “Rebooting...”, puts the system to sleep for 5 seconds and then reboots the system.
- **Globals:** None
- **Arguments:**
 - `$1` (`MSG`): The message to be logged and echoed just before the reboot command. It’s optional, and if it’s null or not defined, it won’t be used.
- **Outputs:** Echoed statements to the standard output for logging and status updates.
- **Returns:** No value is returned as the terminal will be closed upon successful function execution because of the reboot command.
- **Example usage:** `ipxe_reboot "System updates are completed"`

11.6.534 Quality and security recommendations

1. Always make sure to validate the input parameters. Even though this function does not explicitly use user-provided inputs, it’s still a good practice.
2. The `reboot` command is a powerful system command. Ensure this function is only accessible to and executable by authorized users and applications to prevent misuse.
3. There are no command success/failure checks in the function. Consider adding error handling or command result checks to make the function more robust.
4. In an environment where the system log is reviewed or parsed, consider standardizing the log format to make the logs more readable.
5. The function depends on another function, `ipxe_header`. Ensure this dependant function is robust, secure, and available where `ipxe_reboot` is used to avoid runtime errors.
6. Avoid hard-coding values such as the sleep duration. It would be a better practice to make such values configurable.

11.6.535 `ipxe_reboot`

Contained in `lib/functions.d/ipxe-menu-functions.sh`

Function signature: `c2e2fb28eaa8f9898d8b5cb181b2b278c884005d1a98813c38797c3225f12b52`

11.6.536 Function overview

The function `ipxe_reboot` is designed to print a log message indicating a reboot request, followed by the actual reboot request. Initially, it defines a local variable `MSG`

and assigns it the value of \$1, the first argument passed to the function. If the MSG is not null, it is printed, and regardless, a standard rebooting message is printed. Finally, it waits for 5 seconds and then triggers the reboot.

11.6.537 Technical description

- **Name:** ipxe_reboot
- **Description:** This function is used to log a reboot request and then reboot the system. This includes constructing a header using `ipxe_header`, printing a custom reboot message if provided, and echoing a standard “Rebooting...” message. It then instructs the system to sleep for 5 seconds before rebooting.
- **Globals:** No globals are used explicitly in this function.
- **Arguments:**
 - \$1: MSG: An optional message to be displayed during the reboot process.
- **Outputs:**
 - If a message is provided as an argument (MSG is not null), it is printed.
 - A standard “Rebooting” message is always printed.
- **Returns:** The function doesn’t return a value.
- **Example usage:**

```
ipxe_reboot "Scheduled system reboot"
```

This will log a message “Reboot requested Scheduled system reboot”, output the provided message, and then the system will reboot after a pause of 5 seconds.

11.6.538 Quality and security recommendations

1. Buffer Overflow Protection: To avoid any risk of buffer overflow, the function should previously validate the length of the \$1 input argument before assigning it to the MSG variable.
2. Error Handling: If the reboot command fails to execute, this function doesn’t handle it. This can be improved by adding error handling routines or exit codes checking for each command in the function.
3. Secure Logging: The function simply logs the message, “Reboot requested \$MSG”. To ensure secure logging, it should also log the identity of the user or process that initiated the reboot, and timestamp each log entry.
4. Documentation: Each step and the overall purpose of the function should be properly documented in the code.

11.6.539 ipxe_show_info

Contained in `lib/functions.d/ipxe-menu-functions.sh`

Function signature: `cbd14427c29a67a77d52cb0fd250b99f45cb94110f4d4b0a3f1123ec4fe219a4`

11.6.540 Function overview

The `ipxe_show_info()` function is mainly used to display different configurations of a particular host. This function pulls up information such as the host's IP, hostname, platform, UUID, serial number, product, cluster configuration, and HPS (High Performance Storage) paths.

The function achieves this by taking a single argument, which determines the category of information that will be displayed for the host. The categories include `show_ipxe`, `show_cluster`, `show_host`, and `show_paths`.

11.6.541 Technical description

- **Name:** `ipxe_show_info()`
- **Description:** This function displays specific information about a host, based on what category is passed as an argument.
- **Globals:**
 - `HPS_CLUSTER_CONFIG_DIR`: The directory where the cluster configuration is stored
 - `CGI_URL`: Link to CGI (Common Gateway Interface) script which is used to implement the command `process_menu_item`
 - `HPS_CONFIG`: The file where the HPS (High Performance Storage) configuration is stored
- **Arguments:**
 - `$1: category`: The type of information to display about the host. The options are `show_ipxe`, `show_cluster`, `show_host`, and `show_paths`.
- **Outputs:** Information about the host in the requested category, such as IP address, hostname, platform, UUID, serial and product numbers, and more.
- **Returns:** Does not return a value.
- **Example usage:** `ipxe_show_info show_ipxe`

11.6.542 Quality and security recommendations

1. Always ensure variable sanitization before using any variable in command substitution. This removes potential harmful commands being hidden as a value for a variable.
2. Always quote the variables. This will prevent word splitting and pathname expansion.
3. Implement error handling for when file reading or command chain replacement fails, or when an unknown item category is passed as an argument.
4. Implement checking for edge cases where the category argument is not passed at all.
5. The `ipxe_show_info` function is essentially storing, processing, and displaying sensitive information about a host. It is recommended to add necessary security

measures to protect this sensitive data from potential breaches.

6. Use `printf` instead of `echo` for better string handling, especially while displaying file contents.
7. Consider using local variables. This can help in preventing variable clashing and accidental modification of environment variables which can have impact on how the system functions.

11.6.543 `iscsi_targetcli_export`

Contained in `lib/host-scripts.d/common.d/iscsi-management.sh`

Function signature: `1febc38504cbf3c2f9e3c9454db03a4906bf3c248c462581bd42ab50e69d0296`

11.6.544 Function overview

The `iscsi_targetcli_export()` is a Bash function part of the iSCSI protocol used for block-level storage traffic management between server and client in a network. This function creates a block-based backstore, an iSCSI target, a Logical Unit Number (LUN), sets the iSCSI target parameters, creates the portal listening on the specified IP address and port, and saves the `getConfigureduration`.

11.6.545 Technical description

- Name: `iscsi_targetcli_export`
- Description: It executes a sequence of `targetcli` commands to create an iSCSI target with a block-based backstore and certain attributes. The target and portal creation can be conditional based on the `new_target` flag.
- Globals:
 - None
- Arguments:
 - \$1: The IQN (iSCSI Qualified Name) for the iSCSI target
 - \$2: The path to the zvol block device
 - \$3: Backstore name
 - \$4: IP address to bind for this iSCSI target
 - \$5: Port to bind for this iSCSI target
 - \$6: Flag to indicate new target (Value: 1) or existing target (Value: 0)
- Outputs: None
- Returns: None. However, if an error occurs within `targetcli` commands being executed, the error will be written to `stderr` by `targetcli`.
- Example usage:

```
iscsi_targetcli_export "iqn.2022-01.com.example:target1"  
↪  "/dev/zvol1" "backstore1" "192.168.1.10" "3260" "1"
```


11.6.546 Quality and security recommendations

1. Implement error checking and input validation: Currently, the function doesn't check if the provided arguments are valid. Input validation to ensure valid IQN, Zvol path, IP address, etc. can help mitigate errors and enhance the function reliability.
2. Sanitize all inputs: To prevent any command injection or arbitrary command execution, all inputs should be properly sanitized, especially when they're being directly used in command-string construction as in this function.
3. Handle `targetcli` command errors: The function doesn't handle any errors that may arise during the command execution. Error handling or exit codes from all `targetcli` commands can greatly improve the robustness of the function.
4. Logging: Implement extensive logging to capture all activities and errors. This is crucial for troubleshooting and system auditing.
5. Consider explicit return values: Right now there are no explicit return values. Providing explicit return values (return 0 upon success or a unique non-zero value upon each type of error) would make this function more reusable by other scripts/-functions.

11.6.547 `_is_tty`

Contained in `lib/functions.d/system-functions.sh`

Function signature: 53fbf6cc003bc7d9b4614fc9d372f3471ed01447874f4db99da2816eb3cb7e69

11.6.548 1. Function Overview

The Bash function `_is_tty` is designed to check whether the standard input (stdin), standard output (stdout) or the standard error (stderr) of the current terminal is attached to a tty (terminal).

11.6.549 2. Technical Description

- **Name:** `_is_tty`
- **Description:** The function checks if standard input (stdin), standard output (stdout), and standard error (stderr) are currently attached to a tty (terminal). This is achieved by using the `-t` test which returns true if the file descriptor is open and associated with a terminal.
- **Globals:** None
- **Arguments:** None
- **Outputs:** No explicit output. Internally the function returns with a status of 0 if any of the standard I/O streams are attached to a tty, or with a non-zero status otherwise.

- **Returns:** It returns true if at least one of stdin, stdout, or stderr is attached to a terminal. Otherwise, it returns false.
- **Example Usage:**

```
if _is_tty; then
    echo "We're in a tty terminal."
else
    echo "We're not in a tty terminal."
fi
```

11.6.550 3. Quality and Security Recommendations

1. Do not use this function if data privacy is your concern. In shared systems, other users can read or write to this terminal if they know its tty device file.
2. Always check the result of the function to handle the non-interactive environments appropriately.
3. It's highly recommended to handle the return status of this function properly to prevent faults in scripts that depend on a tty terminal.
4. Since the function has no arguments or global side-effects, it's safe to use this in any part of your scripts. But be aware that it only checks the state of the terminal at the time the function is called, not continuously.

11.6.551 `_is_tty`

Contained in `lib/functions.d/system-functions.sh`

Function signature: `53fbf6cc003bc7d9b4614fc9d372f3471ed01447874f4db99da2816eb3cb7e69`

11.6.552 Function Overview

The function `_is_tty()` is a utility function that checks if the standard input (stdin), standard output (stdout), or standard error (stderr) of the current process is connected to a terminal. The function uses the `-t` test to determine if each file descriptor (0, 1, and 2) corresponds to a terminal. It returns a boolean value, with `true` indicating that at least one file descriptor is connected to a terminal, and `false` indicating that none are.

11.6.553 Technical Description

Here is a block definition for pandoc:

- **Name:** `_is_tty`
- **Description:** Checks if standard input (stdin), standard output (stdout), or standard error (stderr) of the current process is connected to a terminal.
- **Globals:** None

- **Arguments:** None
- **Outputs:** None. However, the function will return `true` if at least one file descriptor (0,1,2) is connected to a terminal and `false` if they are not.
- **Returns:** 0 if at least one of `stdin/stdout/stderr` is connected to a terminal, 1 otherwise.
- **Example usage:**

```
if _is_tty; then
    echo "Running in a terminal"
else
    echo "Not running in a terminal"
fi
```

11.6.554 Quality and Security Recommendations

1. Add a function comment that summarizes what the function does. This makes it easier for others to understand the function's behavior without having to read through its code.
2. Secure the function against potential vulnerabilities or bugs by adding error handling and validation checks as necessary.
3. Audit the use of this function in larger scripts, especially those that handle sensitive information. The function checks whether input/output is connected to a terminal, which might be an important security consideration in certain contexts.
4. Continually review and update the function as necessary to reflect current best practices and accommodate changes in the system environment or the Bash shell itself.

11.6.555 `list_cluster_hosts`

Contained in `lib/functions.d/dns-dhcp-functions.sh`

Function signature: `9d0cf700c3819fc0712f3ee2664e742b93ed371801fdb6e35ba382d3f7fdc9bd`

11.6.556 Function overview

The Bash function `list_cluster_hosts()` is used to list the hosts in a cluster. It checks if the host configuration directory exists and is readable. The function then finds all configuration files inside that directory, extracts and normalizes MAC addresses for each file, and returns a list of these addresses. If an invalid MAC address is found, it logs a warning and skips to the next file.

11.6.557 Technical description

- **Name:** `list_cluster_hosts`

- **Description:** This function checks if the host configuration directory exists and is readable. Then, it scans through all configuration files in the directory, normalizes the MAC addresses found, and returns a list of these MAC addresses as a string. In case a configuration file does not exist or an invalid MAC address is found, a warning is logged and the function proceeds with the next file.
- **Globals:** `$HPS_HOST_CONFIG_DIR`: A path to the host configuration directory
- **Arguments:** None
- **Outputs:** A space-separated string of normalized MAC addresses.
- **Returns:** Function returns 0 if the function completes successfully, and 1 if the host configuration directory does not exist or is not readable.
- **Example usage:**

```
hosts=$(list_cluster_hosts)
echo "Host MAC addresses: $hosts"
```

11.6.558 Quality and security recommendations

1. **Error Handling:** Additional error handling could be integrated to capture and mitigate potential issues such as a missing key function like `normalise_mac`.
2. **Parameter Validation:** Apply strict validation rules on parameters like `HPS_HOST_CONFIG_DIR` to ensure they contain valid data before proceeding with the rest of the function.
3. **Security:** Make sure that sensitive data such as MAC addresses are handled securely and not exposed in logs or any other insecure means.
4. **Optimization:** The functionality of iterating over all configuration files and extracting MAC addresses could potentially be optimized to improve the performance of the function.
5. **Testing:** Extensive testing should be carried out to ensure the robustness of the function under various scenarios.

11.6.559 `list_clusters`

Contained in `lib/functions.d/cluster-functions.sh`

Function signature: `3f954969c054c0715b3cb4975cc16f60b59455fed57d76db287682a09738b747`

11.6.560 Function overview

The `list_clusters()` function in Bash is used to list all clusters. It gathers cluster directories using the `_collect_cluster_dirs clusters` function call and stores the directories in a local array. The function also attempts to determine the active cluster name, ignoring any potential errors in case it is not set. If no clusters are found, the function alerts the user and returns an exit status of 0. Finally, it iterates over the array of

clusters, and for each one, it checks if the cluster name matches the active cluster name and appropriately tags the active cluster in the output it echoes.

11.6.561 Technical description

- **name:** `list_clusters()`
- **description:** This bash function lists all clusters. It collects cluster directories, determines the active cluster (if any), and iterates over the cluster list to echo each one while highlighting the active cluster.
- **globals:**
 - `HPS_CLUSTER_CONFIG_BASE_DIR`: The base directory where cluster configurations are stored.
- **arguments:** None.
- **outputs:** Prints names of all clusters, marking the active one as '(Active)', if any.
- **returns:**
 - 0: When no clusters are found in the `HPS_CLUSTER_CONFIG_BASE_DIR`.
 - Other values could be returned if inner functions (`_collect_cluster_dirs` and `get_active_cluster_name`) return them.
- **example usage:** Simply run the function without any arguments as `list_clusters`.

11.6.562 Quality and security recommendations

1. Error messages should be handled adequately. For instance, when no clusters are found, the program could inform the user on what steps to take.
2. The function could return unique exit codes for each type of failure to make debugging easier.
3. Consider sanitizing any user inputs or outputs, to prevent potential security risks.
4. The function should handle the possibility of not being able to collect cluster directories gracefully.
5. To achieve better code readability, consider adding more comments to explain complex code segments.
6. While the function does a good job managing local scope variables, careful attention must be paid to global variable usage. It could potentially cause conflicts with other parts of the program.

11.6.563 `list_local_iso`

Contained in `lib/functions.d/iso-functions.sh`

Function signature: `2551c5d34fea622a4876a48c635391772020f951a812cc21a423cc6b87227a67`

11.6.564 Function Overview

The `list_local_iso` function in `bash` is designed to locate and list ISO files that match a specific naming pattern in a designated directory. The naming pattern is created based on the input arguments detailing the CPU type, manufacturer, operating system name, and optionally, the operating system version. If no matching ISO files are found, the function outputs an alert message and returns a status code of 1. If matches are found, the base name of each matching ISO file is outputted.

11.6.565 Technical Description

- Name: `list_local_iso`
- Description: The function searches for and lists ISO files in a directory that match a specified naming pattern, made up from the arguments for CPU type, manufacturer, operating system name, and potentially, the operating system version.
- Globals: None required.
- Arguments:
 - \$1: The CPU type descriptor.
 - \$2: The manufacturer descriptor.
 - \$3: The operating system name descriptor.
 - \$4: (Optional) The operating system version descriptor.
- Outputs: A status message indicating the search process for local ISOs and the names of any ISO files found that match the naming pattern. If no matches are found, an alert message is outputted.
- Returns: If no matching files are found, the function returns 1.
- Example Usage: `list_local_iso intel dell windows 10`

11.6.566 Quality and Security Recommendations

1. Conduct regular checks on the permissions assigned to the directory referenced in this function to ensure ISO files are not accessible to unauthorized parties.
2. Sanitize inputs to prevent potential code injections or file misdirections.
3. Implement error handling or exception management for situations where the directory does not exist or is not accessible.
4. Improve the clarity of existing function documentation especially with regards to its behavior when handling optional arguments.
5. Consider enhancement to allow for different patterns or multiple directories searching.

11.6.567 `load_cluster_host_type_profiles`

Contained in `lib/functions.d/cluster-functions.sh`

Function signature: `dc36e1e312e63f09fcc2eec91fb9e4bba6f60eb0ff592f0a1648a9a56bd34a80`

11.6.568 Function Overview

This bash function, `load_cluster_host_type_profiles()`, is designed to load the configuration files corresponding to the host types within a particular active cluster. The function operates by reading the key-value pairs from the active configuration file and storing them in an associative array for each declared host type.

11.6.569 Technical Description

- **Name:** `load_cluster_host_type_profiles`
- **Description:** This function reads an active cluster's configuration file and populates an associative array for each declared host type with key-value pairs from the file.
- **Globals:** [`__declared_types`: This global associative array stores the host type declarations.]
- **Arguments:** None
- **Outputs:** If no errors occur, it outputs nothing; however, to `stderr`, it'll output any error encountered during the process if the configuration file doesn't exist. It takes each key-value pair from the file and adds it to an associated array corresponding to the host type.
- **Returns:** If the configuration file does not exist, it returns 1; otherwise, it doesn't return a value.
- **Example Usage:** `load_cluster_host_type_profiles`

11.6.570 Quality and Security Recommendations

1. It is always a good idea to add more error checking. For example, this function will fail in silence if the configuration file does not have the right format. It could be improved with more feedback for incorrect format.
2. Regarding security aspects, it is suggested to add a security layer when opening and reading the file. An attacker could replace the content of the configuration file and this function would process it without any validation.
3. The function will process any key-value pair from configured hosts, including possibly non-expected pairs. It could be a good idea to validate the keys and values are expected and not a possible injection attempt.
4. It's always recommended to limit the global scope of variables. A better approach could be returning the processed data to the caller function instead of storing it in a global array. This would make the function safer and easier to maintain and debug.

11.6.571 `load_opensvc_conf`

Contained in `lib/host-scripts.d/common.d/opensvc-management.sh`

Function signature: `34b85a1047a61b533260c1512950ed5d8805b33c625908de3bcc4caf2d262537`

11.6.572 Function Overview

This function, `load_opensvc_conf()`, is mainly used to fetch, validate, and load the OpenSVC configuration from a specified gateway. The function performs the following actions:

1. Resolves a provisioning node.
2. Creates a directory for the configuration file if it does not already exist.
3. Downloads and temporarily stores the configuration file from the gateway.
4. Checks the integrity of the fetched configuration file.
5. If the fetched file is unchanged, it discards the temporary file and logs a message stating that there is no need to restart.
6. If the file has changed, it makes a backup of the current configuration file, replaces it with the new one, and restarts the daemon.

11.6.573 Technical Description

- Function: `load_opensvc_conf()`
- Description: The function fetches, checks, and loads a new opensvc conf file from a specified gateway. If the file is new or modified, it backs up the old file, installs the new one, and restarts necessary services.
- Globals: `conf_dir`: The directory where the OpenSVC configuration file is stored, `conf_file`: The OpenSVC configuration file.
- Arguments: None.
- Outputs: Logs a message regarding the process completion status and potential error messages to the corresponding system log.
- Returns: 0 if the function successfully completes all its tasks and 1 if any of the tasks encounter an issue that causes it to terminate.
- Example usage: This function is not intended to be invoked with any command-line arguments and thus its usage is simply: `load_opensvc_conf`.

11.6.574 Quality and Security Recommendations

1. Always ensure the validity of the configuration file source (the “gateway”) by using trusted certificates.
2. Make sure that the user running the script has proper permissions for all the necessary operations, like creating directories, fetching data, and restarting services.
3. It would be beneficial to implement more sophisticated error handling for the various stages of the function.
4. Strengthen the validation checks for the fetched configuration file.
5. When running the `curl` command, use the `-S` or `--show-error` flag to ensure error messages are displayed if an error occurs.
6. Use the `-s` or `--silent` with `curl` to not show progress meter but still show error messages.

7. To protect the original file, make a backup copy before making changes.
8. Encrypt sensitive parts of the configuration file to ensure security.
9. Use detailed log entries for any operation in order to trace back in case of any issue.
10. Always check the return status of each command, rather than assuming it will succeed. This will help to capture and handle any errors.

11.6.575 load_remote_host_config

Contained in `lib/host-scripts.d/common.d/common.sh`

Function signature: `7a3ad554ca390fc6c05a7c79d9c40323f388eca53f8e402fc2f7f2adaf358fc0`

11.6.576 Function overview

The `load_remote_host_config` function is used to obtain configuration data for a host from a remote source via HTTP request. This data is then evaluated and applied to the host.

11.6.577 Technical description

- **Name:** `load_remote_host_config`
- **Description:** This function establishes a connection with a remote host and fetches configuration details required for the host. It uses 'curl' to send an HTTP GET request to the remote host, and 'eval' to evaluate and apply the configuration data that is received from the host.
- **Globals:** None
- **Arguments:** None
- **Outputs:** If the function cannot fetch or load the host configuration, it outputs a log message "Failed to load host config". If debug is enabled, it outputs the fetched configuration with the log message "Remote config: \$conf".
- **Returns:** This function will return 1 if there is any error in fetching the configuration.
- **Example usage:**

```
load_remote_host_config
```

11.6.577.1 Code breakdown:

- `local conf` and `local gateway="$(get_provisioning_node)"` are used to declare and initialize local variables.
- `curl -fsSL "http://${gateway}/cgi-bin/boot_manager.sh?cmd=host_get_config"` is used to fetch the configuration information from the remote host.

- `remote_log "Failed to load host config"` logs the error message if there is any issue in getting the configuration.
- `remote_log "Remote config: $conf"` outputs the received configuration for debugging purposes.
- `eval "$conf"` is used to apply the fetched configuration to the host.

11.6.578 Quality and security recommendations

1. Use HTTPS instead of HTTP for the curl request to ensure that the data transmission is secure.
2. Verify the integrity of the received configuration data before evaluating it with eval command. Untrusted data can lead to code execution or security breaches.
3. Add error handling for 'get_provisioning_node' function as the function depends on it.
4. The curl command should have timeouts set to prevent the script from hanging indefinitely in case the remote server does not respond.
5. The function should return different error codes for different types of errors - like failure in getting provisioning node or failure in fetching the configuration. This can help in better troubleshooting.
6. Logging should be improved to detail what type of error occurred - whether it is related to network connectivity, server response, etc. for better issue tracking.

11.6.579 _log

Contained in `lib/host-scripts.d/common.d/zpool-management.sh`

Function signature: `f7f5d16961fc9339682891b49f2fefad3611a12964b1a3cd095d11a46c108959`

11.6.580 Function overview

This Bash function, `_log()`, works as a logging subroutine that outputs a given string to both the console and a remotely configured server. It belongs within a larger script, and its specific use case relates to logging messages within a ZFS pool creation workflow.

11.6.581 Technical description

- **Name:** `_log`
- **Description:** This function performs logging operations to both a remote log and the system console (if `LOG_ECHO` environment variable is set to 1). It is specifically designed for logging operations regarding the creation of ZFS pools on free disks.
- **Globals:** [`LOG_ECHO`: An environmental variable used to control if the logging output should also be printed to console.]

- **Arguments:** [\$*: A variable-length list of arguments to be logged. These arguments are usually messages related to the operations performed in the ZFS pool creation process.]
- **Outputs:** Printed statements are sent to the console (if LOG_ECHO is set to 1) and the `remote_log` function.
- **Returns:** Nothing directly, but it outputs strings to the console and the remote logging service.
- **Example Usage:** `LOG_ECHO=1 _log "ZFS pool created successfully"`

11.6.582 Quality and security recommendations

1. Always sanitize the inputs before logging to prevent log injection attacks.
2. Ensure that the `remote_log` function correctly handles connection failures and other errors to prevent disruption of the main program.
3. Consider adding timestamp and log level (info, warning, error, etc.) information to the log messages to provide better logging context.
4. Regularly rotate and archive logs to prevent them from occupying too much disk space.
5. Encrypt sensitive data in logs to protect them from being exposed to unauthorized users.
6. Consider implementing rate limiting to prevent DOS attacks via rapid, repeated calls to the `_log` function.

11.6.583 log

Contained in `lib/host-scripts.d/common.d/common.sh`

Function signature: `8036c583d5ea28bb7d9be8b86179593acefd5b00fca118ffa37907a177f20aec`

11.6.584 Function Overview

The `log` function is designed to output and forward log messages. It prepends the string “HPS” and current time in the “Hour:Minute:Second” format to the log message, then writes the complete string to the standard output. The function also sends the log message to a remote log function called `remote_log`.

11.6.585 Technical Description

- **Name:** `log`

- **Description:** This function takes any number of arguments, appends them after a prepended string which contains “HPS” and the current time. The final string is then echoed (printed to stdout). It also sends every log message to `remote_log` which is assumed to be another logging function.
- **Globals:** No global variables are used or modified by this function.
- **Arguments:** Any set of strings which need to be logged. Example: `$*: log message`.
- **Outputs:** Outputs to STDOUT. Example: `[HPS:14:20:35] your log message`.
- **Returns:** It does not return any value but calls another function called `remote_log`.
- **Example usage:**

```
log "Application started."
```

This will print:

```
[HPS:14:20:35] Application started.
```

and will also send “Application started.” as a log message to the remote log.

11.6.586 Quality and Security Recommendations

1. **Secure transfer for remote logging:** The function `remote_log` which is called within this function should ensure secure transmission of log messages, especially if these logs are being sent over a network.
2. **Error Handling:** The `remote_log` function is called without any error handling. If it fails for any reason, there won't be any fallback or even notification given to the user. This needs to be ensured for better quality.
3. **Detailed TimeStamp:** The timestamp is currently pretty simple just HH:MM:SS. For better track of logs, it could be enriched to contain more details like the date, timezone, milliseconds etc.
4. **Input Validation:** This function accepts any string without validation. Although this is generally okay for a logging function, it could be an issue if certain types of strings could cause problems with the `remote_log` function. Consider validating inputs or sanitizing them if necessary.

11.6.587 make_timestamp

Contained in `lib/functions.d/system-functions.sh`

Function signature: `7ea1fc9d3621ad0a04879323d75cd0a5f1aa2468bf98b9b3c83ff2b66dfa8e3d`

11.6.588 Function Overview

The function `make_timestamp` is a simple Bash function that returns the current date and time in a specific format (Year-Month-Day Hour:Minute:Second UTC). The `-u` option

makes sure that the command uses Coordinated Universal Time (UTC) instead of the local timezone.

11.6.589 Technical Description

- **Name:** `make_timestamp`
- **Description:** This function uses the `date` command with the `-u` option to get the current date and time in UTC, formatted as: Year-Month-Day Hour:Minute:Second UTC.
- **Globals:** None.
- **Arguments:** None.
- **Outputs:** Prints current date and time formatted as Year-Month-Day Hour:Minute:Second UTC.
- **Returns:** Nothing.
- **Example usage:**

```
$ make_timestamp
2023-01-01 00:00:00 UTC
```

11.6.590 Quality and Security Recommendations

1. This function has no user-input, hence it should be safe from security flaws that could result from unreliable user input.
2. Since the `date` command is a common Unix command, it should be available in most environments. However, in case the environment does not support the `date` command, appropriate error handling could be added.
3. For quality, consider adding checks if UTC timezone is required or if local timezone would suffice. If different timezones are needed, consider making the timezone an optional input to the function.
4. If the function will be used as part of larger scripts, consider returning the value instead of printing to allow better usage of this function.

11.6.591 `make_timestamp`

Contained in `lib/functions.d/system-functions.sh`

Function signature: `7ea1fc9d3621ad0a04879323d75cd0a5f1aa2468bf98b9b3c83ff2b66dfa8e3d`

11.6.592 Function Overview

The `make_timestamp()` function is a Bash shell function that generates the current timestamp in the universal coordination time (UTC). This function returns a for-

matted string representing the current date and time formatted in the following way:
YYYY-MM-DD HH:MM:SS UTC.

11.6.593 Technical Description

- name: `make_timestamp`
- description: Generates a current date and time string in the format YYYY-MM-DD HH:MM:SS UTC.
- globals: None
- arguments: None
- outputs: A string representing the date and time (YYYY-MM-DD HH:MM:SS UTC).
- returns: 0 on successful execution, non-zero error code from the `date` command if there's any error.
- example usage:

```
timestamp=$(make_timestamp)
echo "The current UTC time is $timestamp"
```

11.6.594 Quality and Security Recommendations

1. Error Handling: Currently, there is no error handling mechanism. Consider adding an error handling routine to catch and manage any potential errors from using the `date` command.
2. Return consistency: Currently, if the `date` command fails, it returns a different value. Ensuring the function always return a consistent output or an explicit error value can significantly improve usability.
3. Commenting: To improve maintainability, consider adding brief comments in the code describing what the function is doing.

11.6.595 `mount_distro_iso`

Contained in `lib/functions.d/iso-functions.sh`

Function signature: `9ee707a09f131340e40ead1764695e3c9db4201a8c595ad06c9a5f7028376217`

11.6.596 Function Overview

The `mount_distro_iso` function is used to mount an ISO file of a Linux distribution. The function takes a string as an argument that represents the name of the Linux distribution. It checks for the presence of the ISO file in a predefined directory. If the ISO file exists, the function then proceeds to check if the ISO file is already mounted. If not, the function creates a mount point directory and mounts the ISO file to it.

11.6.597 Technical Description

name: `mount_distro_iso` **description:** This function is used to mount an ISO file of a Linux distribution to a predefined directory based on the passed distribution string. **globals:** `HPS_DISTROS_DIR`: This is a reference to the directory where the distribution ISO files are stored. **arguments:** `$1` – `DISTRO_STRING`: The name of the Linux distribution, `$2` – `iso_path`: The full path to the Linux distribution ISO file **outputs:** Logs information about the status of the ISO file and its mount point. **returns:** Returns 1 if the ISO file is not found, or if the mount point already exists, else it does not return any value. **example usage:** `mount_distro_iso ubuntu`

11.6.598 Quality and Security Recommendations

1. Incorporate comprehensive error handling and validation checks. This will help in ensuring that the mounted ISO file is legitimate and not corrupted.
2. Make sure to properly sanitize the `DISTRO_STRING` input. This will protect against potential code injection risks.
3. Document the expected values for each input variable. This would help prevent errors when the function is used by other team members.
4. Use a more descriptive name for the `hps_log` function. This would make the function more understandable to other developers who might read the code.
5. Write unit tests for this function to ensure it behaves as expected under different scenarios.

11.6.599 `node_get_functions`

Contained in `lib/functions.d/configure-remote-host.sh`

Function signature: `1b666a5e2b77b4bf76b075e77c70af41c70e3996b2c4adb712d2c69e9dbccba1`

11.6.600 Function overview

The `node_get_functions` is a bash function in charge of building a function bundle for a specific Linux distribution (or 'distro'). The function accepts a string that represents a 'distro' and a base directory as parameters. The Distro string is a concatenation of the CPU architecture, the manufacturer, the OS name, and the OS version, divided by dashes. A function bundle is a collected set of bash functions that are relevant and compatible with the specific system represented by the distro string.

11.6.601 Technical description

- **name:** `node_get_functions`

- **description:** This function builds a function bundle for a specified Linux distribution based on the specific CPU, manufacturer, OS name and OS version. It finds the relevant functions from a provided directory base and echoes them out.
- **globals:**
 - [LIB_DIR: an optional global variable that might contain the base directory]
- **arguments:**
 - [\$1: a ‘distro’ string in the format:]
 - [2 : *Optional.isabasedirectorywherefunctionfilesaresearched.Defaultsto"{LIB_DIR}/host-scripts.d” or “/srv/hps-system/lib/host-scripts.d” if LIB_DIR is not presented*]
- **outputs:** Outputs the relevant functions found in the provided base directory.
- **returns:** It doesn’t explicitly return anything. The function has side-effects of echoing the functions it finds and logs info and debug messages.
- **example usage:**

```
node_get_functions intel-dell-ubuntu-20.04 /my/base/dir
```

11.6.602 Quality and security recommendations

1. Always use double quotes around variable expansions to avoid word splitting and pathname expansion.
2. Make sure the provided base directory path is validated and sanitized.
3. Consider adding error messages in case the incorrect number of arguments have been provided.
4. Always validate user inputs. In this function, consider applying validation checks on the distro string and possibly the base directory.
5. If possible, add some unit tests around this function. It’ll help to ensure the function behaves as expected when modifying or adding new code.

11.6.603 node_lio_create

Contained in lib/host-scripts.d/common.d/iscsi-management.sh

Function signature: 98eefecb076b12ce8574e70fbcdc3dbe7da279292f2b27fad96ad89b969a1f6e

11.6.604 Function overview

The node_lio_create function sets up an iSCSI target node using the Linux IO (LIO) Target-core storage abstraction layer. It does so by taking three optional arguments for iqn, device, and acl. The function validates and checks the existence of required arguments (iqn and device), extracts the backstore name from iqn, creates the backstore, the iSCSI target and the LUN, configures the ACL if provided and saves the configuration.

11.6.605 Technical description

- **name:** node_lio_create

- **description:** A bash function to set up an iSCSI target node using the Linux IO (LIO) Target-core storage abstraction layer via targetcli.
- **globals:** None
- **arguments:**
 - \$1: iqn – the iSCSI Qualified Name, an unique identifier for an iSCSI node.
 - \$2: device – the device to be used for the iSCSI target.
 - \$3: acl – Access Control List to control which initiators are granted access.
- **outputs:** Logs messages regarding process status (e.g., success/failed to create backstore/iSCSI target/etc.)
- **returns:**
 - 0: if the iSCSI target has been successfully created.
 - 1: if a failure occurred (unknown parameter, missing required parameters, non-existing device, failure to create backstore/iSCSI target, etc.)
- **example usage:** `node_lio_create --iqn iqn.2003-01.com --device /dev/sdb --acl iqn.1994-05.com.redhat:dhclient`

11.6.606 Quality and security recommendations

1. The function should sanitise all user input data to prevent potential command injection attacks or faulty operations.
2. The function could implement some logging mechanism, preserving a history of operations for further analysis or debugging.
3. Improve error messages by stating exactly which parameter(s) is/are missing instead of the generic message
4. The function could handle different input formats and conversions. For instance, device input could be accepted as device UUID instead of device path only.
5. Including more comprehensive tests for edge cases or exceptional circumstances, such as when the system lacks the necessary resources to create a new target.
6. Encryption should be considered for the transmitted data as this function disables authentication for demo/testing purposes. This feature should be taken out of production versions or controlled by an additional variable.

11.6.607 node_lio_delete

Contained in `lib/host-scripts.d/common.d/iscsi-management.sh`

Function signature: `198653d7f8219c254d0f7d905c865d8994dd5dc9072af9cdbc979296660c9662`

11.6.608 Function Overview

The function `node_lio_delete()` is used to delete an iSCSI target and associated backstore in a Linux system. It accepts an iSCSI Qualified Name (IQN) as an argument. This function safeguards by validating the provided argument, checks if the iSCSI target

exists before attempting deletion, and similarly for the associated backstore. Post-deletion task involves saving the configuration.

11.6.609 Technical Description

Function: `node_lio_delete`

Description: Deletes an iSCSI target and its associated backstore.

Globals: None

Arguments:

- `$1`: IQN (iSCSI Qualified Name) of the target to be deleted.

Outputs: Logs information about the execution process, which

- ↪ includes successful or unsuccessful deletion of the iSCSI
- ↪ target and the backstore.

Returns:

- 0 on successful deletion of the target and backstore, or if
↪ they do not exist.
- 1 if the parameter is unknown or missing, or if deletion of
↪ the target or backstore is unsuccessful.

Example Usage:

```
node_lio_delete --iqn iqn.2003-01.org.linux-  
↪ iscsi.localhost.x8664:sn.4afce5632cfd
```

11.6.610 Quality and Security Recommendations

1. Add a mechanism to validate the structure of the IQN argument. Validating the IQN increases security by preventing injection or incorrect usage.
2. Add error handling for the `targetcli` execution when executing commands like `targetcli /iscsi delete "${iqn}"`. Execution success should not be implicitly assumed.
3. Incorporate additional logging for more insight into function execution process.
4. Include a mechanism to backup current configuration before making changes. Recovery from failure scenarios can be quicker.

11.6.611 `node_lio_list`

Contained in `lib/host-scripts.d/common.d/iscsi-management.sh`

Function signature: `f36324bce9bac497425dd3ef1f0cb414f4b7fe0e640904c41e749f78c8e29a70`

11.6.612 Function overview

The `node_lio_list` is a Bash function that lists out iSCSI targets and block backstores from a system. The function leverages the `targetcli` command for pulling and organizing the required information, thus providing an easy and direct interface for iSCSI and backstore management.

11.6.613 Technical description

Name: `node_lio_list`

Description: The function retrieves and displays the iSCSI targets and block backstores present in the system by using the `targetcli` command. It showcases the results one after the other for ease of visibility and understanding.

Globals: None

Arguments: None

Outputs: The function prints two lists. The first list displays the iSCSI Targets. This segment of the program executes the `targetcli` command to get the iSCSI targets. After a line break, the second list displays the block backstores also using the `targetcli` command.

Returns: It returns 0 after executing the echo commands to signify that the function executed successfully.

Example Usage:

```
node_lio_list
```

11.6.614 Quality and security recommendations

1. The function is relatively simple and does not handle errors. Users should be aware that any errors occurring in the `targetcli` command will break the function. Proper error handling should be considered to ensure resilience of the function.
2. The use of global variables could be considered for more customized output.
3. While not a problem in this specific function, keep in mind that usage of echo to output data can lead to potential command injection if user-provided data is not properly sanitized.
4. Always review your Bash scripts for potential security vulnerabilities. This could include examining how it handles input, reviewing the script for potential command injection vulnerabilities, and checking the script for insecure file operations.
5. Use helper functions to make the code reusable and modular. The function should do one thing and do it well. This function is a good example of that.

11.6.615 `node_lio_manage`

Contained in `lib/host-scripts.d/common.d/iscsi-management.sh`

Function signature: `86138e13e14957c8703d9fb6016f95dbd6eec12c63716fc419c9afcb1c48e6ca`

11.6.616 Function Overview

The `node_lio_manage()` function is a controller function in a node's local I/O (LIO) management system. Using a case statement, it dispatches the given 'action' to the

appropriate functions. Actions include ‘create’, ‘delete’, ‘start’, ‘stop’, ‘status’, and ‘list’. If the input ‘action’ is not recognized, the function logs an error message and then returns with a value of 1.

11.6.617 Technical Description

- **name:** `node_lio_manage()`
- **description:** This function serves as a dispatcher for different actions of a node’s local I/O (LIO) management system. The designated ‘action’ argument is checked against a predefined list (‘create’, ‘delete’, ‘start’, ‘stop’, ‘status’, ‘list’), and upon match, the corresponding function is invoked.
- **globals:** None.
- **arguments:**
 - \$1: Action to be performed on a node’s LIO. Valid: ‘create’, ‘delete’, ‘start’, ‘stop’, ‘status’, ‘list’.
 - \$@: Options that may be provided following the ‘action’, to be passed on to the corresponding function.
- **outputs:** Logs a usage message if no ‘action’ is provided, or if the ‘action’ given is not recognized. Also dispatches tasks to various functions which may have their own outputs.
- **returns:** 1 for invalid ‘action’ or missing argument; the exit status of the last executed command otherwise (\$?).
- **example usage:**
 - `node_lio_manage create`
 - `node_lio_manage delete`

11.6.618 Quality and Security Recommendations

1. Incorporate more stringent input validation (e.g., checking the format and/or value of the ‘action’, or other argument inputs, prior to usage).
2. Use more descriptive error messages and create separate logs for errors and usage to make the system easier to debug.
3. Implement specific security measures to protect against command substitution or code injection attacks.
4. Make sure that the user running the script has the necessary permissions for all functions that this one is dispatching to, without granting more permissions than necessary.
5. Test the function in various scenarios to ensure it behaves as expected, including cases of erroneous input.

11.6.619 `node_lio_start`

Contained in `lib/host-scripts.d/common.d/iscsi-management.sh`

Function signature: 618a6b47c51365d7704455f712c3beca1d289fd4105a628552ebf67c3fbc5573

11.6.620 Function Overview

The Bash function `node_lio_start` is designed to start a “target service” in a Linux environment, using `systemd`’s `systemctl` command. It does this through calling `systemctl start target`, which starts the target service, and `systemctl enable target`, which ensures that the service will be started at boot time. Log messages are sent to a remote server via a function called `remote_log` before starting the service and after either successfully starting the service or failing to do so. A success or failure status is indicated by returning 0 or 1 respectively.

11.6.621 Technical Description

Below is a technical description of the `node_lio_start` function:

- **name:** `node_lio_start`
- **description:** A Bash function designed to start a “target service” using `systemd`’s `systemctl` command. Also ensures said service will be started at boot, and records the result (either success or failure) by sending a log message to a remote server.
- **globals:** None.
- **arguments:** None.
- **outputs:** Logs messages to a remote server indicating the process of starting and enabling the service.
- **returns:**
 - 0 if the service is successfully started and enabled.
 - 1 if either starting or enabling the service failed.
- **example usage:**
`node_lio_start`

11.6.622 Quality and Security Recommendations

1. Always ensure the `remote_log` function is properly defined in the same shell environment where `node_lio_start` is called to avoid errors.
2. Implement error catching for case when `systemctl` command is unknown, i.e., `systemd` isn’t in use.
3. Safeguard this function with user and permission checks to ensure only authorized personnel or scripts can initiate `systemd` services.
4. Incorporate status checks before starting the service to prevent launching an already running service.

5. Add more logging for better troubleshooting, such as logging the status code of the `systemctl` command.

11.6.623 `node_lio_status`

Contained in `lib/host-scripts.d/common.d/iscsi-management.sh`

Function signature: `327a33e59e304acd75f9e12b0cd6fa4aca564708c04f3615da39246c5573ba39`

11.6.624 Function overview

The function `node_lio_status()` is designed for the purpose of querying the status of a Target Service and displaying its LIO configuration. It uses Linux's `systemctl` and `targetcli` commands to retrieve and demonstrate this data.

11.6.625 Technical description

- **Name:** `node_lio_status`
- **Description:** The function echoes a status report for a target service and its LIO configuration utilizing the `systemctl` and `targetcli` commands respectively. The function doesn't take any parameters or global variables.
- **Globals:** None
- **Arguments:** None
- **Outputs:** The target service status report and LIO configuration data are printed to stdout.
- **Returns:** Always returns 0 signifying that the function has successfully completed its task.
- **Example usage:** `bash node_lio_status`

11.6.626 Quality and security recommendations

1. It would be beneficial to make sure that the required commands (`systemctl` and `targetcli`) are available in the system before actually utilizing them in the function.
2. The command outputs might require interpretations which can be hieroglyphic or error-prone to a non-technical user. It might be advantageous to include added verbosity/reasoning to the echoed content for the sake of usability.
3. Remember to check the return values of the commands and handle any errors that might occur. Currently, the function always returns 0 irrespective of whether the commands run successfully or not.
4. It's good to utf-8 sanitize the output for proper display and to prevent possible injection attacks. This is especially crucial if the output is intended to be utilized or displayed elsewhere.

5. Always validate that the function is being run with the appropriate permissions since commands like `systemctl` and `targetcli` often require superuser rights.

11.6.627 node_lio_stop

Contained in `lib/host-scripts.d/common.d/iscsi-management.sh`

Function signature: `971a38bb22f277cd0531cfa279368d8f7e4995e9a385bb6221485f8a90cc3bcd`

11.6.628 Function Overview

The function `node_lio_stop` is a bash script function designed to stop a specific service, named `target`, running on a system. It logs its activity using another function `remote_log` and uses the `systemctl` command to stop the service. If successfully stopped, it generates a success message and returns 0. If the attempt to stop the service fails, it generates a failure message and returns 1.

11.6.629 Technical Description

- **Name:** `node_lio_stop`
- **Description:** This function attempts to stop the ‘target’ service on a system. It employs the ‘`systemctl`’ command to accomplish this, but also provides logging functionality through the `remote_log` function for tracking activity.
- **Globals:** None.
- **Arguments:** None.
- **Outputs:** This function generates log messages indicating the status of the stop operation - either successful or failed.
- **Returns:** This function returns a binary response. It returns 0 if the stop operation is successful and 1 if it fails.
- **Example Usage:**

`node_lio_stop`

11.6.630 Quality and Security Recommendations

1. Consider implementing error logging or notifications, such as sending an alert email, when the service fails to stop.
2. For enhanced security, you could use the `sudo` command to enforce the function to run with root permissions only, thus restricting the usage of the function to authorized users only.
3. Include verbose commenting, particularly for areas handling failures or errors, which can help for easier troubleshooting in future.
4. Consider allowing the function to accept a parameter to specify the service to stop, instead of hardcoding ‘target’. This would make it more flexible and reusable.

5. Implement a check to understand whether the ‘target’ service is active before attempting to stop it.

11.6.631 `node_storage_manager`

Contained in `lib/host-scripts.d/common.d/common.sh`

Function signature: `739f69b5a6036d980bc51405a34e0ec8a34bed5236ca2984db5c0d656a80c5c6`

11.6.632 Function overview

The `node_storage_manager` function manages storage components on a node. By using two main arguments, a component and an action, it decides which storage management function to call. The user can also optionally pass further arguments for specific storage management tasks. Component options include “lio” and “zvol,” corresponding to different subsystems in the node’s storage system. Actions can include operations like ‘start’, ‘stop’, ‘status’, etc. The function validates the arguments, dispatches to the appropriate function based on the given component, logs the execution and handles any errors.

11.6.633 Technical description

`node_storage_manager`

- **Description:** This function manages storage components for a machine node. It logs an error and returns 1 if any of the main arguments is missing. The function then checks the value of the ‘component’ argument and calls the corresponding function for ‘lio’ or ‘zvol’. If the ‘component’ argument doesn’t match any of these, it logs an error and returns 1. The function then captures the exit status of the called function, logs a respective message and returns the result.
- **Globals:** None
- **Arguments:**
 - \$1: The component to manage. Can either be ‘lio’ or ‘zvol’.
 - \$2: The action to execute on the given component.
 - . . . : Optional arguments passed to the dispatched function.
- **Outputs:** Logs messages to the remote log regarding the action being performed, its success or failure, and its return code.
- **Returns:** 0 if the dispatched function was successful, 1 if the arguments are invalid or the dispatched function failed.
- **Example usage:** `node_storage_manager lio start`

11.6.634 Quality and security recommendations

1. Refrain from passing sensitive data as arguments or using them in log messages.
2. Add more detailed error logging to help with troubleshooting potential errors.

3. Add input sanitization to the function to prevent possible shell injection attacks.
4. Consider adding more explicit validation logic for argument values to catch common user input errors.
5. Be sure that called functions 'node_lio_manage' and 'node_zvol_manage' are implemented securely and robustly.
6. Protect all logging data, especially if it contains sensitive information.
7. Use specific return codes for different errors to help users better understand why the function failed.
8. Document any modifications made to the function for clarity.

11.6.635 node_zvol_check

Contained in `lib/host-scripts.d/common.d/zvol-management.sh`

Function signature: 07447af71573f33e1e9896011244fa50252bca03c45d1685a0b8677a4720d46c

11.6.636 Function Overview

The `node_zvol_check` function is a part of the command-line tool intended for use in ZFS (Zettabyte file system). The function checks if a specific ZFS volume (zvol) exists, where the zvol is defined by its pool and name. These two parameters are passed to the function as arguments. The function prints the outcome of the check to the console and returns an error code.

11.6.637 Technical Description

- **Name:** `node_zvol_check`
- **Description:** This function checks the existence of a ZFS volume (zvol) where the zvol is identified by its ZFS pool and name.
- **Globals:** No global variables
- **Arguments:**
 - \$1: Parameter pair `--pool <pool>`.
 - \$2: Parameter pair `--name <name>`.
- **Outputs:**
 - If the zvol exists, it prints: `Zvol <zvol_path> exists.`
 - If the zvol does not exist, it prints: `Zvol <zvol_path> does not exist.`
- **Returns:** 0 if the zvol exists; 1 otherwise, or if the required arguments were not correctly passed.
- **Example usage:** `node_zvol_check --pool myPool --name myZvol`

11.6.638 Quality and Security Recommendations

1. Implement parameter validation, ensuring the input parameters meet the expected format.
2. Add more explicit error codes for different error conditions to aid in troubleshooting.
3. Implement a help message function when parameters are not provided or incorrectly supplied.
4. Consider using more secure methods when processing the output and display messages. This could help against possible Bash Injection attacks.

11.6.639 `node_zvol_create`

Contained in `lib/host-scripts.d/common.d/zvol-management.sh`

Function signature: `ec6e430fb8789091cfe8fdd4b1e197a7a20d2619e9cb9986827d740d74268de1`

11.6.640 Function Overview

The Bash function `node_zvol_create()` is designed to create a ZFS (Zettabyte File System) volume, or zvol. This is done through the use of environment parameters, namely a pool, name and size, which are parsed and passed as variables to the system's `zfs` command. The function validates the provided parameters and checks for the existence of the zvol before attempting to create a new one.

11.6.641 Technical Description

- **Name:** `node_zvol_create`
- **Description:** This function creates a ZFS volume using given parameters. The parameters are parsed for three variables: pool, name, and size. The function then uses these variables in the `zfs create` command to generate the volume, first checking to ensure that it does not already exist.
- **Globals:** None
- **Arguments:**
 - `$1`, `$2`: These are used by the command to parse the given parameters. Each shift command removes the current value of `$1` and assigns the next parameter value to it.
- **Outputs:** conditional logging to a remote destination detailing the function's activity and success or failure to create a volume.
- **Returns:**
 - 0: Success - zvol was created without issue.
 - 1: Failure - Either due to an unrecognized parameter, missing required parameters, or pre-existing zvol.
- **Example Usage:**

```
node_zvol_create --pool tank --name vol1 --size 1G
```

11.6.642 Quality and Security Recommendations

1. Enforce stricter validation for the ‘pool’, ‘name’, and ‘size’ parameters to prevent any potential command injection attacks or inconsistencies.
2. Implement robust error handling to make the function more resilient to unexpected behavior and to give more informative responses when something goes wrong.
3. Use unique and clear logging statements to make the progress of the function easier to follow and troubleshoot.
4. Store sensitive data securely to ensure this data isn’t exposed to unauthorized access.
5. For improved auditability, incorporate logging for all major events such as input validation failures, zvol creation failure, and zvol creation success.
6. Find ways to limit the use of global variables to lessen the chances of unwanted side effects from their modification.
7. Whenever possible, encapsulate complex sections of code into their own functions to improve readability and maintainability.

11.6.643 node_zvol_delete

Contained in `lib/host-scripts.d/common.d/zvol-management.sh`

Function signature: `21fe8b803dc5e01613ebe91d459c8ff30011593e6a87412bc6122f867aeec227`

11.6.644 Function overview

The `node_zvol_delete` function is designed to handle the deletion of ZFS volumes (`zvol`). The function accepts two parameters: the name of the ZFS storage pool and the name of the specific volume. It first checks whether these parameters have been provided or not. If they are missing, it will return an error message and exit. If they are provided, it will construct the path to the ZFS volume and then check whether that volume exists. If it doesn’t, it will return an error message and exit. If it does, it will attempt to delete the volume and return a success message if successful or an error message if not.

11.6.645 Technical description

The following properties describe the function in a technical manner:

- **Name:** `node_zvol_delete`
- **Description:** Deletes a node’s ZFS volume
- **Globals:** None
- **Inputs:**

- `--pool`: The pool parameter (a type of ZFS storage)
- `--name`: The name of the volume to be deleted
- **Outputs:** Messages indicating the process of deletion and its success or failure
- **Returns:**
 - 0 if the volume is successfully deleted
 - 1 if there's an error (either because of missing parameters or failure in deletion)
- **Example usage:**

```
node_zvol_delete --pool mypool --name myvolume
```

11.6.646 Quality and security recommendations

1. Add more in-depth input validation for both `--pool` and `--name` parameters to ensure they conform to expected formats or value ranges.
2. Implement error handling that doesn't merely log and return but also performs some sort of troubleshooting or recovery operation.
3. Consider adding an optional force-delete parameter to allow bypassing certain checks when necessary.
4. Ensure that all output messages, especially error messages, do not leak sensitive system state information.
5. Implement strict access control measures to protect against unauthorized zfs pool and volume deletion.

11.6.647 `node_zvol_info`

Contained in `lib/host-scripts.d/common.d/zvol-management.sh`

Function signature: `80d5ff3413c64a25ffdb5394756a3fcb79290986c8d4a3bfd18cdf837a95a543`

11.6.648 Function Overview

This function `node_zvol_info()` is a bash function primarily used to fetch and display ZFS volume (zvol) information from a specified pool. It accepts themed arguments, `--pool` and `--name`, to determine the pool and the name of the volume, respectively. Once the parameters are provided, it validates their existence, then proceeds to check if the zvol exists in the system. Should all these checks pass, it displays the zvol information—its name, volume size, usage, available space, and reference association—along with its device path.

11.6.649 Technical Description

- **Name:** `node_zvol_info`
- **Description:** This function checks specific zvol in the provided zfs pool and name, returns its existence status and displays the zvol information if exists.

- **Globals:** Remote_log:logs for the remote server.
- **Arguments:**
 - \$1: --pool: The name of the ZFS storage pool that contains the ZFS volumes.
 - \$2: --name: The name of the volume to be checked in the pool.
- **Outputs:** The function prints zvol information to stdout, which includes name, volume size, usage, available space, and references; additionally, it provides the exact device path.
- **Returns:** It returns 1 if there's a failure (invalid/unspecified parameters, or non-existent pool/volume). It returns 0 if it successfully found the zvol and displayed its status.
- **Example Usage:**

```
node_zvol_info --pool tank --name volume1
```

11.6.650 Quality and Security Recommendations

1. Implement more robust error handling. This could include more specific error messages depending on the type of error encountered.
2. Use secure methods when handling the arguments to prevent potential command injection attacks. Escaping or quoting variables can help in this regard.
3. Enhance the logging process. Log all the errors to a specific error log file with timestamps. This will be helpful for debugging purposes if anything goes wrong.
4. Refrain from displaying too much detailed information to unprivileged users. This can be used to gather information for potential attacks.
5. Implement a usage function or comprehensive help menus to assist users in case of syntax confusion or misuse.

11.6.651 node_zvol_list

Contained in lib/host-scripts.d/common.d/zvol-management.sh

Function signature: 95f86f7af006e06fa8df84870e8f7eff728b79dd9bcb448fd864b4d7b350276b

11.6.652 Function overview

The node_zvol_list function is primarily used to list zvols (ZFS volumes) that exist either in a specified ZFS pool or in the entire system.

11.6.653 Technical description

- **Name:** node_zvol_list
- **Description:** This function lists all the zvols present either in a specified pool or in the entire system depending upon the argument provided. It logs any unknown parameters and errors associated with the listing of the zvols.

- **Globals:** `pool`: Holds the name of the pool wherein zvols are to be listed.
- **Arguments:** `[--pool "$2"]`: Optional argument where \$2 is the name of the pool.
- **Outputs:** This function will output the list of zvols along with their volume size on stdout.
- **Returns:** If the function successfully lists the zvols, 0 (zero) is returned, otherwise 1.
- **Example usage:**

```
node_zvol_list --pool my-zfs-pool
```

11.6.654 Quality and security recommendations

1. Sanitize input: Always check and sanitize the inputs provided by the user to avoid any sort of injection or overflow attacks.
2. Error handling: The current function only returns 1 when an error condition is met. However, different error codes can be returned for different error situations to improve error understanding and debugging.
3. Usage of variables: Avoid using global variables as it might lead to unexpected results if multiple functions are modifying them at the same time.
4. Documentation: Maintain proper comments and explanation for the different parts of the function. This will help other developers understand the functionality quickly and accurately.
5. Functional scope: Be clear about the function scope. A function named “`node_zvol_list`” should ideally only list node zvols and not do anything else. This will maintain the function’s purity and make it easier to debug and maintain.
6. Exit early: If a conditional check fails, exiting the function as soon as possible can enhance the performance and readability of the code.

11.6.655 `node_zvol_manage`

Contained in `lib/host-scripts.d/common.d/zvol-management.sh`

Function signature: `fb75475f9ed8b4e90c9fe423b227bbd08f041872fc46c3bee9ac215f09937`

11.6.656 Function overview

`node_zvol_manage()` is a function that acts as a controller for the operations related to managing ZFS volumes (Zvol). Necessary operations such as `create`, `delete`, `list`, `check` and `info` are handled based on the `action` argument passed by the user. In case no `action` is provided or an invalid `action` is passed, the function logs an error message and returns.

11.6.657 Technical description

- **Name:** `node_zvol_manage`
- **Description:** This function manages ZFS volumes (Zvol). It handles different operations like creation, deletion, listing, checking, and access information of Zvols. These operations are carried out based on the `action` argument provided to the function. If an invalid `action` is given, it logs an error message and returns.
- **Globals:** None
- **Arguments:**
 - `$1: action` - specifies the operation to be performed on the Zvols.
- **Outputs:** Executes one of the Zvol operations (`create`, `delete`, `list`, `check`, `info`), or logs an error message in case of invalid action.
- **Returns:**
 - 1 if no action or an invalid action is provided.
 - Exit status of the Zvol operation otherwise.
- **Example usage:** `node_zvol_manage create [options]`

11.6.658 Quality and security recommendations

1. Provide more specific error messages to help identify issues quickly.
2. Safeguard function against invalid number of arguments beyond the `action` argument depending upon the operation.
3. Consider handling operations asynchronously to improve performance and efficiency.
4. Incorporate additional security measures especially while deleting or modifying Zvols to avoid accidental data loss.
5. Document usage of the function in detail to avoid any misuse or misunderstanding.

11.6.659 `normalise_mac`

Contained in `lib/functions.d/network-functions.sh`

Function signature: `597ad94d9559d604fae13c8fd8ef8bde5b6c19fee81c409691213ad492d3d6b3`

11.6.660 Function Overview

The `normalise_mac` function in Bash is used to standardize the format of MAC addresses. It accepts a single MAC address as input, removes all common delimiters (such as `"-"`, `":"`, `."` or white space), converts it to lowercase, and validates the resulting output for being exactly 12 hexadecimal characters, which is the standard MAC address format.

11.6.661 Technical Description

- **name:** `normalise_mac`

- **description:** Normalizes the provided MAC address by removing common delimiters, converting to lowercase, and validating the MAC address format.
- **globals:** NA
- **arguments:** [\$1: A string representing a MAC address. Input MAC address could be with any common delimiters like colon(:), hyphen(-), dot(.) or a whitespace.]
- **outputs:** The function outputs a standardized MAC address in lowercase and without delimiters if the input is valid. If the input doesn't match the standard MAC address format, it outputs an error message to stderr "[x] Invalid MAC address format: \$1".
- **returns:** The function returns with a status of zero when successfully offering a normalized MAC address. It returns a status of 1 if the provided MAC address is invalid.
- **example usage:** `$ normalise_mac "01-23-45-67-89-ab"`
0123456789ab

11.6.662 Quality and Security Recommendations

1. Be cautious while handling the argument and take steps to ensure that it does not carry any harmful inputs such as shell code injections.
2. Always use the function in a safe environment, taking into consideration input sanitation and privilege separation.
3. To enhance the function, consider adding more input validations, such as checking if the input is null or empty before processing.
4. Lastly, always inform the users of the function about the correct input to provide. Since the function expects an argument to be a MAC address, providing this in the function's help or documentation would help avoid confusion and unexpected output.

11.6.663 `osvc_apply_identity_from_hps`

Contained in `lib/functions.d/opensvc-functions.sh`

Function signature: 36b4990e9b99b775a8e6be52ea2f3098da2a0dd4579130e4277b663f40451e06

11.6.664 Function overview

The `osvc_apply_identity_from_hps` function specifies the identity for a node in a cluster by checking the agent nodename in the OpenSVC configuration file. If the nodename is not found, the function logs an error and terminates. If the daemon is not running, the function provides a debug log stating that the identity will be configured upon daemon startup. The function assigns a value to the `node.name` attribute if the agent nodename is "ips". It also attempts to assign a value to the `cluster.name` attribute from a cluster configuration file.

11.6.665 Technical description

- **Name:** `osvc_apply_identity_from_hps`
- **Description:** This function assigns identities to a node and a cluster using values from configuration files.
- **Globals:** [`conf`: A string storing the path to the OpenSVC configuration file]
- **Arguments:** [No command line arguments]
- **Outputs:** Logs error, debug, and warning messages related to the process of assigning identities.
- **Returns:** 1 if the configuration file is unreadable or the agent nodename is not found; 0 otherwise.
- **Example usage:**

`osvc_apply_identity_from_hps`

11.6.666 Quality and security recommendations

1. Ensure that the OpenSVC and cluster configuration files exist and are readable.
2. Handle other potential error situations, such as if the 'om' command (which checks the status of the cluster) fails or if the daemon is not running.
3. Consider parameterizing the function so that the paths to the configuration files can be specified at the command line.
4. Consider validating the cluster and node name to ensure they adhere to expected naming conventions, decreasing the likelihood of issues during later processes.
5. Be mindful of the system permissions required to run commands involving cluster status and name assignments. It's best practice to run these commands with the minimum privileges necessary to reduce the potential security risks.

11.6.667 `osvc_bootstrap_cluster_on_ips`

Contained in `lib/functions.d/opensvc-functions.sh`

Function signature: `ec4bc28c436647a67226fa73910324c86c5355dbbbdbcaf89b48043c9fdb5796`

11.6.668 Function Overview

The `osvc_bootstrap_cluster_on_ips` function acts as a procedure for initializing an OpenSVC cluster within an IPS (Internet Protocol Suite) environment. This function involves several steps including:

1. Configuration generation and enforcing the key.
2. Starting up a daemon with the aid of `supervisord`.
3. Then undertaking further configuration steps with regards to the cluster,
4. Including the setup of a heartbeat.
5. Generation and/or retrieval of the cluster secret.
6. Ultimately resulting in the verification of the daemon.

11.6.669 Technical Description

- **Function Name:** `osvc_bootstrap_cluster_on_ips`
- **Description:** This function bootstraps an OpenSVC cluster on IPS (Internet Protocol Suite). It performs several steps such as generating a configuration, starting a daemon, setting up cluster settings, and verifying the daemon.
- **Globals:**
 - `CLUSTER_SERVICES_DIR`: Directory containing services for the cluster.
 - `cluster_name`: Name of the cluster.
 - `cluster_secret`: Secret passphrase for the cluster.
- **Arguments:** This function does not accept any arguments.
- **Outputs:** This function logs either the success or the failure of each step and prints these messages to the standard output.
- **Returns:** 0 on successful execution, 1 on failure due to configuration issues or failure to set parameters, and 2 on failure due to daemon start issues or the daemon not responding after bootstrapping.
- **Example Usage:**

```
osvc_bootstrap_cluster_on_ips
```

11.6.670 Quality and Security Recommendations

1. Strict handling of the `cluster_secret`: The cluster secret should be securely stored and its read access should be tightly controlled. Logging the secret value should also be avoided.
2. Exception handling: The function should robustly handle any exceptions or failures that may occur during execution especially during critical actions such as creating configurations and starting daemons.
3. Function validation: Additional validation checks could be included to ensure the function is executing in the expected environment that complies with other necessary specifications.
4. Logging: Including more logs in the function would help tracking the progress and tracing any failures more easily.
5. Secure daemon: Ensure that the “om daemon run” command runs securely, especially when exposing IP addresses and other sensitive details.

11.6.671 `osvc_config_update`

Contained in `lib/functions.d/opensvc-functions.sh`

Function signature: 623bd8e4588409e783c922c0cd2b96d347d9f58a1f1e5c0efabb5d906e5719ef

11.6.672 Function overview

The `osvc_config_update` function updates the OpenSVC cluster configuration by setting the entered key=value pairs. This function first checks if at least one key=value pair was entered as argument. If not found, it logs an error message and exits with a return status of 1. If all is well, it sets the entered arguments with each key=value as an argument to `om cluster config update`. The successful or unsuccessful update is logged and the function returns with a status of 0 or 1.

11.6.673 Technical description

Following are the technical details for the function `osvc_config_update`:

- **name:** `osvc_config_update`
- **description:** This function updates the OpenSVC cluster configuration with key=value pairs entered as arguments.
- **globals:** `set_args`: An array to store the `--set` with key=value pair for the `cluster` configuration update.
- **arguments:** `$@`: Key=value pairs for updating the cluster configuration.
- **outputs:** Logs the setting of cluster configurations as well as success or failure of update.
- **returns:** 0 if the cluster configuration updated successfully.
1 if it fails or if no arguments provided.
- **example usage:** `osvc_config_update database=postgres user=admin`

11.6.674 Quality and security recommendations

1. Function should validate the key=value pairs for known configurations before processing to avoid any issues.
2. Proper error handling should be done when the command to update the cluster configuration fails.
3. There should be a limit on the number of key=value pairs that could be processed at once to prevent any resource issues.
4. Log the output of command to update the configuration for debugging in case of any failures.
5. All data, including log messages and return values, should be properly sanitized to prevent any security threats.

11.6.675 `osvc_configure_cluster_identity`

Contained in `lib/functions.d/opensvc-functions.sh`

Function signature: `1f5c78a106d5bed621fe281cceb1a8daa04bf12c433437185fe19c7e31f0192d`

11.6.676 Function Overview

The `osvc_configure_cluster_identity()` function is used to configure the identity of an OpenSVC cluster. It waits for the socket to be ready, checks if the daemon is responsive, retrieves the cluster name and heartbeat type, and applies the configuration updates. It also records these stages in logs. If any stage fails, it records the error in the logs and the function stops. The function returns 0 if the process is successful and 1 if it encounters an error.

11.6.677 Technical Description

- **Name:** `osvc_configure_cluster_identity()`
- **Description:** This function configures the identity of an OpenSVC (open source virtual clustering) cluster.
- **Globals:** [`i`: Iterator for a loop from 1 to 10, `cluster_name`: Name of the cluster, `hb_type`: Cluster heartbeat type]
- **Arguments:** None
- **Outputs:** Informational and error logs throughout the process. Final log provides cluster identity, including name and heartbeat type.
- **Returns:** 0 if the configuration process is successful, 1 if an error is encountered.
- **Example Usage:** `osvc_configure_cluster_identity`

11.6.678 Quality and Security Recommendations

1. The function should have error handling for when the socket is not ready after 10 seconds. Currently, the function would proceed further even in this scenario.
2. The function should check the existence of the cluster, before trying to configure its identity.
3. The function should validate the `cluster_name` and heartbeat type before using them.
4. The unchecked use of globals can be risky. These should be validated or handled locally within the function.
5. Ensure that error logging is detailed enough to troubleshoot potential problems. Sensitive information, however, should not be included in logs to prevent security risks.
6. To prevent potential errors, consider setting default values for the `cluster_name` and heartbeat type if they are not set or invalid.
7. Consider adding more extensive testing to cover all aspects of this function to ensure it behaves as expected.

11.6.679 `_osvc_kv_set`

Contained in `lib/functions.d/opensvc-functions.sh`

Function signature: `e83bb88e314eb7098eb2ed8868cdecf150f3a7c4cec631f0ac14eadabbeaa0d3`

11.6.680 Function overview

The `_osvc_kv_set()` function in Bash is a utility intended for setting key-value pairs in an Omniscale configuration. This function takes two arguments—a key and a value—and passes them into an `om config set` command using a keyword argument.

11.6.681 Technical description

- **Name:** `_osvc_kv_set`
- **Description:** This function takes a key and a value as inputs and sets this key-value pair in an Omniscale configuration using the `om config set` command.
- **Globals:** None.
- **Arguments:**
 - `$1`: A string representing the key of the key-value pair.
 - `$2`: A string representing the value of the key-value pair.
- **Outputs:** No explicit return, but changes will be made in the Omniscale configuration via the `om config set` command.
- **Returns:** Does not return a value.
- **Example usage:**

```
_osvc_kv_set "username" "admin"
```

11.6.682 Quality and security recommendations

1. Always validate input before processing: The function currently does not perform any validation or sanitization of the arguments it receives before processing them, making it open to potential misuse or errors. Consider implementing input validation to ensure that correct and safe values are being passed in as arguments.
2. Handle errors gracefully: While the function does use the `${var:?word}` expression to handle missing arguments, its error handling could be more explicit. For example, it could log a custom error message or use an `exit` call to halt the entire script.
3. Write test cases: To fully ensure the function's reliability and maintainability, consider writing unit tests that inspect the behavior of the function in various contexts and with different inputs.

11.6.683 `_osvc_run`

Contained in `lib/host-scripts.d/common.d/opensvc-management.sh`

Function signature: `b4d91dc406a4ee6f725ae31a5e96d8ba670b6af35d65b7ad9e0416aecb14c186`

11.6.684 Function overview

The `_osvc_run` function is a utility function in Bash that encapsulates the command execution pattern accommodating argument passing, error capture, and logging. This

function accepts command arguments, executes them, logs any error codes or output, and returns the error code.

11.6.685 Technical description

11.6.685.1 Name

`_osvc_run`

11.6.685.2 Description

This script function accepts a description and a command as inputs, and executes the given command. Any resulting output or errors are logged using a custom logger, tagged with a provided description and an exit code. The function then returns the same exit code.

11.6.685.3 Globals

None

11.6.685.4 Arguments

- `$1`: desc - A description tag for the command being run. This gets logged with any output or errors.
- `$@`: The remaining arguments form the command to be executed.

11.6.685.5 Outputs

Logs an information message on execution status including the provided tag, the exit code, and any command output or errors.

11.6.685.6 Returns

The exit status of the command that was run.

11.6.685.7 Example Usage

```
_osvc_run "List running processes" "ps aux"
```

11.6.686 Quality and security recommendations

1. Aside from the facilities provided by Bash, there's no explicit error checking or sanitization. Ensure the commands or arguments that this function wraps are safe to execute.

2. Consider adding checks for the number of parameters provided to the function to avoid undesired behavior.
3. Make sure to properly escape the parameters that are passed to this function to avoid command injection vulnerabilities.
4. Always use the function with trusted inputs, as this function executes commands directly.
5. Ensure that the right permissions are granted at the directory and shell level. The function should not have more permissions than what it needs.
6. Incorporate a controlled logging mechanism to reduce the risk of potential log forgery.

11.6.687 osvc_verify_daemon_responsive

Contained in `lib/functions.d/opensvc-functions.sh`

Function signature: `4ff390712ebba6d37b0e6c9adaeddd2892a0085401cb7e2c2539e52d8ebb211f`

11.6.688 Function Overview

The function `osvc_verify_daemon_responsive()` is used to check the responsiveness of the OpenSVC daemon. It attempts to get the status of the OpenSVC cluster and based on the result it returns either a success or a failure message. It logs the process and the status of the responsiveness of the daemon.

11.6.689 Technical Description

- **Name:** `osvc_verify_daemon_responsive`
- **Description:** This function verifies if the OpenSVC daemon is responsive. It tries to get the status of the OpenSVC cluster and logs the process. It returns 0 if the daemon is responsive and exits with a status of 1 if it isn't.
- **Globals:** None
- **Arguments:** None
- **Outputs:** This function logs the process and the result of the verification of the OpenSVC daemon's responsiveness.
- **Returns:**
 - 0 if the OpenSVC daemon is responsive
 - Exits the program with a status of 1 if the daemon isn't responsive
- **Example usage:**

`osvc_verify_daemon_responsive`

11.6.690 Quality and Security Recommendations

1. Adding error handling mechanisms to manage unexpected errors that might occur while getting the OpenSVC cluster status.

2. It might be useful to customise the exit status codes to indicate different types of errors, instead of using only 1 for any error.
3. This function depends on the `om` and `hps_log` commands. Ensure that these commands are secure and permissions are correctly set.
4. It would be beneficial to implement a log rotation system for the log files to avoid them becoming too large and to maintain the logging system's efficiency.
5. Avoid using global variables as much as possible to enhance security and prevent unintended side-effects.
6. Regularly update and audit the dependencies (`om` and `hps_log` commands) utilised in this function for any known vulnerabilities.

11.6.691 `osvc_wait_for_socket`

Contained in `lib/functions.d/opensvc-functions.sh`

Function signature: `52ca00661cdfd2a34c7a7c9485780d5e482cb5aba723144647d5e1fabb6e9e5d`

11.6.692 Function overview

The function `osvc_wait_for_socket` is designed for waiting until the OpenSVC daemon socket becomes ready. It sends a debug log informing that it's waiting for the OpenSVC daemon socket. In a loop of 10 iterations, it checks if the daemon socket is ready and if so, it sends a debug log to confirm that the socket is ready and returns a success status. If the socket doesn't get ready within 10 seconds, an error log is sent indicating the socket's unavailability and the program is forced to exit with failure status.

11.6.693 Technical description

- **Name:** `osvc_wait_for_socket`
- **Description:** Waits for the OpenSVC daemon socket to be ready.
- **Globals:** None
- **Arguments:** None
- **Outputs:** Sends debug log about the status of the daemon, sends an error log if the OpenSVC daemon socket is not ready after 10 seconds.
- **Returns:** 0 if the OpenSVC daemon socket is ready within 10 seconds, terminates with error status 1 otherwise.
- **Example usage:**

```
osvc_wait_for_socket
```

11.6.694 Quality and security recommendations

1. The function could be improved to accept the socket path and timeout as arguments, instead of hardcoding the values. This will make the function more flexible

and reusable.

2. The function uses `exit 1` in case of an error, which will terminate the whole script and not just the function. This could be inappropriate in some cases where the unavailability of the socket should not result in the termination of the whole script. It would be better to use something like `return 1` to just stop this function and return error status.
3. In terms of security, it is advisable to use `set -e` and `set -u` Bash options to handle unassigned variables and errors.

11.6.695 `parse_apkindex_package`

Contained in `lib/functions.d/repo-functions.sh`

Function signature: `c8272ffc02ed22252a554f027f250eccd763e445bdf8b52ad5784e2b6d338cbc`

11.6.696 Function overview

The Bash function `parse_apkindex_package` is designed to parse APKINDEX files. The function takes a package name and APKINDEX file as arguments, then finds and prints the paragraph associated with the package name. A paragraph ends with a blank line in APKINDEX file. If it does not end with a blank line, then it checks the last record. The function returns 0 if the package details are found or 1 if the package was not found.

11.6.697 Technical description

- **Name:** `parse_apkindex_package`
- **Description:** This function is used to read an APKINDEX file, paragraph by paragraph (blank-line separated), and find package information. The package information is then echoed (printed) to the console. If the file does not end with a blank line, it checks the last record for matching package information.
- **Globals:** None
- **Arguments:**
 - `$1` (`apkindex_file`): The path of the APKINDEX file to be parsed.
 - `$2` (`package_name`): The name of the package for which information is sought.
- **Outputs:** Prints the information of the package found within the APKINDEX file.
- **Returns:** Returns 0 if the package details have been found in the APKINDEX and 1 if the package was not found.
- **Example Usage:**

```
parse_apkindex_package "./apkindex.txt" "mypackage"
```

11.6.698 Quality and security recommendations

1. **Error Handling:** The function currently doesn't handle possible exceptions. For instance, if the provided file is not accessible or does not exist.
2. **Input Validation:** There is currently no validation of inputs. This might lead to problematic behaviour if, for instance, an empty or incorrect APKINDEX filename or an empty package name is provided.
3. **Documentation:** The function could be clearer with more comments explaining any complex parts of the syntax used. Currently, there are no comments explaining what variables like `in_record` serve for, which might confuse other contributors or end users.
4. **Efficiency:** The function reads the file line by line, incurring heavy I/O operations. If the APKINDEX file is very large, this might not be efficient. There's room to optimize this for better performance.

11.6.699 `prepare_custom_repo_for_distro`

Contained in `lib/functions.d/repo-functions.sh`

Function signature: `f41fadc94c0ebb53ebb87776324ccfcc73903f665dfe9bfd0a9d3f4c9cbb827c`

11.6.700 Function Overview

The function `prepare_custom_repo_for_distro` serves to prepare a custom software repository for a given linux distribution represented by the input `dist_string`. It also takes an array of URLs or local files to download and add to the custom repository. The function segregates URLs and local file paths from the other entries stamped as required packages in an iterative manner. Further, it initiates repository creation, downloads the package from each URL or copies from local files, and builds the YUM repository. It verifies that the required packages are available in the repository and logs an error if a package is missing.

11.6.701 Technical Description

- **Name:** `prepare_custom_repo_for_distro`
- **Description:** This function prepares a custom software repository for a given distribution by segregating source links (URLs or local file paths) from required package listings, creating repositories, building YUM repository, and ensuring presence of required packages.
- **Globals:** `HPS_PACKAGES_DIR`: The root directory where repositories are created.
- **Arguments:**
 - `$1`: `dist_string`, the string representation of a Linux distribution.
 - `$@`: An array of URLs, local file paths to the required packages or names of the required packages.

- **Outputs:** Logs various informational and error messages during repository preparation. Copies or downloads package files to the repository.
- **Returns:** The function uses return values varying from 0-6 to indicate success or various types of failures (such as, failure to create repository directory, download or copy a package, etc.).
- **Example usage:**

```
prepare_custom_repo_for_distro    "ubuntu"
                                "https://example.com/package1.deb"          "package2"
                                "/path/to/package3.deb"
```

11.6.702 Quality and Security Recommendations

1. **Error Handling:** At present, the function returns error messages in various places but a more rigorous error handling system should be implemented for each step.
2. **Input Validation:** Currently, the repository name and the package sources are not validated. They should be confirmed to avoid unpredictable behavior.
3. **User Permissions:** Permissions required to create directories or copy files need careful considerations.
4. **Logging & Auditing:** It would be good to maintain consistent logging throughout the function to account for all actions taken.
5. **Dependency Handling:** The function should manage potential dependencies of the required packages.
6. **Use Secure Communication:** If possible, use secure protocols (like HTTPS) for downloading packages from URLs.

11.6.703 print_cluster_variables

Contained in `lib/functions.d/cluster-functions.sh`

Function signature: `c5583d8fd4715e19ec442b98a50c60b43c9b20d0f5892f8d147bbb29263d00fa`

11.6.704 Function Overview

The function `print_cluster_variables()` in Bash is used to read and print the variables from a configuration file of a currently active cluster. This function first checks if the config file exists, and if it doesn't, returns an error message and terminates with a non-zero return value. If the config file exists, this function reads every line from it, ignores blank lines or lines starting with the hash symbol (`#`), and prints the rest after removing surrounding quotes.

11.6.705 Technical Description

- **Name:** `print_cluster_variables()`

- **Description:** This function reads a configuration file for a currently active cluster and prints its variables after removing surrounding quotes. It returns an error and stops execution if the config file doesn't exist.
- **Globals:** No global variables are used.
- **Arguments:** This function doesn't take any arguments.
- **Outputs:** Variables from the cluster's config file. The output is sent to stdout.
- **Returns:** 1 if the config file doesn't exist. Otherwise, the return value is dependent on the final command in the function.
- **Example usage:** `print_cluster_variables`

11.6.706 Quality and Security Recommendations

1. For enhanced security, consider adding additional checks besides the presence of the config file. For example, verifying file permissions or ownership.
2. Always quote variable references to prevent word splitting and globbing. For instance, consider replacing constructs like `$k` with `"$k"`.
3. The function could return a specific non-zero exit code for different types of failures (file not found, permission errors, etc.) to make error handling more specific.
4. To improve readability, consider adding a comment describing what each local variable specifically stores.
5. Consider handling possible errors in command substitutions used in variable assignments (like the call to `get_active_cluster_filename`) for more robust error handling.

11.6.707 `_print_meta`

Contained in `lib/functions.d/cluster-functions.sh`

Function signature: `8f66cc2d9e01b400ddc05a00d6399036188baa213e119c36c4340e95320bcf7f`

11.6.708 Function Overview

The function `_print_meta` is used to obtain and display metadata about a cluster configuration using a specific key and label. The function makes use of two arguments, the key and the label, to fetch the metadata details of a cluster configuration. It then prints the label and the value if the value is not empty. Furthermore, it also handles conditions where there might be multiple clusters, none active cluster and also when there is no active one and exactly one cluster.

11.6.709 Technical Description

- **Name:** `_print_meta`
- **Description:** A shell function to fetch and display metadata information from a cluster configuration.

- **Globals:**
 - `key`: The configuration key used to fetch value from the cluster configuration.
 - `label`: Label to be print before printing the configuration value.
- **Arguments:**
 - `$1`: The first argument, used as the key in the function.
 - `$2`: The second argument, used as the label in the function.
- **Outputs:** The function can output the label and configuration value if the value is not empty.
- **Returns:** The function returns 0.
- **Example usage:** `_print_meta "DESCRIPTION" "Description"`

11.6.710 Quality and Security Recommendations

1. As per good practice, make sure that the necessary sanitization of input parameters is performed before they are used.
2. Error checking should be established to handle cases where the cluster configuration or key provided as function argument might not exist.
3. Clear readability should be maintained throughout the code. Usage of clear variable names and comments can help to ensure a good level of readability.
4. Always consider the security implications of your script; be cautious of the possibility of code injection attacks.
5. Regular updates and security patches should be implemented to ensure your bash shell is up-to-date and secure from known vulnerabilities.

11.6.711 `refresh_node_functions`

Contained in `lib/host-scripts.d/common.d/common.sh`

Function signature: `4d94de7cd8d9dc976cb47f74764f248bf1285856047f02be71ec96d00b4e6d74`

11.6.712 Function Overview

`refresh_node_functions` is a Bash function that updates the node functions in a server's local library. This is done by getting the provisioning node IP address and constructing a URL to download the updated functions from. If the download is successful, the function refreshes the local node functions and reloads them in the current shell.

11.6.713 Technical Description

Name: `refresh_node_functions`

Description: This function is designed to refresh the local node functions on a server by downloading them from a provided URL. It first determines the provisioning node, constructs the URL for the functions file, checks that necessary directory exists and, if

successful, downloads the functions file, and reloads it within the same shell the function was run from.

Globals: [provisioning_node: Stores the provisioning node IP address, functions_url: Holds the fully formed URL from where the functions will be downloaded]

Arguments: No Arguments are taken by this function

Outputs: Status messages and potential error messages are output to the terminal. May output either a success message (“Successfully refreshed node functions from...”) or one of two error messages (“Could not determine provisioning node”, “Failed to download functions from...”)

Returns: The function will return 1 on encountering an error in determining provisioning node or downloading functions. Will return 0 upon successful completion.

Example usage:

```
refresh_node_functions
```

11.6.714 Quality and Security Recommendations

1. Consider adding validation to check if the provisioning node and URL are in the correct format before making the request.
2. Examine the potential for integrating error handling to catch and handle any curl download failures more gracefully.
3. Always ensure that sensitive data, such as IP addresses, are not exposed within error messages.
4. Consider utilizing HTTPS instead of HTTP to enhance the security of your download.
5. Bears ensuring the server you are downloading the functions from is trusted to prevent malicious code execution.
6. Verify if the required directory `/srv/hps/lib` has correct permissions assigned to prevent any unauthorized access.

11.6.715 register_source_file

Contained in `lib/functions.d/prepare-external-deps.sh`

Function signature: `3a47ee217d4c21d4f884f60289c0b97418f2bea9352f1117fe8517c7bb30bb0b`

11.6.716 Function Overview

The function `register_source_file()` has been designed to register a source file to an index within a specific destination directory. This function takes in a filename and

handler as inputs and if any duplicate is avoided, it proceeds to register the file and its handler to the index file located in the destination directory. The function also provides feedback on successful registration or if the file is already registered.

11.6.717 Technical Description

Name: `register_source_file`

Description: This Bash function registers a source file to an index present in the given destination directory.

Globals:

- `HPS_PACKAGES_DIR`: Path to the packages directory.

Arguments:

- `$1` or `filename`: The name of the source file to be registered. - `$2` or `handler`: The handler associated with the source file.

Outputs:

- Prints a message that indicates successful registration or if a source file is already registered.

Returns:

- Returns 0 if file is already registered, effectively preventing any changes.

Example Usage:

```
register_source_file "myfile.txt" "myHandler"
```

11.6.718 Quality and Security Recommendations

1. Validate the inputs: Make sure both filename and handler are not empty or null before proceeding with the registration process.
2. Error handling: Add error catching mechanisms to handle unexpected situations such as issues with directory creation or file writing.
3. Secure Files: Ensure that the permissions for both the index file and directory are set appropriately to prevent unauthorized access.
4. Logging: Introduce detailed logging in each step for easier debugging and traceability.

11.6.719 `reload_supervisor_config`

Contained in `lib/functions.d/supervisor-functions.sh`

Function signature: `d65132fc9d374ab0fcb20731915231186b9e78da61ece048b77813c0e1a57baa`

11.6.720 Function Overview

This is a simple bash shell function named `reload_supervisor_config`. The main aim of this function is to reload the configuration file for the Supervisor process control system. It achieves its objective by first setting the Supervisor configuration file path and then making two calls to `supervisorctl` with the `-c` flag, executing `'reread'` and `'update'` commands respectively. These actions are logged using the `hps_log` function with an "info" flag.

11.6.721 Technical Description

Definition Block:

- **Name:** `reload_supervisor_config`
- **Description:** This function is used to reload a Supervisor server's configuration file. It rereads the configuration file and automatically applies the changes to the services under its management.
- **Globals:**
 - `SUPERVISORD_CONF`: The environment variable used to store the path to the Supervisor configuration file, which will be reread and updated.
 - `CLUSTER_SERVICES_DIR`: This environment variable should hold the directory path where the Supervisor configuration file is located.
- **Arguments:** None in this function.
- **Outputs:** The result of the `'reread'` and `'update'` actions, as well as their respective exit statuses, are logged.
- **Returns:** Not explicitly defined in this function, returns the exit status of the last executed command.
- **Example Usage:** `reload_supervisor_config`

11.6.722 Quality and Security Recommendations

1. Add error checking to ensure the Supervisor configuration file exists before trying to reload it.
2. Handle the potential failure of the `'reread'` or `'update'` commands, such as by checking `"$?"` after each call to `supervisorctl` and responding accordingly.
3. Improve documentation by adding descriptive comments within the function, detailing what each command does.
4. Consider restricting who has permission to execute this script to limit potential security risks associated with unauthorized changes to Supervisor tasks.
5. Use environment variable expansion `${VAR?}` to ensure required environment variables are set before invoking the function.

11.6.723 `reload_supervisor_config`

Contained in `lib/functions.d/configure-supervisor.sh`

Function signature: 1ad96e9487bac5b94d7eca312662f18392454507c7cddf5de26d924e922722

11.6.724 Function overview

The `reload_supervisor_config` function is a bash function created to reload the configuration of `supervisord`. It first specifies the path to the `supervisord` configuration file, then uses the `supervisorctl` command with the `-c` option to read this configuration file, before using the same command to update the `supervisord` system with the new configuration.

11.6.725 Technical description

- **Name:** `reload_supervisor_config`
- **Description:** This function is responsible for reloading the `supervisord` configuration. It utilizes `supervisorctl`, a command line utility provided by `supervisord`, to reread and update based on the given configuration file.
- **Globals:**
 - `HPS_SERVICE_CONFIG_DIR`: This is the directory that contains the `supervisord` configuration file.
 - `SUPERVISORD_CONF`: This is the path to the `supervisord` configuration file, assembled by concatenating `HPS_SERVICE_CONFIG_DIR` with `/supervisord.conf`.
- **Arguments:** No arguments required.
- **Outputs:** No explicit outputs. Function performs operations and sends tasks to `supervisord`.
- **Returns:** Nothing explicitly but based on the tasks it sends to `supervisord` it determines whether `supervisord` configuration has been successfully reloaded or not.
- **Example usage:** `reload_supervisor_config`

11.6.726 Quality and security recommendations

1. Before running the `supervisorctl` commands, the function should check if the `SUPERVISORD_CONF` file actually exists to avoid errors.
2. The function should also check that `HPS_SERVICE_CONFIG_DIR` is set and not empty before proceeding.
3. Use the `-c` option for `supervisorctl` carefully, as it allows overriding the main configuration file (`supervisord.conf`), which in the wrong hands can lead to potential misconfigurations or disruptions.
4. Implement error handling. For example, catch errors from `supervisorctl` and handle them appropriately, such as logging the error and exiting the function with an error status.

5. As a general rule, avoid storing passwords or other sensitive information in environment variables like `HPS_SERVICE_CONFIG_DIR`. Use a secure method to handle these credentials, such as a secret manager.

11.6.727 reload_supervisor_services

Contained in `lib/functions.d/supervisor-functions.sh`

Function signature: 0551d7e290e09fc3d49350121abd16aeb548ae2a69179c18759cb9973b83ff65

11.6.728 Function overview

`reload_supervisor_services` is a Bash function that reloads services managed by `Supervisor`. The function is capable of reloading a specific service or all services based on the provided argument. During the reloading process, it validates the Supervisor configuration file and sends a HUP (Hang UP) signal to the target service(s).

11.6.729 Technical description

Name: `reload_supervisor_services`

Description: This function is used to reload services managed by Supervisor. It sends a HUP signal to the services which informs them to close and reopen their log files. If a specific service name is given, it sends the signal to the specified service otherwise all of the services are signaled. It checks if the Supervisor configuration file exists and if it does not, it logs an error and returns.

Globals: - `SUPERVISORD_CONF`: specifies the path of the supervisor configuration file.

Arguments: - `$1`: name of the service to be reloaded. If this argument is not given, the function will send the signal to all services.

Outputs: Logs either an informational or an error message, depending on whether the HUP signal is sent successfully or not.

Returns: - 0: if the HUP signal is sent successfully. - 1: if the `SUPERVISORD_CONF` global variable is not set, or the supervisor configuration file does not exist. - 2: if the HUP signal fails.

Example Usage:

```
reload_supervisor_services "my_service"
```

11.6.730 Quality and security recommendations

1. Add input validation for the service name parameter to prevent the function from accepting arbitrary user input.

2. Due to indirect use of variables in the command line ("`$target`"), it is advisable to use arrays to avoid misinterpretation of any special characters in the commands.
3. Use absolute file paths for the supervisor configuration file, which makes your script more robust and less error-prone.
4. Avoid using `2>&1`, which combined stdout and stderr because it can cause potential issues with parsing and troubleshooting.
5. Log all relevant function events at the appropriate verbosity level for easier debugging and greater transparency of its activities.
6. Always consider using `-e` bash option to force your script to exit with non-zero automatically when any command fails. This will make your script more reliable.
7. Check if the required commands (`supervisorctl`, `sed`, `tr`) are available before using them.

11.6.731 `remote_cluster_variable`

Contained in `lib/host-scripts.d/common.d/common.sh`

Function signature: `80d48ac6c7e7f4246491425bf9e89a589b854a03598f7fff94cfddfd4f9b7a55`

11.6.732 Function Overview

The function `remote_cluster_variable` is essentially used for manipulating the cluster variable in a remote Bash environment. It enables both the reading (GET) and writing (SET) of a cluster variable over HTTP, through a POST or GET request. The function communicates with the server via curl and utilizes a gateway system as the intermediary node for communicating with the remote cluster. The variable's value is URL-encoded before it's sent, ensuring its safe passage across the network.

11.6.733 Technical Description

- **Name:** `remote_cluster_variable`
- **Description:** The function manages a cluster variable in a remote bash environment. If two arguments are provided, it sets (POSTs) the specified value of the variable. If only one argument is provided, it gets (GETs) the value of the variable.
- **Globals:** None
- **Arguments:** `$1`: name (The name of the variable. If it's not provided, the function raises an error and terminates), `$2`: value (The value to set to the variable. This argument is optional. If it's provided, function sets (POSTs) this value to the variable. If it's not provided, function gets (GETs) value of the variable)
- **Outputs:** The function outputs the response from the POST or GET request.
- **Returns:** Returns 1 if there's an error in `get_provisioning_node` function. Otherwise, nothing is returned explicitly.
- **Example Usage:** `bash remote_cluster_variable username john`
`remote_cluster_variable username`

11.6.734 Quality and Security Recommendations

1. Always URL encode all inputs to prevent potential security threats such as injection attacks.
2. Error checking for all variable assignments and function calls should be implemented for greater resilience.
3. Consider implementing retry logic for the curl command, as network requests are oftentimes flaky and may need to be retried.
4. Implement logging to catch and debug potential errors during runtime.
5. Make sure that the gateway system, being the relay of data, is secure against possible threats.
6. Validation for the user-inputted name and value, such as type or format restrictions, could be considered.
7. Always use HTTPS for such requests to ensure the privacy and data integrity of your communications.

11.6.735 remote_function_lib

Contained in `lib/functions.d/kickstart-functions.sh`

Function signature: `b6d9cd335e4f61b186f614aa07279b2f8bdd77298bf573a231ffd21869c18118`

11.6.736 1. Function Overview

The `remote_function_lib` function is designed for assembling functions to be injected into pre and post sections of certain scripts. It does this by using a technique called “here document” or heredoc, denoted by the `<<EOF` syntax. This is generally used to reduce the need for invoking `echo` repeatedly or for creating a temporary file containing the lines of text. This function in particular does not do anything else yet and waits for the user to fill in the necessary details.

11.6.737 2. Technical Description

- **Name:** `remote_function_lib`
- **Description:** This function is a placeholder for functions to be injected into pre and post sections of a script. It deploys a heredoc approach for adding lines of text.
- **Globals:** None
- **Arguments:** None
- **Outputs:** Outputs a block of text that is written between the EOF markers.
- **Returns:** Depending on the usage (not specified in the provided function details).
- **Example usage:**

`source remote_function_lib`

Note: For actual usage, the function injected within this heredoc should be implemented first.

11.6.738 3. Quality and Security Recommendations

1. Be mindful while using heredocs. If not properly utilized, sensitive data may be inadvertently written to a file and may remain there even after the script has finished.
2. As this function is passive and doesn't perform any actions except print to standard output, ensure that the functionality it provides is actually needed in your script.
3. Consider using functions as external files for better organization, modularity, and error handling.
4. Always validate and sanitize input to functions; although this function doesn't take any inputs, it's a good general practice.
5. Include comments to thoroughly explain the details of the function, its inputs, outputs, and what it specifically does.
6. Also consider handling possible errors and return suitable error codes/messages for better debugging.

11.6.739 `remote_function_lib`

Contained in `lib/functions.d/kickstart-functions.sh`

Function signature: `b6d9cd335e4f61b186f614aa07279b2f8bdd77298bf573a231ffd21869c18118`

11.6.740 1. Function Overview

The `remote_function_lib` function in Bash is a method to output a collection of functions that are supposed to be introduced into the pre and post sections of a script or command sequence. This function can be extremely useful in scenarios where repeated sections of code are needed across various parts of a script, allowing the user to maintain the DRY (Do not Repeat Yourself) principle. While the function currently outputs a predefined set of functions directly into the pre and post sections, future enhancements might permit the functions to be defined within an independent file and imported when needed.

11.6.741 2. Technical Description

- **Name:** `remote_function_lib`
- **Description:** A Bash function that outputs multiple functions to be used in pre and post script/command sequences.
- **Globals:** None.
- **Arguments:** No explicit arguments.

- **Outputs:** A collection of functions that can be injected into the pre and post sections of a script or command sequence.
- **Returns:** The function does not explicitly return any value; it instead outputs multiple functions to be used elsewhere.
- **Example Usage:**

`remote_function_lib`

11.6.742 3. Quality and Security Recommendations

1. To enhance code readability, consider moving the return functions into an independent file. This independent file can then be called within this function.
2. It would be better to include proper namings of the functions as comments to enhance readability.
3. For security purposes, consider adding input validations when the function starts handling command-line arguments.
4. Avoid exporting unnecessary global variables. They can conflict with variables from other parts of the scripts or be manipulated by unauthorized users.
5. Always use the latest version of Bash to get higher security and better new features from the latest updates.
6. Regularly run your code through static code analyzers (like ShellCheck) to catch potential bugs and security issues.

11.6.743 `remote_host_variable`

Contained in `lib/host-scripts.d/common.d/common.sh`

Function signature: `d808c9948262b9c10e2d29dbf91fff37d9e5c1ad4ce5c78af73813e01521b322`

11.6.744 Function overview

This Bash function, `remote_host_variable`, is designed to interact with a remote host's boot manager. The function sets a local variable for the name, and optionally for the value. The gateway is set to the function call of `'get_provisioning_node'`. If an error is found, the function will immediately return with a status of 1.

If a second argument is provided, it is encoded and added to the URL along with the encoded name. A POST request is made to this URL, which presumably interacts with the remote host, potentially setting a variable or performing some other change.

11.6.745 Technical description

- Name: `remote_host_variable`
- Description: This function performs an HTTP POST request to a remote boot manager, potentially setting a variable.

- Globals: None
- Arguments: [\$1: This is the name of the variable to be set or changed, a mandatory input. \$2: This is the optional value of the variable.]
- Outputs: Outputs are dependent on the return of the `curl` call.
- Returns: If the `get_provisioning_node` function call encounters an error, the function returns 1. Otherwise, return value is determined by the `curl` request.
- Example usage: `remote_host_variable "NAME" "VALUE"`

11.6.746 Quality and security recommendations

1. This function relies on a successful return from `'get_provisioning_node'`, but does not validate its output in any way. You should add a check to ensure that `'get_provisioning_node'` returns a valid result.
2. There is a potential risk of URL injection with the \$2 argument. Ensure the validation and sanitization of the \$2 and \$1 parameter, especially if using this function in web applications or other programs where user input might be used as these arguments.
3. Rather than forming your URL by concatenation, consider using an URL-building library function for a more robust and error-resistant approach.
4. Always ensure that the correct headers are used in your curl requests to ensure authenticity and data integrity when communicating with the remote server.
5. The URL base-string used in the function is hardcoded into the code itself, consider moving it out to a configuration file or argument to make the script more versatile.
6. Consider using HTTPS for secure communication instead of HTTP.

11.6.747 remote_log

Contained in `lib/functions.d/kickstart-functions.sh`

Function signature: `1f586056ba1b573eed69d32639c3940c1c742ba73af4aae917a1d65d36c5367c`

11.6.748 Function overview

The `remote_log()` function is created to send log messages from a local machine to a remote gateway server. The function takes one argument, which is the message that needs to be logged. Before the message is sent, it is URL-encoded to ensure that it is correctly received by the remote server. The encoding is done in a loop, character by character. After the encoding, the message is sent to the remote server using a `curl` HTTP POST request.

11.6.749 Technical description

Name: - `remote_log()`

Description: - This is a Bash function created to send log messages from a local machine to a remote gateway server. It takes one argument, the message, which needs to be logged. The function first URL-encodes the message and then sends it to the remote server via a POST request using `curl`.

Globals: - VAR: Not applicable - `macid`: The Mac id of the local machine - `HOST_GATEWAY`: The IP address or URL of the remote server

Arguments: - `$1`: The initial message that needs to be logged

Outputs: - Sends an HTTP POST request to the remote server with the URL-encoded message

Returns: - Null

Example Usage:

```
remote_log "This is a test log message"
```

11.6.750 Quality and security recommendations

1. Input validation: Ensure that the variable message does not contain malicious commands or SQL injections.
2. Error handling: Handle exceptions to avoid function crashes e.g., if the remote server URL is incorrect or unreachable.
3. Logging: Include more logging within the function for debugging purposes.
4. Encryption: Consider encrypting the message content for confidentiality and data integrity.
5. Immutable globals: Since `macid` and `HOST_GATEWAY` are used as globals, careful handling is necessary as their value shouldn't be easily manipulated.

11.6.751 `resolve_alpine_dependencies`

Contained in `lib/functions.d/repo-functions.sh`

Function signature: `689d5ee3849151be4f6f15527c700df7d9c8156b9a6e4b876f2b82c3b22f9d9e`

11.6.752 Function overview

The `resolve_alpine_dependencies()` Bash function takes two arguments, `apkindex_file` and `package_name`. Its primary purpose is to parse an `APKINDEX` file corresponding to an Alpine Linux package repository, extract information relating to a specific package and its dependencies from the file, and print the details to standard output. Specifically, it extracts the package name, its version, and its associated dependencies, skipping over any shared library or file dependencies. It then recursively resolves and lists the dependencies for each of these package dependencies.

11.6.753 Technical description

- **Name:** `resolve_alpine_dependencies`
- **Description:** This function reads an APKINDEX file and recursively prints a list of the input package name's dependencies. It checks if the APKINDEX file exists and whether the package exists in the APKINDEX. For each found package it extracts its name and version and if there are dependencies, it recursively calls itself to resolve these dependencies.
- **Globals:** [`hps_log`: Function for logging]
- **Arguments:**
 - `$1`: APKINDEX file path
 - `$2`: Alpine Linux package name
- **Outputs:** Corresponding filename for each resolved dependency package along with error messages related to the package resolution process. Outputs will be printed on stdout.
- **Returns:** It returns '1' in case of errors, like if the APKINDEX file or the package doesn't exist. There's no successful return value ('0'), function output is meant to be captured from stdout.
- **Example usage:** `resolve_alpine_dependencies "/path/to/APKINDEX" "package_name"`

11.6.754 Quality and security recommendations

1. Implement a method for throwing errors instead of returning '1' when any occur.
2. Add error handlers to check whether `$1` and `$2` have been provided when the function is called.
3. Extend error handling to deal with cases in which a dependency package is not found in the APKINDEX file.
4. Increase the robustness of the function by validating APKINDEX content format.
5. Add more logging messages to provide insight into the function's execution and progression, especially during the recursive traversal of dependencies.
6. Use a more efficient and error-proof method than `echo` and `grep` chains for extracting field value.
7. Implement a limit on recursive depth to avoid potential problems with cyclic dependencies.

11.6.755 `rocky_latest_version`

Contained in `lib/functions.d/iso-functions.sh`

Function signature: `f42e139a07133ae36c4e9783c1fbb144e1167e59537bd374c0b346f853c77a3d`

11.6.756 Function overview

The `rocky_latest_version` function is used to fetch the latest version number of the Rocky Linux distribution. It retrieves the listing of the Rocky Linux distribution repository using `curl`, filters, sorts, and echoes out the version number of the latest release.

11.6.757 Technical description

- **name:** `rocky_latest_version`
- **description:** This function retrieves the HTML of the Rocky Linux distribution repository, filters out the version numbers using regular expressions, sorts them in reverse order, and echoes out the latest release version.
- **globals:**
 - `base_url`: the URL of the Rocky Linux distribution repository
 - `html`: stores the html data fetched from the `base_url`
 - `versions`: array holding the version numbers sorted in reverse order (`-Vr` option for “version sort”)
- **arguments:** None
- **outputs:** The latest Rocky Linux version number.
- **returns:**
 - 1 if either fetching the HTML fails, or no version numbers could be extracted from the HTML
 - otherwise, does not explicitly return anything
- **example usage:**

```
latest_version=$(rocky_latest_version)
echo "The latest Rocky Linux version is ${latest_version}"
```

11.6.758 Quality and security recommendations

1. Since this function heavily relies on the format of the HTML file, it may become brittle with changes to the website structure. It would be more reliable to use an official API, if available.
2. Error handling can be improved. Currently, the function returns 1 in case of either connection failure or when no versions are found in the HTML. Considering these two situations could be differentiated for better troubleshooting.
3. For security reasons, consider validating the HTML content before processing, as maliciously-crafted content might lead to unexpected behavior.
4. Add comments to give context to the regular expression used in the `grep` command. This will improve maintainability for developers not familiar with this pattern.

11.6.759 `rtslib_add_lun`

Contained in `lib/host-scripts.d/common.d/iscsi-management.sh`

Function signature: `b7dcfb5ba5fb6cf52d4e8f8990eea725726e0de361be466ec9a246c28e59ba62`

11.6.760 Function Overview

The function `rtslib_add_lun()` is a bash script function that interacts with a Python script. It adds a Logical Unit (LUN) to a specific target in the `rtslib-fb` Python module. A LUN is a logical reference to a portion of a storage subsystem. This function is used in the configuration of iSCSI (Internet Small Computer System Interface) targets, which allows the sharing of storage resources over a network.

11.6.761 Technical Description

- **Name:** `rtslib_add_lun`
- **Description:** This function adds a Logical Unit (LUN) to an iSCSI target in the `rtslib-fb` Python module using a Python script. The Python script checks to see if the target exists by comparing `iqns`. If the target is discovered, the LUN is appended to the list of that target's LUNs.
- **Globals:** [] There are no global variables.
- **Arguments:**
 - `$1`: `remote_host`, it is the name of the remote host.
 - `$2`: `zvol_path`, refers to the path to the zvol (ZFS volume).
- **Outputs:** If the target is not found, the script will print "Target not found".
- **Returns:** The function doesn't directly return anything since it is used to manipulate state in the Python `rtslib-fb` module rather than produce output within the bash script. However, a side effect is that the Python script will cause an exit with status 1 when the target is not found.
- **Example Usage:** `rtslib_add_lun "remotehost" "/path/to/zvol"`

11.6.762 Quality and Security Recommendations

1. The function could benefit from more sophisticated error handling. For instance, adding more explicit checks for whether the command-line arguments `$1` (`remote_host`) and `$2` (`zvol_path`) were provided.
2. This function relies heavily on Python script, it is recommended to preserve dependencies and ensure that the used python modules are kept up-to-date for the overall system's security.
3. It would be prudent to add validation to ensure that the `zvol_path` provided as an argument leads to a real and accessible path.

4. For better security, the function should handle cases where there are multiple iSCSI targets with the same iqn. This could involve updating the process to accept an additional parameter to uniquely identify targets.

11.6.763 `rtslib_check_available`

Contained in `lib/host-scripts.d/common.d/iscsi-management.sh`

Function signature: 600db02b1741924e983f7cfb8291447bd0b965fffa533f6f2b35a4b9ae08b44

11.6.764 Function overview

The `rtslib_check_available` function checks if the `python3 rtlib_fb` module is available. If the module is not available, the function will output a text message, indicating the absence of the module. If the module is found, nothing happens and the function returns 0, which signifies success. The function is executed in the current shell, meaning any changes made happen within the current shell.

11.6.765 Technical description

Name: `rtslib_check_available`

Description: This function checks whether the module `python3-rtlib_fb` is accessible. If the module is unavailable, it outputs a message notifying the user and returns 1. If the module is available, the function returns 0 and no message is output.

Globals: None

Arguments: None

Outputs: If the module `python3-rtlib_fb` is not available, it outputs “`python3-rtlib_fb is not available`”.

Returns: - 1 if the `python3-rtlib_fb` module is not available. - 0 if the module is available.

Example usage:

```
rtlib_check_available
```

```
# Expect output if rtlib_fb isn't present:  
# python3-rtlib_fb is not available
```

11.6.766 Quality and security recommendations

1. Where possible, ensure that all external modules used in the function are up-to-date. This is to mitigate possible security issues that come with using outdated modules.

2. Error handling helps maintain the quality of the code. Check for potential errors and handle them properly. In this case, the function handles the error that arises from the `python3-rtplib_fb` module not being available.
3. Using a standard success/failure status like this function does (returning 0 for success and 1 for failure) is also a good practice. This makes the function's usage easier for other developers or scripts checking for success.
4. Lastly, do thorough testing on various environments because a module available on one machine or in one environment may not be available on another.

11.6.767 `rtplib_create_target`

Contained in `lib/host-scripts.d/common.d/iscsi-management.sh`

Function signature: `05a02a7bcad8e26b1c77c0b01e3325bd5fd95fd6fc531cb13f7285b33082a1df`

11.6.768 Function Overview

The `rtplib_create_target()` function is primarily used to create an iSCSI target. The function establishes the iSCSI target by invoking Python 3 from within a Bash shell, utilizing the `rtplib_fb` Python library with preset attributes and saving them to a file. The iSCSI target is given a name based on the `remote_host` variable and the current date.

11.6.769 Technical Description

- **Name:** `rtplib_create_target()`
- **Description:** This function creates an iSCSI target by using Python 3 and the `rtplib_fb` Python library. The function takes in a remote host as an argument and creates an iSCSI target name determined by the current date and the provided remote host.
- **Globals:** None
- **Arguments:**
 - `$1`: The name of the remote host
- **Outputs:** Either creates an iSCSI target or returns a failure message.
- **Returns:** Nothing
- **Example usage:** `rtplib_create_target 192.168.1.15` will create an iSCSI target related to the provided IP address.

11.6.770 Quality and Security Recommendations

1. Proper Validation: The function should have checks put in place to ensure that the input, in this case, the remote host, is correctly validated and sanitized. This can help prevent any sort of injection or misconfiguration.

2. **Error Handling:** Increase robustness of the function by putting more specific error handling with clearly defined error messages.
3. **Secure attribute setting:** Be mindful of the iSCSI target attributes being set here. For example, this function disables authentication ("authentication", False), which may not be desirable in a secure setup. Check and review these attribute settings to better suit them to your environment.
4. **Use of globals:** This function does not use any global variables which is a good practice in terms of code clarity and avoiding unexpected side effects. However, if any are to be used in the future, they should be carefully managed.
5. **Logging:** Proper logging techniques need to be followed so that any issue can be debugged effectively. In the absence of effective logging, it will be difficult to trace the issue in the case of any failure.

11.6.771 `rtslib_delete_target`

Contained in `lib/host-scripts.d/common.d/iscsi-management.sh`

Function signature: 318b4dea0022039da24b23ab7ac595a5bfd08e12b0ec24fa41fd9d6813b14435

11.6.772 Function Overview

The Bash function `rtslib_delete_target()` is designed to delete a specific target on a remote host with the help of an iSCN (iSCSI Qualified Name). This function takes a remote host as an argument and generates an iSCN with the remote host embedded in it. Later, it uses python3 to implement the SCSI protocols with the help of the `rtslib` library for Python, deletes the required target, and saves the updated configuration to a file. If the target doesn't exist, it will print a message saying the target is not found.

11.6.773 Technical Description

- **Name:** `rtslib_delete_target`
- **Description:** This function deletes an iSCSI target provided by the remote host. It will print a confirmation message if the deletion is successful or a fail message in case the target is not found.
- **Globals:**
 - [VAR: `iqn`]: It describes the iSCSI qualified name dynamically generated using the current year and month, appending the remote host name provided. It helps as the unique identifier for identifying the iSCSI targets.
- **Arguments:**
 -
- **Outputs:** Prints a confirmation message indicating whether the deletion action was successful. It either shows a checkmark with 'Target deleted' or a cross sign with 'Target not found'.
- **Returns:** Nothing.

- **Example usage:** `rtslib_delete_target my_remote_host`

11.6.774 Quality and Security Recommendations

1. Ensure that the remote host's input is correctly sanitized to avoid code injection or other kinds of security issues.
2. The function should have an error handling mechanism for scenarios where the Python code fails to execute.
3. Although the script currently handles targets not found gracefully, it could be further improved to provide more specific error messages.
4. It is essential to have permission checks so that unauthorized users cannot delete targets.
5. Ensure that date command used to generate iqn is correctly referenced in terms of the timezone, as it may create discrepancies due to timezone differences.

11.6.775 `rtslib_list_luns_for_target`

Contained in `lib/host-scripts.d/common.d/iscsi-management.sh`

Function signature: `2e1e32139249bea08ce91bc7a5e64b08a087499b34dbb1b776cbd4e7d369d1e0`

11.6.776 Function overview

The `rtslib_list_luns_for_target` is a Bash function designed to list the Logical Unit Numbers (LUNs) for a specific iSCSI target. It accepts two arguments, generates an iSCSI Qualified Name (IQN) for the target, and then uses a Python script to iterate through the LUNs. Information on each LUN is then outputted, or an error message is displayed.

11.6.777 Technical description

- **Name:** `rtslib_list_luns_for_target`
- **Description:** This bash function lists the LUNs for a particular iSCSI target. It achieves this by launching an embedded Python script which uses the `rtslib_fb` library to enumerate over the LUNs for a target.
- **Globals:** None
- **Arguments:**
 - `$1`: The remote host. Used to construct the iqn.
- **Outputs:** Depending on the results of its operations, the function either outputs each of the LUNs for the specified iSCSI target or an error message when the target is not found.
- **Returns:** Nothing
- **Example Usage:** `rtslib_list_luns_for_target "localhost"`

11.6.778 Quality and security recommendations

1. Ensure that the remote host argument is appropriately sanitized to prevent potential command injection attacks.
2. Consider handling or logging the captured Python exceptions for debugging and traceability.
3. Look into dealing with situations where the Python environment or the required libraries are not available.
4. Assess whether the function behaves as expected if called with an invalid or unreachable host.
5. Consider how the function will respond if date cannot provide the expected output.
6. Check that the user running this script has the necessary privileges for the script's operations.
7. Enforce strict mode in the bash script to minimize potential errors due to uninitialized variables or unhandled errors.

11.6.779 `rtslib_list_targets`

Contained in `lib/host-scripts.d/common.d/iscsi-management.sh`

Function signature: `6e4bf5d3b011e59cd3a4346f13ded587844ffead7fe935f5f5e08c258f5d737e`

11.6.780 Function Overview

The `rtslib_list_targets` function is a Bash function that is used to list all the target World Wide Names (WWNs) in the fabric module “iscsi”. The function uses Python3 code embedded in a Bash function to interact with the RTS (Runtime Simple) library.

11.6.781 Technical Description

- **Name:** `rtslib_list_targets`
- **Description:** This function is designed to list all the target WWNs in the fabric module named “iscsi”. Internally, it uses a small Python3 script that utilizes the `RTSRoot` class from the `rtslib_fb` (RTS Library Full-Blown) module. This class provides an interface to the root of the RTS configuration model. The function iterates through each target of the RTS root, checks if its fabric module's name is “iscsi”, and if so, prints its WWN.
- **Globals:** None.
- **Arguments:** None.
- **Outputs:** Prints the WWNs of all targets in the “iscsi” fabric module to stdout.
- **Returns:** None. Outputs are sent directly to stdout.
- **Example Usage:** To use this function, simply source it in your bash script and call it without any arguments. The function will print the output to stdout, so

you may wish to capture that in a variable or pipe it to another command for further processing: `sh source rtslib_list_targets.sh rtslib_list_targets`

11.6.782 Quality and Security Recommendations

Here are some suggestions to improve the quality and security of the function:

1. **Input Validation:** Even though this function does not accept any arguments, checking for unexpected inputs is always a good habit. Always make sure that the environment in which you are running your script is safe and secure.
2. **Error Handling:** Add error handling to the function to deal with potential issues that may arise, such as failure to import the required Python module or issues accessing the RTS configuration root.
3. **Documentation:** Enhance the code readability by adding more comments within the function explaining what each line or block of code is meant to do.
4. **Security:** This function operates with the RTS configuration model and may have effects on the network configuration. Make sure it is only operative under necessary permissions and not susceptible to unauthorized execution.

11.6.783 rtslib_save_config

Contained in `lib/host-scripts.d/common.d/iscsi-management.sh`

Function signature: `ac265647656df50229187ef098655cc149e59be37bf19e8244f9356fddcace3e`

11.6.784 Function Overview

The `rtslib_save_config` function creates a Python3 child process and runs an inline (here doc) Python script. This script imports the `RTSRoot` class from the `rtslib_fb` module and calls the `save_to_file()` method on an instance of that class. Therefore, using this function writes the current state of the runtime configuration of the target (presumably a storage target/subsystem) to the configuration file.

11.6.785 Technical Description

- **Name:** `rtslib_save_config`
- **Description:** The function calls a Python3 subprocess, importing the `RTSRoot` class from the `rtslib_fb` module within the Python environment. It then invokes the `save_to_file` function of a `RTSRoot` class instance, thus saving the current state of the target's runtime settings into a file.
- **Globals:** None.
- **Arguments:** None.

- **Outputs:** Writes the target's current runtime settings into a file.
- **Returns:** None. If the Python subprocess encounters an error, it will be written to stderr.
- **Example Usage:**

```
rtslib_save_config()
```

11.6.786 Quality and Security Recommendations

1. Ensure proper exception handling is done. While using the Python code, if there is any exception caused due to environment issues or some runtime exceptions, this script would fail. So, appropriate error handling will make the function more robust.
2. Validate Python version: Although the script specifically asks for Python3, there may be differences between minor versions. It's always a good idea to make sure the right version of Python is available.
3. Document the “rtslib_fb” module and its usage: If the person using this function isn't familiar with the “rtslib_fb” module being imported, they could incorrectly use or modify this function. A brief explanation of the module and its usage would be helpful.
4. Use a more secure way to call Python scripts: While Bash is a flexible language, it also can have security implications when running scripts. Always consider the security of your Python scripts when using them in a Bash environment.

11.6.787 script_render_template

Contained in `lib/functions.d/kickstart-functions.sh`

Function signature: `064ddcb49f3688c1b0ede8ce884eae36c9ac96f5117338bdce1f62d5c6960a67`

11.6.788 Function Overview

The `script_render_template()` function is designed to remap all variable placeholders (`@...@`) with their corresponding values (`${...}`). The function cycles through all the variables using `compgen -v`, assigns the value of the variable to a local variable, and then adds it to an array of values for `awk`. The `awk` utility then processes the array of values, replacing each placeholder in the original string with the corresponding variable value.

11.6.789 Technical Description

- **Name:** `script_render_template`

- **Description:** This function is used to remap variable placeholders with their actual value. It does this using the `awk` utility and a `for` loop iterating over all variables in the scope.
- **Globals:** None
- **Arguments:** The function does not require any arguments. It acts on all variables in its scope.
- **Outputs:** The function outputs a string with all `@var@` placeholders replaced with their corresponding `${...}` values.
- **Returns:** The function does not have a return value.
- **Example Usage:** Assume that the following variables are already defined in the context:

```
script_name='MyScript'  
script_version='1.0'
```

If we call `script_render_template` in a context where a template string like `'This is @script_name@ version @script_version@'` is present, the function will output the string `'This is MyScript version 1.0'`.

11.6.790 Quality and Security Recommendations

1. Always make sure the substitution values (`${...}`) are securely obtained and sanitized to prevent command injection attacks.
2. Consider validating the variable names that `compgen -v` produces to ensure they adhere to expected patterns and rules.
3. Beware of potential performance issues if the function is used in a context with a large number of variables.
4. Handle errors and exceptions gracefully. For instance, what should happen if a placeholder variable does not exist?
5. Ensure that the function fits well within your specific use case, as its current implementation is very general and may not be suitable for more specific tasks.

11.6.791 `select_cluster`

Contained in `lib/functions.d/cluster-functions.sh`

Function signature: `d6887af36fa9b9a8372e021a7e7c69665b0d3eee270caacdddd223f760546494`

11.6.792 Function overview

`select_cluster` is a shell function designed to effectively manage multiple clusters within a shell environment. Depending on the argument passed, it returns either the directory or the name of the selected cluster. If the shell is non-interactive, it picks the

active or first cluster depending on different conditions. If the shell is interactive, it prompts the user to select a cluster from a list of available clusters.

11.6.793 Technical description

- **Name:** `select_cluster`
- **Description:** This function is used to manage multiple clusters in a shell environment. It either returns the directory or the name of the active or first available cluster, depending on specific conditions.
- **Globals:** [`HPS_CLUSTER_CONFIG_BASE_DIR`: Base directory for the cluster configuration]
- **Arguments:** [`$1`: Sets the return mode to 'name' if value is '--return=name']
- **Outputs:** Prints either the directory or the name of the selected cluster.
- **Returns:** Returns 1 if no clusters are found. Returns 0 after successfully selecting a cluster in either interactive or non-interactive mode.
- **Example usage:** `bash select_cluster --return=name`

11.6.794 Quality and security recommendations

1. For better script security, ensure all input data is validated and sanitized to mitigate the risk of injection attacks.
2. Always check return values from functions and handle any potential errors appropriately.
3. Make good use of local variables to limit the scope to the current shell and to avoid possible variable name conflicts.
4. Use double quotes around variable references to avoid word splitting and path-name expansion.
5. Regularly update the script and maintain proper documentation for easier debugging and maintenance.

11.6.795 `set_active_cluster`

Contained in `lib/functions.d/cluster-functions.sh`

Function signature: `ff67ed1fe94af1bc0a6ebd9436efd66e00f5b5f9905b2979af037bdf49fce60e`

11.6.796 Function Overview

This function `set_active_cluster` is responsible for defining the active Kubernetes cluster by name, in a specific environment. It accomplishes this through several steps - assigning the input to a variable, checking to see if the input variable is empty, defining cluster directory and configuration file locations, ensuring both the directory and configuration file exist, and finally, linking to the active cluster.

11.6.797 Technical Description

- **Name:** `set_active_cluster`
- **Description:** This function sets a specific Kubernetes cluster as active based on cluster name provided as an argument.
- **Globals:** None
- **Arguments:**
 - \$1: `cluster_name` - Name of the Kubernetes cluster to be set as active.
- **Outputs:**
 - Echoes an error message and usage suggestion to `stderr` if no argument is supplied, or if cluster directory or configuration file cannot be found.
 - Echoes a success message if the active cluster is successfully set.
- **Returns:**
 - Returns 1 if no argument is supplied or if cluster directory or configuration file cannot be found.
 - Returns 2 if clustered configuration not found.
- **Example usage:**

```
set_active_cluster my_cluster_name
```

11.6.798 Quality and Security Recommendations

1. Input validation: Implement more comprehensive input validation; currently, the function only checks if an argument is supplied, but there might be specific naming rules for cluster names that could also be checked.
2. Error handling: More specific error messages could help with easier problem diagnostics.
3. Security: Review if and where this script allows for problems like symlink attacks; if a potential exists, steps should be introduced to minimize the risk.
4. Robustness: Consider introducing checks for edge cases such as file permission issues or lack of disk space.
5. Testing: Incorporate this function in unit testing to ensure it continues to work correctly as the cluster environment evolves.

11.6.799 `storage_deprovision_volume`

Contained in `lib/host-scripts.d/common.d/storage-management.sh`

Function signature: `668ca90d247ae9b85bd28558e102cecbf0dd119fbbf5fca3b61e894cf67082d7`

11.6.800 Function overview

The `storage_deprovision_volume` function is primarily used to deprovision a given volume in storage. It first parses the arguments for IQN and volume name, then ensures that the host type is suitable for deprovisioning. Additionally, it verifies the

existence of a local zpool name. The function subsequently deletes the iSCSI target and the volume based on the provided IQN and volume name. Failure in deletion results in an error log, and successful deprovisioning returns 0.

11.6.801 Technical description

- **Name:** `storage_deprovision_volume`
- **Description:** Parses arguments and deprovisions a specified storage volume by deleting the iSCSI target and the volume itself. Checks for host type and presence of a local zpool name.
- **Globals:** None.
- **Arguments:**
 - \$1: Action flag. Possible flags are `--iqn` (sets internal `iqn` variable) and `--zvol-name` (sets internal `zvol_name` variable)
 - \$2: Corresponding value for the flag set by \$1 (either `iqn` value or `zvol` name depending on \$1).
- **Outputs:** Logs into a remote system information about the steps made by the function. If errors occur, they are reported in the log.
- **Returns:**
 - 1 if invalid flag is set, required flags are not set, host type is not 'SCH', zpool name could not be determined or deletion of zvol fails.
 - 0 if the volume was successfully deprovisioned.
- **Example usage:**

```
storage_deprovision_volume --iqn
↪ iqn.2003-01.org.linux-iscsi.localhost:x8664.sn.d33fadd1d40
↪ --zvol-name zvol1
```

11.6.802 Quality and security recommendations

1. Make sure only authorized users can run this function to ensure that volumes are not accidentally deprovisioned.
2. Enhance error checking to catch unknown flags and handle them appropriately.
3. Consider adding functionality to backup the volume before deletion to allow recovery in case of accidental deprovisioning.
4. Always use secure credentials when logging into the remote system to protect the integrity of the data.
5. Consider hardening the script by restricting the ability to deprovision based on additional factors, such as the time of day or the load on the system.
6. Regularly review and audit the logs produced by the function for any irregular activities.

11.6.803 `storage_get_available_space`

Contained in `lib/host-scripts.d/common.d/storage-management.sh`

Function signature: 0ecd662d000f3348b1348219cc1c7541ead05b93b29de5edfc244de079f38d03

11.6.804 Function overview

The function `storage_get_available_space` is a bash function used to find out the available space in a specific ZPOOL. The function uses ZFS commands to determine the amount of available space in the ZPOOL. It first fetches the ZPOOL name from a remote host and then checks if the ZPOOL name could be fetched successfully. It then queries ZFS for the available space in that ZPOOL and logs an error if this is not successful. Finally, it echos the available space in bytes and returns 0 on successful execution.

11.6.805 Technical description

- **Function name:** `storage_get_available_space`
- **Description:** This function retrieves the available storage space in a specified ZPOOL residing on a remote host. The ZPOOL name is obtained via the `remote_host_variable` function.
- **Globals:** [`ZPOOL_NAME`: the ZPOOL's name in the remote host]
- **Arguments:** None
- **Outputs:**
 1. Logs errors if the ZPOOL name cannot be determined, or if ZFS fails to retrieve available space.
 2. Prints the available space in bytes if the function executes successfully.
- **Returns:**
 1. If the ZPOOL name cannot be determined or ZFS fails to retrieve available space, the function returns 1.
 2. Upon successful execution, the function returns 0.
- **Example usage:** `storage_get_available_space`

11.6.806 Quality and Security Recommendations

1. Ensure that the `remote_host_variable` and `remote_log` functions are secure and cannot be exploited to execute remote shell commands.
2. Make sure that error logs do not reveal sensitive information about the system, which could be utilized by an attacker.
3. Validate the output of the `zfs get` command to ensure it can't be manipulated to inject malicious code.
4. Check the network connection between the local and remote systems securely to prevent MITM (Man in the Middle) attacks.
5. Consider adding more error handling and logs for network issues or ZFS failures.
6. Regularly update the ZFS command-line tools to benefit from the most recent security patches and improvements.

11.6.807 storage_parse_capacity

Contained in `lib/host-scripts.d/common.d/storage-management.sh`

Function signature: `c67957f8ee4be8442088f8824ef3828b5a9b2d67c4b4f352df2b050bf6c8bbbee`

11.6.808 Function overview

The Bash function `storage_parse_capacity()` takes a string argument indicating a data size (with an optional suffix denoting the scale e.g., K, M, G, or T for Kilobytes, Megabytes, Gigabytes, and Terabytes respectively). The function will parse this string and return the data size in bytes. The function will return 1 and terminate if the input does not match the expected format, i.e., a numerical value optionally followed by a suffix (K, M, G, or T).

11.6.809 Technical description

Definition block for `storage_parse_capacity()` function:

- **name:** `storage_parse_capacity()`
- **description:** Parses a string argument depicting a data size with an optional suffix (K, M, G, T) and returns the equivalent size in bytes.
- **globals:** None
- **arguments:**
 - \$1: The capacity string to be parsed. Could be a plain number (considered as bytes), or suffixed with K, M, G, or T (case-insensitive) to denote Kilobytes, Megabytes, Gigabytes, and Terabytes, respectively.
- **outputs:** The function echoes the parsed capacity (data size) in bytes.
- **returns:**
 - 0: if the processing was successful.
 - 1: if the input string is empty or does not match the expected format.
- **example usage:** `$ storage_parse_capacity 20K` would output `20480`.

11.6.810 Quality and security recommendations

1. Add input validation: There should be some additional error handling to make sure the value before the suffix is a valid number. Currently, non-numeric characters before the suffix can lead to unexpected behaviour.
2. Use consistent error: The function should always echo an error message to `stderr` whenever it returns 1. This would make it easier for users to understand any error that the function encountered during its execution.
3. Create unit tests: To ensure that the function consistently works as expected, create some unit tests that will run the function with different inputs and compare the return values with expected results.

11.6.811 storage_provision_volume

Contained in `lib/host-scripts.d/common.d/storage-management.sh`

Function signature: `5e3861e2975c7a31100612e49933846099de90b2882d09056b15392c4698cf6b`

11.6.812 Function Overview

The `storage_provision_volume()` function is designed to automate the process of provisioning a storage volume within a networked storage solution. The function accepts three parameters to define the fully qualified iSCSI Qualified Name (IQN), the storage capacity, and the zvol name of the volume to be created. The function then creates an iSCSI target, or storage resource, that other iSCSI initiators on the network can access.

11.6.813 Technical Description

storage_provision_volume()

- **Description:** Provisions a storage volume within a networked storage solution by creating an iSCSI target.
- **Globals:**
 - `host_type`: checks to verify this is a storage host.
 - `zpool`: gets local zpool name.
- **Arguments:**
 - `--iqn ($2)`: The IQN of the iSCSI.
 - `--capacity ($2)`: Storage size requirement in appropriate units (Byte, Kilobyte, Megabyte, or Gigabyte).
 - `--zvol-name ($2)`: The name of the volume to be created.
- **Outputs:**
 - Validates required parameters.
 - Verifies that this is a storage host.
 - Calculates and reports available space, ensuring enough space exists for the requested volume.
- **Returns:**
 - 1 if an error occurs, such as missing required parameters, incorrect host type, running out of available space, or failure in creating zvol or iSCSI target.
 - 0 if the volume is successfully provisioned.
- **Example usage:**

```
storage_provision_volume --iqn
↪ iqn.2021-05.com.example:storage:disk1 --capacity 1G
↪ --zvol-name disk1
```

11.6.814 Quality and Security recommendations

1. The function should check the validity of the passed arguments which includes checking if the `iqn` and `zvol-name` are properly formatted and if the capacity is realistically feasible before attempting to provision the storage volume.
2. It may be beneficial to have an additional parameter to choose the type of volume to be provisioned—block, file or object—to add versatility to the function.
3. This function could further be improved by providing a rollback mechanism for partial completions, in case the storage provisioning operation fails midway.
4. Always ensure that error messages do not disclose too much information, which might end up being a security risk. For instance, revealing the host type or zpool name could provide useful information to malicious users.
5. To reduce the risk of injection vulnerabilities, ensure that all parameters within the shell command are appropriately escaped or quoted.
6. For a more robust implementation, consider using a try-catch mechanism to handle unexpected errors. This will prevent the script from crashing and provide a more elegant way of logging the error.

11.6.815 `strip_quotes`

Contained in `lib/functions.d/network-functions.sh`

Function signature: `b2892c33de60887f65c7f2b7946fafee0d873041d3a244be57b49c6339bb1cb5`

11.6.816 Function overview

The `strip_quotes` function is designed to process a string and remove leading and trailing quotes. This includes both single (') and double (") quotes. The function achieves this through the use of local bash string manipulations.

11.6.817 Technical description

- **Name:** `strip_quotes`
- **Description:** The function takes one string argument with either leading or trailing or both types of quotes and remove them.
- **Globals:** None
- **Arguments:**
 - `$1: str` This is a string input sent to the function. It is expected to have leading and/or trailing quotes which are to be removed.
- **Outputs:** The function prints the input string with the quotes removed.
- **Returns:** The function does not explicitly return a value, it only outputs the string without quotes via an `echo` command.
- **Example usage:** `strip_quotes "\"Hello World!\""` or `strip_quotes "'Hello World!'"`

11.6.818 Quality and security recommendations

1. As the function does not account for inner quotes within the string, it is advisable to extend its functionality to handle such cases.
2. While executing scripts, always validate/filter the inputs for any malicious content, as appropriate. In this context, the function could be extended to check and validate the input string.
3. Since the function echoes the output, it might lead to risk of command injection attacks. A better way would be to return the value and let the caller decide what to do with the returned value.
4. Provide more descriptive error messages that would explain the problem if the script fails.
5. Increase resilience of the function by handling edge cases such as empty strings or strings without quotes.

11.6.819 `_supervisor_append_once`

Contained in `lib/functions.d/configure-supervisor.sh`

Function signature: `5f79fe5e7a4a7d3eb591b9e371a140f9bd1de60abf5e87d0e47082bfff056191`

11.6.820 Function Overview

The function `_supervisor_append_once` is designed to modify the settings for Supervisor, a control system for UNIX-based servers. It takes two arguments—`stanza` and `block`. The function checks the supervisor configuration file, which is determined by the environment variable `${SUPERVISORD_CONF}`. If the stated `stanza` is not present in the configuration file, it then appends the corresponding `block` at the end of the file. This function is used multiple times subsequently in the script to ensure Supervisor is configured to securely and efficiently manage various service programs.

11.6.821 Technical Description

- **Name:** `_supervisor_append_once`
- **Description:** This function appends configuration blocks to the Supervisor configuration file for a specified service program, but only if that block does not already exist. It is used in managing UNIX based servers.
- **Globals:** [`${SUPERVISORD_CONF}`]: Path to the Supervisor configuration file]
- **Arguments:**
 - `$1`: `stanza`— the name of the service program, e.g. `program:nginx`
 - `$2`: `block`— the complete configuration string to be appended.
- **Outputs:** Appends a configuration block to the Supervisor configuration file for a specific service program.
- **Returns:** Nothing.

- **Example Usage:** `_supervisor_append_once "program:dnsmasq"`
`"$(cat <<EOF`
 - This usage of the function appends dnsmasq related configurations to the Supervisor configuration file.

11.6.822 Quality and Security Recommendations

1. Enhance error handling to prevent the potential mishandling of non-existent files.
2. Introduce a mechanism to backup the original configuration file before making changes to it.
3. Expand on arguments validation by checking for empty or null arguments.
4. Since shell scripts are susceptible to injection attacks, consider potential sanitization steps for the `stanza` and `block` variables.
5. Avoid hard coding paths and consider converting them into arguments or environment variables to improve scalability.

11.6.823 `sync_alpine_packages`

Contained in `lib/functions.d/repo-functions.sh`

Function signature: `b2fb9504a3d2f06000545f9301edac370a296cec8ec2e779dd1d02496acfeafc`

11.6.824 Function Overview

The function `sync_alpine_packages` is used to synchronize the Alpine software packages. It takes the version of Alpine, the version of the mirror, the name of the repository, and a set of package names as inputs. It locates the desired Alpine packages and their dependencies from the specified repository, downloads them, and finally places all the files at the desired destination. If any issues happen during the process such as failing to create directory, download or extract `APKINDEX`, or resolve dependencies, the function handles errors and cleans up temporary storage to maintain a safe state.

11.6.825 Technical Description

- **Name:** `sync_alpine_packages`
- **Description:** This function synchronizes software packages of given package names from specified Alpine Linux mirror and repository.
- **Globals:**
 - `HPS_DISTROS_DIR`: The base directory for the distros.
- **Arguments:**
 - `$1`: Alpine version.
 - `$2`: Mirror version.
 - `$3`: Repo name.
 - Rest arguments: Package names.

- **Outputs:** Downloaded packages at `${HPS_DISTROS_DIR}/alpine-${alpine_version}/apks/${rep`
- **Returns:**
 - 0: On successful completion.
 - 2: If any error occurred during package synchronization.
- **Example Usage:** `sync_alpine_packages 3.12 3.12 main bash curl`

11.6.826 Quality and Security Recommendations

1. Always check if required arguments are passed before executing the function.
2. All variables inside the function should be local to avoid potential naming collisions.
3. Validate user inputs to prevent potential command injection attacks.
4. Handle all error paths gracefully, such as download failures or file system limitations.
5. Use more secure methods to download packages, if available, such as HTTPS instead of HTTP.
6. Consider verifying packages' authenticity and ensuring data integrity through checksums or digital signatures.
7. Document return codes of a function and their meanings for easier debugging and maintenance.
8. Cleaning up temporary storage areas is a good practice but also consider a more robust cleanup mechanism in case of unexpected termination.

11.6.827 `sync_alpine_repo_arch`

Contained in `lib/functions.d/repo-functions.sh`

Function signature: `1e50c32388b6c1ff2be449bd1f28c57fd1d94b6fd34b27c243dba0e74d19f4e4`

11.6.828 Function Overview

The function `sync_alpine_repo_arch` is designed for synchronizing alpine repositories to a local directory. It uses parameters such as alpine version, mirror version, repository name, and architecture to determine the repository to sync. The destination directory is built using the environmental variable and provided parameters, defaulting to `x86_64` architecture if no other is provided. This function includes numerous error checks such as ensuring the destination directory can be created, and the available disk space is sufficient for downloads.

11.6.829 Technical description

- **Name:** `sync_alpine_repo_arch`

- **Description:** This function synchronizes an Alpine Linux package repository to a local directory. It retrieves the file list from the repository, checks if there's enough disk space, then uses `wget` to download the packages.
- **Globals:** [`HPS_DISTROS_DIR`: The base directory for syncing distros]
- **Arguments:**
 - [\$1: Alpine Linux version (e.g., `v3.12`)]
 - [\$2: Version of the mirror (e.g., `edge`)]
 - [\$3: Name of the package repository (e.g., `main`)]
 - [\$4: Architecture (e.g., `x86_64`), optional, default `x86_64`]
- **Outputs:** Logs info, debug, and error messages. Stores downloaded files in the `dest_dir` defined in the function.
- **Returns:** 2 if an error occurred that prevented successful function execution, return value of the `validate_apkindex` function otherwise.
- **Example usage:** `sync_alpine_repo_arch v3.12 edge main x86_64`

11.6.830 Quality and Security Recommendations

1. Enhance error handling: Presently, the function returns 2 when an error occurs, thus doesn't differentiate between failures. It would be beneficial to have unique error codes for the different failure points.
2. Download verification: After each download, consider adding functionality that verifies the authenticity and integrity of the downloaded files. This not only helps in ensuring that the complete file has been downloaded but also safeguards against any potential security concerns.
3. User input validation: Validate user inputs earlier in the function to prevent possible misuse, such as directory traversal.
4. Secure Temporary file handling: Ensure the security of temporary files by setting appropriate permissions and minimizing their lifespan.
5. Implement error retries: Temporary issues may hamper the download process, therefore implementing reparative measures such as retries with backoff periods could increase reliability.
6. Insert proper cleanup actions: Even if the function fails at any point, it should clean up any temporary files created during its execution. Implementing a trap for `EXIT` signal could be beneficial for this task.

11.6.831 `sync_alpine_repository`

Contained in `lib/functions.d/repo-functions.sh`

Function signature: `d834802852064ff82cad597b7d94e4f49698aa48de1ec1ef61daaaab11190357`

11.6.832 Function overview

The `sync_alpine_repository` function is a Bash shell script used to synchronize Alpine Linux repositories to your local system. The script is characteristic of version management, supporting specific versions of the Alpine Linux distro and selective synchronization based on the given mode.

11.6.833 Technical description

- **Name:** `sync_alpine_repository`
- **Description:** Synchronizes an Alpine Linux repository version to the local system. The function supports selective sync through different modes: `all`, `minimal`, `main`, `community` and `packages`.
- **Globals:** [`HPS_DISTROS_DIR`: The directory where the Alpine Linux distributions are stored]
- **Arguments:** [`$1`: `alpine_version`, `$2`: `sync_mode`, `$3`: `package_list`, `$4`: `arch` (optional, defaults to `x86_64`)]
- **Outputs:** The status of the synchronization process is logged to `stdout`.
- **Returns:** If successful, it returns 0; if unsuccessful, it returns 1 or 2 based on the failure's nature.
- **Example usage:** `sync_alpine_repository 3.12 all`

11.6.834 Quality and security recommendations

1. Implement error checking for other global variables as done for `HPS_DISTROS_DIR`.
2. Soft-code the hard-coded array that determines the repositories' mapping for the `minimal` sync mode. This allows more flexibility and efficiency.
3. Consider handling command line arguments in a more fault-tolerant way, e.g., by using `getopts`.
4. Improve security by isolation. Run the script within a sandbox or container or use a user with limited privileges to execute the script to limit the potential impact of malicious or erroneous scripts operating on unauthorized files/directories.

11.6.835 `tch_apkovl_create`

Contained in `lib/functions.d/tch-build.sh`

Function signature: `a007a0818596cb39af8229e52a0b9936c78e7463a081efc4d73d2c4067076a5c`

11.6.836 Function Overview

The `tch_apkovl_create()` is a Bash function responsible for creating an Alpine Linux `apkovl` (Alpine local backup) file. The function performs this task by obtaining

important configuration information first (like the gateway IP, Alpine version, and name server), logging relevant activities and potential issues/errors, and then building the necessary components of the apkowl. Upon successful creation, information is logged, and a tarball archive is created. If any step fails, it performs necessary cleanups and returns an error status of 1.

11.6.837 Technical Description

- **Name:** `tch_apkowl_create()`
- **Description:** Bash function that creates an Alpine Linux apkowl file.
- **Globals:** None.
- **Arguments:**
 - `$1: output_file`: Represents the name/path of the output file of the tarball archive that will be created.
- **Outputs:** Logging messages to the console regarding the progression of apkowl creation or errors if they occur.
- **Returns:** Function returns 0 if the apkowl and tarball creation is successful or 1 if it encounters an error.
- **Example Usage:**

```
tch_apkowl_create "my_archive.tar.gz"
```

11.6.838 Quality and Security Recommendations

1. **Error handling:** The function already does a good job of returning 1 when an error is encountered, but it might be beneficial to also return distinct error codes for different types of errors.
2. **Input validation:** The function does not currently validate the input `$1`. It should check that the output file path provided as an argument is both valid and writable before attempting to create a tar archive.
3. **Temporary directory:** The usage of `tmp_dir` can be better managed with a trap on EXIT signal to clean it up instead of cleaning in several places.
4. **Security:** Ensure that the files which are written to `tmp_dir` have proper file permissions set, to avoid unauthorized access.
5. **Code Commenting:** Commenting is important to understand the function, variables, and logic used in the code. Detailed comments explaining complex parts of the function would improve readability and maintainability.

11.6.839 `tch_configure_alpine`

Contained in `lib/host-scripts.d/alpine.d/tch-configure.sh`

Function signature: `04d22f41c2428713b77df149ab5bd5ab40f9e84f845cf0bdc82a7b0437b1dfb6`

11.6.840 Function overview

The function `tch_configure_alpine` is a bash function used for printing configuration messages. The function is named to suggest that it is meant to configure something related to 'TCH' on an 'Alpine' system, yet it currently serves as a placeholder giving informative echoes about the start and completion of TCH configuration. At the end it returns a 0, indicating a successful execution.

11.6.841 Technical description

- **Name:** `tch_configure_alpine`
- **Description:** This function, as of current state, acts as a placeholder for configuring TCH on an Alpine system by printing messages regarding the commencement and completion of the configuration process.
- **Globals:** None
- **Arguments:** This function does not accept any arguments.
- **Outputs:** The function outputs two lines of text indicating the start and completion of the supposed configuration.
 - [HPS] `tch_configure_alpine`: Starting TCH configuration
 - [HPS] `tch_configure_alpine`: Placeholder - configuration complete
- **Returns:** The function returns a 0, indicating successful execution.
- **Example Usage:** `tch_configure_alpine`

11.6.842 Quality and security recommendations

1. Provide a meaningful implementation: Currently, the function simply outputs some text and does not perform any configuration tasks. The function should be improved by adding meaningful logic that accomplishes the intended task of configuring TCH on an Alpine system.
2. Validation: If the function is updated to accept parameters, ensure that appropriate validation checks are in place to secure against injections or misuse.
3. Logging: Instead of using `echo` for logging, use a reliable logging mechanism to record important information, errors, warnings etc.
4. Error Handling: Implement error handling mechanisms to anticipate the unexpected and handle any exceptions in a graceful manner.
5. Comments: Provide comprehensive comments for complex coding blocks to ensure maintainability and easy understanding of the code by other developers.
6. Secure Return Status: Choose correct status codes based on the function's execution, rather than simply returning 0, to make the function more reliable.

11.6.843 `ui_clear_screen`

Contained in `lib/functions.d/cli-ui.sh`

Function signature: 0493ca3490c6e95475fb08d2f387298f827857616ec67718e0dd7bc3268a31d2

11.6.844 Function Overview

The `ui_clear_screen()` function is a simple Bash utility function used to clear the terminal screen. If the `clear` command is available and successful, it will be used. Otherwise, it will fall back to outputting an escape sequence that should also perform the same function, this is especially handy in various system environments where the `clear` command may not be available.

11.6.845 Technical Description

- Name: `ui_clear_screen`
- Description: This function is responsible for clearing the terminal screen. It first tries to use the `clear` command and falls back to an escape sequence (`\033c`) if `clear` is not successful.
- Globals: Not applicable as this function does not use or modify any global variables.
- Arguments: Not applicable as this function does not take any arguments.
- Outputs: A cleared terminal screen.
- Returns: If the `clear` command is successful, this function will return the exit status of the `clear` command which is expected to be 0 indicating success. If the `clear` command fails (non-zero exit status), 0 will be returned after printing the escape sequence.
- Example Usage: `bash ui_clear_screen()`

11.6.846 Quality and Security Recommendations

1. Always ensure that functions correctly handle unexpected or erroneous input. For `ui_clear_screen()`, that's not necessary as it doesn't take any arguments.
2. The function doesn't provide or log any error messages which might be useful in some scenarios to debug issues related to the terminal or environment. You could consider adding error logging.
3. Evaluate the fallback method of using an escape sequence for screen clearing, it might not work or could lead to unexpected results in certain terminal environments. A more robust way of handling the unlikely event of `clear` failing could be devised.
4. Make sure to check the return values of commands you are running within your functions, incorporate error checking mechanisms wherever possible.
5. Consider outputting a warning or notice to the user when the fallback method is employed letting them know `clear` command wasn't successful.
6. This function should be safe from code injection as it does not process external input or environmental variables.

11.6.847 ui_clear_screen

Contained in `lib/functions.d/cli-ui.sh`

Function signature: 0493ca3490c6e95475fb08d2f387298f827857616ec67718e0dd7bc3268a31d2

11.6.848 Function Overview

The `ui_clear_screen` function in Bash is simply meant to clear any terminal output previously written. It firstly tries to use the `clear` command, but if that fails, it uses an ASCII escape sequence to clear the terminal.

11.6.849 Technical Description

- **Name:** `ui_clear_screen`
- **Description:** The function clears the terminal screen.
- **Globals:** None.
- **Arguments:** No arguments needed.
- **Outputs:** None.
- **Returns:** 0 if the screen is cleared successfully. If not successfully, the corresponding error code will be returned.
- **Example usage:** To use this function within your Bash script, you simply need to call it as `ui_clear_screen`. No arguments need to be passed.

```
#!/bin/bash
```

```
# ... Some commands ...
```

```
ui_clear_screen
```

```
# ... More commands ...
```

This will clear the terminal at that point in the script's execution.

11.6.850 Quality and Security Recommendations

1. It is recommended to validate the calling environment before executing commands that manipulate the terminal directly.
2. Excessive usage of clear screen function can lead to loss of important information in terminal history. Make sure to use this function only when required.
3. It is recommended to implement error handling mechanisms (like checking the last command return status) to provide a smooth and secure functionality.
4. Consider using `echo` with newline characters for small amounts of screen clearances instead of using terminal-specific commands. This would improve compatibility with different terminal emulators.

5. Where possible, do not depend on specific environment variables. This function should fail gracefully, with well-defined behaviors, in the absence of environment variables.
6. This function is reliant on system-level commands (`clear` and ASCII escape sequence). Ensure the system these commands are being invoked on supports them to avoid any exceptions or errors.

11.6.851 `__ui_log`

Contained in `lib/functions.d/cli-ui.sh`

Function signature: `d3d49681e2b691a5ff908ab2cfc80feb43c6f2c7f2aa6c60521377957fc9244b`

11.6.852 Function Overview

The function `__ui_log` in Bash is a utility for logging or displaying messages with a distinctive prepend label. It can be utilized to show system-level or user interface related informational, warning, or error messages. The function outputs the passed messages to the standard error stream.

11.6.853 Technical Description

- **Name:** `__ui_log`
- **Description:** This is a logging function used to display messages with a distinctive label '[UI]'. It's mainly intended for outputting system or user interface messages. The function employs the `echo` command to output the messages, which are passed in as arguments.
- **Globals:** None
- **Arguments:** `$*` - A list of arguments to be logged or displayed. They're concatenated into a single string by the `echo` command.
- **Outputs:** The function redirects its output to the standard error output stream (`>&2`). Hence, any message passed to this function would appear in the `stderr` stream, with '[UI]' as a prefix.
- **Returns:** None. The function doesn't explicitly return a value.
- **Example Usage:**

```
# Example to log a message
__ui_log "This is a UI log message"
```

11.6.854 Quality and Security Recommendations

1. For added clarity and readability, it could be beneficial to add comments within the function to explain what each component does.

2. The function might be enhanced with the addition of explicit return values, aiding in error handling and flow control in scripts using this function.
3. To prevent command injection attacks, ensure that all variable data is properly escaped before it is included in the log message.
4. Implement a mechanism to control the log level. Thereby, control what kind of messages (debug, info, warning, error) should be directed to the output.
5. Always make sure that no sensitive data is logged in order to maintain information security and privacy.
6. Consider directing the logs to a specific file or a log management system, for easier troubleshooting and better performance in large systems.

11.6.855 __ui_log

Contained in `lib/functions.d/cli-ui.sh`

Function signature: `d3d49681e2b691a5ff908ab2cfc80feb43c6f2c7f2aa6c60521377957fc9244b`

11.6.856 Function Overview

The `__ui_log` function is a Bash function designed for logging and debugging purposes. The function is prefixed by two underscores, a convention sometimes used to indicate a function is ‘private’ or expected not to be used directly by users or other scripts. This function is responsible for echoing the provided arguments `"$*"` to the standard error (`>&2`), prefixed with the string `"[UI]"`. It provides a clean and easy way to standardize the logging output format and centralize the logging functionality.

11.6.857 Technical Description

- **Name:** `__ui_log`
- **Description:** This function echoes its input arguments to the standard error, with a `[UI]` prefix, which could indicate that this log is related to user interface operations.
- **Globals:** None.
- **Arguments:**
 - `$*`: All positional parameters joined together. These are the messages to be logged by this function.
- **Outputs:** The output is a string formatted as `[UI] your input` where `your input` is the value of the provided arguments. Note, output is always to standard error.
- **Returns:** Does not return a value, only produces output to standard error.
- **Example Usage:**

```
__ui_log 'This is a test log.'  
# Output: "[UI] This is a test log."
```

11.6.858 Quality and Security Recommendations

1. The current function uses the `$*` variable to capture all input arguments. However, it's generally safe to use `$@` in quotes ("`$@"`") instead of `$*` to prevent argument splitting or word splitting.
2. The function echoes the messages to `stderr`, which is a great practice to distinguish between error messages and regular output. However, consider providing a provision for users to redirect the error messages to a file to keep the console clean and review the logs later.
3. Since this function is used for logging, it may be beneficial to add a timestamp to the logged messages for better traceability of the events.
4. Although the function's name starts with two underscores, indicating it may be a 'private' function, Bash does not have true private functions. Therefore, clearly document the intended use of this function.
5. If this function is going to deal with sensitive information, make sure to sanitize or mask the data before logging it to prevent leaking sensitive information.

11.6.859 `ui_menu_select`

Contained in `lib/functions.d/cli-ui.sh`

Function signature: `228d39d076d4308652684c61dd46d71781022466ed70155a37a899acdc69da71`

11.6.860 Function Overview

The `ui_menu_select()` function in Bash is used to create a user interface menu that accepts input from the user and presents a list of selectable options. On selection of a valid option, the function outputs the chosen value and gracefully exits. In case of an invalid selection, it prompts the user for a valid selection until a valid option is selected.

11.6.861 Technical Description

- **Name:** `ui_menu_select`
- **Description:** A bash function that prints a list of options (menu) to the console, takes user input, and validates the input. If the input is valid, it prints the selected option and returns. If the input is invalid, it asks for a new input from the user.
- **Globals:** None
- **Arguments:**
 - `$1`: The prompt string for the UI menu.
 - `$@`: An array containing the selectable options for the UI menu.
- **Outputs:** Prints to `stdout` either the prompt and options for the UI menu, the selected option on valid input, or an error message on invalid input.
- **Returns:** Returns 0 on success, no explicit return on failure.
- **Example Usage:**

```
options=("option1" "option2" "option3")
ui_menu_select "Please select an option:" "${options[@]}"
```

11.6.862 Quality and Security Recommendations

1. Add checks to validate the input arguments—especially check if the supplied options are not empty.
2. Handle signal interrupts for better robustness.
3. It's generally a good practice to avoid the use of `echo -n` since its behavior might be different across different systems.
4. Sanitize error messages to avoid misleading information or potential injection vulnerabilities.
5. Consider adding a timeout for user input to prevent potential denial of service if the script is being run as a server-side script.
6. Use `unset` to destroy variables that store sensitive data after their use to prevent unintentional exposure or leakage of such data.
7. Exit with an error code on failure instead of just printing an error message to indicate error status to the calling script or function.

11.6.863 `ui_pause`

Contained in `lib/functions.d/cli-ui.sh`

Function signature: `1636563cd29eae7fb011743bbb84e1bd4b67567168166cfc3cd0cf46472383ec`

11.6.864 Function overview

The Bash function `ui_pause()` is designed to introduce a pause in the script execution, requiring a user intervention to continue. This behavior is accomplished by using the `read` command with the `-rp` option, which reads a line from standard input and prompts with a message until input is received.

11.6.865 Technical description

- **Name:** `ui_pause`
- **Description:** This function uses the `read` command to pause the execution of a script and display a prompt to the user, requesting them to press the [Enter] key to continue.
- **Globals:** None.
- **Arguments:** None.
- **Outputs:** Outputs a message to the user prompting them to press the [Enter] key.
- **Returns:** Does not return a value.
- **Example Usage:**

```
ui_pause
```

```
# The script will pause here until the user presses [Enter].
```

11.6.866 Quality and security recommendations

1. Lack of user input validation: In this function, any key strike will be interpreted as a signal to continue execution. It would be advisable to include some level of user input validation to ensure that only the specific [Enter] key press is given the ability to continue.
2. Error handling: Unexpected errors or exceptions during the function execution are not accounted for, introducing the potential for unexpected behavior and script crashes. A recommended improvement would be to include error handling mechanisms within the function code.
3. Usability: The current function prints a static message which may be unclear to some users or not applicable in all scenarios. Allowing customization of the pause message could improve usability.
4. Return value: Even though this function does not need to return a specific value, for consistency with other Bash functions, it may be beneficial to return a success status (0) after successful execution.
5. Security: As this function does not make use of any user-provided data or values, there are no apparent security risks. However, it's always a good practice to keep security in mind and follow safe coding practices throughout.

11.6.867 ui_pause

Contained in `lib/functions.d/cli-ui.sh`

Function signature: 1636563cd29eae7fb011743bbb84e1bd4b67567168166cfc3cd0cf46472383ec

11.6.868 1. Function overview

The `ui_pause` function is aimed to provide a ways of adding a pause or break point in the execution of the script. This function doesn't take any parameters. When the function is called, the script will stop at the place it has been called and will not continue until the user manually presses the [Enter] key.

11.6.869 2. Technical description

- Name: `ui_pause`
- Description: The function is used to create a pause in the execution of your Bash script until the user manually presses the [Enter] key. This serves as a beneficial breakpoint utility for troubleshooting or pacing the script execution.
- Globals: None
- Arguments: None

- Outputs: “Press [Enter] to continue...” prompting the user to manually press [Enter].
- Returns: None. The function doesn’t return any value. It waits for the user input (pressing [Enter]) to continue the script execution.
- Example usage:

```
echo "This will display first."  
ui_pause  
echo "This will display after the user presses [Enter]."
```

In this example, after printing “This will display first.”, the script will pause until the user presses [Enter]. After that, “This will display after the user presses [Enter].” will be printed.

11.6.870 3. Quality and security recommendations

1. For security purposes, be cautious while using user input in your scripts to prevent any potential injection attacks.
2. Always validate and sanitize user input, this helps to improve code quality and security.
3. In order to ensure the response from user is indeed an [Enter] key, check ASCII value of the user input.
4. Ensuring that written scripts are easily readable and maintain efficiency, as this function can be used throughout, creating a standalone utility function in a different bash file for such repetitive tasks can be beneficial.
5. Always ensure to comment and document your code properly, to ensure others can understand and maintain it.

11.6.871 ui_print_header

Contained in `lib/functions.d/cli-ui.sh`

Function signature: `57e2fede290b133fd0e31c859a63c119ad15a0eea0326adcc61d2c65983dfecc`

11.6.872 Function Overview

The `ui_print_header()` function in Bash is a utility function built for printing headers in terminal-based user interfaces. The function itself is quite straightforward - it accepts a string as an argument which is then displayed as a title within a border of equals signs (=) before and after the title.

11.6.873 Technical Description

Name: `ui_print_header()`

Description: This function prints a passed title surrounded by a set of equals signs (=) on the lines before and after the title. This creates a clear and visually distinct header within

a terminal or console.

Globals: None

Arguments:

- \$1: `title` - This is a string placeholder that the function expects. This value is what the function will print out as the header text.

Outputs: This function prints to stdout. The printout consists of an empty line, then a line of equals signs, followed by the title, another line of equals signs, and finally, another empty line.

Returns: Returns nothing.

Example Usage: `ui_print_header "Welcome to My Program"` - Will print:

```
=====
Welcome to My Program
=====
```

11.6.874 Quality and Security Recommendations

1. Be aware that there are no sanity checks on the supplied argument. The function will print whatever is supplied as an argument, making it susceptible to potentially handling unexpected or rogue inputs. Therefore, consider validating or sanitizing the input on a higher level of function call.
2. This function does not check the length of input strings. If an excessively long string is supplied as an argument it can lead to inconsistent formatting and potentially unreadable headers.
3. The function doesn't use the locale settings to determine the orientation of the symbols. This may cause issues when it is used in locales that use right-to-left writing systems.
4. As the function doesn't return anything it would not be suitable for scenarios where error handling or feedback would be required based on the output of the function.

11.6.875 `ui_print_header`

Contained in `lib/functions.d/cli-ui.sh`

Function signature: `57e2fede290b133fd0e31c859a63c119ad15a0eea0326adcc61d2c65983dfecc`

11.6.876 Function Overview

The function `ui_print_header()` is used to print a shell output that presents a sort of header, divided by lines of equals ("`=`") signs, with a given title text, provided as an argument, displayed centrally. The function makes use of the `echo` command to print the header to the terminal. The header has a consistent format which makes the console log outputs legible and organized.

11.6.877 Technical Description

- **name:** `ui_print_header`
- **description:** A simple bash function which prints out a header text encapsulated in rows of equal signs. This function can be used for making terminal outputs more clear and readable.
- **globals:** N/A
- **arguments:**
 - `$1: title` The text that will appear as the title of the header.
- **outputs:** The function outputs a header in the following format:

```
=====
      title
=====
```

- **returns:** N/A
- **example usage:**

```
ui_print_header "Initialization Process"
```

11.6.878 Quality and Security Recommendations

1. Always pass a non-empty string as the `title` parameter to avoid generating headers with no titles.
2. Keep the usage of the function within safe environments since it directly outputs to the terminal without any security check.
3. For better readability, avoid using very long strings as `title`. The title should be concise and clear.
4. The function doesn't handle errors or exceptions, so ensure the inputs are well formatted and appropriate.
5. In order to increase function versatility, consider implementing additional formatting parameters, such as underline, bold, or the inclusion of timestamp.
6. Considering globally specifying a maximum length for the `title` to keep all headers uniform and distinct.

11.6.879 `ui_prompt_text`

Contained in `lib/functions.d/cli-ui.sh`

Function signature: `698fb0f6a52fc5fb7d346cb2755ab85d4e44cef66d1ac20f8f40870457b2970a`

11.6.880 Function Overview

The `ui_prompt_text` function in Bash is designed for presenting an interactive prompt to users. This function receives two arguments - a prompt message and a default value. The function first displays the prompt, afterwards if the default value is nonempty,

it is also displayed in brackets. A colon and a space character are appended to prepare the text for an expected user input. The user's input is read and stored in a variable, `result`. In case of no user input, the default value is returned, otherwise the user's input is returned.

11.6.881 Technical Description

Feature	Details
Name	<code>ui_prompt_text</code>
Description	A Bash function to prompt users for input with the option of a default response.
Globals	None
Arguments	<code>\$1</code> (prompt): The message prompt to present to the user. <code>\$2</code> (default): The default value that will be used in case of absence of user input.
Outputs	Prompts the user with a message and optional default value.
Returns	The user's input if provided, otherwise the default value.
Example Usage	<pre>ui_prompt_text "Please enter your name" "John Doe"</pre>

11.6.882 Quality and Security Recommendations

1. Always use `read -r` to prevent interpreting backslashes as escape characters.
2. Beware of potential security risks of command injection if the result is used in further commands without sanitization.
3. You should always quote your variable substitutions like so: `"$var"`. This is to prevent issues with multi-word strings.
4. Remember to initialize local Bash variables. This can help avoid problems if there's a global variable with the same name.
5. Provide clear and user-friendly prompts to facilitate the operation for end users.
6. Default values should be carefully chosen to prevent problems in case of user misuse or misunderstanding.
7. When handling sensitive data, ensure that input is hidden or obscured to protect it from unauthorized access or exposure.

11.6.883 ui_prompt_yesno

Contained in `lib/functions.d/cli-ui.sh`

Function signature: `5b56e1af7169ad6721f24f1f595cee61d779f0b22963936d53073e284b177e9c`

11.6.884 Function Overview

The `ui_prompt_yesno` function in Bash is a user interface function that prompts the user with a yes/no question. It loops until the user provides a valid response. The first argument is the prompt text and the second optional argument sets the default answer.

11.6.885 Technical Description

- **Name:** `ui_prompt_yesno`
- **Description:** This function prompts the user with a question and loops till it gets a valid “yes” or “no” answer from the user. It ensures the users interaction in a script where a binary input is necessary for further execution.
- **Globals:** No global variables are used.
- **Arguments:**
 - \$1: This is the prompt text that is to be displayed to the user.
 - \$2: This optional argument specifies the default answer.
- **Outputs:** Outputs the prompt question with an optional default value and user response.
- **Returns:**
 - Returns 0 if the response is “yes”.
 - Returns 1 if the response is “no”.
- **Example Usage:** If we need user’s confirmation for proceeding further, we can use the function as follows:

```
ui_prompt_yesno "Do you want to continue?" "n"
if [ $? == 0 ]
then
    echo "User wants to continue"
else
    echo "User doesn't want to continue"
fi
```

11.6.886 Quality and Security Recommendations

1. Validate that \$1 (the prompt text) exists and is non-empty before proceeding.

2. Regularly audit and update third-party dependencies to keep them up to date with the latest security updates.
3. Avoid using raw user inputs without validation. In this case though, the input is controlled to “y” or “n” only.
4. Use case-insensitive matching to allow inputs as ‘Y’, ‘n’, etc.
5. Provide more detailed information about valid inputs for prompt. Do not assume users know they need to enter ‘y’ or ‘n’.
6. Test the function rigorously with a variety of different inputs and edge cases to ensure it handles those correctly.
7. Consider implementing a limit on the number of invalid attempts before automatically selecting the default option.
8. Follow secure code practices and maintain a regular schedule for reviewing/updating the function.

11.6.887 unmount_distro_iso

Contained in `lib/functions.d/iso-functions.sh`

Function signature: 8736c022aa1702dff2efeb53cc382894bbe29e4a44b302d6fe6c2f67439871c6

11.6.888 Function Overview

The function `unmount_distro_iso` is used to unmount a distribution ISO. Given the string name of a distribution, it attempts to unmount the ISO if it is currently mounted. The ISOs are stored in the directory specified by the global variable `HPS_DISTROS_DIR`. If the ISO is not mounted, it returns 0, otherwise it logs the attempts to unmount the ISO and returns 0 if it is successfully unmounted and 1 otherwise.

11.6.889 Technical Description

- **name:** `unmount_distro_iso`
- **description:** Attempts to unmount a given distribution ISO.
- **globals:**
 - `HPS_DISTROS_DIR`: The directory of where the ISO’s are stored.
- **arguments:**
 - `$1`: string name of the distribution (`DISTRO_STRING`)
 - `$2`: none
- **outputs:** Logs info level messages documenting the unmounting attempt.
- **returns:**
 - Returns 0 if the ISO is not mounted or successful in unmounting.
 - Returns 1 if it fails to unmount the ISO.
- **example usage:** `unmount_distro_iso ubuntu`

11.6.890 Quality and Security Recommendations

1. Add error checking for the input arguments. Currently the function assumes that the input will be correct and does not check if the distribution string is empty or if it contains illegal characters.
2. Improve logging. On failure, consider logging the error message output from the mount command.
3. Use absolute paths. Relying on relative paths can lead to unexpected behaviour, depending on the working directory when the script is executed.
4. Check that HPS_DISTROS_DIR value is set before using it to avoid errors.
5. Make sure that appropriate permissions and ownership settings are in place. This can prevent unauthorized access or modifications.

11.6.891 update_distro_iso

Contained in `lib/functions.d/iso-functions.sh`

Function signature: `2cc9510e14e685d9bd46ee12b37cf92d19a872fbfd40ed5f9884b79c76dc02d6`

11.6.892 Function overview

The `update_distro_iso` function is a Bash function used to update an ISO image of a given Linux distribution. This function takes a single string argument `DISTRO_STRING` representing the distribution to update. It then unmounts the respective distribution ISO, prompts users to update the ISO file manually, checks if the ISO file exists and is properly updated, and finally re-mounts the ISO. If any step fails, the function returns with an exit code of 1.

11.6.893 Technical description

- **Name:** `update_distro_iso`
- **Description:** This bash function is designed to update the ISO image of a desired Linux distribution. It unmounts the currently mounted ISO, prompts user to manually update the ISO file, verifies if the updated ISO file exists, and remounts the updated ISO. If any of these steps fails, it outputs an error message and exits with a 1 status.
- **Globals:** [`HPS_DISTROS_DIR`: Directory path where the Linux distributions ISO are stored]
- **Arguments:** [`$1`: `DISTRO_STRING`: String representing the Linux distribution to be updated]
- **Outputs:** Messages detailing the status and results of each operation, including potential errors.
- **Returns:** The function returns 1 if any step of the process fails. If everything proceeds smoothly, there is no explicit return statement, so the function's status is the exit status of the last command executed (per bash's standard behavior).

- **Example usage:** `shell update_distro_iso "x86-Ubuntu-16.04"`

11.6.894 Quality and security recommendations

1. Validate that `DISTRO_STRING` is not empty at the beginning of the function to avoid unnecessary operations.
2. Consider checking that the manual ISO update was successful after the `Press ENTER when ready to re-mount...` statement.
3. Implement file existence and readability checks before trying to unmount or mount ISO files to avoid potential errors.
4. Consider using explicit return and exit codes to help clarify the function's behavior during debugging.
5. It's a good practice to use proper sanitization to free shell commands from potential code injection attacks. Ensure that user inputs are sanitized and safe before constructing paths based on them.

11.6.895 update_dns_dhcp_files

Contained in `lib/functions.d/dns-dhcp-functions.sh`

Function signature: `03499dd691020e6c8242d0a4a804e259ee89585ebf7a058cf6c54eefbd9da8f7`

11.6.896 Function Overview

The Bash function `update_dns_dhcp_files` is designed to update DNS (Domain Name System) and DHCP (Dynamic Host Configuration Protocol) configuration files in a network environment. It does this by utilizing other internally defined functions, logging information, and error messages to keep track of the process. If either of the file update operations (DNS or DHCP) fails, the function will return an error. If both operations succeed, the `dnsmasq` service is reloaded to apply the new configurations, and the function successfully exits.

11.6.897 Technical Description

- **Name:** `update_dns_dhcp_files`
- **Description:** This bash function updates DNS and DHCP configuration files. The update is considered successful when the DHCP addresses file and DNS hosts file both build without errors. If either file fails to build, the function logs an error message and a failure status is returned. If both files build successfully, `dnsmasq` is reloaded to pick up the new configuration.
- **Globals:** None.
- **Arguments:** None.
- **Outputs:** Logs [INFO] and [ERROR] messages using the inbuilt `hps_log` function.
- **Returns:**

- 0 If updating both DNS and DHCP files is successful.
- 1 If building either DNS hosts file or DHCP addresses file fails.
- **Example usage:** The function is used without any arguments i.e
`update_dns_dhcp_files`

11.6.898 Quality and Security Recommendations

1. Consider using more detailed and unique error messages to assist debugging process and increase maintainability.
2. Techniques such as input validation and data sanitization should be implemented to increase security.
3. Always prefer using local variables inside a function to avoid side effects and accidental modification of global variables.
4. For long running processes, consider using methods to keep the user informed of the progress instead of simply running in the background.
5. Implement logging levels in `hps_log` function to control the verbosity of the logs in different environments (production, staging, development etc).
6. It might be helpful to wrap the `dnsmasq` service reload in a try-catch block to handle unexpected errors with the service restart.

11.6.899 `url_decode`

Contained in `lib/functions.d/hps_log.sh`

Function signature: `5975a81b62dcd9f26bb0618375a64252646f6ad04c8687e2bce4da8abea4dae4`

11.6.900 Function overview

The `url_decode` function is a bash function that decodes URL-encoded strings. The function operates by first replacing any '+' characters with spaces, then replaces any '%' characters with their corresponding ASCII values. This decoded string is then printed out.

Function's primary usage is to work with decoded URL data. It's mostly used to handle and interpret data gotten from URLs or to make such data more humanly readable.

The other part of the code is not part of the `url_decode` function, but provides example usage. This code decodes a message, sends it to syslog, and writes it to a file if possible.

11.6.901 Technical description

- **Name:** `url_decode`
- **Description:** This function decodes a URL-encoded string.
- **Globals:** None

- **Arguments:** [\$1: String to be decoded]
- **Outputs:** Decoded string
- **Returns:** None. The function directly prints to stdout.
- **Example Usage:**

```
msg="[ $(hps_origin_tag) ] ( $(detect_client_type) ) $(url_decode  
↪ "$raw_msg")"  
logger -t "$ident" -p "user.${level},," "${FUNCNAME[1]}"  
↪ $msg"
```

11.6.902 Quality and security recommendations

1. Escape User Inputs: Always escape user-supplied inputs in the logging systems as failure to do so can lead to injection attacks or the printing of sensitive information.
2. Secure the Log Files: If sensitive information is being logged, make sure the log files are secure and are not readable by unauthorized users.
3. Error Handling: When the function fails to write to `$logfile`, it should handle such a failure more gracefully than just logging the failure message.
4. Avoid Globals: While no global variables are used here, always minimize the use of them in bash scripts.
5. Return Values: Even though this function is printing to the stdout, it could also return the decoded string for more flexibility.

11.6.903 url_encode

Contained in `lib/host-scripts.d/common.d/common.sh`

Function signature: 244a7c143b13a17ef50725eef748c82a2bd2df462a35d3727b81937558bae3ec

11.6.904 Function Overview

The `url_encode` function is used to encode URLs in Bash. The function accepts a string as an input and returns a URL-encoded string as output. It scans through each byte of the input string and performs a case-by-case encoding. Specifically, all alphanumeric characters and the special characters `.`, `_` and `-` are left as they are. Spaces are replaced with `'%20'` hex value for fast processing, and all other characters are replaced with their respective `'%'` followed by their hex value.

11.6.905 Technical Description

- **Name:** `url_encode`

- **Description:** This function is used for URL encoding in Bash. It accepts a string, iterates through its characters and replaces each character based on particular rules.
- **Globals:** N/A
- **Arguments:** \$1: input string for url encoding.
- **Outputs:** URL-encoded string.
- **Returns:** 0
- **Example usage:** `s="HELLO World!/?/#" url_encode "$s"`

11.6.906 Quality and Security Recommendations

1. Consider using a well-established library or built-in function for URL encoding if available, as there might be edge cases we are not considering in this implementation.
2. Make sure you always encode URLs when using them as part of commands or requests, as failure to do so could lead to command injection or other security vulnerabilities.
3. It is essential to write unit tests for this function to make sure it works as expected. Consider testing with all types of characters including alphanumeric and special characters.
4. Avoid the use of the variables with generic names like 's', 'out', 'c', 'i' to improve readability.
5. Consider handling the case of receiving more than one parameter by showing an error message.
6. Document the purpose of disabling SC2039 ShellCheck warning at the beginning of the function.

11.6.907 `urlencode`

Contained in `lib/functions.d/cgi-functions.sh`

Function signature: `f758d39e7a343eef82fc4e92ae1358118cf9a79d8accf9f5013313b5448282ac`

11.6.908 Function Overview

The `urlencode` function is a utility function used to encode a string by converting certain characters into their hexadecimal representation prefixed by `%`. The function operates by iterating over each character in the source string and checking if it falls within certain ranges. If it doesn't, that character is replaced by its encoded form.

11.6.909 Technical Description

- **Name:** `urlencode`

- **Description:** The `urlencode` function is designed to encode a string by replacing certain characters with their URL-encoded form based on the ASCII character set. For every character not in the set `[a-zA-Z0-9.~_-]`, it is replaced with its hexadecimal ASCII value prefixed by `%`.
- **Globals:** None
- **Arguments:**
 - `$1`: This is the string that needs to be URL-encoded.
- **Outputs:** The function prints the URL-encoded version of the input string.
- **Returns:** None
- **Example Usage:**

```
$ urlencode "Hello, World!"  
Hello%2C%20World%21
```

11.6.910 Quality and Security Recommendations

1. Validate the inputs: Before processing, validate that the input is in fact a string and not any other data type to avoid errors during and unexpected results from processing.
2. Error handling: The function currently lacks error handling. Add an error message or an error code to handle situations where an invalid input is given.
3. Unit testing: Ensure each piece of this function is adequately tested, both with expected and unexpected inputs to ensure it behaves as expected in all scenarios.
4. Use of `local`: This function appropriately uses `local` variables to ensure that they do not clash with variables outside of the function. This should be continued for any new variables introduced into the function.
5. Security: The function as is does not have any direct security concerns. However, always be aware of potential security risks when dealing with URL encoding, especially in web development contexts where URL encoded strings can sometimes be manipulated by malicious actors.

11.6.911 `validate_alpine_repository`

Contained in `lib/functions.d/repo-functions.sh`

Function signature: `40cdbdfa69d128c0e53275f522f4913827561f596df427e85b5d208433153364`

11.6.912 Function overview

The `validate_alpine_repository` function is used to validate alpine repositories set on a given directory (defined by `HPS_DISTROS_DIR`). It first checks if the environment variable `HPS_DISTROS_DIR` is set; if not, an error is logged and the function returns. If the Alpine version is not provided as an argument, it attempts to auto-detect one. Next, it builds the repository directory based on provided or auto-detected Alpine version and `HPS_DISTROS_DIR`. It then validates the existence of the

repository directory and its APKINDEX file. Finally, it checks if the number of packages in the repository exceeds a defined minimum expectation.

11.6.913 Technical description

- Name: `validate_alpine_repository`
- Description: Validates a set Alpine repository.
- Globals: [`HPS_DISTROS_DIR`: Location of the Alpine distribution directories]
- Arguments: [`$1`: Alpine version. If not provided, `get_latest_alpine_version` is used to find it, `$2`: name of repository, defaults to `main` if not provided]
- Outputs: Logs detailing the validation process and validation outcome.
- Returns: 0 if validation is successful; 1 otherwise.
- Example usage: `validate_alpine_repository 3.9 main`

11.6.914 Quality and security recommendations

1. Enforce strict argument checking to ensure all necessary inputs are provided, and they are in the correct format.
2. Implement a feature that handles unexpected errors or exceptions to prevent potential security risks or system crashes.
3. Add more detailed logging for each step in the function to facilitate easier debugging and maintenance.
4. Consider using more secure methods for file and directory checking to avoid potential security vulnerabilities associated with file traversals.
5. Protect the function against possible command injection attacks by validating and sanitizing input arguments.
6. Use more robust error handling techniques to provide useful feedback when the function encounters any errors during execution.
7. Secure the repositories with digital signatures or hashes to ensure their authenticity and integrity.

11.6.915 `validate_apkindex`

Contained in `lib/functions.d/repo-functions.sh`

Function signature: `8b01a88e6110ddb2342de0b3ad1c7ac2c81b45e6b739b8d49cb390eeb12b97d9`

11.6.916 Function Overview

The `validate_apkindex` function is designed to check an `APKINDEX.tar.gz` file in a given directory. This file is a compressed tag-separated file used by Alpine Linux package management to store metadata about packages in an Alpine repository. It first checks whether such a file exists within the given directory, and then checks if said file is corrupt

or not. If the file does not exist or has been corrupted, it will return an error. Once the file has been validated successfully, the function returns 0.

11.6.917 Technical Description

- ****Name****: ``validate_apkindex``
- ****Description****: Validates the availability and integrity of the APKINDEX.tar.gz file
- ****Globals****: None
- ****Arguments****:
 - ``$1``: ``repo_dir`` - The directory where the APKINDEX.tar.gz file is expected to be.
- ****Outputs****: Logs either a successful validation of APKINDEX.tar.gz, or logs errors
- ****Returns****:
 - ``2`` if the APKINDEX.tar.gz file is not found or is corrupted.
 - ``0`` if the APKINDEX.tar.gz file is successfully validated.
- ****Example usage****:

```
```bash
validate_apkindex "/path/to/directory"
```

### ### Quality and Security Recommendations

1. Provide clear and explicit error messages that can be acted upon without revealing s
2. Consider adding file permissions checks to ensure that the file can be accessed by th
3. Where possible, avoid using global variables to avoid potential conflicts and increa
4. Explicitly declare input expectations to help prevent potential manipulation and mi
5. Always exit with a non-zero status code when a failure occurs to allow other scripts
6. Regularly check for and manually handle potential errors and exceptions in your scri

### ### `validate\_hostname`

Contained in ``lib/functions.d/network-functions.sh``

Function signature: `1b48bafa4d6a9144287a36e810a96b33fe5fc9b78c467e672aa1e7c0186f54b`

### ### Function overview

This bash function, ``validate_hostname``, undertakes the task of validating a hostname

### ### Technical description

#### **\*\*Function details:\*\***

- **\*\*Name\*\***: ``validate_hostname``
- **\*\*Description\*\***: Verifies if the given hostname complies with the permissible conditions.
- **\*\*Globals\*\***: None
- **\*\*Arguments\*\***:
  - ``$1: hostname`` - Hostname to be validated.
- **\*\*Outputs\*\***: No explicit output; all outputs are implied through return codes.
- **\*\*Returns\*\***:
  - ``0`` - If the hostname complies with all conditions.
  - ``1`` - If the hostname does not comply with any condition.
- **\*\*Example usage\*\***: ``validate_hostname google.com``

### ### Quality and Security Recommendations

1. Consider using explicit error messages to elaborate on the reason for validation failure.
2. Use principles of least privilege for any direct access to system level resources, such as files or network interfaces.
3. Keep an eye on performance and compliances with large input values.
4. Implement unit tests to ensure the functionality of this function is as expected.
5. Include a logging mechanism to audit the program's activity which helps in troubleshooting.

### ### ``validate_ip_address``

Contained in ``lib/functions.d/network-functions.sh``

Function signature: `e46beef7993583a0fd0da40d09a03ba29d295b6eea7552ec53c3a8434a1b7eb`

### ### Function Overview

The function ``validate_ip_address`` is a bash function that is used for validating that a string is a valid IP address.

### ### Technical Description

- **\*\*name:\*\*** `validate_ip_address`
- **\*\*description:\*\*** Validates that a string is a valid IP address.
- **\*\*globals:\*\*** [ No global variables used ]
- **\*\*arguments:\*\*** [ ``$1: String to validate as IP address`` ]
- **\*\*outputs:\*\*** None
- **\*\*returns:\*\*** Returns 1 if the IP is not valid, and 0 if the IP is valid.
- **\*\*example usage:\*\***

``bash``

```
validate_ip_address "127.0.0.1" # Valid; returns 0
validate_ip_address "999.0.0.1" # Not valid; returns 1
```

### 11.6.918 Quality and Security Recommendations

1. **Input Validation:** The function should handle edge cases where the input is either empty or an invalid data type. This prevents unexpected behavior and potential script vulnerabilities.
2. **Error Messaging:** Instead of simply returning 0 or 1, consider including an informative error message if the provided IP address is invalid. This would make the script more user-friendly and easier to troubleshoot.
3. **Documentation:** Each section of the code should be well commented for easier maintenance and readability. This includes the function's purpose, its input and output, and how it handles different conditions.
4. **Unit Tests:** Create unit tests to cover different edge cases and function behaviors. This will ensure the function works as expected and helps identify bugs.

### 11.6.919 verify\_checksum\_signature

Contained in `lib/functions.d/iso-functions.sh`

Function signature: 6927fb684cf8e6a1fa690734435cbce8b702cdeef1affb7a199413ba1ed2e406

### 11.6.920 Function Overview

The `verify_checksum_signature` function takes four positional arguments to define the type of ISO file to check and the location from where to fetch the ISO's CHECKSUM and GPG key. This function is designed to verify the integrity and authenticity of a downloaded ISO file according to the checksum and GPG signature fetched from the file's distribution server.

Currently, the function only supports the Rocky Linux distribution.

### 11.6.921 Technical Description

- **Name:** `verify_checksum_signature`
- **Description:** This function verifies a specified ISO checksum and GPG signature to ensure the authenticity and integrity of the file.
- **Globals:**
  - `HPS_DISTROS_DIR` (Defaults to `/srv/hps-resources/distros`): Specifies the directory of the distros.
- **Arguments:**
  - `$1 (cpu)`: Specifies the CPU architecture.
  - `$2 (mfr)`: Specifies the manufacturer.
  - `$3 (osname)`: Specifies the operating system name.



- `$4 (osver)`: Specifies the operating system version.
- **Outputs:** Various status updates printed to the console. Error messages are redirected to standard error.
- **Returns:** 0 on success, 1 if there is an error (such as the ISO not being found, a failed download, checksum mismatch, or signature verification failure).
- **Example Usage:**

```
bash verify_checksum_signature "x86_64"
"rocky" "rockylinux" "8"
```

### 11.6.922 Quality and Security Recommendations

1. The function could benefit from input validation to ensure the provided `cpu`, `mfr`, `osname`, and `osver` arguments are in the expected formats before they are concatenated into URLs.
2. The function may need to be modified to handle ISO files for operating systems other than Rocky Linux.
3. Consider the implementation of stronger error handling rather than simply returning 1 on an error. Detailed and distinct exit codes may help to better identify specific issues that could occur.
4. The echoing of status updates could be optionally silenced for running the script in a quiet mode.
5. Temporary directories and files should be securely deleted to prevent sensitive information from being exposed on the server.
6. Implement HTTPS download error handling to improve security and stability. This could be accomplished within the `curl` command statements.

### 11.6.923 `verify_required_repo_packages`

Contained in `lib/functions.d/repo-functions.sh`

Function signature: `a16392408f8b201975834d3d38029e3859302d45f825b35156dca0ae85e71609`

### 11.6.924 Function Overview

The `verify_required_repo_packages` function checks for the presence of certain required packages in a specified repository path. It is intended for use with RPM-based package repositories. The function takes the repo path as the first argument, then a list of required package names. If the function cannot find a required package in the repository, it logs an error message and returns a value of 1 or 2, depending on the type of error. If all required packages are present, it logs an informational message and returns a zero value.

### 11.6.925 Technical Description

- **Name:** `verify_required_repo_packages`

- **Description:** This function inspects a specified package repository and verifies the presence of required packages. It is used for checking the availability of certain essential packages in a RPM repository.
- **Globals:** None.
- **Arguments:**
  - \$1: `repo_path` - The path to the package repository.
  - \$2...: `required_packages` - An array of package names that the function will check for in the repository.
- **Outputs:** Logs error messages through `hps_log` function if required packages are missing or repository path is not provided. Logs an informational message if all required packages are present.
- **Returns:**
  - 0 if all required packages are present in the specified `repo_path`.
  - 1 if `repo_path` not provided or does not exist.
  - 2 if any of the required packages are missing.
- **Example Usage:** `verify_required_repo_packages "${HPS_PACKAGES_DIR}/${DIST_STRING}/RHEL7/`  
`zfs opensvc`

### 11.6.926 Quality and Security Recommendations

1. It is recommended to use absolute paths for `repo_path` to avoid any ambiguity or errors derived from relative paths usage.
2. Ensure the proper access rights are in place for the directory path specified by `repo_path`, so the function can process the commands effectively.
3. Always sanitize input given to the function to prevent potential security vulnerabilities, such as command injection.
4. Consider adding further error checking mechanisms, like checking if each package name in `required_packages` is a non-empty, non-null string.
5. Implement a feature to handle version-specific packages in the array `required_packages`. Currently, it assumes that the requirement is met if any version of the package exists.
6. Leverage the return status of the function to handle error situations in the script that calls this function.

### 11.6.927 `verify_rocky_checksum_signature`

Contained in `lib/functions.d/iso-functions.sh`

Function signature: `63d87412289beb590b23dead3edfbc2afb85d4ee141526d80e8ad9104dcfbcb`

### 11.6.928 Function overview

The `verify_rocky_checksum_signature` is a Bash function designed for verifying the checksum signature of a specified Rocky Linux ISO image. It accomplishes this

by first downloading the checksum and signature associated with the provided version of the Rocky Linux ISO image, then importing the GPG key from the Rocky Linux server, and finally, verifying the GPG signature. If all passes, the function will then calculate and compare the sha256 checksum of a specified ISO file to ensure the ISO image has not been tampered with.

### 11.6.929 Technical description

- Name: `verify_rocky_checksum_signature`
- Description: This function downloads and verifies the checksum and its signature of a Rocky Linux ISO image. Additionally, it validates the SHA256 hash of the ISO file against the downloaded checksum.
- Globals:
  - `HPS_DISTROS_DIR`: A directory path where Rocky Linux distribution files are stored.
- Arguments:
  - `$1 (version)`: The version of Rocky Linux ISO image to verify.
- Outputs: The function will output various status messages indicating the status of download, GPG key import, and verification steps.
- Returns:
  - 0: Success. The checksum signature has been verified and matches the checksum of the specified ISO image.
  - 1: Failure. The GPG key import process was unsuccessful.
  - 2: Failure. The checksum signature verification failed.
  - 3: Failure. No checksum was found for the specified ISO image.
  - 4: Failure. The checksum of the ISO file does not match the downloaded checksum.
- Example Usage:
  - `verify_rocky_checksum_signature "8.4"`

### 11.6.930 Quality and security recommendations

1. Always ensure to use secure and up-to-date versions of all tools and packages used in the function.
2. Instead of hardcoding the architecture (`x86_64`), consider making it an argument or an environment variable that is configurable by the user.
3. Consider checking whether the `curl` and `gpg` commands succeed immediately rather than using a mixture of exit status checks and error redirections.
4. Employ better error handling and provide more specific output messages in case of failures.
5. Avoid reassigning constants in the middle of the function, such as `checksum_path` and `sig_path`.

6. Uncomment the downloading commands for CHECKSUM and CHECKSUM.sig or explain why they are commented out.

### 11.6.931 write\_cluster\_config

Contained in `lib/functions.d/cluster-functions.sh`

Function signature: `cf6f48116ee4f6ae2fa075ed74a4476f6c22ffe594564ebf704d40d3dce09039`

### 11.6.932 Function overview

The `write_cluster_config` function is designed to write a series of values (array) into a targeted configuration file. The function starts by checking the length of the array of values, and if empty, outputs an error message and returns 1 (indicating an error occurred). If the array is not empty, the function prints the values to the terminal and then writes the inter-space-separated values into the targeted file.

### 11.6.933 Technical description

- **Name:** `write_cluster_config`
- **Description:** Writes an array of values to a targeted configuration file. Reports an error and returns 1 if the array is empty. Otherwise, prints the array of values to the screen and writes them to the file.
- **Globals:** None
- **Arguments:** [ \$1: Target file for writing the array, \$2: The array of values ]
- **Outputs:** “[x] Cannot write empty cluster config to \$target\_file” to stderr if the array is empty; otherwise, “Writing: \${values[\*]}” and “[OK] Cluster configuration written to \$target\_file” to stdout.
- **Returns:** Returns 1 if the array is empty. Does not explicitly return a value otherwise.
- **Example usage:**

```
write_cluster_config "config.txt" "value1" "value2" "value3"
```

### 11.6.934 Quality and security recommendations

1. Consider validating file write operations: While the function currently reports whether a configuration file is written, it could potentially add error checking for the file write operation to catch and report errors.
2. Input validation: More robust validation of input arguments (such as checking if \$1 is a valid file path) will help prevent accidental misuse of the function.
3. Atomic writes: Consider using atomic write operations to prevent potential race conditions or half-written files in case of errors or interruptions during write operations.

4. Secure handling of error messages: Rather than writing error messages to stderr, consider logging them securely in a way that would not expose potentially sensitive information.
5. Sanitization of inputs: Always sanitize inputs especially if they are used as part of a command to be executed to avoid command injection vulnerabilities.

### 11.6.935 `zfs_get_defaults`

Contained in `lib/host-scripts.d/common.d/zpool-management.sh`

Function signature: `7075829ac859fa1a3d095289b6f87eeae5bfd9df08c9d1cfd66df7de132cf128`

### 11.6.936 Function Overview

The function `zfs_get_defaults` is designed for setting sensible defaults for pool options (`_POOL_OPTS`) and ZFS properties (`_ZFS_PROPS`). This function helps to customize and optimize your ZFS filesystem according to your needs. Pool options include sector size, while ZFS properties cover settings such as compression type, access time, extended attribute style, Access Control List (ACL) type, mode, inheritance, node size, and log bias.

### 11.6.937 Technical Description

- Name: `zfs_get_defaults`
- Description: This function sets default values for pool options (`_POOL_OPTS`) and ZFS properties (`_ZFS_PROPS`). The pool options are safest for SSD/NVMe/HDD with 4K-sectors. The ZFS properties include features like compression, access time, ACLs, and others.
- Globals:
  - `_POOL_OPTS`: An array to store pool options.
  - `_ZFS_PROPS`: An array to store ZFS properties.
- Arguments:
  - `$1`: A reference to a variable intended to hold pool options.
  - `$2`: A reference to a variable intended to hold ZFS properties.
- Outputs: There are no explicit outputs besides the modified `_POOL_OPTS` and `_ZFS_PROPS` variables.
- Returns: This function does not have any explicit return values and does not generate exit status.
- Example usage: `zfs_get_defaults POOL_OPTS ZFS_PROPS`

### 11.6.938 Quality And Security Recommendations

1. Consider making the function's name more descriptive, such as `set_zfs_defaults`, to specify the action that the function performs.

2. For improved security, validate the variables that are passed into the function to ensure they allow item assignment.
3. When setting sensible defaults, remember to base it on the actual use case scenario and the nature of the data handled.
4. To enhance transparency and ease of debugging, consider implementing a logging system to record any changes made by the function.
5. The comments flagged with 'TODO' should be addressed. Consider implementing the -O props within the function as these can provide an additional level of customization for the user.
6. Always include error checking in your functions to catch unexpected scenarios and improve robustness.

### 11.6.939 `zpool_create_on_free_disk`

Contained in `lib/host-scripts.d/common.d/zpool-management.sh`

Function signature: `ff0bf00dc7c2ed8816d8a88a6b148a5993ba6e4c5d1ae70415c5960612576a80`

#### 11.6.940 1. Function overview

The `zpool_create_on_free_disk` function is a Bash utility for working with data storage in a Unix-like operating system. It initializes a zpool on free disk space using a specified strategy. It has preset values for variables such as `strategy` (defaults to "first"), `mpoint` (defaults to "/srv/storage"), `force`, `dry_run`, and `apply_default` all of which can be adjusted according to user requirements.

#### 11.6.941 2. Technical description

- **Name:** `zpool_create_on_free_disk`
- **Description:** A bash function designed to initialize a zpool on available disk space. This function allows users to set their preference with a number of preset options to customize the setup according to their needs. The possible options include specifying a storage strategy, the mount point for the storage, and whether to force the operation, do a dry run or apply default settings.
- **Globals:**
  - `strategy`: Defines the strategy for initializing the zpool storage. Default is "first".
  - `mpoint`: The directory in which the initialized storage will be mounted. Default is "/srv/storage".
  - `force`: Defines whether or not to force the initialization. Default is 0 (do not force).
  - `dry_run`: Defines whether or not to perform a dry run. Default is 0 (do not dry run).

- `apply_defaults`: Defines whether to apply the default settings. Default is 1 (apply defaults).

- **Arguments:**

- \$1: The first positional parameter is not explicitly used in this function.
- \$2: The second positional parameter is not explicitly used in this function.

- **Outputs:** Outputs the status of the zpool creation process.

- **Returns:** Returns a status code indicating the success or failure of creation.

- **Example usage:** To use this function, you would typically include in a Bash script like this:

```
zpool_create_on_free_disk
```

### 11.6.942 3. Quality and security recommendations

1. Validate input: Although this function does not take arguments, it is always good practice to ensure any input or sources from which input is derived are valid.
2. Error handling: Include error handling mechanisms for situations where the disk space is not available or the formatted storage cannot be mounted at the specified mount point.
3. Security: Ensure that the storage's mount point has appropriate permissions set, and sensitive data is securely managed.
4. Code clarity: Some variables are initialized but not used, these could be removed for improved code clarity.
5. Logging: Add comprehensive logging for tracking the sequence of operations. Logs would aid in debugging and resolving any issues that might arise.
6. Code Comments: Adding comments to explain the purpose of complex code blocks would make the function more maintainable.

### 11.6.943 `zpool_name_generate`

Contained in `lib/host-scripts.d/common.d/zpool-management.sh`

Function signature: `5995908b1c71b7b0931db1a09cf94c2257d6d0ed783b7e45ca8926f62b975cf6`

### 11.6.944 Function Overview

The `zpool_name_generate` function is a Bash function that generates a ZFS pool (zpool) name based on certain parameters. It specifically takes an input, "class," and produces a zpool name incorporating information about the type of storage (like NVMe, SSD, HDD, ARC, or MIX), the cluster name, current Unix timestamp, and a random 3-byte hexadecimal identifier.

### 11.6.945 Technical Description

- **Name:** `zpool_name_generate`

- **Description:** This function receives a “class” parameter, which represents the type of storage, and generates a unique zpool name that includes the storage type, cluster name, Unix timestamp, and a random hexadecimal identifier.
- **Globals:** None
- **Arguments:**
  - `$1 (class)`: an indicator of the storage class. It could be either of `nvme`, `ssd`, `hdd`, `arc`, `mix`. If it’s not set or doesn’t match these, the function will return an error.
- **Outputs:** This function outputs a zpool name to stdout.
- **Returns:**
  - Return code 2 if the class argument is not given or doesn’t match the expected values.
  - Return code of the `zpool_slug` function call if it fails.

- **Example Usage:**

```
$ zpool_name_generate nvme
```

### 11.6.946 Quality and Security Recommendations

1. The function doesn’t validate the cluster name from the `remote_cluster_variable` function. Such validation could be beneficial to avoid creating zpools with incorrect or malicious names.
2. The random three-byte value is generated using both `/dev/urandom` and the `RANDOM` variable. To maintain consistency, consider using only one of these methods. Using `/dev/urandom` would make it more random and secure.
3. Consider adding error checking for critical operations like `od` and `printf`, which could fail due to various reasons.

### 11.6.947 `zpool_slug`

Contained in `lib/host-scripts.d/common.d/zpool-management.sh`

Function signature: `18bbbffc19d9426de745e5b45fe837f9d50ec9443122924bb19421975f877cbc`

### 11.6.948 Function overview

The `zpool_slug()` function is designed to convert a given string into a slug, suitable for use in URL paths or IDs. This function takes up to two arguments: the string to convert and (optionally) the maximum length of the slug. It converts the string to lower case, replaces any non-alphanumeric characters with a hyphen and removes any consecutive or trailing hyphens. It then truncates the slug at the specified maximum length or at 12 characters if no length was provided.



### 11.6.949 Technical description

- **name:** `zpool_slug()`
- **description:** Converts a string into a slug of a specified or default length.
- **globals:** None.
- **arguments:** [ \$1: The string to be transformed into a slug, \$2: Maximum length of the slug. This argument is optional, with a default length of 12 if left unspecified.]
- **outputs:** Prints the slug to stdout
- **returns:** None.
- **example usage:** `zpool_slug "Example String" 10` will output `example-str`

### 11.6.950 Quality and security recommendations

1. Right now, there is no explicit handling of invalid input, such as non-string values. Adding type checking and error handling for these scenarios would improve robustness.
2. Consider adding input sanitization to prevent any potential security issues (although current implementation is already reasonably safe due to removal of all non-alphanumeric characters).
3. Ensure that maximum slug length is not overly restrictive and take possible unicode characters into account.
4. Write unit tests for this function to guarantee it behaves as expected and validate the slug creation logic.
5. Specify locale in the script to ensure consistent character conversion, as the current implementation uses the locale setting of the running environment which can lead to unexpected results.

